

University of Surrey
Faculty of Engineering and Physical Sciences
School of Computer Science and Electronic Engineering

MSc Dissertation

Evaluating Reinforcement Learning Algorithms for Portfolio Trading

A Reinforcement Learning Study on the S&P 500 ETF (SPY)

Fiyinfoluwa Akano

URN: 6962514

Supervisor: Dr. Nguyen

Submitted for the Degree of
Master of Science in Artificial Intelligence



September 2026

©2026 Fiyinfoluwa Akano

Statement of Originality

I confirm that this report and the work it presents are my own. Any ideas, data, images or text taken from the work of others, whether published or unpublished, and including any content generated by a deep learning/artificial intelligence tool, have been clearly identified and appropriately acknowledged in the text and bibliography.

I confirm that this report has not been submitted, either in whole or in part, for any other academic degree or professional qualification.

I understand and agree that the University may submit this work to TurnitinUK to check its originality.

Fiyinfo luwa Akano

September 2026

Abstract

An investor holding a stock-index fund faces the same decision each day: buy more, sell part of the position, or do nothing. This dissertation asks whether reinforcement learning, which learns by trial and error from rewards, can make these decisions more effectively than the simple strategy of buying and holding permanently.

Three algorithms are compared: REINFORCE and Deep Q-learning, implemented from scratch, and Proximal Policy Optimisation, from an open-source library. Each agent trades the S&P 500 exchange-traded fund (SPY) in monthly episodes with transaction fees, action masking, and a trade size calibrated on validation data. Every configuration is trained with six random seeds, and every experimental decision is made on validation data before the test data is used.

The work consists of two studies. The first trained the agents on real daily data alone: 60 monthly episodes from 2018–2022, tested on 2024–2025. No agent beat buy-and-hold, which earned +\$180.25 per month. A behavioural analysis showed why the returns could not be trusted: the two policy-gradient methods earned their returns from fixed always-buy patterns that never responded to market information, on twelve seeds out of twelve, while the only method whose actions did depend on its state, Deep Q-learning, earned the least. The cause is a data bottleneck: 60 mostly-rising episodes cannot train networks of 18,000 parameters beyond memorisation.

The main study removed that constraint. A market simulator with three controlled conditions — rising, falling, and flat — was calibrated on the same training years and generated 3,000 balanced episodes, and the networks were reduced to 2,800 parameters. Under these conditions the ranking inverted. Deep Q-learning earned +\$126.94 per episode, on all six seeds, on a balanced test set where no fixed trading pattern is profitable, and remained profitable at ten times the assumed transaction fee. The risk-aware reward, which had only suppressed trading in the first study, now cut falling-market losses by 73% at a 1.5% cost in mean return. Transferred unchanged to 180 real months that no training or tuning decision ever saw (2006–2017 and 2023–2025), the learned policies remained profitable on every seed, held zero exposure through October 2008, and earned a higher Sharpe ratio than buy-and-hold (0.224 against 0.198).

The dissertation contributes a verified evaluation framework, a behavioural test of whether a policy uses its state, a calibrated market simulator for controlled evaluation in each market condition, and evidence that the amount and balance of training data — not the choice of algorithm — set the limit on what reinforcement learning can learn in this setting.

Acknowledgements

First and foremost, I would like to thank God for giving me the idea for this project, the strength, wisdom and perseverance to see it through. What started as an initial idea developed significantly throughout this journey, and I am grateful to everyone who played a part in bringing it to where it is today. It is only by His Grace to Start, Continue, and Finish that this was possible.

I would like to thank the professors who helped me narrow down the original idea to shape it into a research problem that could be properly investigated. I am especially grateful to my supervisor, Dr. Nguyen, for his guidance throughout the project. His feedback helped me streamline the work, question my assumptions, and keep the research focused on what evidence could support. More importantly, his guidance helped me develop a much stronger understanding of reinforcement learning and how it can be applied to a problem as challenging as financial trading.

I would also like to express my deepest gratitude to my family, especially my mum and dad, for their constant encouragement and support. They kept pushing me to keep going until the final stage. I am equally grateful to my uncle, aunties, grandma, church circle, and everyone else who took the time to check in on me throughout the journey. Their encouragement and help, even in small moments, but especially in sick, tired moments, all meant a great deal to me.

Finally, I would like to give a special shoutout to my friends from the cohort: Katherine, Joel, Yumna, Akeeb, David, and Dr. Tim. The laughs, outings, and support throughout the year are unforgettable. Having people around who were going through the same experience made the long parts of the journey so much fun, shorter, and easier to get through.

To everyone who contributed to this journey in one way or another, thank you. I am genuinely grateful.

Use of Generative AI. Generative AI tools (ChatGPT and Cursor) were used during the preparation of this dissertation to assist with drafting and restructuring prose, literature organisation, LaTeX formatting, and debugging of experimental code. All research design choices, experimental results, analysis, and the final wording of this report are my own. AI-generated suggestions were reviewed, edited, and verified before inclusion. No AI tool is cited as a scholarly source.

Contents

Statement of Originality	i
Abstract	ii
Acknowledgements	iii
List of Figures	vi
List of Tables	vii
1 Introduction	1
1.1 Background and Motivation	1
1.2 Problem Statement	2
1.3 Research Questions and Objectives	2
1.4 Scope	4
1.5 Contributions	4
1.6 Dissertation Structure	5
2 Background and Literature Review	7
2.1 Reinforcement Learning for Sequential Decisions	7
2.2 Two Families of Learning Algorithms	8
2.3 Reinforcement Learning for Trading	11
2.4 Evaluating Trading Agents Honestly	14
2.5 The Gap Addressed In This Study	15
3 Methodology & Mathematical Formulation	18
3.1 Overview: Trading as a Markov Decision Process	18
3.2 Episode Definition and Data Structure	19
3.3 State Representation Design	21
3.4 Action Space Design	30
3.5 Reward Function Design	32
3.6 Objective Function	35
3.7 REINFORCE: Monte-Carlo Policy Gradient	37
3.8 Deep Q-Learning	38
3.9 PPO: Proximal Policy Optimisation	40
3.10 Comparison of the Three Learning Approaches	41
4 System Design and Experiment Framework	42
4.1 Introduction	42
4.2 Data Pipeline	43

4.3	Simulated Market Data	45
4.4	Training, Validation and Testing Procedure	47
4.5	Experiment Configuration	49
4.6	Implementation Verifications	51
4.7	Implementation Corrections	52
4.8	Complete Implementation Pipeline	53
5	Experimental Results and Analysis	54
5.1	Introduction	54
5.2	Experiment Protocols and Evaluation Criteria	55
5.3	The First Study: Real Data Alone	56
5.4	The Data Bottleneck	59
5.5	Fixed Trade Size Under the Simulated Data	59
5.6	Main Results: The Controlled Comparison Across Market Conditions	60
5.7	Effect of the State Representation	62
5.8	Effect of the Risk-Aware Reward	63
5.9	Transaction-Fee Sizing	66
5.10	Transfer to Real Markets	67
5.11	Discussion of the Main Findings	69
5.12	Chapter Summary	71
6	Conclusion and Future Work	72
6.1	Summary of the Work	72
6.2	Conclusions Against the Research Questions	73
6.3	Critical Appraisal of the Project	75
6.4	Limitations	77
6.5	Future Work	79
6.6	Closing Remarks	80
	References	81
A	Verification Tests	84
B	Algorithm Listings	87
C	Detailed Test Results	89
C.1	First Study: Wealth Change in Each Held-Out Month	89
C.2	Main Study: Wealth Change by Transfer Year	90
D	Glossary and Notation	91
D.1	Glossary of Terms	91
D.2	Notation	93

List of Figures

3.1	The agent–environment interaction loop for the trading MDP.	19
4.1	How the two studies use the real SPY months	45
4.2	System architecture showing the data flow described in Section 4.7.	53
5.1	First study: behaviour of each seed.	58
5.2	Main study: results by market condition and real-data transfer.	62
5.3	Risk-penalty sweep on the simulated validation episodes.	64
5.4	Transaction-fee sizing on the simulated validation episodes.	66
5.5	Cumulative wealth change over the 180 real transfer months.	69

List of Tables

3.1	Full State Definition	26
4.1	Generator parameters fitted on the 60 real training months.	47
5.1	First-study test results over the 24 held-out real months (2024–2025), benchmarks and the three learned agents in their final configuration, six seeds each.	57
5.2	Attainable exposure for each fixed trade amount on the simulated validation episodes. The policy buys whenever a purchase is valid.	59
5.3	Main study: mean wealth change per episode on the 600 balanced test episodes, the mean within each market condition, and the mean over the 180-month real transfer set.	61
5.4	Nine- against ten-feature state on the simulated test episodes (wealth-change reward, six seeds), with the real-transfer mean.	62
5.5	Risk-penalty size on the simulated validation episodes (Deep Q-learning, six seeds). The rule requires $\Delta W \geq +\$79.14$ and exposure ≥ 0.299 ; the largest value meeting both is $\lambda = 0.25$	64
5.6	Wealth-change versus risk-aware reward ($\lambda = 0.25$) on the simulated test episodes, ten-feature state, six seeds.	65
5.7	Transaction-fee sizing on the simulated validation episodes; six seeds. Trades are per episode. Fees are shown in basis points.	66
5.8	Transfer test on the 180 real months: overall and per window, with risk measures. Learned agents are six-seed means.	67

Chapter 1

Introduction

1.1 Background and Motivation

A private investor holding an S&P 500 asset faces a similar decision each trading day: buy more, sell part of the position, or do nothing. On most days, the choice may have little effect on the overall outcome. During a heavy market move, however, the decision can have a much larger impact.

March 2025 is one of such periods studied in this dissertation. Under the experimental setup used later in this study, it was the worst month for a simple buy-and-hold strategy, which lost \$398.41 from an initial \$10,000 investment, approximately 4% of the starting capital. An investor who reduced their position before the decline could have avoided part of this loss. However, reducing positions too often would also mean missing gains during the other months, which produced an average monthly gain of more than \$200.

The problem is therefore not just when to invest, or deciding when the market will rise or fall, but how to make a series of decisions about an asset's position as market conditions change, as each decision made affects the position long-term and the consequences of that action will only become clear several days or weeks later.

This makes trading a sequential decision-making problem, which is one reason reinforcement learning (RL) [1] is appropriate for the task. RL allows an agent to learn by interacting with an environment, observing the results of its actions, and receiving rewards over time based on those results. In essence, it learns a policy/strategy that makes better decisions to receive more rewards.

Trading has several characteristics that make it suitable for this machine learning approach. The agent observes information about the market and its portfolio, selects from a set of possible actions such as buying, selling, or holding, and receives a measurable reward based on the change in portfolio wealth. As a result, reinforcement learning has been increasingly applied to financial trading [2, 3, 4].

However, another important problem arises. The Stock/Equity markets generally rise over the years rather than fall. Meaning that, in a rising market, a strategy that simply remains invested without any pulling out can generate a strong return without using any market information. A learned agent that produces a positive return therefore has not necessarily demonstrated, by itself, that it has learned a useful trading strategy in such markets.

Two questions become relevant. First, can the learned agent outperform the simple buy-and-hold strategy with which it is being compared? Second, does the agent actually use the information available to it when deciding what action to take? These questions are important when evaluating on data not used during training.

This dissertation therefore focuses on both the financial performance of reinforcement learning agents and whether their decisions show evidence of using the information provided to them.

1.2 Problem Statement

The main problem this dissertation addresses is whether reinforcement learning algorithms can learn a trading policy for a portfolio — containing only one asset in this case — that beats a simple benchmark strategy on unseen data.

However, financial performance alone is not enough to determine that an agent has actually learned a useful trading strategy. For example, an agent that buys whenever it can, would perform well in a generally rising market without using any of the market information provided to it. A positive return from such a policy could be incorrectly interpreted as proof that the agent learned when to buy, sell, or remain out of the market, whereas that was not the case.

Evaluating a learned trading agent therefore requires a second perspective alongside financial performance: **behaviour**. The study must analyse whether the agent's actions are a result of changes in the observed market and portfolio state. This helps to differentiate a policy that genuinely responds to the information it receives from one that simply follows a fixed trading pattern.

This dissertation therefore addresses two interrelated problems: whether reinforcement learning can produce better financial performance than simple benchmarks, and whether that performance stems from a learned policy that actually uses the information available to it.

A third problem sits underneath both of these: the data itself. A network of thousands of parameters trained on a few dozen historical months can memorise those months instead of learning a rule that carries over to new ones, and a market that mostly rises rewards the same behaviour in every agent: constant buying. Whether the training data can support learning at all is therefore examined as a question in its own right rather than assumed.

1.3 Research Questions and Objectives

The aim of this dissertation is to evaluate three reinforcement learning algorithms — REINFORCE [5], Deep Q-learning (DQN) [6], and Proximal Policy Optimization (PPO) [7] — and determine whether they can learn profitable and genuinely intelligent trading strategies for the SPY exchange-traded fund.

The evaluation is carried out as two studies. The **first study** trains the three algorithms on the real market months alone and evaluates them on held-out months; its negative results are then traced to the size and balance of that training data. The **main study** corrects the data: it trains the same algorithms on thousands of balanced simulated episodes, evaluates them under controlled rising, falling, and flat market conditions, and finally tests the trained policies, unchanged, on 180 real months they have never seen. The two studies are reported together because each explains the other: the first shows what fails at real-world data volumes, and the second shows that the failure belonged to the data rather than to the algorithms.

The study is structured around five research questions:

1. **Performance:** Do the learned agents outperform simple benchmarks on unseen data after transaction costs are taken into account?
2. **Behaviour:** Do the learned agents use the market and portfolio information provided in their state representation when making trading decisions?
3. **State representation:** Does adding a drawdown feature to the main nine-feature state change the behaviour learned or the financial performance?
4. **Reward:** Does a risk-aware reward improve the balance between return and drawdown, rather than simply encouraging the agent to trade less?
5. **Data:** Do the amount and balance of the training data explain the learned behaviour, and does training on large, balanced simulated data produce the state-dependent trading the real data could not?

These research questions are touched upon through eight measurable objectives, each traceable to a specific part of the dissertation

O1. Build a reproducible daily trading environment for SPY using monthly episodes, transaction fees, and action masking, and verify the environment (Chapter 4).

O2. Implement REINFORCE, DQN, and PPO and verify their implementations before running the trading experiments (Chapters 3 and 4).

O3. Carry out the first study: train the three algorithms on the real training months, compare them with the benchmark strategies on held-out months using six independently trained seeds, and apply behavioural analysis to determine whether the learned policies change their actions in response to the observed state (Section 5.3).

O4. Diagnose the negative results of the first study and trace them to the size and balance of the training data (Section 5.4).

O5. Build a market simulator with controlled up, down, and flat conditions, calibrated on the real training months only, and generate thousands of balanced simulated episodes (Section 4.3).

O6. Retrain the three algorithms on the simulated episodes and evaluate them against the benchmark strategies in each market condition, using the same behavioural analysis and the same six-seed protocol (Section 5.6).

O7. Repeat the supporting experiments under the corrected data: the trade-size calibration, the drawdown state feature, the risk-aware reward, and the transaction-fee sensitivity (Sections 5.5, 5.7, 5.8, and 5.9).

O8. Test the trained policies, unchanged, on 180 real months that no training or tuning decision ever saw, including the 2008 financial crisis (Section 5.10).

These objectives create a structured way to evaluate both the financial performance and the learned behaviour of the three reinforcement learning methods.

1.4 Scope

The study is deliberately narrowed to one asset to make the experiments more controlled and easier to analyse.

All experiments use the SPY asset, an exchange-traded fund that tracks the S&P 500 index, with daily data from 2006 to 2025. The first study uses the 2018–2025 period, divided chronologically into training (2018–2022), validation (2023), and held-out testing (2024–2025). The main study trains on a simulated dataset generated from a market model with up, down, and flat conditions, calibrated on the 2018–2022 training years alone; it supplies thousands of balanced episodes for the controlled study, while every real month outside the calibration window — 2006–2017 and 2023–2025, 180 months including the 2008 financial crisis — serves as its transfer test.

Each monthly episode starts with \$10,000 of capital and no asset holdings, and each transaction to open a position incurs a 0.05% fee. The available actions are discrete: the agent can buy a fixed fraction of the initial capital, sell the same amount, or hold. The action masking prevents invalid actions, such as attempting to sell when no position is held.

Three reinforcement learning algorithms are evaluated. REINFORCE and Deep Q-learning are implemented from scratch, while PPO is implemented using a popular library.

Several areas are intentionally outside the scope of this study. These include multi-asset portfolios, continuous position sizing, intraday trading data, and live deployment. Keeping the study within these boundaries limits the number of factors that can affect the results, making it easier to identify which specific factor causes a change in performance or behaviour during testing rather than multiple changes at once.

1.5 Contributions

This dissertation makes five main contributions:

1. A framework for evaluating trading agents.

The study uses a controlled process in which experiment variables, such as trade size and transaction cost, are set using validation data and finalized before moving on to the held-out test data. The trading environment, state features, rewards, and benchmark strategies are also verified through tests to identify and correct problems before final evaluation,

2. A behavioural analysis for testing state utilization.

A separate analysis was developed to determine whether a policy learned actually responds to the information in its state, rather than simply following a fixed pattern of available actions.

In this evaluation, steps are grouped by their action mask. When a policy selects different actions at steps with the same mask pattern, the difference cannot be explained by the available actions alone, thereby demonstrating a change in its action based on the current observed state. The analysis does not identify which individual state feature drives this dependence, nor does it make claims about the learned network's internal representations.

This reveals a difference between algorithms that would not be seen from financial return figures alone: all six Deep Q-learning seeds were state-dependent, while none of the twelve policy-gradient seeds (REINFORCE and PPO) were (Section 5.3).

3. An explanation for the first study's negative results.

None of the three learning algorithms outperformed buy-and-hold across the first study's 24 held-out test months. The behavioural analysis helps explain this outcome. The policies that achieved returns closest to the benchmark, REINFORCE, generally did so through a fixed pattern of high exposure, while the algorithm that did use its state information, Deep Q-learning, reduced market participation rather than successfully understanding market conditions to perform better.

Although the overall result is negative, evidence from the six training seeds per configuration, behavioural analysis, and peripheral experiments on state representation, reward function, and transaction costs supports it as a finding rather than a failure.

4. A calibrated market simulator for controlled evaluation.

The study builds a market generator with three controlled conditions — up, down, and flat — whose drift, volatility, and intraday range are fitted on the real training years alone. The generator supplies thousands of balanced training episodes and, more importantly, labelled test conditions: every test episode is known to be a rising, falling, or flat market, so the evaluation can state how each agent behaves in each condition rather than in whatever mixture one historical period happens to contain.

5. Evidence that the data, not the algorithms, set the limit.

Retraining the same three algorithms on 3,000 balanced simulated episodes, with a smaller network, inverts the real-data ranking. Deep Q-learning, the worst method on real data, becomes profitable and state-dependent on all six seeds, earning \$126.94 per episode on a test set where no fixed trading pattern is profitable. Transferred unchanged to 180 real months, its policies stay profitable on every seed and hold zero exposure through October 2008, the worst month of the financial crisis, earning a higher risk-adjusted return than buy-and-hold over the window. The policy-gradient methods collapse to fixed corner policies once always-buying stops paying. The comparison shows that the real-data failure was a property of 60 imbalanced training episodes, not of the learning methods.

1.6 Dissertation Structure

The dissertation is organized as follows. Chapter 2 reviews existing research on reinforcement learning for financial trading and portfolios and explains how this study's evaluation approach relates to previous

work.

Chapter 3 presents the mathematical formulation of the portfolio trading problem, describing the trading environment as a Markov decision process; the state, action, and reward formulas; the three reinforcement learning algorithms; and their optimisation functions.

Chapter 4 describes the system implementation, data pipeline, verification process, trading environment, and the corrections made during development.

Chapter 5 presents the experiment results and analyses them in relation to the five research questions. It reports the first study and diagnoses the data bottleneck behind its negative results. It then reports the main study: the controlled comparison across market conditions on the simulated data, the state, reward, and transaction-cost experiments repeated under the corrected data, and the transfer of the trained policies to 180 real months including the 2008 crisis.

Finally, Chapter 6 concludes the dissertation by summarising the main findings, critically evaluating the work, discussing its limitations and possible directions for future research.

Chapter 2

Background and Literature Review

This chapter introduces the main concepts and prior work on which the rest of the dissertation is based. It begins with an overview of reinforcement learning and explains why the portfolio trading problem can be formulated using the reinforcement learning framework. It then presents the two algorithm families used in this study: value-based methods, from Q-learning extending to Deep Q-learning (DQN), and policy-gradient methods, from REINFORCE to Proximal Policy Optimisation.

The second half of the chapter focuses on reinforcement learning in financial trading. It reviews findings from previous studies, particularly how researchers have evaluated trading agents and the limitations that commonly appear in this area of research.

The chapter also considers the market-efficiency perspective, which provides an important baseline for measuring trading performance—judging whether a learning algorithm can consistently find profitable trading opportunities.

The chapter concludes by bringing these points together and identifying the two gaps addressed by this dissertation. The first is whether financial performance alone is sufficient to establish that a trained trading policy actually uses the information available to it. The second is whether the historical record of a single asset contains enough data to train such a policy at all — a question the reviewed studies answer implicitly through their workarounds, and this study answers by experiment.

2.1 Reinforcement Learning for Sequential Decisions

Most machine learning taught at the master's level is based on supervised learning, where a model is given inputs together with the correct answers and learns to reproduce the desired output. Reinforcement learning (RL) works differently. Instead of being given the correct action, the learner, known as the **agent**, interacts with an environment and learns from the results of its own decisions [1].

This interaction is a repeated process. At each step, the agent observes the current state, selects an action, and receives a reward, then moves to the next state. The reward indicates how useful the action was, but it does not directly tell the agent whether the action was correct. The agent must therefore learn through experience which actions are likely to produce higher rewards over time.

Trading fits naturally into this framework. A trader looks at the markets, decides whether to buy or sell more units of shares, and then waits for the effect of that decision on his portfolio account. A decision made early in a month affects the portfolio permanently from that point onwards. The agent must therefore learn not only whether a decision was profitable, but also how earlier actions contribute to later outcomes. This is known as the **credit-assignment problem** and is an important part of the

mathematical formulation presented in Chapter 3.

The absence of explicit labelled answers also creates a practical issue. In supervised learning, an incorrect prediction can be compared directly with a known target during training. In reinforcement learning, however, there is no equivalent answer to compare against during training. An experiment can therefore run without producing an obvious error in the environment, reward calculation, state representation, or action handling and still learn the wrong behaviour. This was an important consideration in this project and is one reason the verification tests were developed for use before the main experiments, as described in Chapter 4 and Appendix A.

The broader history of reinforcement learning systems shows this approach can learn complex sequential tasks. TD-Gammon, for example, learned to play backgammon strongly through self-play [8]. Later, Deep Q-learning achieved human-level performance across several Atari games using raw image inputs [6]. The analogy with trading is not that financial markets are equivalent to games but that in both cases, the agent is not given the correct action for each situation. Instead, they learn from a numerical reward or game score. In a trading environment, financial outcomes such as changes in portfolio wealth play this role.

The standard mathematical framework for representing sequential problems is the **Markov Decision Process (MDP)**. Its foundations are closely tied to Bellman's work on dynamic programming [9], and Sutton and Barto present a modern reinforcement learning formulation [1]. An MDP represents a problem through states, actions, transitions to the next state, rewards, and a discount factor, assuming that the current state contains all the information needed to make the next decision. In financial markets, however, it is difficult to show that any collection of market observations can fully capture the information relevant to future prices.

For this reason, Chapter 3 uses a practical state representation rather than claiming that the Markov assumption is perfectly satisfied: it captures market, portfolio, and timing information rather than claiming to reconstruct the complete market state.

The study also examines this design experimentally by testing whether adding drawdown information to the state changes the behaviour or performance of the learned agents. The next section discusses the two main families of reinforcement learning algorithms considered in this dissertation.

2.2 Two Families of Learning Algorithms

The algorithms used in this dissertation belong to two different families of reinforcement learning methods. This section discusses the main ideas behind each family and explains why comparing them on the same trading problem is useful.

Value-based methods estimate the value of taking an action and derive a policy from those estimates, whereas policy-gradient methods select the action with the highest probability to optimise the policy itself. These differences matter because they offer two different learning mechanisms for the same trading problem.

2.2.1 Value-based methods: from Q-learning to DQN

The Value-based family approaches the question of what action to take indirectly. Rather than learning the policy directly, the agent learns how valuable different actions are in a given state and then selects the action with the highest estimated value.

This approach started from **temporal-difference learning** [10], which updates value estimates using the difference between sequential predictions rather than waiting until the end of an episode to observe the final outcome. **Q-learning** turned this idea into an off-policy method that directly estimates the action-value function [11]. Under appropriate conditions, including continued exploration of state-action pairs and suitable learning-rate decay, tabular Q-learning converges to the optimal action values for finite state spaces. This convergence property does not depend on the particular actions used while learning.

In tabular Q-learning, a separate value is stored for every state-action pair. This works well when the number of possible states is small, such as in a simple grid-world, but becomes impractical when the state contains continuous financial information. Two very similar market states would also be treated as completely separate entries in the table. As a result, experience gained in one state would not automatically help the agent in a similar state.

Deep Q-learning (DQN) [6] addresses this problem by replacing the table with a neural network that estimates the action values. The network can therefore generalise across similar states, allowing information to be shared, but it introduces instability. However, combining neural networks with bootstrapping and off-policy learning can make training unstable [12, 1].

DQN then introduces two important mechanisms to improve this stability. **Experience replay** [13] stores previous transitions in a replay buffer and samples them later in random minibatches during training. This reduces the strong correlation present in consecutive experiences, allowing each transition to contribute to multiple updates. A **target network** is also used as a periodically updated copy of the value network to construct the learning target, preventing it from changing with every network update as the parameters are being optimised.

These mechanisms are highly useful for financial data, where consecutive daily observations are naturally correlated. Both were implemented from scratch in this project, as described in Chapter 4. As a result, the implementation shows the underlying learning mechanics rather than treating the algorithm as a library component.

Several Other improvements have since been proposed to address the limitations of the original DQN formulation. For example, **Double Q-learning** [14] was developed partly to reduce overestimation bias in action values that can arise in standard Q-learning when the same estimates are used to select and evaluate actions.

Such modifications are relevant to the broader DQN literature or discussed as possible future improvements but are not included in this study's present implementation, allowing the core DQN method to remain relatively simple and easier to verify.

2.2.2 Policy-gradient methods: from REINFORCE to PPO

Policy-gradient methods take a more direct approach to the decision problem. Instead of first estimating the value of each action to derive a policy from those estimates, the agent learns a policy that maps states to actions and adjusts its parameters to increase the expected return. The policy is represented by a differentiable model, typically a neural network, whose parameters are updated for the objective.

This study uses the foundational policy-gradient algorithm **REINFORCE** [5]. One of its main advantages is its simplicity, as it makes the relationship between the policy, observed states, selected actions, and expected return relatively transparent. It does not require bootstrapping, experience replay, or a separate target network. The policy is updated using returns from complete episodes, making the learning process straightforward to understand. For these reasons, REINFORCE is used as the main policy-gradient implementation in this dissertation.

The main weakness of REINFORCE is the high variance of its updates. The learning signal is based on Monte Carlo returns, which can vary considerably between episodes. This can make learning slow and unstable when rewards are noisy.

Later policy-gradient methods introduced techniques to address this problem. Actor-critic methods use a learned value estimate as a baseline to reduce the variance of the policy update [15]. The actor represents the policy while the critic estimates value, allowing the learning signal to use an advantage rather than the raw return. **A2C** is one example of this approach [16]. **Generalised Advantage Estimation (GAE)** [17] is another way of controlling the bias-variance in the advantage estimate.

A second problem is the size of policy updates. An update that changes the policy too aggressively can cause training to become unstable and degrade performance sharply. Trust-region methods [18] address this by restricting how far the policy can change between successive updates. **Proximal Policy Optimisation (PPO)** [7] uses a simpler approach by clipping the policy update when it moves too far from the previous policy.

PPO is widely used in reinforcement learning applications, including trading frameworks such as FinRL [19], and is therefore included in this study as a library-based policy-gradient method for comparison while implementing REINFORCE from scratch. This allows for both a transparent baseline policy-gradient method and a more modern, stabilised method.

The trading environment also contains actions that are not always valid. For example, an agent should not be allowed to sell when it has no position or buy when there is insufficient cash. All three learning methods use **invalid-action masking** to prevent such actions from being selected while preserving only valid actions. The use of action masking in reinforcement learning, including its interaction with policy-gradient methods, has been examined by Huang and Ontañón [20]. In this dissertation, masking is part of the common environment rather than a difference between algorithms, ensuring that the comparison is made under the same set of trading constraints.

2.2.3 Why compare the two families?

The two algorithm families have different implementation trade-offs.

Value-based methods such as DQN can learn from individual transitions and reuse previously collected experience through replay. However, they rely on bootstrapping value estimates, meaning the agent learns partly from estimates produced by other learned estimates in the same learning process.

Policy-gradient methods instead optimise the policy directly with respect to the objective. In its basic form, REINFORCE provides a balanced policy-gradient estimate, but it requires new experiences for each update from fresh trajectories, and gradient estimates can suffer from high variance. PPO improves the stability of this family through constrained policy updates.

Neither family can therefore be assumed to be better for the portfolio trading problem considered here. Comparing them in the same environment and under the same evaluation procedure is therefore useful. If both families develop similar behaviour, this may suggest that the behaviour is mainly caused by the structure of the problem, reward, state representation, or environment. If their behaviour differs, the difference becomes evidence that the choice of learning method influences the resulting strategy.

This comparison becomes an important part of the experimental design in Chapter 5, as it is not limited to only asking which algorithm produces the highest financial return. It also asks whether the resulting policies learn fundamentally different ways of responding to the trading environment. The next section discusses how previous research applies these algorithms to financial markets.

2.3 Reinforcement Learning for Trading

The use of reinforcement learning (RL) for trading has been studied for many years, but the application comes with several well-known challenges. One of such early influential work is Moody and Saffell on direct reinforcement methods [2], which uses gradient-based methods to learn trading policies directly. Their work raised an important point: the reward should account for risk. If an agent is rewarded only for increasing profit, it may have a strong incentive to remain highly exposed to the market without explicitly rewarding control of downside risk.

This observation is highly relevant to this dissertation, giving us reason to run two experiments in Chapter 5: first, to examine what happens when the agents are trained using a simple wealth-change reward and the objective is tied only to portfolio growth; and second, whether adding a drawdown penalty to the reward changes their behaviour.

Drawdown is also introduced separately as an additional state feature. These two uses of drawdown act differently. Adding drawdown to the state changes the information available to the agent, whereas adding a drawdown penalty in the reward changes what the agent is encouraged to optimise. One changes what the agent can see, while the other changes what it is rewarded for. However, the specific drawdown penalty used in this dissertation is also an experimental design choice to be evaluated rather than a reward formulation from previous literature.

Keeping these changes separate allows their effects to be examined separately. The experiments determine whether either change affects the learned behaviour, rather than assuming what the resulting behaviour should be.

Later work extended these ideas using more complex models. Deng et al. [3], for example, combined direct reinforcement with deep recurrent representations. Jiang et al. [21] applied deep reinforcement learning to multi-asset portfolio management using cryptocurrency data, using a different formulation for the action, which represents portfolio allocation rather than a discrete buy, sell, or hold decision.

More recently, frameworks such as FinRL [19] have made it easier to build trading environments and apply these reinforcement learning algorithms to financial experiments, but it often defaults to PPO-based methods.

These frameworks are useful for experimentation, but they can also hide parts of the underlying learning process. This was one reason for implementing REINFORCE and DQN from scratch in this dissertation, allowing their behaviour and training mechanisms to be examined directly.

The study most closely related to the present work is Théate and Ernst [22]. They evaluate a DQN-based trading agent on a test portfolio set of 30 stocks, comparing it with buy-and-hold and traditional active strategies. Their results report that the DQN performs better on average than the benchmark strategies, although they include several active strategies rather than only passive buy-and-hold, which is more relevant to this dissertation. For SPY, which is also used in this dissertation, their reported Sharpe ratio for the DQN was 0.834, the same as buy-and-hold.

Across multiple stocks, they describe the agent's performance as identical or very close to passive strategies and suggest that the agent can tend towards a passive strategy when active trading becomes more uncertain. The variations between training runs showed that the agent traded less as transaction costs increased, eventually becoming passive. They also report non-negligible variance across identical training runs. These findings make Théate and Ernst study relevant to this study's evaluation.

There are, however, important differences between their experimental design and the one used here. Their study trains a separate agent for each stock, allows short positions, and uses a continuous test period, whereas this dissertation studies one ETF, does not permit shorting, and uses independent monthly episodes with a fixed trade size. Their benchmarks also differ from this study's buy-and-hold strategies, preventing a direct numerical comparison between the studies.

More importantly, although their results show that the learned agent can become increasingly passive, they do not directly test how the learned policy generates its decisions: whether the policy responds to its observations or simply follows a fixed trading pattern. Both behaviours could produce similar trading trajectories.

This difference is fundamental to this dissertation, as a policy can achieve a strong return without making meaningful use of market information, particularly when the asset always goes up over time. Chapter 5 therefore includes a specific behavioural analysis to determine whether the actions of each trained agent actually change when its observed state changes while the available actions remain the same.

Read together, the previous studies reviewed here vary widely in how RL-based trading is formulated and evaluated. Some focus on individual assets (Moody and Saffell [2], Deng et al. [3], and Théate and Ernst [22]), while others use multi-asset portfolios (Jiang et al. [21] and Guan and Liu [23]). Action representations range from discrete trading decisions [2, 3, 22] to continuous portfolio weights [21]. Reward functions also vary, from direct profit or return objectives [3, 21, 22] to risk-adjusted objectives such as Moody and Saffell's differential Sharpe formulation [2]. Most papers include transaction costs, but their assumptions and representations differ. Benchmark strategies also vary, from traditional trading strategies [22] to portfolio-selection methods [21].

There is also variation in the results across seed runs. Some studies acknowledge differences between training runs [22], but reported results are often based on individual runs. None of the reviewed studies also describe a procedure in which all experimental design decisions are selected using a validation period before being used on an unseen test period. Behavioural analysis is similarly limited, with previous work using methods such as trading trajectories [22] or feature attribution [23].

These differences make direct comparison between studies difficult and highlight an important issue: the evaluation procedure can be just as important as the choice of algorithm. The studies reviewed in this section were selected because they are particularly relevant to this dissertation's design, especially their use of RL algorithms, transaction costs, and benchmarks. Fischer's survey [4] provides a broader review of RL in financial markets.

2.3.1 Data requirements and simulated markets

There is one further issue the trading studies above rarely confront directly: how much data reinforcement learning needs. The landmark results of the field were all obtained with effectively unlimited experience. The Atari agents consumed tens of millions of frames per game [6], and AlphaGo sidestepped the shortage of human games entirely by generating its own training data through self-play [24]. A single asset's daily history offers nothing comparable. Eight years of daily bars is roughly 2,000 observations, and grouping them into monthly episodes leaves only dozens of episodes to learn from.

The common response in the literature is to train on generated data. Théate and Ernst, whose setting is the closest to this study's (a deep Q-network trading single stocks on daily bars), trained their agent "entirely based on the generation of artificial trajectories from a limited set of stock market historical data" [22]. The earlier studies solved the same problem differently: Moody and Saffell matched a tiny single-layer policy to scarce monthly data [2], while Deng et al. [3] and Jiang et al. [21] moved to minute and 30-minute bars, where a few years supply hundreds of thousands of observations. Open-source frameworks such as FinRL [19] and the Gym trading environments [25] structure the task so that any price series, real or synthetic, can drive the environment. Generators range from classical stochastic-process models, including regime-switching models — models that alternate between market conditions — built on hidden Markov chains [26], to learned generators such as time-series GANs [27]. A generator gives the experimenter two things a historical record cannot: as many training episodes as the learner needs, and control over which market conditions appear in training and in testing. The price is a domain gap, since no generator reproduces real markets exactly, so studies that train on synthetic data normally

keep a real held-out period as a transfer test.

This dissertation follows that pattern. Its first study trains on real monthly episodes and shows what goes wrong at that data volume; its main study trains on episodes from a calibrated generator with controlled market conditions and tests under those conditions, with 180 real months — every month outside the generator’s calibration window, including the 2008 financial crisis — retained as a transfer test. Chapters 3 and 4 define the generator, and Chapter 5 reports both studies.

The next section therefore focuses more specifically on how to evaluate trading agents and on the weaknesses that can arise when financial performance is considered without sufficient behavioural and experimental controls.

2.4 Evaluating Trading Agents Honestly

A review of reinforcement learning in finance [4], together with research on financial machine learning [28, 29], identifies several common problems that can make trading results appear stronger than they are. Four of these issues are relevant to this dissertation, and each one influenced the experiment design. These are not claims that every reviewed study states as problems; they are risks that an experiment should account for when being designed.

The first issue is **weak benchmarks**. A learned trading strategy can appear successful if it is compared with a benchmark that invests less capital or does not represent a realistic alternative. For this reason, this study uses buy-and-hold as the main benchmark, with the same initial capital as the learned agents. The trade size was also calibrated in Chapter 5 so the agents could reach a comparable level of market exposure. This keeps the comparison fair, as the agents are not disadvantaged simply by being unable to invest enough capital.

The second issue is **look-ahead bias**. Financial data are ordered in time, so treating it like randomly ordered observations and shuffling the data before splitting could place future information in the training set and produce overly optimistic results. The first study’s data are therefore divided chronologically: 2018–2022 for training, 2023 for validation, and 2024–2025 as the held-out test period. The main study keeps the separation stricter still: its transfer months made no training, calibration, or tuning decision at all, as described in Chapter 4.

The third issue is **ignoring transaction costs**. A strategy that trades frequently may look profitable when every trade is assumed to be free, but the same strategy could lose much of its return once real trading costs are included. In this dissertation, transaction fees are part of the trading environment and are therefore included while the agents are learning rather than only after training. Chapter 5 also varies the assumed fee to examine whether the learned strategies trade less when trading becomes more expensive.

The fourth issue is **backtest overfitting**. Repeatedly adjusting a strategy after looking at its test results gradually turns the test set into another training set. Bailey et al. [28] discussed this problem for investment backtests, and López de Prado [29] discussed methodology for controlling leakage and data reuse. The defence here is that all design choices were selected using the training and validation data, then

evaluated once on the held-out test months, without further tuning.

A related concern comes from deep reinforcement learning itself. Henderson et al. [30] show that reported **results can vary substantially across random seeds**, making a single successful training run a weak basis for comparison. Each configuration in this study is therefore trained with six different seeds, and the spread is reported alongside the mean rather than presenting one favourable run as representative.

Financial economics provides more context for understanding the results. Under the **efficient market hypothesis**, current prices should already reflect the information contained in past prices [31]. If this assumption holds, an agent whose market information consists only of price history should not be expected to consistently earn higher risk-adjusted returns than the market itself. In this study, buy-and-hold on the single fund provides the practical benchmark against which this expectation is tested.

Portfolio theory also emphasises that return cannot be considered separately from risk [32]. A strategy that earns less may still be useful if it achieves a large enough reduction in risk. The evaluation therefore reports mean return together with drawdown and the Sharpe ratio [33]. No single metric decides whether a strategy is successful; the results are read together so that the trade-off between return and risk remains visible.

This approach does not aim to guarantee a positive result for the learning algorithms. If an agent does not beat buy-and-hold, that outcome is treated as a finding to explain rather than a failure to hide. Understanding why a strategy does not improve on a simple benchmark is therefore just as important as reporting successful results. However, even a carefully controlled financial experiment only establishes what happened to the portfolio; it does not establish what the agent learned to do. That distinction leads to the gap considered in the next section.

2.5 The Gap Addressed In This Study

Financial performance alone does not tell us what a trading agent has actually learned. An agent can follow a simple rule, such as buying whenever possible, and still perform well during a rising market without using any market information. At the same time, an agent that genuinely responds to market conditions may earn less over the same period. For this reason, financial performance and agent behaviour need to be evaluated separately, and the second is needed to interpret the first.

This distinction matters because most of the trading studies reviewed in this chapter focus mainly on financial measures such as return, profit, Sharpe ratio, and drawdown. More recent work also examines how trained models make their decisions, including feature-attribution methods by Guan and Liu [23]. These methods can identify which inputs are related to a model's output. However, they do not establish whether the policy changes its actions as its inputs change. A policy that produces the same action regardless of its observations leaves no variation in action for an attribution method to explain.

Checking whether a learned agent responds to its observations is already used in reinforcement-learning research and in the wider evaluation of machine-learning systems. This dissertation therefore does not claim to introduce sensitivity testing or behavioural analysis as new concepts, as there are related ideas

that exist in input sensitivity.

The more specific issue addressed here is how to test dependence on state information in a trading environment with constraints on actions. The agent's cash and holdings determine which actions are legal: it cannot buy when it has insufficient cash and cannot sell when it has no position. The set of available actions can therefore change during an episode, and the selected action may change simply because an action became unavailable rather than because the agent responded to market information, and this would risk counting these forced changes as evidence of state use.

To account for this, this dissertation defines state dependence intentionally: **trained policies must select different actions in situations where the same actions are legally available, with the difference occurring in the observed state.** Conditioning the analysis on action masking reduces the chance of misinterpreting changes caused by the constraints as evidence that the agent is using its observations.

The diagnostic, however, answers a limited question — does the policy show any evidence of changing its decisions in response to its observed state? It does not identify which feature caused the change, whether the response was economically sensible, or whether it improved financial performance.

Among the trading studies reviewed here, policy responsiveness is not reported as a quantitative evaluation result alongside financial performance. This is a narrower claim than saying previous research has never used sensitivity testing or behavioural analysis. Established practices such as chronological splits, transaction-cost modelling, benchmark comparisons, and multi-seed reporting are also not presented as new contributions here.

Théate and Ernst [22] show why this distinction matters. Their DQN moved towards passive trading under some conditions, but passive behaviour could come from a policy that responds to market information or from one that follows a fixed trading rule. Without directly testing the relationship between the observed state and the chosen action, both cannot be explained separately.

This dissertation addresses this gap through a controlled comparison of three different learning algorithms. REINFORCE, Deep Q-learning, and PPO are trained using the same observed state, the same masked action space, the same fee-inclusive reward, and the same chronological splits. Their financial performance is compared against benchmark strategies with the same capital, and their behaviour is examined using state-dependence analysis. The diagnostic is reported for each seed alongside the financial results rather than being treated as a technical implementation check.

Two further experiments examine whether changing the information available to the agent or changing its incentive affects the learned behaviour. The first adds drawdown to the state, allowing the agent to observe its recent portfolio losses. The second adds a drawdown penalty to the reward, changing the incentive to take risks. One changes what the agent can see, while the other changes what the agent is encouraged to optimise.

The contribution of this dissertation is therefore not the invention of a new evaluation method, but the controlled combination of experimental practices established from prior works to measure the agent's state-dependent behaviour along with financial performance. This allows the evaluation to distinguish

between an agent that earns a strong return by following a simple fixed rule and one whose decisions actually respond to the information provided.

A second gap follows from the discussion of data requirements earlier in this chapter. The single-asset trading studies reviewed here train on historical data alone and rarely ask whether that history is large enough for the networks being trained on it. This dissertation makes the question explicit: it measures what the algorithms learn from 60 real episodes, diagnoses the failure, and then repeats the comparison on thousands of balanced generated episodes where the data volume is no longer the binding constraint.

These points lead directly to the five research questions introduced in Chapter 1. The first asks whether the learned agents can outperform meaningful benchmarks after transaction costs. The second asks whether their decisions depend on the information in their states. The third asks whether adding drawdown to the state changes behaviour or performance. The fourth considers whether penalising drawdown through the reward produces a different risk-sensitive outcome. The fifth asks whether the amount and balance of training data explain the learned behaviour, and whether fixing the data fixes the behaviour. The next chapter formalises these questions through the MDP formulation, state representation, trading rules, reward functions, the simulated market model, and the algorithms used in the experiments.

Chapter 3

Methodology & Mathematical Formulation

This chapter explains how the portfolio problem can be treated as a Markov Decision Process (MDP) [9, 1], i.e the trading process is treated as a sequence of step-by-step decision-making problems. The objective is to develop an environment in which an autonomous trading agent can learn an optimal decision-making behaviour by interacting with a simulated market. At each step, the portfolio system observes current market and portfolio conditions, takes an action such as buying, selling, or holding, receives a result based on how the portfolio's wealth changed, and then moves to the next step.

The chapter then explains the two main learning methods used in the project: **REINFORCE** [5]—the Monte-Carlo policy gradient, which learns by improving its strategy based on the rewards it receives—and **Deep Q-learning** [11, 6], which learns which actions are likely to give the best results. Finally, **PPO** is introduced as an additional method in the ablation row to compare against the other approaches, as it is designed to make the policy learning process more stable.

Appendix D collects every symbol used in this and the following chapters, for reference.

3.1 Overview: Trading as a Markov Decision Process

Reinforcement Learning (RL) trains an agent to make sequential decisions by interacting with an environment. The environment is modelled as an MDP where decisions are made at regular time steps, represented by the tuple $(\mathcal{S}, \mathcal{A}, \mathcal{P}, \mathcal{R}, \gamma)$: at step t the agent observes the current state $s_t \in \mathcal{S}$, selects an action $a_t \in \mathcal{A}$, receives a numerical reward based on the outcome r_t , and the environment moves to the next state s_{t+1} according to the current state and the action selected by the agent $P(s_{t+1}|s_t, a_t)$.

In this project, the trading environment is modelled in the same way, providing a formal description of how the trading agent interacts with the market over time; one time step t corresponds to one trading day.

The state s_t represents all information available to the agent on day t to fully capture its position within the trading episode. This includes the selected market features, the agent's current portfolio position, and its available cash.

Based on this state, the agent selects one of three possible actions $a_t \in \{0, 1, 2\}$: *Hold the current position*, *Buy a fixed amount*, or *Sell a fixed amount*. The selected trade is then executed, with transaction fees c applied where appropriate.

After taking an action, the environment calculates the change in the agent's total portfolio value and returns it as the reward r_t .

The environment advances to the next trading day s_{t+1} , and the process repeats throughout the episode

τ , allowing the agent to update its policy π_θ or value estimates Q_ϕ based on the experience collected.

The interaction can therefore be represented as: $s_t \rightarrow a_t \rightarrow r_t, s_{t+1}$

Figure 3.1 illustrates the interaction between the trading agent and the environment.




Figure 3.1: The agent–environment interaction loop for the trading MDP. One cycle corresponds to one trading day.

The formulation can be made easier to understand by comparing it with a familiar game-playing task, Flappy Bird, which was used initially to understand reinforcement learning. Where the game’s state is the bird’s position and speed together with the next pipe, the trading state is the market information and the agent’s current portfolio position. Where the game’s actions are flap or do not flap, the trading actions are Hold, Buy, or Sell, and the reward is the change in portfolio wealth after trading fees rather than a survival score. One complete game session corresponds to one calendar month of daily trading, and where the game ends when the bird crashes, the trading episode ends at the final trading day, with any remaining position sold.

This analogy also highlights a design principle used throughout the project: as the Flappy Bird state contains the information relevant to the next decision, the trading state includes the relevant information available to the agent at each trading step, without being given future information. Section 3.3 discusses the design of the trading state in more detail.

3.2 Episode Definition and Data Structure

3.2.1 Episode definition (τ)

Each episode (τ) represents one calendar month of daily closing prices, from P_0 to P_T ($P_0, P_1, \dots, P_T$) where T is typically between 18 and 23 trading days per month, excluding public holidays when markets are not open and weekends of market inactivity ($T \approx 18 - 23$). At the beginning of each episode, the agent starts with cash and no asset holdings. It can then trade throughout the month based on the available information. At the end of the final trading day t , any remaining position is automatically sold. This ensures that no position is carried over from one episode to the next.

3.2.2 Independence of Trading Episodes

For the purposes of the experiment, each monthly episode is treated as a separate and independent trading period, similar to treating each game session in Flappy Bird as a separate episode. This is a simplification, because events and trends in one month do influence the market in the following month, so the resulting data are not independent over time. Modelling these temporal relationships in greater detail would require time-series or partially observable approaches, which are outside the scope of this project. The implications of this assumption are discussed further in the conclusion.

3.2.3 Choice of Episode Length

Monthly episodes were selected to provide a balance between the length of each trading period and the total number of training examples available.

A single year of daily data provides approximately 12 monthly episodes and around 250 daily trading decisions. Using eight years of data from 2018 to 2025 gives 96 monthly episodes and approximately 2,000 daily transitions. For the first study, these episodes are divided chronologically into 60 months for training, 12 months for validation, and 24 months for the held-out test period, as detailed in Chapter 4. The data are not shuffled, which prevents information from later periods from being used during training and helps avoid look-ahead leakage. The main study additionally evaluates on real months from 2006 to 2017, which no training or calibration step ever reads.

The same environment could be adapted to shorter intervals, but intraday trading is outside the scope of this study.

3.2.4 Selection of the Historical Period

The training and calibration data use daily market data from 2018 to 2025 as a fixed period. Earlier data were not included in the training period because the study aims to train the learning algorithms under relatively recent and varied market conditions. The selected period contains a range of market conditions, including relatively stable periods in 2018, the market decline associated with the COVID-19 pandemic and its subsequent recovery, the later period of increased inflation and interest rates, and the war in Ukraine in 2022, which caused weaker market conditions.

Older data, from 2006 to 2017, is used differently: as evaluation-only transfer data for the main study. Because no agent trains on it and no design decision reads it, market conditions that differ from the training period — older microstructure, different fee structures, and above all the 2008 financial crisis — become an asset rather than a contamination risk. They test how well policies trained on recent, calibrated conditions generalise to markets they have never seen.

3.2.5 The Simulated Market Model

Even the full 2018–2025 period yields only 60 training episodes, and Chapter 5 shows this is far too few for the networks to learn rather than memorise. The study therefore also defines a simulated market,

and uses it as the primary source of training data. The idea follows a standard pattern in reinforcement learning: when real experience is scarce, generate experience from a model of the task and keep the real data for a final test [24].

The simulated market has three **market conditions**: up, down, and flat. An episode is a sequence of 21 daily bars drawn entirely from one condition m . Within the condition, daily log-returns are independent draws from a normal distribution,

$$r_t = \ln \frac{P_t}{P_{t-1}} \sim \mathcal{N}(\mu_m, \sigma_m^2),$$

so the price follows a geometric random walk whose drift μ_m and volatility σ_m belong to the condition. High and low prices extend each day's move by random wicks sized to match the real average true range, and volume is drawn from a lognormal distribution. The parameters of each condition are estimated from the real training months only; Chapter 4 reports the fitted values.

The model is deliberately the simplest one that supports a controlled experiment. Because the generator's condition is known for every episode, the evaluation can report how a policy behaves in rising, falling, and flat markets separately, which no historical test period can guarantee to provide. Because episodes are generated, the training set can be as large and as balanced as the experiment requires. And because consecutive generated episodes share nothing, the independence assumption of Section 3.2.2, which is a simplification for real months, holds exactly in simulation.

The cost of this simplicity is realism. Real returns are not independent across days, volatility clusters within months, and market conditions change mid-month. A policy trained on the simulated market therefore faces a domain gap when applied to real data, and Chapter 5 measures that gap directly by evaluating the simulation-trained agents on the real held-out months.

Real financial markets are also more complex than any MDP and cannot be assumed to be fully observable; the information the agent cannot observe, and the consequences of this, are discussed with the state representation in Section 3.3.

3.3 State Representation Design

In reinforcement learning, how information is presented to the agent plays an important role in the quality of decisions it can learn. The state representation must provide enough information for the agent to understand the current situation and choose an appropriate action.

In this portfolio trading problem, this requires capturing information about both the market and the agent's own portfolio. The agent needs to know how the asset has been moving, but it also needs to know how much cash is available, whether it currently holds the asset, whether that position is profitable, and how much time remains before the trading period ends.

A state that leaves out important information can make it difficult for the agent to learn effective strategies, even if the learning algorithm is suitable for the task. This is because the agent can make decisions only

using the information provided. Therefore, the state representation in this study was developed gradually by first examining simpler representations and then adding information as needed to address specific limitations of the previous representation.

The objective was not to create an unnecessarily complex state representation with many variables, but to design a compact representation containing all relevant information needed for decision-making within the defined portfolio trading environment, while avoiding features that simply duplicate information already available from other variables

3.3.1 Information Available to the Agent

As discussed in Section 3.1, the state should see the whole by providing the agent with information relevant to its decision, as the Flappy Bird state does. In a trading environment, this information can be divided into two halves: **the agent's own portfolio** and **the wider market**. These two areas behave very differently and should therefore be considered separately when designing the state.

The first area is the agent's own portfolio. This includes information such as the amount of cash available, how much of the asset it owns, what it paid to enter a trade position, and the time remaining in the trading period. The agent controls this half; it is not hidden or estimated, and it can be completely represented. Given the sequence of the asset prices and the agent's actions, the portfolio can be updated directly using the trading environment's rules.

The second area is the market. Here, complete information about the market is impossible in principle and cannot be provided to the agent because many factors affecting an asset's price are not directly observable from the available price data. For example, the agent cannot observe other market participants' decisions, expectations, or intentions. Similarly, information such as news, market sentiment, and macroeconomic events may influence prices without appearing directly in the closing price at the end of the episode.

Adding more features calculated from past prices can give the agent additional information about recent market behaviour, but it cannot reveal everything happening in the market. The attainable goal is therefore a state that is complete on the portfolio side and informative enough on the market side: **the portfolio information can be represented completely because the environment directly controls and records it, whereas market information can only approximate the wider market**. Whether the additional market features provide useful information is ultimately an empirical question, which is considered when evaluating the state representation later in the dissertation.

3.3.2 Criteria for a Suitable Trading State

It is tempting to design a state by asking which features look useful, but this is a weak approach as it says nothing about whether the object being trained is an MDP at all.

Before deciding on the final state representation, three conditions were considered. These provide a basis for deciding whether the proposed state contains the information required for the portfolio trading problem.

1. Reward measurability

The reward should be calculable from information available to the agent and the resulting transition. If the reward depends on information not contained in the state or the transition, the agent is effectively being evaluated using information it cannot observe.

The primary reward used in this study is the change in portfolio wealth, $r_t = W_{t+1} - W_t$. The final state includes cash and asset exposure as fractions of the initial capital, and these two features combine to give the current wealth, so the reward can be calculated from the observed transition (s_t, a_t, s_{t+1}) with no hidden portfolio variable. Section 3.5 defines the reward and shows this derivation in full.

2. Transition consistency

The state should also contain enough information for the next state to be determined from the current state, the selected action, the next observed market price, and the trading environment rules.

This matters because the agent should not need to access hidden historical variables to describe how its portfolio changes after an action.

For example, the current portfolio exposure can be used to determine the position's value, while the current unrealised profit/loss contains information about the position's entry price. Together with the current price and the environment's trading rules, these quantities provide the information required to update the portfolio after an action.

This keeps the state representation consistent with the defined trading environment rather than requiring the environment to retrieve additional portfolio history that is not represented in the state.

3. Decision sufficiency

The state should contain enough information for the agent to make a decision based on the current trading situation without ambiguity. If two different situations look identical to the agent but require different actions, then the state is missing important information, making the representation incomplete.

For example, consider the following two situations:

- the asset price has increased by 2% today, and the agent does not hold the asset;
- the asset price has increased by 2% today, but the agent holds a position that is currently making a loss.

In both cases, the observed price movement is the same. However, the most appropriate action may be different. In the first situation, the agent may consider opening a position, whereas in the second, it may consider reducing or closing its existing position to minimize its losses.

The state representation must therefore include information about both the market and the agent's own portfolio.

3.3.3 Development of the State Representation

The state was developed in stages, starting from the smallest representation that could describe a trading decision and adding information only where the previous stage left the agent unable to distinguish situations that require different actions.

The starting point was a price-only state:

$$s_t^{(0)} = [\Delta P_t], \quad \Delta P_t = \frac{P_t - P_{t-1}}{P_{t-1}}$$

where ΔP_t is the one-day return. This is insufficient because the same market movement can occur in very different trading situations. A positive daily return means one thing when the agent holds nothing and may consider entering a position, and another when it already holds a large position and may consider reducing it. The agent cannot see its own portfolio.

Portfolio, position, and time information were therefore added in turn, producing the following features:

1. Normalised cash (c_t):

$$c_t = \frac{C_t}{C_0}$$

This is the proportion of the initial capital that remains available for future decisions. Using a normalised value keeps the information consistent regardless of the starting account size.

2. Normalised exposure (v_t):

$$v_t = \frac{h_t P_t}{C_0}$$

This is the current market value of the open position relative to the initial capital. It replaces the raw unit count h_t because the same number of units can represent a small or large investment depending on the asset price. Together, the two portfolio features recover the agent's current wealth as $c_t + v_t = W_t/C_0$.

3. Unrealised profit/loss (PnL_t):

$$PnL_t = \frac{P_t - \bar{P}_t^{entry}}{\bar{P}_t^{entry}}$$

where $\bar{P}_t^{entry}$ is the average price at which the current position was opened. This tells the agent whether its position is in profit or loss — two situations that look identical from the market side but may call for different actions. Holding an asset after a price increase may be profitable if it was bought at a lower price, or losing money if the position was entered at a higher price before the price declined.

4. Episode clock (τ_t):

$$\tau_t = \frac{t}{T}$$

where T is the total length of the episode. A value close to 0 means the episode has just started, while a value close to 1 means the forced close at the final trading day is approaching. This matters because the same action can have different consequences depending on when it is taken: buying early leaves the position time to grow, whereas buying shortly before the episode ends does not.

3.3.4 Final State Representation

Although the previous stages successfully introduced information about the agent's portfolio and position, they provide limited understanding of the broader market conditions. A trading decision should not depend only on the most recent price change, as different market situations exist, producing similar short-term movements.

For this reason, the agent also needs information that helps it distinguish between market conditions: whether the asset is in a short-term upward or downward trend, whether the market is stable or highly volatile, and whether the current price is above or below its recent average level. Additional market-related features are therefore included to give the agent a more complete view of the trading environment.

The final state representation is defined as:

$$s_t^{(4)} = [\Delta P_t, mom_t, vol_t, gap_t, rng_t, \tau_t, c_t, v_t, PnL_t] \in \mathbb{R}^9$$

The following variables were added:

1. Momentum

$$mom_t = \frac{P_t - P_{t-k}}{P_{t-k}}$$

where $k = 5$ previous days (about one trading week). Momentum gives the agent a longer view of the recent price direction, helping it differentiate a temporary movement from a more consistent trend developing over several timesteps.

2. Volatility

$$vol_t = \sigma(\Delta P_{t-k+1}, \dots, \Delta P_t)$$

where σ is the standard deviation of the last $k = 5$ daily returns. This helps the agent recognise differences in market uncertainty: a 2% price increase during a relatively stable period is different from a 2% increase during a period of large and frequent price movements, which indicates a riskier situation.

3. Price trend position

$$gap_t = \frac{P_t}{MA_{m,t}(P)_t} - 1$$

where $MA_{m,t}(P) = \frac{1}{m} \sum_{i=0}^{m-1} P_{t-i}$ is the simple moving average of the closing price over the previous ($m = 20$) days (about one trading month). A positive value means the current price is above its recent average level and a negative value means it is below, helping the agent identify potential trends or deviations from normal price behaviour.

4. Normalised Trading Range

The normalised trading range is based on the Average True Range (ATR) [34]:

$$rng_t = \frac{ATR_{n,t}}{P_t}$$

The True Range at timestep t is:

$$TR_t = \max(H_t - L_t, |H_t - C_{t-1}|, |L_t - C_{t-1}|)$$

where H_t and L_t are the current high and low prices and C_{t-1} is the previous closing price. The ATR is the rolling arithmetic mean of the last ($n = 14$) True Range values (about two trading weeks):

$$ATR_{n,t} = \frac{1}{n} \sum_{i=0}^{n-1} TR_{t-i}$$

The ATR therefore measures the size of recent price movement using the asset's high, low, and closing prices, and dividing it by the current price creates a normalised measure that can be compared across different price levels.

The nine features and what each tells the agent are summarised in Table 3.1.

Table 3.1: Full State Definition

Table 3.1: The nine-feature state representation ($s_t \in \mathbb{R}^9$) used in the main experiments. C_0 denotes the initial cash, h_t the number of asset units held, and $W_t = C_t + h_t P_t$ the portfolio wealth at time t . the mark-to-market wealth. The feature timestep lengths are set to ($k=5$) for momentum and volatility, ($m=20$) for the moving-average reference, and ($n=14$) for the Average True Range.

#	Feature	Definition	What it tells the agent
Market			
1	ΔP_t	$P_t - P_{t-1} / P_{t-1}$	The most recent price movement

#	Feature	Definition	What it tells the agent
2	mom_t	$P_t - P_{t-k} / P_{t-k}$	whether that recent price movement is part of a short-term trend
3	vol_t	$\sigma(\Delta P_{t-k+1}, \dots, \Delta P_t)$	how much the asset price has been varying recently
4	gap_t	$P_t / MA_m(P) - 1$	how far the current price is from its recent average
5	rng_t	$ATR_n(P) / P_t$	The size of the recent price ranges, which closing prices alone cannot show
Portfolio			
6	c_t	C_t / C_0	how much of the initial cash remains available as cash
7	v_t	$h_t P_t / C_0$	The current value of the open position relative to the initial capital
Position			
8	PnL_t	$(P_t - \bar{P}_{tentry}) / \bar{P}_{tentry}$	whether the open positions are in profit or loss
Time			
9	τ_t	$t / T \in [0, 1]$	how much time remains before the trading episode ends

3.3.5 Feature Redundancy and Excluded Information

The state was not designed by adding every potentially useful market indicator. Features were included only where they provide information that cannot already be recovered from the other variables in the state.

This is important because adding a feature that simply repeats existing information increases the size of the input without providing the agent with anything new.

The following features were excluded:

1. Running portfolio return

For example, the portfolio's current return relative to its starting capital does not need to be included as a separate feature because it can already be calculated from the normalised cash and asset exposure as $c_t + v_t$. Adding it would duplicate information already available in the state.

2. Raw position size

The raw number of units held, h_t , is also not included in the final state. The financial value of the position is more directly represented by:

$$v_t = \frac{h_t P_t}{C_0}$$

This provides information about the actual exposure rather than only the number of units.

3. Raw asset price

The raw asset price is also not required as a separate feature. The state instead uses returns and normalised values, which describe price movements without depending on the standard price level.

The general principle is therefore to include features that provide useful information while avoiding variables that can already be reconstructed from information available elsewhere in the state.

3.3.6 The Risk-Aware Extension: Drawdown State

The main state representation contains the information needed for the change in wealth reward. However, it does not describe the history of the portfolio's losses: current wealth alone does not show how the portfolio reached its current value. For example, two portfolios can have exactly the same wealth at the current timestep but may have followed very different paths:

- one portfolio may have increased steadily over time;
- another may have suffered a large loss before recovering to the same level.

To address this, an additional feature representing portfolio drawdown is introduced as a separate risk-aware extension:

$$dd_t = \frac{\max_{u \leq t} W_u - W_t}{\max_{u \leq t} W_u}$$

where dd_t represents the current decline in portfolio wealth from the highest wealth level reached so far during the episode.

A value of zero means that the portfolio is currently at its highest wealth level reached so far in the episode. A larger value indicates that the portfolio has fallen further below its previous peak.

For this reason, drawdown is treated separately from the main state representation. This makes it possible to evaluate whether explicitly providing information about portfolio loss history affects the agent's behaviour and performance.

The resulting risk-aware state is therefore:

$$s_t^{risk} \in R^{10}$$

where:

$$s_t^{risk} = [s_t, dd_t]$$

Since the original state s_t contains nine features, adding drawdown produces a **ten-dimensional** risk-aware state.

3.3.7 Feature Calculation and Information Timing

The timing of feature calculation matters because the agent should receive only the information available when making its trading decision. Using future information would give the agent an unfair advantage and could lead to overly optimistic results.

For this reason, the market features are calculated using information available at or before timestep (t). The current price and past prices can be used to create the state, but the agent must not use future prices when it decides its action (a_t).

The same principle applies to the features that use historical windows: momentum uses the previous ($k = 5$) timesteps, the moving average the previous ($m = 20$) prices, and the normalised trading range the ($n = 14$)-day ATR, as reported in Table 3.1.

Therefore, each state is based only on information that is available at the time the decision is made. This prevents future market information from entering the agent's observation and helps avoid **look-ahead bias**, where a model appears to perform well because it has indirectly been given information that would not have been available during actual trading.

3.3.8 State Representation Limitations

Although the final state representation gives the agent a wider range of information, the trading environment is still a simplified representation of a real financial market.

The agent does not have access to several types of information that can influence real-world trading decisions, including order flows, market sentiment, financial news, macroeconomic conditions, relationships between the asset and other financial assets, and the actions of large financial institutions.

These limitations mean that the state should not be interpreted as a complete description of the real financial market. Because this relevant information cannot be represented in the environment, it is better to describe this as a *partially observable Markov decision process (POMDP)* rather than a fully observable one.

However, this dissertation does not aim to replicate every aspect of real-world financial markets. The purpose is to compare and evaluate reinforcement learning algorithms within a controlled trading environment for an asset.

The proposed state representation therefore provides sufficient information for the defined task and can be **considered complete with respect to the information used by the defined trading environment**, which includes recent market behaviour, the current portfolio position, the performance of its current position, and the amount of time remaining in the trading period, while acknowledging that the model does not capture all factors present in real-world trading.

3.3.9 Summary

The final state was not selected because it contained a particular number of variables, but because its components collectively describe the market conditions, portfolio exposure, position performance, and remaining trading episode available within the defined environment. The staged development above records the limitation each stage resolved, from the price-only state S_0 to the nine-feature main state S_4 . The separate tenth feature, drawdown, is used only for the risk-aware experiments described in Section 3.3.6.

The main state representation provides the basis for defining the available actions and applying the reinforcement learning algorithms described in the following sections.

3.4 Action Space Design

3.4.1 Overview

The action space defines the choices available to the trading agent at each timestep. Its design matters because it determines how the agent interacts with the market and what its learned policy means.

A poorly designed action space can introduce unnecessary complexity or create unrealistic trading behaviour. For example, allowing the agent to choose any number of shares to buy or sell adds learning challenges: the agent must simultaneously learn when to trade, which direction to take, and how much to trade. This can make it difficult to determine whether poor performance results from an ineffective trading strategy or from difficulty learning appropriate position sizes.

To keep the focus on the trading decisions made by the reinforcement learning algorithms, this study uses a **simplified discrete action space**. The agent decides on the action (trading direction), while the environment determines the amount traded.

This separates the trading problem into two decisions:

1. **When and in which direction to trade** — this is the trading policy being learned by the reinforcement learning algorithm.
2. **How much capital to trade** — this is fixed by the environment and treated as an experimental setting.

This approach keeps the action space simple and makes it easier to compare the learning algorithms on the same portfolio trading task. It also means that differences in performance can be interpreted mainly in terms of the agents' ability to learn trading decisions, rather than their ability to learn position sizing.

3.4.2 Trading Actions and Trade Size

The action space is defined as $A = \{0, 1, 2\}$, where action 0 **holds** the current position, action 1 **buys** a fixed amount, and action 2 **sells** a fixed amount. At each timestep, the agent selects one of these three actions using the information available in the current state.

The **hold** action leaves the portfolio unchanged. No transaction takes place: the available cash and the number of units held remain the same, and no transaction fee is charged. This allows the agent to remain in its current position when it does not consider buying or selling to be appropriate.

The **buy** and **sell** actions change the agent's exposure to the asset by a fixed monetary amount:

$$\Delta = f C_0$$

where C_0 is the initial capital and f is the fraction of that capital allocated to one transaction. A fixed monetary amount is used rather than a fixed number of shares because the cost of one share depends heavily on the price level: buying one share of an asset priced at £10 is a small investment, while buying one share priced at £1,000 is a large one, even though both would be described as buying one share. Fixing the monetary size keeps every trade at the same proportion of the initial capital throughout the dataset, regardless of the asset price.

The fraction f is treated as an experimental setting rather than a fixed design constant. It is selected using the validation data before the final experiments, as reported in Section 5.5.

The number of units bought or sold at price P_t is:

$$\Delta h_t = \pm \frac{\Delta}{P_t}$$

with the positive sign for a buy and the negative sign for a sell, and with fractional ownership allowed, similar to a real-world portfolio. A sell is valid only when the agent holds enough units for the transaction. Because the trade size is fixed, the agent can build or reduce its position gradually over several decisions while still operating within a simple three-action framework.

3.4.3 Action Constraints

Although the agent has three possible actions, some actions may not be possible at a particular timestep because of the current portfolio position.

For example, the agent cannot buy if it does not have enough available cash, and it cannot sell if it does not hold enough of the asset to withdraw.

To handle these situations, the environment uses a **legal action mask**. The mask does not provide any additional market information to the agent. It simply indicates which actions can be carried out given the current portfolio status.

The mask is represented as:

$$M_t = [m_{\text{hold}}, m_{\text{buy}}, m_{\text{sell}}] \in \{0, 1\}^3$$

where $M_t(a) = 1$ indicates that action a is available and $M_t(a) = 0$ indicates that it is not. The availability of each action is determined from the portfolio state: available cash determines whether a **buy** action can be carried out, the current asset holding determines whether a **sell** action can be carried out, and **hold** remains available throughout the episode.

Because the mask is derived directly from information already in the portfolio state, it is not treated as an additional feature of the state representation. Instead, it acts as a constraint on the agent, preventing it from selecting actions the environment cannot execute.

3.4.4 Transaction Costs

A transaction fee of $c = 0.05\%$ is charged whenever the agent carries out a buy or sell action, so a transaction of value Δ incurs a fee of Δc . The same fee calculation is applied to both purchases and sales.

Including transaction costs matters because it prevents the agent from overly buying and selling repeatedly. Without these fees, the agent could benefit from making many trades even when price movements are very small, making frequent trading appear more profitable than it would be in practice.

The inclusion of fees therefore introduces a basic cost of trading into the environment and provides a more realistic basis for evaluating learned trading strategies.

3.4.5 Relationship Between the State and Actions

The action space was designed to match the state representation of Section 3.3. Deciding whether to buy or sell depends on the agent's current exposure, its available cash, whether an existing position is in profit or loss, the time remaining in the episode, and the current market conditions — exactly the information the state provides. The agent decides whether to hold, buy, or sell, while the environment determines the amount traded and prevents invalid actions from being carried out.

3.5 Reward Function Design

3.5.1 Purpose of the Reward Function

The reward function tells the reinforcement learning algorithm how well it performed after taking an action and guides the agent toward more desirable actions. Unlike supervised learning, where the model is given the correct answer, the agent learns by taking actions and receiving feedback on the results of those actions.

In a trading environment, the reward function needs to accurately reflect the actual objective of the investment strategy. If it is poorly designed, the agent may learn behaviours that are not desirable, such as trading too frequently, taking unnecessary risks, or performing well during training but poorly under realistic conditions.

The main goal of this dissertation is to investigate whether reinforcement learning algorithms can learn profitable trading strategies for the selected asset while accounting for transaction costs. For this reason,

the primary reward is based on the **change in portfolio wealth** from one timestep (trading day) to the next.

3.5.2 Wealth-Based Reward

The main reward used in this study is based on the change in the portfolio's total value:

$$r_t = W_{t+1} - W_t$$

where portfolio wealth at timestep t is given by:

$$W_t = C_t + h_t P_t$$

where C_t is the cash available in the portfolio, h_t the number of asset units held, and P_t the current asset price.

The reward therefore measures how much the portfolio's total value changes from one timestep (trading day) to the next after the agent's action.

A **positive reward** means the portfolio increased in value, while a **negative reward** means it decreased. This gives the agent a direct measure of the financial outcome of its decisions.

3.5.3 Including Transaction Costs in the Reward

Transaction fees are included when calculating the reward so that the result reflects the actual cost of carrying out a trade.

When a buy or sell action is taken at timestep t , the portfolio value is updated after the transaction fee is deducted:

$$W_{t+1} = C_{t+1} + h_{t+1} P_{t+1} - Fees$$

The reward is then calculated from the resulting change in portfolio wealth. This means that the agent is rewarded or penalised based on the value of the portfolio **after trading costs have been taken into account**.

3.5.4 Wealth Change As the Primary Reward

The change in portfolio wealth was chosen as the main reward for several reasons.

First, it directly reflects the main objective of the trading strategy: **increasing the value of the portfolio**. This also makes the reward closely related to the metric used to evaluate investment performance in the real world.

Second, it satisfies the requirements of the MDP defined in Section 3.2 because the reward can be obtained from portfolio information in the state representation.

The current portfolio wealth can be recovered from the cash and exposure features in the state representation:

$$c_t + v_t = \frac{W_t}{C_0}$$

Therefore, the change in wealth can be written as:

$$W_t = C_0 [(c_t + v_t)]$$

$$r_t = C_0 [(c_{t+1} + v_{t+1}) - (c_t + v_t)]$$

This shows that the reward depends only on the observed transition (s_t, a_t, s_{t+1}) , without requiring additional information unavailable to the agent.

3.5.5 Risk-Aware Reward Extension

Although the main objective is to increase portfolio wealth, focusing only on returns may encourage the agent to take more risk than desired. For example, a strategy may achieve a high final return but experience large losses along the way, which is not a desirable trading behaviour.

Therefore, a risk-adjusted reward that accounts for changes in portfolio drawdown is considered as an additional experimental condition:

$$r_t^{\text{risk}} = \Delta W_t - \lambda C_0 \Delta dd_t$$

where ΔW_t is the change in portfolio wealth, Δdd_t is the increase in portfolio drawdown at step t (zero when drawdown does not rise), C_0 is the initial capital, and λ controls the penalty applied to increases in drawdown.

Drawdown measures how far the current portfolio wealth has fallen from the highest level reached so far:

$$dd_t = \frac{\max_{u \leq t} W_u - W_t}{\max_{u \leq t} W_u}$$

where W_t is the portfolio wealth at time t , and W_u is the wealth at any earlier time u up to t . $\max_{u \leq t} W_u$ is therefore the highest wealth reached so far in the episode.

Δdd_t is therefore the increase in drawdown from one trading day to the next. Multiplying it by the initial capital C_0 keeps the reward in dollars, so a penalty of $\lambda = 0$ lets the wealth-change reward remain the same, while a larger λ discourages actions that worsen the drawdown.

However, this reward requires access to the portfolio's previous performance because the previous highest wealth reached up to time t must be known in order to calculate drawdown.

For this reason, the risk-aware reward is used only with the **risk-aware state extension** described in Section 3.3.6. This ensures that the information needed to calculate the reward is available to the agent.

The wealth-change reward remains the main reward used in the baseline experiments. The risk-aware version is treated separately so that the effect of explicitly including information about portfolio loss history can be examined.

3.5.6 Limitations of the Reward Function

Although the wealth-change reward provides a clear and direct measure of trading performance, it does not capture every aspect of a desirable trading strategy.

For example, it does not directly penalise large falls in portfolio value, unstable performance in returns, or high levels of volatility.

A strategy could therefore achieve a strong overall return while experiencing substantial losses during the trading period. The limitations are addressed by the risk-aware reward in Section 3.5.5 by adding a penalty when drawdown increases.

However, adding more objectives changes the learning problem and what the agent is being trained to optimise. This can make comparisons between algorithms less straightforward because improvements may result from the different reward definition rather than from the learning algorithm itself.

Therefore, the **wealth-change reward is used as the common baseline**, while the risk-aware reward is considered separately as an additional experiment.

3.6 Objective Function

The objective function describes what the reinforcement learning algorithms are trying to achieve during training. In this study, the aim is to learn a trading policy that produces the highest possible expected return over each trading episode. The reward function introduced in Section 3.5 measures the result of an individual decision; the objective function combines the rewards received throughout the episode.

3.6.1 Episode Return

An episode can be represented by the sequence $\tau = (s_0, a_0, r_0, s_1, a_1, r_1, \dots, s_T)$, where s_t is the state observed at timestep t , a_t is the action selected by the agent, and r_t is the reward resulting from that action. The episode contains T trading decisions, and any position that remains open at the end is closed

according to the rules of the trading environment. The discounted return from timestep t , commonly referred to as the **reward-to-go**, is:

$$G_t = \sum_{t'=t}^{T-1} \gamma^{t'-t} r_{t'},$$

where $\gamma \in (0, 1]$ is the discount factor and G_0 is the return of the whole episode. This links each action to the rewards that follow it, rather than assigning the episode's total return to every action taken in that episode.

A discount factor of $\gamma = 0.99$ is used throughout the experiments. The trading episodes contain approximately 18–23 daily observations, so this value results in relatively mild discounting over the length of an episode: $0.99^{20} \approx 0.82$, so a reward received near the end of the monthly episode still contributes around 82% of its original value. Using $\gamma = 0.99$, therefore, maintains the standard reinforcement-learning formulation for discounted returns while still ensuring the objective is maximising end-of-month wealth.

3.6.2 Expected Return

The outcome of a single trading episode is not enough to judge the quality of a policy, because each episode reflects the market conditions of one month. A policy that performs well in one month may not perform equally well in another. The objective is therefore defined using the expected return the policy earns across different trading episodes. The policy $\pi_\theta(a|s)$, with trainable parameters θ , assigns probabilities to the available actions in each state, and the learning objective is:

$$J(\theta) = \mathbb{E}_{\tau \sim p_\theta(\tau)} [G_0(\tau)], \quad \theta^* = \arg \max_{\theta} J(\theta).$$

In simple terms, the aim is not for the agent to find a policy that performs well in only one trading period, but parameters that produce good expected performance across the different monthly episodes used during training. The policy parameters affect the objective only through the actions the agent selects: changing θ changes the probabilities of choosing Buy, Hold, or Sell in each state, and these decisions affect the portfolio and, in turn, the rewards collected over the episode.

3.6.3 Objective and the Trading Environment

The formulation above also separates the trading environment's role from the learning algorithm's. The environment determines the market data, the current portfolio state, the transaction costs, which actions are currently valid, the resulting portfolio wealth, the reward, and whether the episode has ended. The learning algorithm determines only how the agent updates its policy or value function using that experience.

This separation is important for comparing the different learning approaches in this dissertation. **REINFORCE**, **DQN**, and **PPO** use the same state representation, action space and masking, market data, and

reward function, so differences in their results stem from how they learn rather than from the conditions they operate under.

3.7 REINFORCE: Monte-Carlo Policy Gradient

REINFORCE [5] is the first policy-gradient method used in this study. Unlike value-based methods, which estimate how useful different actions are in a given state, **REINFORCE directly learns a policy** that gives probabilities to the available actions. It provides a useful baseline because its main idea is straightforward: the agent selects actions according to its current policy, observes the resulting rewards, and then adjusts the policy so that actions leading to better returns become more likely.

3.7.1 Policy Representation

The policy is represented by a neural network with trainable parameters θ . Given the current state s_t , the network produces one output, or logit, for each of the three possible actions. Actions that are not permitted at the current timestep are removed using the action mask described in Section 3.4, and the remaining logits are passed through a SoftMax function to produce the policy probabilities $\pi_\theta(a_t|s_t)$.

During training, the agent samples an action from this probability distribution, which lets it try different actions and explore alternative trading decisions. During evaluation, the agent selects the valid action with the highest probability without exploring. The policy network used in the first study has **two hidden layers, each with 128 units**, and uses the hyperbolic tangent tanh as the activation function. The main study keeps the same structure but narrows the layers to 64 and 32 units, for the reasons of model capacity discussed in Chapter 4.

3.7.2 Policy-Gradient Update

REINFORCE improves the policy by estimating the gradient of the expected return $J(\theta)$ with respect to the policy parameters. The probability of a complete episode depends on the initial state, the actions selected by the policy, and the market's transitions, but only the action probabilities depend on θ . The likelihood-ratio identity [5, 1] therefore gives a gradient that can be estimated from sampled episodes without any model of the market:

$$\nabla_\theta J(\theta) = \mathbb{E} \left[\sum_{t=0}^{T-1} \nabla_\theta \log \pi_\theta(a_t|s_t) G_t \right].$$

Each action is weighted by the reward-to-go G_t rather than the whole episode return, because an action taken at time t cannot influence rewards already received before that time. In simple terms, the update adjusts the probability of an action based on the outcome that follows it: actions resulting in higher returns are made more likely, and actions associated with lower returns are made less likely.

The expectation cannot be calculated exactly, because the full distribution of possible market trajectories is unknown. REINFORCE therefore estimates it as an average over sampled episodes. This is why it

is described as a **Monte-Carlo policy-gradient method**: it does not require the agent to know how the market will move or to build a model of future price changes, only complete episodes produced by its interaction with the environment.

3.7.3 Loss Function Used for Training

Most neural-network frameworks are designed to minimise a loss function, while the objective here is maximised. The update is therefore implemented as a loss by introducing a negative sign:

$$L(\theta) = -\frac{1}{N} \sum_{n=1}^N \sum_{t=0}^{T_n-1} \log \pi_{\theta} \left(a_t^{(n)} | s_t^{(n)} \right) G_t^{(n)},$$

where n indexes the N sampled episodes. The logarithm appears as a direct result of the likelihood-ratio derivation; it is not an additional trading assumption. The implementation uses the **Adam optimiser** [35] with a learning rate of 10^{-3} , and the sampled returns are normalised before being used in the update, which controls the gradient magnitude across training batches and improves numerical stability.

3.7.4 Strengths and Limitations

REINFORCE is a suitable baseline for this project because its formulation is relatively straightforward: it directly learns the policy and does not require a separate model of the market, so performance differences against more advanced algorithms are easier to interpret.

Its main limitation is that **its gradient estimates can have high variance** because they depend completely on episode returns, and different episodes can produce very different outcomes even in similar trading situations. This can make learning less stable, particularly when training data is limited. Reward-to-go and standardised returns reduce some of this variation, but they do not remove the algorithm's underlying high-variance nature.

For this reason, this provides motivation for comparing REINFORCE with the other methods, **DQN and PPO**, to see whether they can achieve a more stable or effective learning in the same trading environment.

3.8 Deep Q-Learning

Deep Q-learning provides a **value-based alternative** to REINFORCE. Rather than directly learning the probability of selecting each action, it estimates the value of each possible action in the current state. This comparison is useful because both methods address the same portfolio trading problem but learn in different ways.

3.8.1 Action-Value Function

The action-value function, commonly referred to as the **Q-function**, is defined as $Q^{\pi}(s, a)$: the expected discounted return from taking action a in state s and then continuing to follow policy π . The optimal action-value function Q^* satisfies the Bellman optimality equation [9, 11]:

$$Q^*(s_t, a_t) = \mathbb{E} \left[r_t + \gamma \max_{a'} Q^*(s_{t+1}, a') \right].$$

The value of an action therefore depends on the reward received immediately and on the value of the opportunities available afterwards.

Tabular Q-learning, which stores a separate value for every state-action pair, is not suitable here because the state contains continuous market and portfolio variables; discretising them would create a very large state space and lose information. Deep Q-learning instead approximates the Q-function with a neural network $Q_\phi(s, a)$ [6], where ϕ represents the trainable parameters. The network takes the state as input and estimates the Q-value of each action, and the agent selects the most valuable *legal* action — it cannot, for example, buy without enough cash or sell more than it holds. The network used in the first study has **two hidden layers with 128 units each** and ReLU activations; the main study narrows the layers to 64 and 32 units.

3.8.2 Experience Replay and the Target Network

Successive decisions within a trading episode are naturally related because they come from consecutive trading days, and training on them in their original order would produce highly correlated updates. DQN addresses this with an **experience replay buffer** [13, 6]. Each interaction is stored as a transition $(s_t, a_t, r_t, s_{t+1}, M_{t+1}, d_t)$ — the state, action, reward, next state, next action mask, and an end-of-episode indicator — and training samples random minibatches of previously stored transitions. The buffer has a capacity of **50,000 transitions**, with **64 transitions sampled per minibatch**, so the network learns from experience collected at different points during training rather than only the most recent transition.

Using the same network to calculate both the current Q-values and the targets can make training unstable, because the targets change with every update. DQN therefore uses a second network known as the **target network** [6], $Q_{\phi^-}(s, a)$, whose parameters are periodically copied from the online network rather than updated at every step. In this implementation, the target network is synchronised with the online network every **500 training steps**, which provides a more stable reference for the learning target.

3.8.3 Temporal-Difference Target and Loss

Unlike REINFORCE, which waits until the end of an episode to calculate the return, DQN updates its estimates using a one-step temporal-difference target calculated from the target network:

$$y_t = r_t + \gamma(1 - d_t) \max_{a' \in \mathcal{A}_{\text{legal}}(s_{t+1})} Q_{\phi^-}(s_{t+1}, a').$$

Action masking is applied inside the maximum because the agent should not assign value to actions it cannot take in the next state. The end-of-episode indicator d_t matters in this environment because any remaining position is closed at the end of each monthly episode: when $d_t = 1$, the future term vanishes and the target reduces to the final reward alone. Including a future Q-value beyond that state would

incorrectly treat the episode as if further decisions were possible.

The online network is trained by minimising the mean squared error between its predicted Q-values and these targets over each minibatch. The main difference from REINFORCE is the learning signal: REINFORCE updates the policy using returns from complete sampled episodes, whereas DQN updates its Q-values using one-step targets that include an estimate of the future value.

3.8.4 Exploration

If the agent always selected the action with the highest current Q-value, it could settle on actions before exploring the alternatives, because the Q-values are initially poor estimates. To encourage exploration, the DQN uses an ϵ -greedy strategy: with probability ϵ it selects a random action from the set of valid actions, and otherwise it selects the valid action with the highest estimated value. The exploration rate starts at $\epsilon = 1.0$ and is gradually reduced to $\epsilon = 0.05$ over the first half of training, so the agent explores frequently early on and increasingly relies on its learned values later. The agent considers only valid actions when exploring; it cannot select an invalid action simply because it is in its exploratory phase.

3.8.5 Strengths and Limitations

DQN has several advantages for this problem. Experience replay lets the agent reuse transitions collected at different stages of training, the target network stabilises learning, and the separate value estimates make it possible to directly compare the expected value of Buy, Hold, and Sell.

Its limitations are the mirror of these strengths. Training can be affected by the scale of the reward and target values, unlike the REINFORCE implementation, where returns are standardised within each episode. DQN also uses bootstrapping: the target partly relies on Q-value estimates given by the network rather than observed rewards, so inaccurate estimates can carry errors into later updates.

These differences make DQN a useful comparison to the policy-gradient algorithms used in this study rather than using another implementation of the same learning architecture.

3.9 PPO: Proximal Policy Optimisation

Proximal Policy Optimisation (PPO) [7] is used as the third learning method in this study to provide a more stable policy-gradient algorithm that can be compared with REINFORCE. In REINFORCE, the policy can change considerably between updates, specifically when the sampled returns are large, which makes training unstable. PPO reduces this problem by limiting how much the policy changes during each update.

3.9.1 Actor-Critic Structure and the Clipped Objective

PPO uses two neural networks: an **actor**, which represents the trading policy $\pi_{\theta}(a|s)$, and a **critic**, which estimates the value $V_{\phi}(s)$ of the current state. The critic provides a reference against which the outcome of each action is compared, giving the **advantage** $A_t = G_t - V_{\phi}(s_t)$. A positive advantage

means the action performed better than expected for that state, so its probability should increase. The implementation uses Generalised Advantage Estimation [17] to calculate the advantages, which reduces the variance of the updates compared with using full episode returns as REINFORCE does.

The size of each update is controlled by a clipped objective. The probability ratio $\rho_t(\theta) = \pi_\theta(a_t|s_t) / \pi_{\theta_{old}}(a_t|s_t)$ measures how much the new policy has moved from the policy that generated the sampled data, and the objective

$$L^{CLIP}(\theta) = \mathbb{E}_t [\min(\rho_t(\theta)A_t, \text{clip}(\rho_t(\theta), 1 - \epsilon, 1 + \epsilon)A_t)]$$

stops rewarding changes once the ratio leaves the range $[1 - \epsilon, 1 + \epsilon]$. This project uses the standard value $\epsilon = 0.2$. Clipping does not prevent the policy from changing by more than 20%; it removes the incentive for further movement in that direction, which keeps the learning process smaller and more controlled than updating directly from sampled returns, as REINFORCE does.

3.9.2 PPO in the Experiment

The action masking described in Section 3.4 also applies to PPO [20]: the agent should not be able to select an action that the trading environment cannot execute. For this reason, the implementation uses **MaskablePPO** from sb3-contrib [36] instead of the standard PPO implementation. The mask is applied after the policy determines the probability of each action, so it excludes actions that are not currently possible, such as buying when the available cash is too small.

The policy uses the same two-hidden-layer architecture as the other neural-network agents, 128 units per layer in the first study and 64 and 32 units in the main study, and the remaining PPO settings use the library defaults. PPO also operates under the same trading environment and evaluation procedure as the other methods, as established in Section 3.6.

3.10 Comparison of the Three Learning Approaches

The three learning procedures, as implemented, are listed in full as Algorithms 1–3 in Appendix B. REINFORCE and Deep Q-learning are built from scratch following the formulations in Sections 3.7 and 3.8; MaskablePPO is the library comparator of Section 3.9 and applies the same action mask.

REINFORCE and DQN are the primary learning methods because they directly represent the two main reinforcement learning algorithms being investigated: policy-gradient learning and value-based learning. PPO is included as an additional policy-gradient comparison, setting the basic Monte Carlo policy-gradient algorithm against a more stable policy-gradient algorithm.

Using all three algorithms provides a stronger comparison than relying on a single one. If similar trading behaviour appears across the three, it gives greater confidence that the behaviour comes from the trading environment or state representation rather than a specific algorithm; if they produce different results, the differences show how the choice of learning method affects the resulting trading strategy.

Chapter 4

System Design and Experiment Framework

This chapter explains how the mathematical formulations and design choices from Chapter 3 were implemented in the portfolio trading system using reinforcement learning. The state representation, action space, reward function, and learning algorithms described previously were converted into a working experimental system for evaluation.

4.1 Introduction

The system is a Python-based system, divided into separate components for market data preparation, the trading environment, the learning algorithms, experimental configurations, and the evaluation process. This structure allows easy modification of each component without affecting the rest of the system. More importantly, it allows the same trading environment to be used by different algorithms, REINFORCE, DQN, and PPO, with little difference in their experimental conditions.

REINFORCE and Deep Q-Learning (DQN) were implemented from scratch using PyTorch [37]. PPO was implemented using the MaskablePPO implementation from sb3-contrib [36]. The Gymnasium interface [38] simulates the trading environment and is used across all the learning algorithms.

The implementation provides two data pipelines. The first prepares real SPY market data. It serves three roles: it supplies the training months for the first study, the calibration statistics for the market generator, and the real months on which the final policies are transfer-tested. The second pipeline generates the simulated market episodes on which the main study trains, and exists because the real data turned out to be too small to train the networks without memorisation. Both pipelines produce episodes in the same format, so the environment, the agents, and the evaluation code do not know which kind of data they are running on.

The implementation also addressed several practical issues stated in the methodology that the mathematical formulation could not capture. These include limiting the agent to only valid actions, ensuring proper separation between training and testing data, closing any remaining position at the end of each monthly episode, and ensuring the risk-aware states and rewards are used separately. This helps ensure that the experiment setup follows the constraints assumed in Chapter 3 to the letter. These details matter heavily because even correct mathematical formulations can produce unreliable results if implemented improperly.

4.2 Data Pipeline

4.2.1 Market Data Preparation

The first stage of the implementation converts historical market data into the observations the simulation trading environment needs. The experiment uses daily **OHLCV market data (open, high, low, close, and volume) of the SPDR S&P 500 ETF Trust (SPY)** to construct features.

SPY was selected because it is a highly traded ETF that represents the broader US equity market. Using an index ETF rather than an individual company reduces the effect of company-specific events and provides a universal asset for the experiments.

Although most formulations use closing prices, the implementation keeps the open, high, low, close, and volume data. The high and low prices are necessary to calculate the range for the ATR feature used in the state representation.

The main dataset period covers **January 2018 to December 2025**. This period was stated in Chapter 3 to capture a mix of market conditions while keeping the main experiment within a significant historical period.

The data are obtained using the **yfinance interface** and are cached locally after processing. This ensures that later experiments use the same prepared dataset rather than retrieving and processing the data again.

4.2.2 Feature Construction & Continuous Calculation

The state features are calculated before separating the data into monthly episodes. This is necessary because some of the features depend on previous data.

For example, momentum uses a five-day window, while the moving average and ATR features use longer windows. Their calculations at the beginning of a month should still use data from previous months and should not restart simply because the environment starts a new monthly episode.

Calculating these features only in each monthly episode would mean that the first few days of every month cannot access or use data from the previous month. This would artificially shorten that feature's historical window at the beginning of the experiment.

To avoid this, the pipeline first creates a continuous time series dataset and calculates all features that roll from the past months across the full dataset. It then creates the monthly episodes from the resulting feature data.

More observations are used before the start of the main experiment period to provide the historical data needed to calculate rolling features. For instance, the longest feature window, ATR, uses 20-day periods, so the pipeline includes 21 additional previous day bars before the first actual observation. These past data are used only for calculating the initial features and are not included in the main experiments.

For this reason, the implementation does not begin feature calculation exactly at the first observation of January 2018. Instead, additional market data is loaded before the experimental period begins. This

warm-up data provides the previous observations necessary to calculate the first valid features with rolling window calculations.

The final pipeline uses a warm-up period beginning in September 2017. The data from this period are used to initialise the rolling calculations and are then removed before the experimental episodes are created.

This approach avoids using the dataset's first observations with incomplete rolling statistics and treats the first experiment's observation the same as later observations, with the required historical context already available.

The pipeline also checks the processed data for missing values before creating monthly episodes. If the data is not available, the data preparation process raises an error rather than silently producing incomplete features.

Overall, the process ensures that each feature at time t is calculated using information available at or before that time, while still allowing rolling features to use historical observations.

4.2.3 Dataset Split & Monthly Episodes

After feature construction, the continuous dataset is divided into monthly episodes by calendar month. Months containing fewer than 15 trading days are excluded from the dataset. After this filtering, the 2018–2025 dataset is left with 96 monthly episodes.

These 96 episodes are separated chronologically into three periods: 2018–2022 (60 months), 2023 (12 months), and 2024–2025 (24 months). The boundary is chronological rather than random because the observations form a time series; shuffling before splitting would let future information leak into training [29], which would not reflect how a trading strategy is developed in practice.

This one boundary serves both studies, in different ways. In the first study, the three periods play the classic roles: the agents train on the 60 months of 2018–2022, decisions made during development, such as the fixed trading amount, use the 2023 validation months, and the 24 months of 2024–2025 are the held-out test, touched once for the final evaluation. In the main study, the same boundary decides what the market generator of Section 4.3 is allowed to read. Its parameters are calibrated on the 60 training months only; simulated episodes take over the training and validation roles; and every real month outside the calibration window — 2023–2025 together with the earlier 2006–2017 months — is reserved for the transfer test described next. Figure 4.1 shows both uses.

The 2018–2022 window is used for the calibration because it includes most market conditions: the volatility of 2018, the COVID-19 crash and recovery in 2020, and the 2022 downtrend. Chapter 5 nevertheless shows that 60 episodes are far too few to train on directly; the consequences of that, and the simulated data that fix it, are the subject of Section 4.3.

The processed episodes and their metadata are cached locally so that the same dataset can be used consistently across all experiments.




Figure 4.1: How the two studies use the real SPY months. The first study trains, validates, and tests inside 2018–2025 and never touches the earlier years. The main study reads real data only to calibrate the market generator (2018–2022); it trains and validates on simulated episodes, and every other real month — 2006–2017 and 2023–2025, 180 in total — becomes its transfer test.

4.2.4 The Extended Transfer Set

The main study does not train on real months at all, and that changes what the real data can be used for. The first study had to spend most of its history on training and tuning, leaving only 24 months for testing. The main study needs real data only for two things: the 60 training months calibrate the generator, and everything else can test the learned policies.

The transfer set therefore contains **every available SPY month outside the calibration window**: 144 months from January 2006 to December 2017, and 36 months from January 2023 to December 2025, giving **180 real test months** in total. The 2018–2022 calibration months are excluded so that no evaluated month overlaps the data the generator’s parameters were fitted on. A separate warm-up lead starting in September 2005 initialises the rolling features, mirroring the 2017 warm-up of the first study’s pipeline.

The early window matters most. The 2024–2025 months the first study tested on rose almost throughout, so they could never show what a defensive policy is for. The 2006–2017 window contains the 2008 financial crisis, the 2011 correction, and the 2015–16 drawdowns: the real market declines that the simulated down condition stands in for. A policy that claims to manage downside risk can now be checked against actual downside months that no part of the training or calibration process ever saw.

4.3 Simulated Market Data

This section describes the second data pipeline, which produces the training data for the main study reported in Chapter 5. It generates simulated market episodes under controlled conditions, and it exists because the real dataset cannot supply enough training episodes for the learning algorithms to do more than memorise.

4.3.1 Why the Real Dataset Is Not Enough

The real training split contains **60 monthly episodes**. The network used by the first study has two hidden layers of 128 units, which is more than 18,000 trainable parameters for the nine-feature state. A model with 18,000 free parameters fitted to 60 examples can match every training month exactly without learning any rule that transfers to new months. The first study in Chapter 5 shows exactly this behaviour.

The 60 episodes are also far below what reinforcement learning normally uses. The Atari agents of Mnih et al. trained on millions of frames [6], and AlphaGo generated its own training games through self-play precisely because no collection of human games was large enough [24]. When the real data is too small, the standard remedy is to train on generated data and keep the real data for a final transfer test.

The real training years also lean upward. Of the 60 training months, 29 rose by more than 2% and only 16 fell by more than 2%. The held-out test years 2024–2025 rise almost throughout. An agent that always buys therefore scores well on both periods without reading its state once. This is a class-imbalance problem: the data rewards one constant answer, so the constant answer is what gets learned.

4.3.2 Generator Design

The generator creates monthly episodes under three controlled market conditions: up, down, and flat. Each episode contains 21 daily bars, and all of them are generated from one condition only. This means that every simulated episode is labelled with the condition that produced it, so the evaluation can state exactly how an agent behaves when the overall market is rising, falling, or moving sideways.

For each condition, the daily log-returns are sampled independently from a normal distribution with that condition's drift and volatility. The daily high and low prices are then created by adding random wicks to the close-to-close movement, and the size of these wicks is based on the average true range observed in the real data. Volume is generated from a lognormal distribution. Since the feature set only uses relative volume, the exact scale of the generated volume does not affect the observations.

The generated data is passed through the same feature pipeline as the real SPY data. A 21-bar warm-up period is generated before each episode so that the rolling features can be calculated before the first episode day. These warm-up bars are then removed, in the same way that the real pipeline uses its 2017 data to initialise the rolling features before discarding it.

The generator is intentionally simple. More complex approaches, such as hidden Markov models [26] and adversarial time-series generators [27], can produce more realistic financial time series. However, realism is not the main purpose of this experiment. The purpose is to control the market condition, so that the behaviour of each agent can be examined separately in rising, falling, and sideways markets. This comparison is only possible when the condition of every generated episode is known.

4.3.3 Calibration

The generator's parameters come from the **real training months only**. Each of the 60 real training months is labelled by its return: up if it gained more than 2%, down if it lost more than 2%, and flat

otherwise. The daily log-returns of the months in each class are pooled, and their mean and standard deviation become that condition’s drift and volatility. The wick size is fitted so the generated average true range matches the real mean for each condition. Table 4.1 shows the fitted values.

Table 4.1: Generator parameters fitted on the 60 real training months.

Condition	Months	Daily drift	Daily volatility	Mean ATR/close
Up	29	+0.244%	0.967%	1.38%
Down	16	−0.358%	2.000%	1.97%
Flat	15	+0.012%	1.095%	1.25%

The fitted values behave as market data should. Down months carry twice the volatility of up months, which is the volatility clustering seen in real crashes. As a calibration check, the generated episodes reproduce the target average true range to within 4% in every condition. The real validation and test months are never read by the generator, so the transfer test in Chapter 5 stays untouched by any generation decision.

4.3.4 Simulated Splits

The simulated dataset contains **3,000 training episodes, 300 validation episodes, and 600 test episodes**, each split holding equal numbers of up, down, and flat episodes. The three splits are independent draws from the same generator under different random seeds.

This design fixes two problems of the real split at once. First, the volume problem: 3,000 episodes of 20 decisions each give 60,000 transitions per pass over the data, against 1,200 for the real training split. Second, the distribution problem: train and test are now drawn from the same distribution by construction, so a policy that fails on the test episodes cannot blame a change in market conditions between the periods.

The balance also removes the advantage of a constant answer. An agent that always buys earns well in the up third, loses more in the down third, and roughly breaks even in the flat third. Buying regardless of state now loses money on average, so a profitable policy must read its observations. The training procedure for this study is described in the next section alongside the real-data procedure.

4.4 Training, Validation and Testing Procedure

The experiments use a fixed sequence that separates training, validation, and test data. This separation ensures that the test data are not used when making decisions about the model or experimental setup. The same sequence applies to both studies; only the data differs. The first study trains on the 60 real training months, validates on the 12 months of 2023, and tests on the 24 months of 2024–2025. The main study trains on the 3,000 simulated episodes, validates on the 300 simulated validation episodes, and tests on the 600 simulated test episodes plus the 180-month real transfer set.

4.4.1 Training

During training, each learning algorithm interacts with the monthly episodes in the training split. Each algorithm is given the same market data, state representation, portfolio conditions, and trading constraints. The main difference between the algorithms is how they use the collected experience to update their parameters.

REINFORCE updates the policy using the reward-to-go calculated from completed monthly episodes. DQN updates its Q-network using transitions sampled from the replay buffer. PPO collects fixed-length rollouts and uses the resulting batches for several optimisation epochs.

Each configuration is trained using six random seeds. Using multiple seeds reduces the dependence on using one particular network initialisation or a sequence of sampled experiences and allows the variation between runs to be reported [30].

The trained model from each seed is saved with its configuration so that it can be evaluated later using the same settings.

4.4.2 Validation

The validation split carries every decision that has to be made before the final test. It determines the fixed transaction amount traded for each Buy or Sell action, the weight of the risk-aware penalty, and the robustness checks on the transaction fee. In the first study, these decisions used the 12 real months of 2023; in the main study, they use the 300 simulated validation episodes.

The candidate transaction sizes are 0.05, 0.10, 0.25, and 0.50 of the initial capital.

The selection is based on the maximum portfolio exposure that can be achieved from each action rather than on trading performance during the test period. This prevents the action from being chosen simply because it happens to perform well on the final evaluation data.

Once the required configuration decisions have been made, the selected settings are fixed for testing.

4.4.3 Testing

The final evaluation is run once, on data that made no configuration decision. For the first study this is the held-out real months from January 2024 to December 2025. For the main study it is the 600 simulated test episodes, reported per market condition, together with the 180 real transfer months described in Section 4.2. This ensures that results provide an independent assessment of the methods.

Each trained model is evaluated in the same trading environment and using the same evaluation procedure, which includes the initial capital, transaction costs, action constraints, and monthly episode structure.

For the learned agents, actions are selected intentionally by choosing the highest-probability action for the policy-based methods or the highest-valued action for DQN.

The portfolio outcomes are then recorded for each monthly test episode and each random seed, so the stored results contain sufficient information to compare algorithms, state representations, reward functions, and benchmark strategies, which are presented in Chapter 5.

4.5 Experiment Configuration

The implementation was designed so that different experiments could be run by changing their settings without changing the overall trading environment or learning algorithm settings.

For this reason, experiments are defined through configuration files or parameters rather than by manually changing the source code for each run. Each experiment is therefore linked to a specific configuration that records the conditions used to produce it.

This separates experiment settings from the environment's implementation and learning algorithms, allowing the same code to run different state, reward, and action configurations.

4.5.1 Configuration Parameters

The configuration records the main settings that can affect an experiment's outcome: the learning algorithm, the state representation and its dimensionality, the reward formulation, the transaction fee, the fixed transaction amounts, the risk-penalty coefficient where applicable, the action model, the random seed, and the training and evaluation settings.

Using this approach reduces the need to manually change the source code between experiments. It also makes the experimental process easier to track, since each result can be linked directly to the configuration that produced it.

This is particularly important for the experiments reported in Chapter 5, where several state, reward, and action configurations are compared.

4.5.2 Algorithm Configuration

The experiment uses the same starting capital and transaction cost across all three learning algorithms. The initial capital is set to $C_0 = \$10,000$, and a transaction fee of 0.05% is applied to each trade opened.

The main settings used for REINFORCE, DQN, and PPO are shown below.

Parameter	REINFORCE	DQN	PPO
Network	Pyramid MLP $d \rightarrow 64 \rightarrow 32 \rightarrow 3$	Pyramid MLP $d \rightarrow 64 \rightarrow 32 \rightarrow 3$	<i>MlpPolicy</i> $d \rightarrow 64 \rightarrow 32$
Activation	tanh	ReLU	
Optimiser	Adam	Adam	Adam
Learning rate	10^{-3}	10^{-3}	3×10^{-4}

Parameter	REINFORCE	DQN	PPO
Discount factor γ	0.99	0.99	0.99
Replay buffer	—	50,000	—
Batch size	Episode	64	64
Target network update	—	Every 500 steps	—
Exploration	Policy sampling	$\epsilon - greedy$	Policy sampling
ϵ schedule	—	1.0 $\rightarrow$ 0.05	—
PPO clipping	—	—	0.2
GAE	—	—	0.95
PPO epochs	—	—	10
Entropy coefficient	—	—	0
Training budget	300,000 steps	300,000 steps	300,000 steps
Random seeds	42–47	42–47	42–47

All three methods use the same training budget and the same six random seeds to keep the computational budget similar throughout when comparing performance. The pyramid network has about 2,800 trainable parameters for the nine-feature state, against the 3,000 training episodes of the main study. The 300,000-step budget corresponds to roughly 15,000 episode draws over those episodes. Validation sweeps, which only rank configurations rather than report final numbers, run at half this budget (150,000 steps).

The first study used the same configuration except for two settings. Its network was larger, two hidden layers of 128 units ($d \rightarrow 128 \rightarrow 128 \rightarrow 3$, more than 18,000 parameters), and its training budget was 80,000 steps, sized to its 60 training months. The first study therefore had six times more parameters than training decisions; the main study has more training episodes than parameters. Everything else, including the activation functions, optimiser, learning rates, discount factor, and the six seeds, is identical, so the two studies differ only in the data and in the network sized to match it.

The experiments consider the different state and reward configurations described in Chapter 5. These include the standard nine-feature state, the ten-feature state with drawdown, the primary wealth-change reward, the risk-aware reward, and the alternate action configurations.

The implementation also treats each trading position size as a configurable parameter. Its final value is selected using the predefined validation procedure. Chapter 5 reports the different values and their effect on trading results as part of the experimental findings.

4.5.3 Benchmark Implementation

The learned agents are compared with benchmark strategies to provide clear grounds for assessing their performance. These benchmarks are simple trading behaviours that do not require reinforcement learning.

The main benchmark is a buy-and-hold strategy where the initial capital is used to purchase the asset at the start of the episode. The position is then held until the final trading step, when it is closed. The same transaction fee and position closure used for the learned agents apply to this benchmark.

This approach is different from gradually purchasing the asset using the same fixed trade amount available to the learning agents. A gradual purchasing strategy may take several trading days to build a position and, if the fixed trade size is small, it may not become fully invested before the monthly episode ends.

For this reason, the gradual purchase strategy is treated as a separate reference strategy rather than as buy-and-hold. The conventional buy-and-hold strategy gives the learned agents a straight performance target.

The benchmark strategies use the same portfolio conditions, transaction costs, and evaluation procedures as the learned agents. This keeps the comparison consistent and reduces the chance that differences in results stem from inaccuracies in calculating portfolio performance rather than from the strategies themselves.

4.6 Implementation Verifications

Before using the experimental results for analysis, the implementation was tested to ensure that it followed the definitions and constraints established in Chapter 3. The final implementation includes 21 automated tests covering all major components, all of which pass. Appendix A lists every test and the condition it verifies; this section summarises what each group of tests protects.

4.6.1 Environment and Portfolio Tests

The environment tests focus on the correctness of the portfolio calculations during trading. They verify that portfolio wealth matches the cash balance plus the market value of current holdings, that Buy and Sell actions update the portfolio correctly with transaction fees included, that invalid actions are handled correctly, and that each monthly episode ends without an open position. These checks matter because an error in portfolio accounting would affect the reported rewards and performance measures for all three learning algorithms, and could therefore lead to incorrect conclusions even if the algorithms were implemented correctly.

4.6.2 Feature and Data Tests

The feature tests check that the state representation follows the intended information structure defined in Chapter 3. They verify that features do not use future information, that rolling features carry historical

information across episode boundaries, that the state dimensions and feature ordering match the reported configurations, and that the training, validation, and test periods remain separate. This group of tests identified the rolling-feature issue described in Section 4.6.

4.6.3 Reward and Risk Tests

The reward tests verify that the primary wealth-change reward corresponds to the change in portfolio wealth with no risk penalty applied, and that under the risk-aware reward an increasing drawdown does not lead to an increasing reward. An additional test enforces that the risk-aware reward cannot be paired with a state that lacks the drawdown feature, so an inconsistent combination of state and reward settings cannot run accidentally. These tests ensure that the relationship between state representation and reward design described in Chapter 3 is maintained in the implemented system.

4.6.4 Action-Masking Tests

The action-masking tests verify that the environment correctly identifies actions that cannot be executed under the current portfolio conditions: Buy is unavailable when cash is insufficient, and Sell is unavailable when the agent does not hold enough of the asset. They also verify the different configurable fixed trading amounts. A defensive check in the environment converts an invalid action supplied externally to Hold rather than executing it. Logs monitor these checks during the experiments; after the final experiment, the recorded rate of invalid actions during evaluation was zero.

4.7 Implementation Corrections

The verification process identified three implementation issues before the final experiments were run. First, the rolling market features were initially recalculated separately for each monthly episode, so the rolling windows restarted at the beginning of every month instead of continuing from previous observations; this left incomplete feature windows in approximately **43.1% of episode steps**. The correction calculates the rolling features on the continuous dataset before dividing it into monthly episodes. Second, an earlier version of the buy-and-hold benchmark used the same transaction size as the learning algorithms rather than investing fully at the start of each episode, which understated the benchmark by approximately **\$38.95 per month**, or **21.6%**. These values describe only the incorrect implementation, which has been discarded, and are reported to show the potential impact of the error. Third, the risk-aware reward could initially be used without requiring the drawdown feature in the state, creating a mismatch between the information needed by the reward and the information supplied to the agent.

All three issues were treated as implementation errors, because leaving any of them in place would affect the validity of comparisons between the learning algorithms and the benchmark strategies. Each was corrected, verification tests were added or updated where necessary, and the complete automated test batch was run again with all 21 tests passing. The final experiment results in Chapter 5 come only from the corrected implementation.

4.8 Complete Implementation Pipeline

The complete implementation pipeline is a sequence of stages, from collecting historical market data to storing the final evaluation test results.

The first stage loads the SPY OHLCV data, including an initial warm-up period needed to calculate the rolling features. The market features are then calculated on the continuous dataset using only information available up to said observation. Once feature construction is complete, the data are divided chronologically into the training, validation, and test datasets. For the main study, the generator is calibrated on the real training months and produces the simulated splits; the remaining real months form the transfer set.

The observations are then grouped into monthly episodes. At the start of each episode, the portfolio resets to the initial capital with no asset holdings. At each timestep, the learning algorithm receives the current state and selects one of the valid available actions; the environment executes the action with the transaction fee, updates the portfolio, returns the reward, and moves to the next state. REINFORCE, DQN, and PPO all interact with this same environment, and their main difference is how each algorithm processes the experience to update its parameters.

After training and validation, the final experiment settings are fixed, and the trained models are evaluated on the held-out test data: the 2024–2025 months for the first study, and the 600 controlled simulated episodes plus the 180-month real transfer set for the main study. The resulting metrics, monthly results, training information, configuration settings, and performance measures are stored together with information about the software environment used to produce them.

This pipeline clearly links the methodological decisions defined in Chapter 3, their implementation in the system, and the experiment results presented in Chapter 5.

Figure 4.2 summarises the flow.




Figure 4.2: System architecture showing the data flow described in Section 4.7. All learned algorithms and benchmark strategies are evaluated using the same episode runner, ensuring that the comparison uses the same evaluation procedure.

Chapter 5

Experimental Results and Analysis

5.1 Introduction

This chapter reports and analyses the experiment results from the trading environment and learning algorithms described in Chapter 4. The experiments examine the financial performance of the learned agents, the effects of two state representations, their reward functions, and the role of the training data itself.

The experiments are structured around five questions that directly relate to the objectives of the dissertation to increase portfolio wealth:

1. **Performance:** Do the learned agents outperform appropriate simple benchmarks on unseen data after accounting for transaction costs?
2. **Behaviour:** Do the learned agents make use of the market and portfolio information provided in the state representations?
3. **State representation:** Does adding the drawdown feature to the main nine-feature state affect the learned behaviour or financial performance?
4. **Reward:** Does the risk-aware reward improve the balance between return and drawdown, rather than simply encouraging the agent to trade less?
5. **Data:** Do the size and balance of the training data explain the behaviour of the learned agents, and does fixing the data fix the behaviour?

The chapter answers these questions with two studies. The **first study** trained the agents on the 60 real training months and evaluated them on the held-out 2024–2025 months. Its results were negative, and Section 5.4 shows that the negative results share one cause: 60 episodes cannot train networks of this size. The **main study** corrects this. The agents are trained on 3,000 balanced simulated episodes from the generator of Chapter 4, evaluated under controlled market conditions, and finally transferred, unchanged, to 180 real months that include the 2008 financial crisis. Every experimental question the first study asked is asked again in the main study, under data that can support an answer.

A key part of the analysis is determining whether the learned strategies outperform simple strategies. Buy-and-hold, for example, is an important benchmark here. A learned agent producing a positive return does not necessarily indicate that it has learned a useful trading strategy if a simple investment strategy achieves a similar or better result.

For this reason, the analysis considers both **financial performance** and **agent behaviour**. Financial performance is evaluated using measures including wealth change, exposure, trading activity, drawdown,

and Sharpe ratio. Behavioural analysis examines whether an agent's actions change as the state changes. This helps to identify a policy that responds to state information rather than one that follows a fixed trading pattern. The next section sets out the protocols shared by both studies.

5.2 Experiment Protocols and Evaluation Criteria

This section describes how the experiments in this chapter were conducted and how the reported results should be interpreted.

5.2.1 Experimental Configurations

Each configuration trained REINFORCE, Deep Q-learning (DQN), and Proximal Policy Optimisation (PPO) separately in the portfolio environment. The main study trains each seed for 300,000 steps on the simulated episodes; the first study trained for 80,000 steps on the real months, a budget sized to its far smaller dataset.

To account for randomness in reinforcement learning [30], six random seeds were used, with an initial capital of $C_0 = \$10,000$ at the start of every monthly episode. This means six independent training runs all with the same configuration, allowing us to analyse the consistency of each algorithm's results rather than relying on a single training run.

During evaluation, all policies acted deterministically: policy-gradient algorithms select the action with the highest probability, while value-based algorithms select the action with the highest estimated value.

Every experimental decision in this chapter was made on validation data. In the first study, these decisions were made on the 2023 months. In the main study, they are made on the 300 simulated validation episodes: the trade size in Section 5.5, the risk-penalty weight in Section 5.8, and the transaction-fee robustness check in Section 5.9. The simulated test episodes and the real transfer months made no configuration decision and are each evaluated once.

5.2.2 Evaluation Metrics

There are five financial measures used in each experiment report. They are averaged over the episodes and, for learned agents, across the six training seeds:

- Mean wealth change (ΔW) reported in dollars per episode;
- Mean exposure (Δv_t), which represents the proportion of initial capital invested in the asset and is averaged across trading days;
- Trades per episode;
- Mean episode drawdown (Δdd_t);
- and the monthly Sharpe ratio [33], a financial metric that calculates the average return relative to the variation of those daily returns during the month.

The tables also report the **seed spread**, which is the difference between the best and worst result across the six seeds. This is included because there might be large differences between training runs which the average alone may hide [30]. A large spread indicates that the algorithm did not produce a consistent trading behaviour across the different runs.

Financial performance is followed by a behavioural analysis, as returns alone cannot show whether a policy actually used the information in its state. For example, a policy that buys simply because buying is allowed can perform well during a rising market without responding to any market information. The **state-dependence analysis** tests whether a policy’s actions changed when its observed state changed while the valid actions remained the same; if a policy selected different actions at different steps with the same actions available, the difference can be attributed to the observed state.

Three further columns appear in the main study’s tables. “State-dep.” counts how many of the six seeds passed the state-dependence test above; a seed that ignores its state fails no matter how much it earns. “Real” is the mean wealth change per month when the trained policy is run, unchanged, on the 180 real transfer months of Section 5.10. “Down condition” is the mean wealth change per episode on the falling-market test episodes alone, because these are the episodes a risk-aware reward is meant to improve.

5.2.3 Benchmark Strategies

Two non-learning strategies are used as benchmarks. They use the same trading environment, portfolio conditions, and transaction fees as the learned agents.

Buy-and-hold invests the full initial capital on the first trading day and holds the position until the end of the month, when the position is closed. This is the main benchmark as it is a simple passive investment strategy.

Always-buy ($0.25C_0$) purchases a fixed trade amount equal to 25% of the initial capital whenever buying is available, continuing until it has fully invested in that asset. This strategy is useful for comparing against policies that achieve their returns through repeated buying rather than by using information from its observed state, as it also does not use any market or portfolio information when deciding whether to buy.

With the protocols fixed, the next section reports the first study and what it revealed.

5.3 The First Study: Real Data Alone

This section reports the first study: the three algorithms trained on the 60 real training months and tested on the 24 held-out months of 2024–2025. The first study ran the full experimental programme — trade-size calibration, behavioural analysis, the state and reward comparisons, and a fee sweep — and this section reports what each part found. The findings are presented compactly because Section 5.4 shows they share a single cause, and the main study re-examines every one of them under corrected data.

5.3.1 Setup of the First Study

The first study used the larger network of two 128-unit hidden layers and a budget of 80,000 training steps per seed. Its trade size was calibrated on the 2023 validation months by the exposure-ceiling rule described in Section 5.5: a Buy commits a fixed fraction of initial capital, so a small trade size caps the exposure an agent can reach within a month. On the validation months, the $0.25C_0$ size was the smallest that reached at least 90% of buy-and-hold’s exposure and was fixed for all experiments.

5.3.2 Results of the First Study

Table 5.1 reports the first study’s held-out test results in the final configuration: nine-feature state, wealth-change reward, $0.25C_0$ trade size, 0.05% fee, six seeds.

Table 5.1: First-study test results over the 24 held-out real months (2024–2025), benchmarks and the three learned agents in their final configuration, six seeds each.

Policy	ΔW (\$)	Spread	Exposure	Trades	DD	Sharpe	State-dep.
Buy-and-hold	+180.25	—	0.913	2.0	0.0310	0.638	—
Always-buy ($0.25C_0$)	+179.75	—	0.841	5.0	0.0274	0.684	—
REINFORCE	+171.51	24.7	0.810	10.3	0.0266	0.675	0/6
Deep Q-learning	+74.68	71.4	0.312	8.2	0.0124	0.507	6/6
PPO	+127.35	179.8	0.598	3.7	0.0195	0.562	0/6

The main result is negative: no learned agent beat buy-and-hold’s +\$180.25 per month. REINFORCE came closest at +\$171.51, PPO earned +\$127.35, and Deep Q-learning earned +\$74.68. Always-buy at $0.25C_0$, which needs no market information at all, earned +\$179.75 — within \$0.50 of buy-and-hold — so the closest learned agent did not even match a strategy that buys unconditionally.

The behavioural analysis explains where those returns came from. REINFORCE’s seeds converged on always-buying or on fixed trading patterns: none of the six changed its action when the state changed while the same actions stayed valid. PPO split between the always-buy corner and total inactivity, with a seed spread (\$179.80) larger than its mean return. Deep Q-learning was the opposite: all six seeds were state-dependent, yet it earned the least, with the lowest exposure (0.312). Figure 5.1 shows every seed: the policy-gradient seeds sit at the corners, and only the Deep Q seeds lie between them.




Figure 5.1: First study: behaviour of each seed on the validation months. Hollow-marker seeds show actions that do not depend on their state information; filled-marker seeds are state-dependent ones. Policy-gradient methods (REINFORCE, PPO) cluster at high exposure or collapse towards zero, while all six Deep Q-learning seeds fall between those corners.

The remaining first-study experiments reinforced the same picture, and three results summarise them. First, adding the drawdown feature to the state changed mean wealth by less than the seed spread for every algorithm (REINFORCE $-\$4.12$, Deep Q $+\$3.13$, both against spreads above $\$24$), so the ten-feature state offered no measurable advantage under the wealth-change reward. Second, the risk-aware reward ($\lambda = 0.10$, selected on validation) cut Deep Q's drawdown by 41% but only by cutting its exposure roughly in half; its Sharpe ratio fell from 0.505 to 0.416, so the penalty bought less risk with less participation, not better risk management. Third, the fee sweep from 0 to 50 basis points showed only Deep Q-learning reducing its trading in response (10.9 to 1.9 trades per month, an 83% reduction), consistent with it being the only method whose policies read their state.

Even where its state-dependence should have paid, it did not. March 2025 was the worst test month: buy-and-hold lost $\$398.41$. Deep Q-learning held a mean exposure of 0.490 that month, higher than its overall average of 0.312, and lost $\$184.04$ — about the same one-third scaling of the benchmark it showed everywhere else. Its lower drawdown came from holding less overall, not from recognising when to step aside.

The first study therefore ends with a contradiction. The only method that demonstrably used its state earned the least, and the methods that ignored their state earned the most. The next section explains where this contradiction comes from.

5.4 The Data Bottleneck

This section explains why the first study’s results look the way they do. The explanation is not about any single algorithm. It is about how little data the agents were given.

The training split contains 60 monthly episodes. The first study’s networks have more than 18,000 trainable parameters. A model with that much capacity can reproduce 60 training months exactly, in the same way that a student who memorises 60 practice problems scores perfectly on those problems and still fails a new exam. Nothing forces the network to learn a rule that transfers, because memorising the answers is easier.

The training months are also imbalanced. Of the 60 months, 29 rose by more than 2% and only 16 fell by more than 2%, and the test years 2024–2025 rise almost throughout. Under that distribution, “always buy” is a nearly optimal answer that requires no state information at all. This is the same problem as training a classifier on a dataset that is 99% one class: predicting the majority class scores well and learns nothing. The always-buy policies of the policy-gradient agents in Section 5.3 are exactly this artefact.

The same reasoning explains why Deep Q-learning performed worst on the real test months. A value-based method must estimate the value of each state it visits, which needs many visits to similar states under a stable state distribution. The 60 episodes of a shifting market provide neither. The behavioural finding stands — six of six DQN seeds did condition on their state — but the values behind those decisions were estimated from too little data, so conditioning on them did not pay.

None of this can be tested on the real data, because the real data is the problem. Testing it needs a setting where the number of episodes, the balance of market conditions, and the match between training and test distributions are all controlled. That is what the main study provides: 3,000 balanced simulated episodes, a network shrunk from 18,000 parameters to 2,800, and controlled testing in each market condition, with the real months kept back as a final transfer test. The remainder of this chapter reports it, repeating each of its experiments under the corrected data.

5.5 Fixed Trade Size Under the Simulated Data

Before the algorithms are compared, the fixed transaction size is re-calibrated on the simulated validation episodes, exactly as the first study calibrated it on the 2023 months. The trade size matters because each Buy commits a fixed fraction of initial capital, so the size sets a hard ceiling on the exposure an agent can reach within a 21-bar episode. A policy judged against buy-and-hold must first be able to reach buy-and-hold’s exposure in principle; otherwise a poor result would only reflect the exposure limit rather than the quality of the learned policy.

The calibration needs no training. For each candidate size, a policy that buys whenever buying is valid measures the highest exposure that size permits. The selection rule is the first study’s rule: the smallest size whose ceiling reaches 90% of buy-and-hold’s exposure. Table 5.2 shows the outcome on the 300 simulated validation episodes.

Table 5.2: Attainable exposure for each fixed trade amount on the simulated validation episodes. The policy buys whenever a purchase is valid.

Trade amount	Buy days to full investment	Mean maximum exposure	Reaches 90% target
$0.05C_0$	20	0.452	no
$0.10C_0$	10	0.691	no
$0.25C_0$	4	0.834	yes
$0.50C_0$	2	0.882	yes
Buy-and-hold	—	0.906	—

The ceilings are almost identical to the first study’s (0.834 against 0.838 for the $0.25C_0$ size), which is expected: the ceiling is a property of the episode length and the action model, not of which market generated the prices. The $0.25C_0$ size again clears the 90% threshold of 0.815 as the smallest sufficient size and is fixed for every experiment that follows. The $0.10C_0$ size again cannot reach the target even in principle, and $0.50C_0$ adds little ceiling while halving the agent’s control over position changes.

One column from the first study’s version of this table is deliberately not repeated. On the balanced validation episodes, the wealth change of the always-buy policy is close to zero at every size (buy-and-hold itself earns only +\$5.66), so return no longer separates the candidates. This is the first sign of the property the simulated data was built for: no fixed trading pattern is profitable on balanced data. The main comparison relies on this property.

5.6 Main Results: The Controlled Comparison Across Market Conditions

This section reports the central experiment of the dissertation: the three algorithms trained on the 3,000 balanced simulated episodes and evaluated on 600 fresh episodes, 200 in each market condition. Six seeds per algorithm, 300,000 steps each, with the pyramid network of about 2,800 parameters.

5.6.1 Benchmarks Under Balanced Conditions

The balanced test set changes what the two benchmarks earn, and this change is intentional. Buy-and-hold now loses \$60.29 per episode on average. It gains \$510.17 in the average up episode, loses \$710.06 in the average down episode, and roughly breaks even when the market is flat. Always-buy at $0.25C_0$ behaves the same way at slightly smaller magnitudes, losing \$58.70 per episode. A policy that never trades ends every episode with exactly its starting capital, a wealth change of \$0.00.

On the first study’s real test months, buying and holding earned \$180.25 because the period only rose. On the balanced episodes, no fixed pattern is profitable: constant buying loses money and never trading earns nothing. Any agent that earns money here must be reading its observations and behaving differently in different conditions. The benchmarks are therefore no longer a difficult target to beat; instead, they test whether the agent conditions on its state.

5.6.2 Main Results

Deep Q-learning earned **+\$126.94 per episode** on average across the six seeds, with a seed standard deviation of \$7.30. Every other strategy, learned or fixed, lost money or broke even on the balanced test set. All six DQN seeds were classified as state-dependent, and all six were profitable. Table 5.3 summarises the study.

Table 5.3: Main study: mean wealth change (\$) per episode on the 600 balanced test episodes, the mean within each market condition, and the mean over the 180-month real transfer set (Section 5.10). Mean and standard deviation are across six seeds.

In the table, “Sim test” is the mean wealth change per episode over all 600 simulated test episodes, and “Std” is how much that mean varies across the six seeds. “Up”, “Down”, and “Flat” break the same result down by the condition the episode was generated from. “Real” is the mean wealth change per month when the same trained policy is run, unchanged, on the 180 real transfer months.

	Sim test	Std	Up	Down	Flat	Real
DQN	+126.94	7.3	+473.72	−103.35	+10.45	+60.28
REINFORCE	−5.71	52.2	+228.58	−248.32	+2.60	+31.66
PPO	−11.16	27.3	+67.69	−99.04	−2.14	+8.32
Buy-and-hold	−60.29	—	+510.17	−710.06	+19.00	+80.24
Always-buy ($0.25C_0$)	−58.70	—	+477.22	−663.91	+10.58	+76.48

The columns for each market condition show what the Deep Q-learning agents actually learned. In up episodes they earn \$473.72 of the \$477.22 that always-buying earns, so they capture 99% of the available upside. In down episodes they lose \$103.35 where always-buying loses \$663.91, an 84% reduction in the loss. In flat episodes they roughly match the market. That is the state-dependent behaviour the first study looked for and could not find: stay invested while the market rises, step aside while it falls. Figure 5.2 shows both views: the mean for each market condition on the left, and every seed’s simulated result against its real transfer result on the right.

5.6.3 The Policy-Gradient Methods Collapse to Corners

REINFORCE and PPO did not reach that policy reliably. Across the six REINFORCE seeds, three converged to never trading (a wealth change of exactly \$0.00), two converged to always buying (−\$58.70, matching the always-buy row to the cent), and one learned a genuinely conditioned policy. PPO was similar: five seeds never traded and one adopted a fixed buying pattern. No PPO seed was state-dependent.

There is a reason the seeds settle on these corners. On balanced data, unconditional buying loses money, so never trading is a locally safe answer: it earns nothing, but it cannot lose. A policy-gradient method that drifts toward either deterministic corner stops exploring, because the gradient of a near-deterministic policy carries almost no signal. The value-based method does not have this failure mode, since epsilon-greedy exploration keeps visiting both actions regardless of what the current policy prefers.




Figure 5.2: Main study. Left: mean wealth change per episode by test condition for each algorithm and for always-buy. Deep Q-learning keeps almost all of the up-market profit while cutting the down-market loss by 84%. Right: each seed's simulated test result against its real transfer result; filled markers are state-dependent seeds, hollow markers are mask-only seeds. All six Deep Q-learning seeds cluster in the profitable upper-right; the policy-gradient seeds sit at the never-trade origin or the always-buy corner.

One REINFORCE seed did learn a conditioned policy, and its result shows what the method could have achieved. It earned \$83.13 on the simulated test, keeping \$417.02 of the up-market profit while cutting the down-market loss to \$162.11. The desired policy is therefore within REINFORCE's reach, but the training process does not find it dependably. One seed in six shows that the policy can be learned, not that the method learns it reliably.

The comparison so far uses the nine-feature state and the plain wealth-change reward. The next three sections repeat the first study's remaining experiments — the state representation, the risk-aware reward, and the transaction fee — under the corrected data.

5.7 Effect of the State Representation

Chapter 3 defines the nine-feature state as the main state representation and treats drawdown as an additional risk-related feature. The first study found that adding drawdown changed nothing measurable, but that finding was suspect for the same reason as every first-study finding: an agent that memorises 60 months has no use for any feature. This section asks the question again under the corrected data: does adding drawdown to the state change performance or behaviour when the reward stays wealth-change? The comparison keeps everything else fixed: the same simulated training and test episodes, the same 0.25 C_0 trade size, the same six seeds. Table 5.4 reports both configurations. As in the rest of the chapter, "State-dep." counts the seeds whose actions respond to the market information, and "Real" is the mean monthly wealth change over the 180 real transfer months.

Table 5.4: Nine- against ten-feature state on the simulated test episodes (wealth-change reward, six seeds), with the real-transfer mean.

Method	State	ΔW (\$)	Std	Exposure	DD	State-dep.	Real
Deep Q	9-feature	+126.94	7.3	0.585	0.0291	6/6	+60.28
Deep Q	10-feature	+131.31	6.5	0.571	0.0276	6/6	+58.47
REINFORCE	9-feature	−5.71	52.2	0.356	0.0246	1/6	+31.66
REINFORCE	10-feature	−45.55	18.3	0.774	0.0569	3/6	+69.72
PPO	9-feature	−11.16	27.3	0.123	0.0092	0/6	+8.32
PPO	10-feature	+4.73	11.6	0.050	0.0030	1/6	+3.70

For Deep Q-learning, the first study’s conclusion survives the correction. Adding drawdown moves the mean from +\$126.94 to +\$131.31, a change of +\$4.37 against a seed standard deviation near \$7, and every behavioural measure stays the same: six of six seeds state-dependent, similar exposure, similar trading rate, and a slightly smaller down-market loss (−\$87.28 against −\$103.35). The extra feature neither helps nor hurts an agent whose objective gives it no reason to use its loss history. This matches the principle the first study suggested: information in the state only matters when the objective makes it matter, which the next section tests directly.

The policy-gradient rows mainly reflect the corner instability of Section 5.6 rather than an effect of the feature. REINFORCE’s six ten-feature seeds all traded this time, with five settling near the always-buy corner and none finding the conditioned policy that its nine-feature run found once. PPO barely changed: five seeds still never trade, and the sixth learned only a weakly conditioned policy. Because the seeds split between the two corner policies, the mean differences between the state configurations mainly show how many seeds landed on each corner, not what the feature contributed. Only Deep Q-learning, whose six seeds agree, can answer the feature question, and its answer is that the feature has no measurable effect under the wealth-change reward.

5.8 Effect of the Risk-Aware Reward

The first study’s strongest negative finding was about the risk-aware reward: penalising drawdown reduced risk only by reducing participation, and the Sharpe ratio fell. This section repeats the experiment on the corrected data, and the conclusion changes. The reward is the one defined in Chapter 3,

$$r_t^{risk} = \Delta W_t - \lambda C_0 \Delta dd_t,$$

where Δdd_t is the increase in portfolio drawdown at trading day t (zero when drawdown does not rise) and λ controls the penalty. It is paired with the ten-feature state, which carries the drawdown information the penalty refers to.

5.8.1 Selection of the Penalty Size

The penalty weight was selected on the 300 simulated validation episodes, using Deep Q-learning as the selector and the first study’s rule: the largest λ that keeps at least half of the unpenalised return and at least half of the unpenalised exposure. The unpenalised agent earned +\$158.29 at exposure 0.598, so the floors were $\Delta W \geq +\$79.14$ and exposure ≥ 0.299 . Table 5.5 and Figure 5.3 show the sweep.

Table 5.5: Risk-penalty size on the simulated validation episodes (Deep Q-learning, six seeds). The rule requires $\Delta W \geq +\$79.14$ and exposure ≥ 0.299 ; the largest value meeting both is $\lambda = 0.25$.

λ	Mean ΔW	Exposure	Drawdown	DD change	Passes
0 (unpenalised)	+\$158.29	0.598	0.0296	—	—
0.05	+\$166.58	0.575	0.0271	−8.4%	yes
0.10	+\$159.78	0.562	0.0270	−8.8%	yes
0.25	+\$99.13	0.338	0.0158	−46.6%	yes
0.50	+\$15.65	0.039	0.0017	−94.3%	no
1.00	+\$0.61	0.007	0.0005	−98.3%	no
2.00	−\$0.04	0.000	0.0000	−100%	no




Figure 5.3: Risk-penalty sweep on the simulated validation episodes (Deep Q-learning, six seeds). Left: mean wealth change against λ with the +\$79.14 requirement floor; the dotted line marks the selected $\lambda = 0.25$. Right: exposure and drawdown relative to their unpenalised values. Exposure collapses between $\lambda = 0.25$ and 0.5, one grid step further out than in the first study.

The sweep has the same shape as in the first study, but the collapse point moves. Exposure still collapses once the penalty dominates the return signal — from 0.338 at $\lambda = 0.25$ to 0.039 at $\lambda = 0.5$ — and large penalties still stop trading entirely. But the collapse now happens one grid step later, and below that point the penalty reduces drawdown at little cost to return: at $\lambda = 0.25$ drawdown falls 47% while the return stays well above the floor. The rule therefore selects $\lambda = 0.25$, a stronger penalty than the first

study could afford ($\lambda = 0.10$), because agents with reliable value estimates tolerate more penalty before abandoning the market.

5.8.2 Effect on the Test Set

The selected reward was then evaluated once on the 600 simulated test episodes. Table 5.6 compares it with the wealth-change reward at the same ten-feature state. The “Down condition” column reports the falling-market episodes on their own, because these are the episodes the penalty is meant to improve. “Real” is the mean monthly wealth change over the 180 real transfer months.

Table 5.6: *Wealth-change versus risk-aware reward ($\lambda = 0.25$) on the simulated test episodes, ten-feature state, six seeds.*

Algorithm	Reward	ΔW (\$)	Exposure	DD	Sharpe	Down condition	Real
Deep Q	wealth	+131.31	0.571	0.0276	0.285	−87.28	+58.47
Deep Q	risk	+129.38	0.432	0.0190	0.332	−23.26	+36.58
REINFORCE	wealth	−45.55	0.774	0.0569	−0.063	−603.39	+69.72
REINFORCE	risk	0.00	0.000	0.0000	0.000	0.00	0.00
PPO	wealth	+4.73	0.050	0.0030	0.018	−21.82	+3.70
PPO	risk	+0.04	0.001	0.0001	0.001	−0.27	+0.05

For Deep Q-learning, the risk-aware reward now does what it was designed to do. The mean return is essentially unchanged (+\$129.38 against +\$131.31, a 1.5% difference), while drawdown falls 31% and the Sharpe ratio rises from 0.285 to 0.332. The results by market condition show where the improvement comes from: the down-market loss shrinks from −\$87.28 to −\$23.26, a 73% reduction, while the up-market gain keeps 87% of its size (+\$411.01 against +\$470.01). The agent did not simply hold less of the asset; it held less of the asset *when the market was falling*. This is precisely the selective risk management the first study found no evidence of.

The first study’s conclusion is therefore overturned, and the overturning has the same explanation as everything else in this chapter. Penalising drawdown can only produce selective behaviour if the agent’s value estimates can distinguish a falling market from a rising one, and value estimates trained on 60 months could not. The first study measured the penalty’s effect on an agent whose estimates could not make that distinction; this study measures its effect on an agent whose estimates can.

Two qualifications remain. First, the policy-gradient methods collapse completely under the penalty: every REINFORCE seed and effectively every PPO seed stops trading, because for an unreliable policy the penalty makes abstention strictly safer than participation. The risk-aware reward is only usable by the method whose estimates support it. Second, the defensiveness costs return on transfer: the risk-trained Deep Q agents earn +\$36.58 per month on the 180 real months against +\$58.47 for their wealth-trained counterparts, with a lower transfer Sharpe (0.156 against 0.218). The real window rose more often than

the balanced simulation assumes, so a level of caution chosen on balanced conditions turns out to be more than the real market rewards. The wealth-trained agent remains the better real-market configuration; the risk-trained one demonstrates that the reward mechanism works when the data supports it.

5.9 Transaction-Fee Sizing

Transaction costs are an important part of the trading environment because frequent trading can reduce the returns produced by a strategy, even when individual decisions appear reasonable. The main experiments assume a fee of 0.05% per trade, which is defensible for a liquid instrument like SPY but is still an assumption. A conclusion that only holds under an optimistic cost assumption would not be reliable, so the fee was swept from 0 to 50 basis points on the simulated validation episodes with the nine-feature state and everything else held fixed. Six seeds were freshly trained for each algorithm at every fee level. Table 5.7 and Figure 5.4 report the outcome.

Table 5.7: Transaction-fee sizing on the simulated validation episodes; six seeds. Trades are per episode. Fees are shown in basis points.

Fee	B&H ΔW	REINF. ΔW	REINF. trades	Deep Q ΔW	Deep Q trades	PPO ΔW	PPO trades
0 bps	+\$15.67	+\$13.58	6.1	+\$172.25	5.5	+\$6.85	11.3
5 bps	+\$5.66	+\$6.44	2.6	+\$159.63	4.9	−\$7.11	7.0
10 bps	−\$4.34	−\$1.52	0.8	+\$152.30	4.8	−\$1.52	0.8
25 bps	−\$34.29	\$0.00	0.0	+\$142.24	4.1	−\$6.50	0.8
50 bps	−\$83.99	\$0.00	0.0	+\$95.27	3.7	+\$2.72	0.3




Figure 5.4: Transaction-fee sizing on the simulated validation episodes. Left: mean wealth change as the fee rises; Deep Q-learning declines but remains profitable at every level, while buy-and-hold turns negative by 10 basis points. Right: trades per episode; the policy-gradient methods abandon trading as costs rise, while Deep Q-learning trims its rate only slightly.

The most important number is Deep Q-learning’s +\$95.27 at 50 basis points, a fee ten times the study’s

assumption. Its profit declines as costs rise, as expected, but it remains positive at every level while buy-and-hold turns negative by 10 basis points (the balanced episodes give it no upward drift to offset the cost). The agent’s profitability therefore does not depend on the optimistic fee assumption.

The trading columns explain why. The Deep Q agents make about five trades per episode, and those trades are condition-selective: entering when the features say the market is rising, exiting when they say it is falling. Each trade is worth far more than even a punitive fee, so the agents keep trading — the rate falls only from 5.5 to 3.7 — and keep most of their profit. This differs from the first study, where DQN cut its trading 83% as fees rose. The first study’s agent traded on noisy value estimates, so its trades were marginal and a higher fee tipped them into unprofitability; this study’s agent trades on reliable estimates, so its trades absorb the fee instead.

The policy-gradient methods respond by giving up trading. REINFORCE’s trading falls from 6.1 trades per episode to zero by 25 basis points, and PPO’s from 11.3 to 0.3: rising costs make the abstention corner more attractive, and the corner behaviour described in Section 5.6 completes the collapse. Fees, like the balanced data itself, separate the method that evaluates each trade from the methods that settle on a fixed pattern. The final experiment evaluates the trained policies on real market data.

5.10 Transfer to Real Markets

The last question is whether policies trained purely on simulated episodes still perform when they are evaluated on real data. Each trained policy was evaluated, unchanged, on the 180-month transfer set of Chapter 4: every real SPY month from 2006 to 2017 and from 2023 to 2025, excluding only the 2018–2022 calibration window. No agent, and no tuning decision, ever saw any of these months. The window includes the 2008 financial crisis, the 2011 correction, and the 2015–16 drawdowns, so it tests exactly the conditions the simulated down condition stands in for.

5.10.1 Overall Transfer Results

Table 5.8 reports the transfer test. The Deep Q-learning agents earned +\$60.28 per month on average, profitable on every seed (range \$55.15 to \$66.32), against +\$80.24 for buy-and-hold. They did so at a mean exposure of 0.709 against the benchmark’s 0.908, with roughly two thirds of its drawdown (0.0240 against 0.0367).

Table 5.8: Transfer test on the 180 real months: overall and per window, with risk measures. Learned agents are six-seed means.

	ΔW (\$)	2006–17	2023–25	Exposure	DD	Sharpe
Buy-and-hold	+80.24	+56.34	+175.81	0.908	0.0367	0.198
Always-buy ($0.25C_0$)	+76.48	+51.69	+175.63	0.835	0.0341	0.213
DQN	+60.28	+40.14	+140.84	0.709	0.0240	0.224
REINFORCE	+31.66	+21.34	+72.94	0.378	0.0150	0.092
PPO	+8.32	+4.28	+24.45	0.124	0.0052	0.026

This comparison has two halves, and both are needed. On raw return, buy-and-hold still wins: the market rose over these two decades, and a defensive agent gives up some of that rise. On risk-adjusted return, the ordering flips. The Deep Q agents' mean monthly Sharpe ratio is 0.224 against buy-and-hold's 0.198, and five of the six seeds match or beat the benchmark on this measure. In the first study, Deep Q-learning earned less return per unit of risk than buy-and-hold; on the extended window, it earns more. The difference is that this window contains real declines, so the downside protection the agent learned in simulation finally has something to protect against.

5.10.2 Worked Example: The 2008 Crisis

October 2008 is the worst month in the transfer set. Buy-and-hold lost \$1,665.24 that month, 16.7% of its starting capital. All six Deep Q-learning agents held **zero exposure for the entire month** and lost nothing. None of them had ever seen October 2008, or any real month: their training data was simulated, and their calibration months (2018–2022) lie a decade later. The agents stepped aside because the month's volatility features matched the simulated down condition they had learned to avoid.

The same pattern holds across the crisis. From September 2008 to March 2009, buy-and-hold lost \$3,045.66 in accumulated monthly wealth changes; the agents lost \$180.34. In September 2008 (the Lehman collapse) they held 0.08 exposure and lost \$178.73 against the benchmark's \$894.23; from October through February their exposure was zero or negligible in every seed.

The cost side of the same behaviour is equally visible. Buy-and-hold gained \$1,068.28 in the December 2008 rebound and \$1,332.73 in March 2009 when the market turned; the agents, still observing crisis-level volatility, sat both months out. This is the trade-off of the defensive policy in concrete form: the agent gives up part of the recovery in exchange for avoiding part of the crash. Over the whole window the missed gains add up to more than the avoided losses — that is why buy-and-hold keeps the higher raw return — but the agent takes far less risk to earn what it earns, which is what the Sharpe comparison summarises. Figure 5.5 shows the whole window: the agents track the benchmark's rises at reduced scale, and the gap between the two lines through 2008–09 is where the drawdown advantage is earned.

5.10.3 What Transfers and What Does Not

The transfer results also measure the domain gap. The generator produces simpler price paths than the real market, so a policy trained on it cannot be expected to transfer at full strength [24]. Three observations describe the gap. First, profitability transfers: every Deep Q seed is profitable over 180 real months. Second, the behaviour transfers: the agents reduce exposure in real declines exactly as they learned to do in simulated ones. Third, precision does not fully transfer: the agents' timing around real turning points is imperfect — sitting out the December 2008 and March 2009 rebounds is the clearest case — because real markets change condition less cleanly than the generator's fixed ones.

The algorithm ordering on real data matches the simulated ordering (Deep Q, then REINFORCE, then PPO), which is further evidence that the simulated results measure something the real market confirms. The next section discusses what the two studies establish together.




Figure 5.5: Cumulative wealth change over the 180 real transfer months: buy-and-hold against the Deep Q-learning agents (six-seed mean, with the min-max band shaded). The dotted line marks the excluded 2018–2022 calibration window. The agents track the benchmark’s rises at reduced scale and flatten through the 2008–09 crisis, which is where most of their drawdown advantage is earned.

5.11 Discussion of the Main Findings

The experiments answer the five research questions of Section 5.1. This section takes each in turn, using the first study and the main study together, because on three of the five questions the first study alone would have given the wrong answer.

5.11.1 Performance

On raw return against buy-and-hold, the learned agents lose on real data. In the first study, the best agent finished \$8.74 per month behind the benchmark on the 2024–2025 months. On the 180-month transfer set, the Deep Q agents earned +\$60.28 per month against buy-and-hold’s +\$80.24. A passive index position over a period that mostly rises is very hard to outrun, which is consistent with the market-efficiency expectation set out in Chapter 2 [31].

On risk-adjusted return, the main study’s agents win. Their transfer Sharpe ratio of 0.224 exceeds buy-and-hold’s 0.198, with five of six seeds at or above the benchmark, because the agents avoid most of the benchmark’s worst losses at a modest cost in upside. And on the balanced simulated test, where no fixed pattern is profitable, Deep Q-learning is the only strategy of any kind that makes money (+\$126.94 per episode). The performance answer is therefore conditional: the agents do not beat the market’s average direction, but they earn their return with measurably less risk, and they demonstrably possess the state-conditional skill that the raw-return comparison obscures.

5.11.2 Behaviour

All three methods condition on the same state, maximise the same reward, and see the same data, so the first study’s behavioural split — six of six Deep Q seeds state-dependent, none of twelve policy-gradient

seeds — cannot mean that one method “uses the state” and the others cannot. The difference lies in what each method must estimate. A value-based method estimates the value of every state it acts in, which takes many visits to similar states under a stable distribution; the first study’s 60 shifting months could not support those estimates, so Deep Q-learning conditioned on noise and earned the least. The policy-gradient methods only need a policy that scores well on the training months, and on data from a mostly rising market, always buying scores well without using any state information.

The main study confirms this account by inverting the ranking. Given 3,000 balanced episodes, the value estimates become reliable and Deep Q-learning wins on every seed, while the policy-gradient methods lose that easy answer and collapse to corner policies: across all main-study configurations, their seeds land overwhelmingly on never-trade or fixed-buying policies. The single conditioned REINFORCE seed shows the optimum is reachable by policy-gradient training but not reliably found by it, at least not at these settings. The ranking on either dataset alone would mislead; together, the two rankings identify data volume and balance as the deciding variable.

5.11.3 State Representation

Adding the drawdown feature under the wealth-change reward changed nothing measurable in either study. Adding the same feature under a reward that penalises drawdown produced selective downside protection and a higher Sharpe ratio. Together, the two results support one conclusion: information in the state only helps when the learning objective gives the agent a reason to use it. The choice of features and the choice of objective therefore cannot be treated as separate decisions.

5.11.4 Reward Formulation

The first study concluded that the drawdown penalty was a participation tax. The main study shows that with reliable value estimates the same penalty produces exactly the intended behaviour: a 73% smaller down-market loss at a 1.5% cost in mean return. Both studies agree on the failure mode — past a threshold weight, abstention becomes the best reply and exposure collapses — but the corrected study locates the threshold further out and shows genuine risk management below it. The remaining weakness is calibration across domains: a level of caution chosen on balanced simulated conditions is too high for real markets that mostly rise, which is why the wealth-trained agents transfer better than the risk-trained ones.

5.11.5 Transaction Costs

The fee sweep separates fragile advantages from robust ones. The first study’s agents, trading on noisy estimates, lost most of their activity and profit as fees rose. The main study’s Deep Q agents remain profitable at a fee ten times the main assumption, because condition-selective trades are worth far more than they cost. An advantage that survives a tenfold cost increase is unlikely to be an artefact of the cost assumption.

5.11.6 Seed Consistency

The clearest single diagnostic across both studies is the across-seed spread. Deep Q-learning's main-study seed standard deviation is \$7.30 on a mean of +\$126.94; REINFORCE's is \$52.17 on a mean near zero, because its seeds split between corners. A method whose six runs agree is showing a consistent property of the method; a method whose runs disagree is showing the effect of chance. Single-seed evaluations, still common in the applied literature reviewed in Chapter 2, could have reported the one successful REINFORCE seed (+\$83.13) as a working trading system.

5.11.7 Implications for the Research Objectives

Overall, the two studies together answer the main research question of whether reinforcement learning algorithms can learn profitable, state-dependent trading strategies for a portfolio.

On real data at the volumes a single asset provides, they cannot: the 60 training months support memorisation, not learning. Given enough balanced data, they can, and the capability the main study revealed was not missing from the method; it was missing from the data. The main constraint in this problem was not the reward function, the state design, or the trade size, although each had a smaller effect. It was the amount and balance of training data, and once that constraint was removed, the remaining differences between algorithms were differences in how they explore and what they must estimate.

5.12 Chapter Summary

The chapter reported two studies. The first study trained the agents on the 60 real training months and tested them on 2024–2025: no agent beat buy-and-hold, the policy-gradient methods ignored their observations, and the one state-dependent method earned the least. The diagnosis attributed all of it to data: 60 episodes cannot train networks of 18,000 parameters beyond memorisation, and a market that mostly rises makes constant buying a nearly optimal answer.

The main study removed those conditions and re-ran everything. On 3,000 balanced simulated episodes with a 2,800-parameter network, Deep Q-learning earned +\$126.94 per episode on a test set where no fixed trading pattern is profitable, on all six seeds, capturing 99% of the rising-market profit and cutting the falling-market loss by 84%. The trade-size rule re-selected $0.25C_0$; the drawdown feature again changed nothing under the wealth objective; the risk-aware reward, which had only reduced participation in the first study, now cut down-market losses by 73% at almost no cost in return; and the profit survived a tenfold fee increase. Transferred unchanged to 180 real months, the policies stayed profitable on every seed, held zero exposure through October 2008, and finished with a higher Sharpe ratio than buy-and-hold. The policy-gradient methods collapsed to fixed corner policies throughout. The conclusion that survives both studies is that the data, not the algorithms, set the limit on what was learnable — and that value-based learning with sustained exploration was the only method able to exploit the data once it was available. Chapter 6 draws the conclusions.

Chapter 6

Conclusion and Future Work

6.1 Summary of the Work

This dissertation examined whether reinforcement learning algorithms can learn profitable trading policies that utilize market and portfolio information from the observed states when trading assets in a portfolio.

The study designed a complete experimental system for trading the SPY exchange-traded fund. The system uses daily market data organised into monthly episodes, includes transaction costs and action masking, and uses a nine-feature state representation with an optional drawdown feature. Two reward formulations were tested: the primary wealth-change reward and an optional risk-aware reward that penalised drawdown, along with three learning algorithms: REINFORCE, Deep Q-learning (DQN), and Proximal Policy Optimisation (PPO).

Experimental decisions were made on validation data throughout: the 12-month 2023 dataset in the first study, and 300 simulated validation episodes in the main study. Test data, real or simulated, was used only for final evaluations.

The experiments followed a fixed sequence. First, the trade size was adjusted to ensure the agents could reach a reasonable level of exposure relative to the buy-and-hold benchmark. The behavioural experiment then determined whether the learned policies actually respond to the information in their states. Once these environment, calibration and behavioural choices had been decided, the main performance of the learned agents was evaluated and compared against the benchmark strategy on the unseen test period.

Additional experiments then tested whether adding drawdown to the state changed the results, whether a risk-aware reward improved loss management in the portfolio trade, and how the different learning methods responded to changes in transaction costs.

The real-data results then motivated a second study. The failures of the first were diagnosed as a data problem: 60 training episodes cannot train networks of 18,000 parameters to do anything except memorise. A market simulator with three controlled conditions — up, down, and flat — was calibrated on the training years, 3,000 balanced episodes were generated, the network was reduced to about 2,800 parameters, and all three algorithms were retrained and evaluated under controlled rising, falling, and flat conditions. The trained policies were then transferred, unchanged, to 180 real months that no training or tuning step ever read — 2006 to 2017 and 2023 to 2025, including the 2008 financial crisis.

6.2 Conclusions Against the Research Questions

The experiments in Chapter 5 address the five questions asked in 5.1.

6.2.1 Performance

The first question considered whether the reinforcement learning agents could outperform the buy-and-hold benchmark on unseen data. The answer depends on which data, and on which measure.

The main study gives the positive half of the answer. On the balanced simulated test set, where buy-and-hold loses \$60.29 per episode, Deep Q-learning earned +\$126.94 and was the only strategy of any kind to make money. On the 180-month real transfer set, buy-and-hold kept the higher raw return (+\$80.24 against +\$60.28 per month), but the Deep Q agents earned it at three-quarters of the exposure and two-thirds of the drawdown, and their mean Sharpe ratio of 0.224 exceeded the benchmark's 0.198. On raw return over a rising market, buy-and-hold remains unbeaten; on risk-adjusted return over a window containing real crashes, the learned agents finished ahead.

The first study supplies the negative half. Trained on the 60 real months alone, no learned agent outperformed buy-and-hold: the benchmark produced a mean wealth change of +\$180.25 per month on the 2024–2025 test months, while the best learned agent, REINFORCE, achieved +\$171.51. PPO achieved +\$127.35 and DQN +\$74.68. The test period favoured remaining invested, with positive market performance in 17 of the 24 months, so buy-and-hold was difficult to beat.

Transaction costs and the trade size cannot explain these differences, since they were the same throughout the experiments, providing a basis for fair comparison.

6.2.2 Behaviour

The second question examined whether the learned policies genuinely use the state information provided to them.

The behavioural split is the finding on which the two studies agree most strongly. In the main study, all six Deep Q-learning seeds were state-dependent and earned +\$126.94 per episode, while the policy-gradient seeds overwhelmingly collapsed to fixed corner policies — never trading, or buying constantly — once the balanced data removed the free always-buy answer. In the first study, the same split appeared without the reward: all six DQN seeds changed their actions in response to the observed state, while none of the twelve policy-gradient seeds from REINFORCE and PPO did so.

The first study shows why the analysis matters. REINFORCE came close to the benchmark in terms of return, at +\$171.51 per month, but its policies largely shifted to buying rather than using the market information provided. An agent can perform well for the wrong reason: during a generally rising market, a policy that remains invested can achieve a strong return without needing to learn when market conditions require entering, reducing, or exiting a position. DQN produced policies that consistently responded to the state yet earned the least, so state dependence and financial performance were not the same thing

at that data volume. The main study resolves the tension: once the data could support reliable value estimates, the state-dependent method was also the profitable one, on every seed.

Looking only at return would have completely hidden the distinction between these behaviours. The behavioural analysis is therefore an important part of evaluating learned trading policies, as it revealed a clear difference between the value-based and policy-gradient methods.

6.2.3 State Representation

The third question examined whether adding a drawdown feature to the state representation improved agent performance or behaviour.

Both studies give the same answer, and the main study gives it cleanly. In the first study, adding drawdown changed each algorithm's mean return by less than its seed variation. In the main study, where the seeds agree and the comparison is therefore meaningful, Deep Q-learning moved from +\$126.94 to +\$131.31 per episode — a +\$4.37 change against a seed standard deviation near \$7 — with identical behaviour on every measure. The extra information neither helped nor hurt while the objective remained plain wealth change.

Overall, the results show that providing additional information does not automatically produce a better policy or improve learning unless the objective specifies it. This is consistent with an important principle noticed from the experiments: **information in the state is only useful when the learning objective and optimisation process give the agent a reason to use it**. The risk-aware reward experiment below is the positive half of the same principle: once the objective referred to drawdown, the drawdown feature became useful.

6.2.4 Reward Formulation

The fourth question considered whether adding drawdown into the reward could encourage the agents to reduce losses and generate a more risk-aware behaviour. This is the question on which the two studies disagree, and the reason they disagree matters as much as the answer.

In the first study, the risk-aware reward reduced drawdown mainly by reducing market participation. With the selected penalty of $\lambda = 0.10$, DQN's drawdown fell 41% but its return fell 32%, its exposure halved, and its Sharpe ratio dropped from 0.505 to 0.416. The reward appeared to tax participation rather than teach risk management.

The main study overturns that conclusion. With $\lambda = 0.25$, selected on the simulated validation episodes by the same rule, Deep Q-learning kept its return almost unchanged (+\$129.38 against +\$131.31 per episode) while cutting drawdown by 31% and raising its Sharpe ratio from 0.285 to 0.332. The results for each market condition show the improvement is selective, not proportional: the down-market loss fell 73% while the up-market gain kept 87% of its size. The agent reduced its holdings when the market fell, which is exactly the behaviour the reward was designed to produce.

The first study's negative finding was therefore another data artefact. A drawdown penalty can only

produce selective behaviour if the value estimates behind the policy can tell a falling market from a rising one, and estimates trained on 60 months could not. Two qualifications remain. The penalty pushed every policy-gradient seed into complete abstention, so the mechanism is only usable by a method with reliable estimates and sustained exploration. And the risk-trained agents transfer more conservatively than the real market rewards: on the 180 real months they earned +\$36.58 per month against +\$58.47 for the wealth-trained agents, so the wealth-trained configuration remains the better real-market choice.

6.2.5 Data

The fifth question asked whether the amount and balance of the training data explain the learned behaviour, and whether fixing the data fixes the behaviour. The simulated-data study answers yes to both.

On 3,000 balanced simulated episodes, Deep Q-learning earned +\$126.94 per episode across six seeds (standard deviation \$7.30) on a test set where buy-and-hold loses \$60.29 and doing nothing earns \$0.00. All six seeds were profitable and all six were state-dependent. The results in each market condition show the behaviour the real-data study looked for: the agents captured 99% of the profit available in rising episodes and cut the falling-episode loss by 84% relative to always buying. The same policies remained profitable on every seed when transferred unchanged to the 180 real transfer months, and the defensive behaviour they learned in simulation appeared where it mattered: all six agents held zero exposure through October 2008, the month buy-and-hold lost \$1,665.24, and across the seven crisis months from September 2008 to March 2009 they lost \$180.34 where the benchmark lost \$3,045.66.

The study also explains the real-data algorithm ranking. Deep Q-learning was worst on real data not because it is unsuited to trading but because value estimation needs many visits to similar states, which 60 shifting months cannot provide. Once the data supported the estimates, the ranking inverted. REINFORCE and PPO, which looked strong on real data by adopting the majority always-buy action, collapsed to fixed corner policies on balanced data, where no fixed action is profitable. One REINFORCE seed in six did learn the conditioned policy, so the optimum is within reach of policy-gradient methods, but the training process does not find it reliably.

The overall conclusion of the dissertation therefore has two halves. At real single-asset data volumes, none of the tested algorithms learns a genuine trading policy. With enough balanced data, value-based learning does, on every seed, and behaves sensibly in every market condition the simulator can produce.

6.3 Critical Appraisal of the Project

The most important lessons from this project were from the experimental methodology rather than the choice of reinforcement learning algorithm alone. Several of the most useful findings came from implementation problems identified during development and testing.

One of the strongest aspects of the project was the verification tests. The automated tests identified several implementation problems during development, **three of which could have affected the experiments' validity**. These included issues with rolling feature calculation at the end of monthly episodes and an implementation error in the buy-and-hold benchmark that reduced its performance by \$38.95

per month, about 21.6%. If this error had remained, the dissertation's main conclusion could have been misleading, as the learned agents would have appeared to outperform the benchmark.

The verification tests also noticed the experimental design problem in the trade size. Using the original $0.10C_0$ trade size in Chapter 3 would make the maximum attainable exposure only 0.690, compared with 0.910 for buy-and-hold. This meant that the learning agents were at a disadvantage before training even began. Selecting $0.25C_0$ allowed a fairer comparison while still letting the agents control their positions in smaller steps than the $0.50C_0$ alternative.

The complete set of 21 verification tests passed before the final experiments were accepted, giving a higher confidence that the reported results reflect an appropriate experimental design rather than implementation errors.

The behavioural analysis was another important part of the study. Without it, the REINFORCE result of +\$171.51 could have been interpreted as evidence that the agent had learned an intelligent, effective trading strategy. The analysis showed this interpretation was not justified, as the policies acted like a consistent buyer because the market was consistently rising, completely changing the interpretation of the result.

The project also highlighted the importance of random seed variation. The results, particularly for PPO, varied considerably between independent training runs. PPO had a seed spread of \$179.80 in one configuration, which was greater than its mean return. This made it difficult to draw conclusions based on individual training runs, as that would have been unreliable. Reporting multiple seeds widened the evaluation, providing a more consistent view of how each algorithm behaved, although it increased the computational and experimental workload; however, six seeds are still lacking for more definite statistical inference.

The largest single lesson, however, was about data. The project spent most of its effort on the state representation, the reward, the trade size, and the verification of the environment, and all of that effort was necessary; none of it was sufficient. The binding constraint was 60 training episodes feeding 18,000 parameters, and no amount of feature or reward engineering could compensate for it. Diagnosing that constraint and rebuilding the study around simulated data changed the results more than every other design decision combined. The lesson for similar projects is to budget the data before budgeting the model: count the episodes, count the parameters, and if the second number is larger, change the data or the model before running anything else.

Overall, the project demonstrated that careful experimental design is as important as the choice of reinforcement learning algorithm. The main lessons came from testing the assumptions behind the experiments, rather than from simply comparing model returns. State dependence, action constraints, reward design, seed variation, and above all data volume affect how the final results could be interpreted. This made the experimental process more demanding but also produced a more reliable basis for evaluating the learned trading policies.

6.4 Limitations

The findings of this study should be considered within the scope of the experimental design, but there are several limitations.

6.4.1 One asset

The experiments use SPY as the only asset. The main study’s transfer set does cover a wide range of real market conditions, including the 2008 crisis, so the earlier concern that no real declining period was tested no longer applies. What remains untested is breadth across instruments: a single, highly liquid index fund cannot show how the approach behaves on individual stocks, other asset classes, or portfolios of several assets, where correlations and idiosyncratic events matter. The findings cannot therefore be automatically generalised beyond broad index trading without further testing.

6.4.2 Simulator realism and the domain gap

The market generator draws each day’s return independently within a condition. Real returns are not independent: volatility clusters, trends persist, and market conditions change mid-month. The generator also holds one condition for a whole episode, so the agents never face a mid-month crash. These simplifications are the price of a controlled experiment, and the transfer test measures their cost directly: the Deep Q-learning agents kept roughly three-quarters of buy-and-hold’s return over the 180 real transfer months while being trained on simulated data alone, and their timing around real turning points was imperfect — they sat out the December 2008 and March 2009 rebounds because real markets change condition less cleanly than the generator’s fixed ones. Conclusions about absolute profitability on real markets should therefore rest on the transfer numbers, not the simulated ones.

6.4.3 Policy-gradient training was not separately tuned

REINFORCE and PPO collapsed to corner policies on the balanced simulated data. Both were trained with the same hyperparameters as the real-data study, including PPO’s default entropy coefficient of zero. Entropy regularisation or other exploration incentives might let the policy-gradient methods find the conditioned policy more reliably, as one REINFORCE seed shows it exists. The study reports the collapse as observed behaviour under shared settings rather than as a bound on what tuned policy-gradient methods could achieve.

6.4.4 Monthly Episode Structure

The trading environment treated each calendar month as a separate episode. Every episode started with the same initial capital of \$10,000, and any remaining position was closed at the end of the month. As a result, the results do not represent a single portfolio that continuously compounds over the full evaluation period as a real portfolio would.

This design made the experiments easier to control by keeping the episode length and starting conditions consistent. However, it also means that gains or losses from one month did not affect the starting capital

or portfolio position in the following month, which does not represent a real-world portfolio. A longer continuous portfolio environment could therefore lead to a different behaviour, especially for policies that require accumulated wealth or changes in portfolio value over time.

6.4.5 Fixed Trade Size

The agents trade fixed fractions of the initial capital rather than choosing their position size continuously. This was calibrated to create a controlled, reasonable environment for comparison with the benchmark, but it restricts how much the agents can express their decisions. A continuous position-sizing model could potentially respond differently based on its exposure in a particular market condition. Rather than repeatedly applying fixed buy or sell actions, such an agent could select its desired portfolio exposure, producing different results in returns and trading behaviour.

6.4.6 Limited Number of Seeds and Lack of Significance Testing

Six seeds were used for each configuration, and the seed-level results and variation between them are reported throughout the results. This is stronger than relying on a single training run, but it does not provide sufficient statistics for formal claims about small differences. For example, the +\$4.37 change observed for Deep Q-learning when drawdown was added to the state sits within a seed standard deviation of about \$7 and is not large enough to be treated as a meaningful improvement.

6.4.7 Risk Penalty Selection was based on DQN

The reward's penalty weight λ was selected using DQN because it was the algorithm that consistently showed state-dependent behaviour. This does not guarantee that the same penalty value is appropriate for REINFORCE or PPO, as it was not independently optimised for every algorithm. A separate calibration for each algorithm could produce different results. This limitation is particularly relevant because the algorithms responded differently to the state and reward formulation.

6.4.8 The risk penalty only targets drawdown

The risk-aware reward penalises drawdown but does not directly measure whether the return achieved is sufficient for the level of risk taken. While drawdown is directly relevant to losses, it does not fully represent the portfolio risk. A reward that relates return to risk, rather than penalizing risk alone, was not tested.

The main study shows the penalty working as intended at the selected weight, but the sweep in Section 5.8 also shows the failure mode is still there: one grid step past the selected weight, exposure collapses and the agents leave the market. The penalty's usefulness therefore depends on a calibration that the study performed on validation data; a reward whose optimum did not sit next to a collapse point would be more robust.

6.5 Future Work

The limitations above suggest several directions for furthering the research. Each stems directly from a limitation or an unresolved result in the experiments. The behavioural analysis should also be applied to these future works to determine the agent's behaviour to changes alongside performance improvements.

6.5.1 Realistic Market Generators

The simulated-data study used the simplest generator that supports a controlled experiment. The natural next step is a generator that keeps the control but adds the structure the current one omits: changes of market condition within an episode, driven by a hidden Markov chain, and volatility clustering within a condition. Learned generators such as time-series GANs offer more realism again, at the cost of transparency. Each step up in realism should shrink the domain gap that the transfer test measures, and the transfer test itself provides the yardstick for whether it does.

6.5.2 Walk-Forward Evaluation and Other Markets

The transfer test now covers two decades of real months, including the 2008 crisis, but it evaluates fixed policies calibrated once on 2018–2022. A walk-forward design, in which the simulator is recalibrated and the agents retrained as conditions change, would test whether the approach holds up as an operating procedure rather than as a one-shot experiment. Extending the evaluation to other index funds, individual stocks, or non-equity assets would test whether the learned response to market conditions is specific to SPY's statistics or carries across markets.

6.5.3 Predictive and Uncertainty-Aware State Information

The drawdown feature describes the portfolio's previous losses but does not provide information about future market movements. Future work could therefore include predictive information, such as probabilistic forecasts containing estimates of forecast uncertainty.

This would give a state-dependent agent more useful information for making decisions about future market conditions rather than relying only on historical data. It would also directly test whether DQN's ability to use the state leads to more useful market understanding and trading decisions.

6.5.4 Risk-Efficiency Objectives

The most direct refinement of the drawdown penalty would be a reward that considers return relative to risk, such as a Sharpe-based or drawdown-relative objective [2]. Two results motivate this. The penalty sweep in Section 5.8 shows a collapse one step past the selected weight, where reducing exposure to zero becomes the best reply; a risk-relative objective would remove abstention as a way of scoring well. And the transfer results show the risk-trained agents giving up more return for downside protection than the real market's declines justify; an objective that weighs risk against return, rather than penalising risk alone, should calibrate that trade-off automatically.

6.5.5 More Seeds for Better Statistical Analysis

More random training seeds would improve the reliability of comparisons [30] between algorithms and configurations. Future work should combine this with appropriate statistical analysis to determine whether observed differences between configurations are likely to represent genuine results rather than random noise.

6.5.6 More flexible trade sizes

Future experiments could allow the agent to choose its position size continuously rather than using fixed trade amounts. This would remove the exposure limit created by the current environment and allow the agent to express different levels of confidence by adjusting its position size, potentially enabling smoother portfolio adjustments.

6.6 Closing Remarks

This dissertation examined whether standard reinforcement learning could learn a genuine trading policy for a single asset. On real data it could not: the best-performing learned agent finished \$8.74 per month below buy-and-hold, and it did so by ignoring its market observations to follow a simple trading behaviour.

The study's contribution is what it did with that failure. The behavioural analysis showed the strong-looking returns were not evidence of learning; the diagnosis traced the failure to 60 training episodes feeding 18,000 parameters; and the main study removed that constraint and showed the same algorithms succeeding. Deep Q-learning became profitable and state-dependent on every seed, and its policies, trained on simulated prices alone, stayed profitable across 180 real months they had never seen — stepping fully out of the market through October 2008 and finishing the two decades with a better risk-adjusted return than buy-and-hold. The negative result and the positive one are the same finding viewed from two sides: what these methods learn is decided by the data they are given.

The implementation, the market simulator, the experimental framework and the behavioural analysis developed in this study provide a useful basis for future work. More advanced models and more realistic generators can build on this framework while being evaluated against the same fair standards.

References

- [1] R. S. Sutton and A. G. Barto, *Reinforcement Learning: An Introduction*, 2nd ed. MIT Press, 2018.
- [2] J. Moody and M. Saffell, “Learning to trade via direct reinforcement,” *IEEE Transactions on Neural Networks*, vol. 12, no. 4, pp. 875–889, 2001.
- [3] Y. Deng, F. Bao, Y. Kong, Z. Ren, and Q. Dai, “Deep direct reinforcement learning for financial signal representation and trading,” *IEEE Transactions on Neural Networks and Learning Systems*, vol. 28, no. 3, pp. 653–664, 2017.
- [4] T. G. Fischer, “Reinforcement learning in financial markets — a survey,” FAU Discussion Papers in Economics, Friedrich-Alexander University Erlangen-Nürnberg, Tech. Rep. 12/2018, 2018.
- [5] R. J. Williams, “Simple statistical gradient-following algorithms for connectionist reinforcement learning,” *Machine Learning*, vol. 8, no. 3–4, pp. 229–256, 1992.
- [6] V. Mnih, K. Kavukcuoglu, D. Silver, A. A. Rusu, J. Veness, M. G. Bellemare, A. Graves, M. Riedmiller, A. K. Fidjeland, G. Ostrovski, S. Petersen, C. Beattie, A. Sadik, I. Antonoglou, H. King, D. Kumaran, D. Wierstra, S. Legg, and D. Hassabis, “Human-level control through deep reinforcement learning,” *Nature*, vol. 518, no. 7540, pp. 529–533, 2015.
- [7] J. Schulman, F. Wolski, P. Dhariwal, A. Radford, and O. Klimov, “Proximal policy optimization algorithms,” *arXiv preprint arXiv:1707.06347*, 2017.
- [8] G. Tesauro, “Temporal difference learning and TD-Gammon,” *Communications of the ACM*, vol. 38, no. 3, pp. 58–68, 1995.
- [9] R. Bellman, *Dynamic Programming*. Princeton University Press, 1957.
- [10] R. S. Sutton, “Learning to predict by the methods of temporal differences,” *Machine Learning*, vol. 3, no. 1, pp. 9–44, 1988.
- [11] C. J. C. H. Watkins and P. Dayan, “Q-learning,” *Machine Learning*, vol. 8, no. 3–4, pp. 279–292, 1992.
- [12] J. N. Tsitsiklis and B. Van Roy, “An analysis of temporal-difference learning with function approximation,” *IEEE Transactions on Automatic Control*, vol. 42, no. 5, pp. 674–690, 1997.
- [13] L.-J. Lin, “Self-improving reactive agents based on reinforcement learning, planning and teaching,” *Machine Learning*, vol. 8, no. 3–4, pp. 293–321, 1992.
- [14] H. van Hasselt, A. Guez, and D. Silver, “Deep reinforcement learning with double Q-learning,” in *Proceedings of the 30th AAAI Conference on Artificial Intelligence*, 2016, pp. 2094–2100.
- [15] V. R. Konda and J. N. Tsitsiklis, “Actor-critic algorithms,” in *Advances in Neural Information Processing Systems 12*, 2000, pp. 1008–1014.

- [16] V. Mnih, A. P. Badia, M. Mirza, A. Graves, T. Lillicrap, T. Harley, D. Silver, and K. Kavukcuoglu, “Asynchronous methods for deep reinforcement learning,” in *Proceedings of the 33rd International Conference on Machine Learning (ICML)*, 2016, pp. 1928–1937.
- [17] J. Schulman, P. Moritz, S. Levine, M. Jordan, and P. Abbeel, “High-dimensional continuous control using generalized advantage estimation,” in *Proceedings of the 4th International Conference on Learning Representations (ICLR)*, 2016.
- [18] J. Schulman, S. Levine, P. Abbeel, M. Jordan, and P. Moritz, “Trust region policy optimization,” in *Proceedings of the 32nd International Conference on Machine Learning (ICML)*, 2015, pp. 1889–1897.
- [19] X.-Y. Liu, H. Yang, Q. Chen, R. Zhang, L. Yang, B. Xiao, and C. D. Wang, “FinRL: A deep reinforcement learning library for automated stock trading in quantitative finance,” *arXiv preprint arXiv:2011.09607*, 2020.
- [20] S. Huang and S. Ontañón, “A closer look at invalid action masking in policy gradient algorithms,” in *Proceedings of the 35th International FLAIRS Conference*, 2022.
- [21] Z. Jiang, D. Xu, and J. Liang, “A deep reinforcement learning framework for the financial portfolio management problem,” *arXiv preprint arXiv:1706.10059*, 2017.
- [22] T. Théate and D. Ernst, “An application of deep reinforcement learning to algorithmic trading,” *Expert Systems with Applications*, vol. 173, p. 114632, 2021.
- [23] M. Guan and X.-Y. Liu, “Explainable deep reinforcement learning for portfolio management: An empirical approach,” in *Proceedings of the 2nd ACM International Conference on AI in Finance (ICAIF)*, 2021.
- [24] D. Silver, A. Huang, C. J. Maddison, A. Guez, L. Sifre, G. van den Driessche, J. Schrittwieser, I. Antonoglou, V. Panneershelvam, M. Lanctot, S. Dieleman, D. Grewe, J. Nham, N. Kalchbrenner, I. Sutskever, T. Lillicrap, M. Leach, K. Kavukcuoglu, T. Graepel, and D. Hassabis, “Mastering the game of Go with deep neural networks and tree search,” *Nature*, vol. 529, no. 7587, pp. 484–489, 2016.
- [25] G. Brockman, V. Cheung, L. Pettersson, J. Schneider, J. Schulman, J. Tang, and W. Zaremba, “OpenAI Gym,” *arXiv preprint arXiv:1606.01540*, 2016.
- [26] L. R. Rabiner, “A tutorial on hidden Markov models and selected applications in speech recognition,” *Proceedings of the IEEE*, vol. 77, no. 2, pp. 257–286, 1989.
- [27] J. Yoon, D. Jarrett, and M. van der Schaar, “Time-series generative adversarial networks,” in *Advances in Neural Information Processing Systems*, vol. 32, 2019.
- [28] D. H. Bailey, J. M. Borwein, M. L. de Prado, and Q. J. Zhu, “Pseudo-mathematics and financial charlatanism: The effects of backtest overfitting on out-of-sample performance,” *Notices of the American Mathematical Society*, vol. 61, no. 5, pp. 458–471, 2014.

- [29] M. L. de Prado, *Advances in Financial Machine Learning*. Wiley, 2018.
- [30] P. Henderson, R. Islam, P. Bachman, J. Pineau, D. Precup, and D. Meger, “Deep reinforcement learning that matters,” in *Proceedings of the 32nd AAAI Conference on Artificial Intelligence*, 2018, pp. 3207–3214.
- [31] E. F. Fama, “Efficient capital markets: A review of theory and empirical work,” *The Journal of Finance*, vol. 25, no. 2, pp. 383–417, 1970.
- [32] H. Markowitz, “Portfolio selection,” *The Journal of Finance*, vol. 7, no. 1, pp. 77–91, 1952.
- [33] W. F. Sharpe, “The Sharpe ratio,” *The Journal of Portfolio Management*, vol. 21, no. 1, pp. 49–58, 1994.
- [34] J. W. Wilder, *New Concepts in Technical Trading Systems*. Greensboro, NC: Trend Research, 1978.
- [35] D. P. Kingma and J. Ba, “Adam: A method for stochastic optimization,” in *Proceedings of the 3rd International Conference on Learning Representations (ICLR)*, 2015.
- [36] A. Raffin, A. Hill, A. Gleave, A. Kanervisto, M. Ernestus, and N. Dormann, “Stable-baselines3: Reliable reinforcement learning implementations,” *Journal of Machine Learning Research*, vol. 22, no. 268, pp. 1–8, 2021.
- [37] A. Paszke, S. Gross, F. Massa, A. Lerer, J. Bradbury, G. Chanan, T. Killeen, Z. Lin, N. Gimelshein, L. Antiga, A. Desmaison, A. Köpf, E. Yang, Z. DeVito, M. Raison, A. Tejani, S. Chilamkurthy, B. Steiner, L. Fang, J. Bai, and S. Chintala, “PyTorch: An imperative style, high-performance deep learning library,” in *Advances in Neural Information Processing Systems 32 (NeurIPS)*, 2019.
- [38] M. Towers, A. Kwiatkowski, J. Terry, J. U. Balis, G. De Cola, T. Deleu, M. Goulão, A. Kallinteris, M. Krimmel, A. KG *et al.*, “Gymnasium: A standard interface for reinforcement learning environments,” *arXiv preprint arXiv:2407.17032*, 2024.

Appendix A

Verification Tests

This appendix lists the 21 automated verification tests referenced in Section 4.5, grouped as in that section. Every test was run and passed before the final experiments in Chapter 5 were accepted. The test names use the dissertation's current terminology rather than the older working names kept by some functions in the code.

Environment and portfolio tests

Test	Property checked
episode always ends flat	Any remaining position is sold on the final trading day, leaving no units held and no exposure.
wealth can be calculated from the current state	The portfolio wealth can be calculated from the cash and exposure features at every step, which is why wealth is not a separate state feature.
drawdown cannot be calculated from the current state alone	Two price paths can arrive at the same cash, exposure, and wealth but have different drawdowns, which is why drawdown must be added to the state as a separate feature.
share action model buys one unit	An alternative action model, in which a buy purchases one whole unit, behaves as defined. The main experiments use the fixed-amount model below.
trade action model spends a fixed sum	A buy spends exactly the configured fraction of the initial capital, excluding the fee.
true buy-and-hold beats gradual buying when prices rise	The corrected buy-and-hold benchmark invests once at the start of the month and earns more than gradual buying in a rising market (the benchmark correction of Chapter 4).
a month with no trades leaves cash unchanged	If no trades are made during a month, the wealth is unchanged and no fees are charged.

Feature and data tests

Test	Property checked
features use no future data	Removing the later prices from the series does not change any earlier feature value, so no feature is calculated using future information.

Test	Property checked
rolling features survive the episode boundary	Momentum on the first day of a month reaches back into the previous month's trading days rather than restarting from zero.
no NaNs or infinities in any episode	Every feature value in every episode is a finite number, and every price is positive.
canonical order is independent of input order	Requesting the same features in a different order still produces the same state layout.
state dimensions are as documented	Each state configuration contains exactly the number of features reported in Chapter 3: nine for the main state and ten when drawdown is added.
every state configuration is a subset of the next	Each larger state configuration contains all the features of the smaller ones, so a comparison between two states isolates the added features.
split is chronological and disjoint	The training, validation, and test periods are ordered in time and do not share any months.

Reward and risk tests

Test	Property checked
accounting identity holds for every policy	The sum of all daily rewards in a month equals the month's true change in wealth, for every policy and state configuration tested.
reward equals wealth change when $\lambda = 0$	When the risk penalty is switched off, the risk-aware reward gives exactly the same values as the wealth-change reward.
risk penalty requires drawdown in the state	Configuring the risk-aware reward without the drawdown feature raises an error, so the reward can always be calculated from what the agent observes.
risk penalty never increases the reward	For any sequence of actions, the penalised reward is never larger than the unpenalised reward.

Action-masking tests

Test	Property checked
mask blocks infeasible trades	Selling is not allowed with no position, buying is not allowed without enough cash, and selling becomes available while a position is held.
illegal action is forced to hold and flagged	An invalid action supplied from outside the environment is executed as Hold, flagged, and no trade is recorded.

Test	Property checked
mask is a function of observed features	The legal-action mask can be predicted exactly from the cash and exposure features; masking adds no extra information.

Appendix B

Algorithm Listings

This appendix lists the three learning procedures used in this dissertation as they were implemented. REINFORCE and Deep Q-learning were built from scratch, while MaskablePPO is the library implementation used for comparison. All three apply the same action mask and operate in the same trading environment, as described in Chapter 3.

Algorithm 1: REINFORCE for the trading MDP (as implemented)

```
1: initialise policy network  $\pi\theta$ 
2: for iteration = 1, 2, ... until step budget exhausted do
3: sample a training month; reset environment ( $C_0$  cash, no position)
4: for  $t = 0, \dots, T - 1$  do
5: compute legal-action mask; set illegal logits to  $-\infty$ 
6: sample  $a_t \sim \pi\theta(\cdot | s_t)$ ; execute; store  $(\log \pi\theta(a_t | s_t), r_t)$ 
7: end for
8: compute rewards-to-go  $G_t$ ; standardise within the episode
9:  $L(\theta) \leftarrow -\sum_t \log \pi\theta(a_t | s_t) G_t$ 
10: one Adam step on  $\nabla\theta L$ 
11: end for
```

Algorithm 2: Deep Q-learning for the trading MDP (as implemented)

```
1: initialise  $Q\phi$ ; target  $Q\phi^- \leftarrow Q\phi$ ; replay buffer  $D$ 
2: for environment step = 1, 2, ... until step budget exhausted do
3: with prob.  $\epsilon$ : random legal action; else  $a_t = \arg \max_a \text{legal } Q\phi(s_t, a)$ 
4: execute  $a_t$ ; store  $(s_t, a_t, r_t, s_{t+1}, \text{mask}_{t+1}, d_t)$  in  $D$ 
5: sample minibatch; form masked temporal-difference targets  $y$  (Section 3.8); Adam step on  $L(\phi)$ 
6: every  $K$  steps:  $\phi^- \leftarrow \phi$ 
7: end for
```

Algorithm 3: MaskablePPO for the trading MDP (as implemented)

```

1: initialise actor  $\pi\theta$  and critic  $V\psi$  (MaskablePPO, MLP 128–128)
2: for iteration = 1, 2, ... until step budget exhausted do
3:   for step = 1, ...,  $N$  do
4:     if episode terminated: sample a training month; reset ( $C_0$  cash, no position)
5:     observe  $s_t$  and legal-action mask; sample  $a_t \sim \pi\theta(\cdot | s_t)$  over legal actions
6:     execute; store  $(s_t, a_t, r_t, V\psi(s_t), \log \pi\theta(a_t|s_t))$ 
7:   end for
8:   estimate advantages  $\hat{A}_t$  by GAE; form returns
9:   for epoch = 1, ...,  $K$  do
10:    for each minibatch from the rollout do
11:       $\rho_t \leftarrow \pi\theta(a_t|s_t) / \pi\theta_{\text{old}}(a_t|s_t)$ 
12:      Adam step on  $L\text{CLIP}(\theta) + \text{value loss} - \text{entropy bonus}$ 
13:    end for
14:  end for
15: end for

```

Notes. $N = 2048$ is the Stable-Baselines3 rollout length used here, and $K = 10$ is the number of training epochs over each rollout. Month sampling and action masking use the same environment interface as REINFORCE and DQN, so all three algorithms face the same action space at every step.

Appendix C

Detailed Test Results

This appendix reports the detailed test results behind Chapter 5. The first section covers the first study’s held-out months; the second covers the main study’s transfer test on the real months.

C.1 First Study: Wealth Change in Each Held-Out Month

This section reports the wealth change ΔW in dollars for each of the first study’s 24 held-out test months (2024–2025), using its final configuration ($C_0 = \$10,000$, $0.25 C_0$ trade size, 0.05% fee, wealth-change reward). The columns for the learned agents show the mean of the six training seeds, so they match the seed means reported in Chapter 5. The final row is the total over the 24 months.

Month	Buy-and-hold	Always-buy ($0.25 C_0$)	REINFORCE	DQN	PPO
2024-01	+206.23	+282.35	+271.95	+72.70	+196.83
2024-02	+375.56	+305.79	+290.50	+54.06	+219.51
2024-03	+220.74	+267.59	+249.89	+31.05	+187.59
2024-04	-396.08	-325.47	-311.74	-95.75	-233.48
2024-05	+529.49	+366.11	+350.33	+291.96	+266.14
2024-06	+334.03	+264.80	+243.42	+127.65	+190.45
2024-07	+90.21	+3.05	+9.27	+72.89	+5.79
2024-08	+370.28	+651.47	+626.02	+363.21	+449.74
2024-09	+414.15	+487.66	+464.38	+105.42	+342.37
2024-10	-9.64	-26.18	-54.57	+96.12	-17.86
2024-11	+541.25	+431.90	+414.48	+132.34	+310.49
2024-12	-267.80	-297.22	-297.55	-32.13	-209.30
2025-01	+283.57	+188.10	+203.74	+18.17	+137.22
2025-02	-69.99	-148.04	-126.82	+46.50	-101.61
2025-03	-398.41	-320.38	-302.86	-184.04	-230.19
2025-04	-124.50	+243.72	+297.84	+79.27	+157.30
2025-05	+543.11	+479.42	+414.44	+390.80	+342.25
2025-06	+444.52	+414.11	+408.93	+69.76	+294.59
2025-07	+223.40	+167.96	+150.83	+51.18	+121.28
2025-08	+364.88	+255.00	+215.99	+98.64	+185.21
2025-09	+423.09	+345.46	+328.23	+119.11	+247.93
2025-10	+193.41	+175.59	+189.11	-76.82	+125.12
2025-11	-9.26	+90.48	+73.42	-32.31	+59.94

Month	Buy-and-hold	Always-buy (0.25 C_0)	REINFORCE	DQN	PPO
2025-12	+43.83	+10.78	+7.10	-7.54	+9.01
Total	+4326.04	+4314.08	+4116.31	+1792.25	+3056.31

C.2 Main Study: Wealth Change by Transfer Year

This section reports the main study's transfer test in more detail. The 180 transfer months are grouped by calendar year, because a table of 180 individual months would be difficult to read; the total wealth change within each year is shown instead. The columns for the learned agents are the mean of the six training seeds, so the totals are consistent with the monthly means reported in Chapter 5. The years 2018–2022 do not appear because they calibrated the market generator and are excluded from the transfer set.

The 2008 row shows the behaviour discussed in Chapter 5 at the yearly scale. Buy-and-hold lost \$3,887.84 over the year, while the Deep Q-learning agents lost \$1,478.59 — about 62% less — because they reduced their exposure as the volatility features rose and then held none of the asset through October.

Year	Buy-and-hold	Always-buy (0.25 C_0)	REINFORCE	DQN	PPO
2006	+1,108.55	+1,000.57	+453.32	+993.39	+90.99
2007	+270.25	+295.76	+119.33	-204.50	+20.16
2008	-3,887.84	-3,142.23	-1,303.53	-1,478.59	-542.97
2009	+2,307.88	+1,895.91	+539.10	+898.19	+182.28
2010	+258.36	-152.11	-22.25	-82.20	-47.32
2011	+888.78	+767.16	+260.69	-17.76	+108.79
2012	+1,161.48	+1,051.53	+458.31	+979.51	+95.36
2013	+2,122.32	+1,997.68	+986.39	+1,985.39	+247.50
2014	+1,681.49	+1,330.85	+445.00	+1,316.61	+181.40
2015	-178.45	-282.68	-140.82	-768.27	-72.75
2016	+912.53	+1,239.24	+588.62	+718.57	+185.88
2017	+1,468.11	+1,441.89	+689.14	+1,439.32	+167.15
2023	+2,003.04	+2,008.68	+872.75	+1,935.73	+260.21
2024	+2,408.41	+2,411.87	+895.87	+1,710.25	+324.22
2025	+1,917.63	+1,902.21	+857.13	+1,424.35	+295.90
Total	+14,442.52	+13,766.32	+5,699.03	+10,849.99	+1,496.81

Appendix D

Glossary and Notation

D.1 Glossary of Terms

This glossary explains the recurring technical terms of the dissertation in plain language. The definitions here are informal; the precise definitions are given in Chapter 3.

Agent The learning program that makes the trading decisions. In this dissertation, it is a small neural network that decides each day whether to hold, buy, or sell.

Environment The simulated world the agent interacts with. Here, it is a simulated trading month built from real historical prices. It executes the agent’s trades, charges the transaction fees, and updates the portfolio.

State The information the agent observes before making a decision. Here, it is nine numbers covering recent market behaviour, the agent’s own cash and holdings, and how far the month has progressed. A tenth number is added when the drawdown feature is used.

Action One of the choices available to the agent. There are exactly three: hold the current position, buy a fixed amount of the asset (a quarter of the initial capital, written $0.25 C_0$), or sell the same amount.

Reward The feedback the agent receives after each action. Here, it is the day’s change in total wealth, with trading fees already deducted.

Episode One complete round of the task, after which everything resets. Here, it is one calendar month of trading, ending with a forced sale of any remaining holdings.

Policy The agent’s decision rule. Formally, it is a function from the state to action probabilities; informally, it describes what the agent would do in any situation.

Policy gradient / REINFORCE A family of methods that improve the policy directly. After each month, decisions that came before a good outcome are made more probable, and decisions that came before a poor outcome are made less probable. REINFORCE is the original and simplest member of this family.

Q-value / Deep Q-learning The value-based family of methods. The agent learns an estimate of how much future profit each action is worth in each situation, then selects the action with the highest estimate. “Deep” means the estimates come from a neural network rather than a lookup table.

PPO A more recent policy-gradient method that limits how much the policy can change in a single update, which makes training steadier.

Neural network A flexible mathematical function with many adjustable parameters, tuned gradually during training so that the outputs improve against an objective.

Training / validation / test split The data is divided into three parts so that no decision is tested on the data that shaped it. In the first study the parts are chronological: 2018–2022 for training, 2023 for every tuning decision, and 2024–2025 for the final evaluation. In the main study the training and validation parts are simulated episodes, and the final evaluations use fresh simulated episodes and the real transfer months.

Benchmark / buy-and-hold A simple non-learning strategy the agents are compared against. The most important one here is *buy-and-hold*: invest everything at the start of the month and wait. Previous research shows it is genuinely difficult to beat using price history alone, so a result that beats it easily should be examined carefully rather than accepted immediately.

Always-buy ($0.25 C_0$) The strongest strategy available to a policy that ignores its observations: buy one fixed trade every day buying is allowed, until fully invested.

Transaction fee / basis points (bps) The cost of trading, charged as a fraction of the traded amount. 5 bps = 0.05%, which is 50 cents per \$1,000 traded.

Sharpe ratio The return earned per unit of variation in returns. Higher is better, but a strategy that barely invests can report a high Sharpe ratio simply because its wealth barely moves.

Drawdown The fall from the highest wealth reached so far in the episode, as a fraction of that peak.

Exposure The proportion of the capital invested in the asset at a given moment, from 0% (all cash) to 100% (fully invested).

Seed The starting value of the random number generator. Six seeds are used per experiment, to check that a result does not depend on one particular random initialisation.

Overfitting Learning patterns that are specific to the training data rather than patterns that hold in general, which leads to poor performance on new data.

Action masking Removing impossible actions from the agent’s available choices, such as selling with no position or buying with no cash. The agent can neither attempt nor learn from invalid actions.

Market condition The overall direction of the market during an episode: up, down, or flat. The market generator produces each simulated episode under one known condition, so every result can be reported for each condition separately.

Market generator The program that produces the simulated monthly episodes used by the main study. Its parameters are calibrated on the real training months only, and every episode it produces is labelled with its market condition.

Transfer test The final evaluation of the main study. The policies trained on simulated episodes are run, unchanged, on 180 real months (2006–2017 and 2023–2025) that no training or tuning decision

ever used.

State-dependent policy A policy whose actions change with the observed state even when the same actions are legal. The opposite is a mask-determined policy, which follows a fixed pattern such as buying whenever buying is allowed.

Accounting identity The check that the sum of all daily rewards in a month equals the month’s true change in wealth. It is verified in every experiment.

Markov property The requirement that the current state contains everything from the past that is relevant to the next decision. This is the reason drawdown had to be added to the state before it could be used in the reward.

D.2 Notation

This section collects every mathematical symbol used in Chapters 3–5 for reference.

Symbol	Meaning
$s_t \in \mathbb{R}^9$	state observed by the agent at trading day t , containing the nine features used by the trading system (Table 3.1)
$a_t \in \{0, 1, 2\}$	action selected by the agent: Hold, Buy, or Sell.
P_t	asset price at day trading t ;
r_t	reward received after an action, calculated from the change in portfolio value after transaction costs.
C_t	amount of cash available to the agent at trading day t
h_t	number of (fractional) units of the asset held at trading day t .
$W_t = C_t + h_t P_t$	Total portfolio value, calculated as untouched cash plus the current value of the asset held;
C_0	initial cash available at the beginning of an episode.
c	proportional transaction fee for taking a trade, set to 0.05% per transaction
$\Delta = f C_0$	fixed monetary size of each trade; the fraction f is selected on validation data (Section 5.5)
γ	discount factor used to determine the importance of future rewards
τ	a trajectory (one complete trading episode), corresponding to one month of daily trading
T, t	episode length in days ($T \approx 18$ – 23) and one trading day within it
$G(\tau), G_t$	the total return of a complete episode, and the accumulated future reward from time t onwards
$\pi_\theta(a s)$	policy network that determines the likelihood of choosing action a given the current state s .
$Q_\phi(s, a)$	The Q-network, which estimates the value of taking a particular action in a given state
Q_ϕ^-	A fixed copy of the Q-network used as a target during Deep Q-learning.
$J(\theta)$	objective function representing the expected return of the policy
ε	exploration rate used by the (ε -greedy) action-selection strategy