

University of Surrey
Faculty of Engineering and Physical Sciences
School of Computer Science and Electronic Engineering

MSc Dissertation

Evaluating Reinforcement Learning Algorithms for Portfolio Trading

A Reinforcement Learning Study on the S&P 500 ETF (SPY)

Fiyinfoluwa Akano

URN: 6962514

Supervisor: Dr. Nguyen

Submitted for the Degree of
Master of Science in Artificial Intelligence



September 2026

©2026 Fiyinfoluwa Akano

Statement of Originality

I confirm that this report and the work it presents are my own. Any ideas, data, images or text taken from the work of others, whether published or unpublished, and including any content generated by a deep learning/artificial intelligence tool, have been clearly identified and appropriately acknowledged in the text and bibliography.

I confirm that this report has not been submitted, either in whole or in part, for any other academic degree or professional qualification.

I understand and agree that the University may submit this work to TurnitinUK to check its originality.

Fiyinfoluwa Akano

September 2026

Abstract

An investor holding a stock-index fund faces the same decision each day: buy more, sell part of the position, or do nothing. This dissertation asks whether reinforcement learning, which learns by using rewards via trial and error, can make these decisions more effectively than the simple strategy of buying and holding permanently.

Three algorithms are examined: REINFORCE and Deep Q-learning (DQN), implemented from scratch, and Proximal Policy Optimisation (PPO), from an open-source library. Each agent trades the S&P 500 exchange-traded fund (SPY) in a simulated environment using monthly episodes based on daily data from 2018 to 2025. The data is split chronologically into training data from 2018–2022, validation data from 2023, and held-out test months from 2024–2025.

The environment includes the trade size, transaction fees, and action validity. Experimental decisions, such as using a trade size equal to one quarter of the initial \$10,000 capital, were selected using the validation data before being applied to the test data. Each configuration was trained using six different random seeds.

None of the learned agents beat buy-and-hold on the held-out test period. REINFORCE earned the highest mean return at +\$171.51 per month, compared with +\$180.25 for buy-and-hold. PPO earned +\$127.35, while DQN earned +\$74.68.

Behavioural analysis provides an important explanation for these results. All six DQN seeds changed their actions in response to the observed state, whereas none of the twelve policy-gradient seeds (REINFORCE and PPO) did so. The returns earned by these policy-gradient algorithms were therefore largely tied to fixed buying patterns that benefited from the generally rising test period, rather than clear evidence of depending on market information to make informed decisions.

Two further experiments supported this interpretation. Adding drawdown to the state produced little change in either behaviour or performance for any of the algorithms. A risk-aware reward reduced DQN's drawdown by 41%, but only by cutting earnings by 32% and approximately halving market exposure.

The main contribution of the study is therefore an evaluative framework that uses calibrated trade sizes, simple benchmarks, behavioural analysis, and verified implementations, creating a reliable basis for interpreting both successful and negative results from reinforcement learning in trading.

Acknowledgements

First and foremost, I would like to thank God for giving me the idea for this project, the strength, wisdom and perseverance to see it through. What started as an initial idea developed significantly throughout this journey, and I am grateful to everyone who played a part in bringing it to where it is today. It is only by His Grace to Start, Continue, and Finish that this was possible.

I would like to thank the professors who helped me narrow down the original idea to shape it into a research problem that could be properly investigated. I am especially grateful to my supervisor, Dr. Nguyen, for his guidance throughout the project. His feedback helped me streamline the work, question my assumptions, and keep the research focused on what evidence could support. More importantly, his guidance helped me develop a much stronger understanding of reinforcement learning and how it can be applied to a problem as challenging as financial trading.

I would also like to express my deepest gratitude to my family, especially my mum and dad, for their constant encouragement and support. They kept pushing me to keep going until the final stage. I am equally grateful to my uncle, aunties, grandma, church circle, and everyone else who took the time to check in on me throughout the journey. Their encouragement and help, even in small moments, but especially in sick, tired moments, all meant a great deal to me.

Finally, I would like to give a special shoutout to my friends from the cohort: Katherine, Joel, Yumna, Akeeb, David, and Dr. Tim. The laughs, outings, and support throughout the year are unforgettable. Having people around who were going through the same experience made the long parts of the journey so much fun, shorter, and easier to get through.

To everyone who contributed to this journey in one way or another, thank you. I am genuinely grateful.

Use of Generative AI. Generative AI tools (ChatGPT and Cursor) were used during the preparation of this dissertation to assist with drafting and restructuring prose, literature organisation, LaTeX formatting, and debugging of experimental code. All research design choices, experimental results, analysis, and the final wording of this report are my own. AI-generated suggestions were reviewed, edited, and verified before inclusion. No AI tool is cited as a scholarly source.

Contents

Statement of Originality	i
Abstract	ii
Acknowledgements	iii
List of Figures	ix
List of Tables	x
1 Introduction	1
1.1 Background and Motivation	1
1.2 Problem Statement	2
1.3 Research Questions and Objectives	2
1.4 Scope	3
1.5 Contributions	4
1.6 Dissertation Structure	5
2 Background and Literature Review	6
2.1 Reinforcement Learning for Sequential Decisions	6
2.2 Two Families of Learning Algorithms	7
2.2.1 Value-based methods: from Q-learning to DQN	8
2.2.2 Policy-gradient methods: from REINFORCE to PPO	9
2.2.3 Why compare the two families?	10
2.3 Reinforcement Learning for Trading	10
2.3.1 Data requirements and simulated markets	12
2.4 Evaluating Trading Agents Honestly	13
2.5 The Gap Addressed In This Study	14
3 Methodology & Mathematical Formulation	17
3.1 Overview: Trading as a Markov Decision Process	18
3.2 Episode Definition and Data Structure	20
3.2.1 Episode definition (τ)	20
3.2.2 Independence of Trading Episodes	20
3.2.3 Choice of Episode Length	20
3.2.4 Selection of the Historical Period	20
3.2.5 The Simulated Market Model	21
3.3 State Representation Design	22
3.3.1 Information Available to the Agent	22

3.3.2	Criteria for a Suitable Trading State	23
3.3.3	Development of the State Representation	24
3.3.4	Final State Representation	25
3.3.5	Feature Redundancy and Excluded Information	27
3.3.6	The Risk-Aware Extension: Drawdown State	28
3.3.7	Feature Calculation and Information Timing	29
3.3.8	State Representation Limitations	29
3.3.9	Summary	30
3.4	Action Space Design	31
3.4.1	Overview	31
3.4.2	Trading Actions and Trade Size	32
3.4.3	Action Constraints	32
3.4.4	Transaction Costs	33
3.4.5	Relationship Between the State and Actions	33
3.5	Reward Function Design	33
3.5.1	Purpose of the Reward Function	33
3.5.2	Wealth-Based Reward	34
3.5.3	Including Transaction Costs in the Reward	34
3.5.4	Wealth Change As the Primary Reward	35
3.5.5	Risk-Aware Reward Extension	35
3.5.6	Limitations of the Reward Function	36
3.6	Objective Function	37
3.6.1	Episode Return	37
3.6.2	Expected Return	38
3.6.3	Objective and the Trading Environment	39
3.7	REINFORCE: Monte-Carlo Policy Gradient	39
3.7.1	Policy Representation	39
3.7.2	Deriving the Policy-Gradient Objective	40
3.7.3	Reward-to-Go	41
3.7.4	Monte-Carlo Estimation	41
3.7.5	Loss Function Used for Training	41
3.7.6	Strengths and Limitations	42
3.8	Deep Q-Learning	42
3.8.1	Action-Value Function	42
3.8.2	From Q-Learning to Deep Q-Learning	43
3.8.3	Experience Replay	44
3.8.4	Target Network	44
3.8.5	Temporal-Difference Learning	45
3.8.6	DQN Loss	45
3.8.7	Exploration	46
3.8.8	Strengths and Limitations	46

3.9	PPO: Proximal Policy Optimisation	46
3.9.1	Actor-Critic Structure & Advantage Function	47
3.9.2	Policy Ratio	47
3.9.3	Clipped Objective	48
3.9.4	Action Masking in PPO	48
3.9.5	PPO in the Experiment	48
3.10	Algorithms	49
3.11	Comparison of the Three Learning Approaches	49
4	System Design and Experiment Framework	50
4.1	Introduction	50
4.2	Data Pipeline	50
4.2.1	Market Data Preparation	50
4.2.2	Feature Construction & Continuous Calculation	51
4.2.3	Dataset Split & Monthly Episodes	52
4.3	Simulated Market Data	53
4.3.1	Why the Real Dataset Is Not Enough	53
4.3.2	Generator Design	53
4.3.3	Calibration	54
4.3.4	Simulated Splits	54
4.4	Training, Validation and Testing Procedure	55
4.4.1	Training	55
4.4.2	Validation	55
4.4.3	Testing	55
4.5	Experiment Configuration	56
4.5.1	Configuration Parameters	56
4.5.2	Algorithm Configuration	57
4.5.3	Benchmark Implementation	58
4.6	Implementation Verifications	58
4.6.1	Environment and Portfolio Tests	59
4.6.2	Feature and Data Tests	59
4.6.3	Reward and Risk Tests	59
4.6.4	Action-Masking Tests	60
4.7	Implementation Corrections	60
4.7.1	Problems Identified During Development	60
4.7.2	Correction and Re-verification	61
4.8	Complete Implementation Pipeline	61
5	Experimental Results and Analysis	63
5.1	Introduction	63
5.2	Experiment Protocols and Evaluation Criteria	64
5.2.1	Experimental Configurations	64

5.2.2	Evaluation Metrics	64
5.2.3	Benchmark Strategies	65
5.3	Fixed Trade Size Tuning	65
5.3.1	Exposure available under different transaction sizes	66
5.4	Do the Learned Policies Use Their Observations?	67
5.4.1	Behaviour of the learned policies	68
5.4.2	State-dependence analysis	69
5.4.3	Interpretation of the learned trading behaviour	69
5.5	Effect of the State Representation	70
5.5.1	Nine-Feature vs Ten-Feature State	70
5.5.2	Analysis When Drawdown Was Added	72
5.5.3	Interpretation	72
5.6	Held-Out Test Results	72
5.6.1	Overall Comparison with the Benchmark Strategies	73
5.6.2	Return and market exposure	74
5.6.3	Risk and risk-adjusted performance	74
5.6.4	Behaviour During the March 2025 Decline	75
5.6.5	Trading behaviour and Seed consistency	76
5.7	Effect of the Risk-Aware Reward	77
5.7.1	Selection of the Penalty size	77
5.7.2	Effect on the Test Set	79
5.7.3	Interpretation of the risk-aware return	79
5.8	Transaction-Fee Sizing	80
5.8.1	Effect of transaction costs on returns & trading behaviour	80
5.8.2	Interpretation	81
5.9	The Data Bottleneck	82
5.10	Simulated-Data Study: Controlled Regime Results	82
5.10.1	Benchmarks Under Balanced Regimes	82
5.10.2	Main Results	83
5.10.3	The Policy-Gradient Methods Collapse to Corners	83
5.10.4	Transfer to the Real Test Months	84
5.10.5	What the Controlled Study Establishes	84
5.11	Discussion of the Main Findings	85
5.11.1	The Benchmark Remains Difficult to Beat on Real Data	85
5.11.2	Additional State Features Did Not Lead to Better Performance	85
5.11.3	Why the Algorithms Ranked the Way They Did	86
5.11.4	Lower Risk Was Mainly Associated with Lower Exposure	86
5.11.5	The Trade Size Affects Performance	86
5.11.6	Transaction Costs Provided a Useful Behavioural Test	87
5.11.7	Variation Between Seeds Is Important	87
5.11.8	Implications for the Research Objectives	87

5.12 Chapter Summary	88
6 Conclusion and Future Work	89
6.1 Summary of the Work	89
6.2 Conclusions Against the Research Questions	89
6.2.1 Performance	90
6.2.2 Behaviour	90
6.2.3 State Representation	91
6.2.4 Reward Formulation	91
6.2.5 Data	92
6.3 Critical Appraisal of the Project	92
6.4 Limitations	93
6.4.1 One asset and a limited market period	93
6.4.2 Simulator realism and the domain gap	94
6.4.3 Policy-gradient training was not separately tuned	94
6.4.4 Monthly Episode Structure	94
6.4.5 Fixed Trade Size	94
6.4.6 Limited Number of Seeds and Lack of Significance Testing	95
6.4.7 Risk Penalty Selection was based on DQN	95
6.4.8 The risk penalty only targets drawdown	95
6.5 Future Work	95
6.5.1 Richer Market Generators	95
6.5.2 Longer and More Challenging Real Market Periods	96
6.5.3 Predictive and Uncertainty-Aware State Information	96
6.5.4 Risk-Efficiency Objectives	96
6.5.5 More Seeds for Better Statistical Analysis	96
6.5.6 More flexible trade sizes	96
6.6 Closing Remarks	96
References	98
A Verification Tests	101
B Algorithm Listings	104
C Per-Month Held-Out Test Results	106
D Glossary of Terms	107

List of Figures

3.1	The agent–environment interaction loop for the trading MDP.	19
4.1	Chronological split of the 96 monthly episodes	52
4.2	System architecture showing the data flow described in Section 4.7.	62
5.1	Exposure ceilings from Table 5.1.	66
5.2	Results for each seed from Table 5.2.	68
5.3	Cumulative sum of average monthly wealth changes across the test months.	74
5.4	Average wealth change compared with exposure (left) and daily drawdown (right) for every learned configuration in this chapter.	75
5.5	Monthly wealth changes of buy-and-hold and Deep Q-learning over the 24 test months. .	76
5.6	Results from the penalty size of Table 5.5. Left: mean wealth change against λ with the requirement floor of +\$39.17. Right: exposure and drawdown relative to their values without penalties.	78
5.7	Transaction Fee sizing from Table 5.7. Left: average wealth change falls for every strategy as fees rise. Right: average trades per month as fees rise.	81
5.8	Simulated-data study: per-regime results and real-data transfer.	84

List of Tables

3.1	Notation used throughout Chapters 3–5.	17
3.2	Mapping a popular game MDP to the portfolio trading MDP.	19
3.3	Full State Definition	26
3.4	The staged development of the state representation. Each stage resolves the limitation named in the row above it.	30
4.1	Generator parameters fitted on the 60 real training months.	54
5.1	Attainable exposure for each fixed trade amount on the validation months. The policy buys whenever a purchase is valid.	66
5.2	Behavioural Results on the validation months (0.25 C_0 trade size, nine-feature state, wealth-change reward, six freshly trained seeds). The always-buy strategy shows the maximum asset exposure this trade size can reach.	68
5.3	Results of the nine- and ten-feature states on the test data (wealth-change reward, 0.25 C_0 trade size, six seeds). DD shows the mean drawdown within each episode, while State-dep. shows the number of seeds where actions changed with the observed state. . .	71
5.4	test results over the 24 held-out months (2024–2025): benchmarks and the three learned agents in their final configuration, six seeds each.	73
5.5	Risk-penalty size on the validation months (Deep Q-learning, six seeds). The rule re- quires $\Delta W \geq +\$39.17$ and exposure ≥ 0.223 ; the largest passing value that meets both requirements is $\lambda = 0.10$	78
5.6	Wealth-change versus risk-aware reward ($\lambda = 0.10$) on the held-out months, ten-feature state, six seeds.	79
5.7	Transaction Fee sizing on the validation months; six seeds. Trades are per month. Fees are shown in basis points and as a percentage of trade value.	80
5.8	Simulated-data study: mean wealth change per episode on the 600 balanced test episodes, per-regime means, and the real-data transfer test.	83

Chapter 1

Introduction

1.1 Background and Motivation

A private investor holding an S&P 500 asset faces a similar decision each trading day: buy more, sell part of the position, or do nothing. On most days, the choice may have little effect on the overall outcome. During a heavy market move, however, the decision can have a much larger impact.

March 2025 is one of such periods studied in this dissertation. Under the experimental setup used later in this study, it was the worst month for a simple buy-and-hold strategy, which lost \$398.41 from an initial \$10,000 investment, approximately 4% of the starting capital. An investor who reduced their position before the decline could have avoided part of this loss. However, reducing positions too often would also mean missing gains during the other months, which produced an average monthly gain of more than \$200.

The problem is therefore not just when to invest, or deciding when the market will rise or fall, but how to make a series of decisions about an asset's position as market conditions change, as each decision made affects the position long-term and the consequences of that action will only become clear several days or weeks later.

This makes trading a sequential decision-making problem, which is one reason reinforcement learning (RL) [1] is appropriate for the task. RL allows an agent to learn by interacting with an environment, observing the results of its actions, and receiving rewards over time based on those results. In essence, it learns a policy/strategy that makes better decisions to receive more rewards.

Trading has several characteristics that make it suitable for this machine learning approach. The agent observes information about the market and its portfolio, selects from a set of possible actions such as buying, selling, or holding, and receives a measurable reward based on the change in portfolio wealth. As a result, reinforcement learning has been increasingly applied to financial trading [2, 3, 4].

However, another important problem arises. The Stock/Equity markets generally rise over the years rather than fall. Meaning that, in a rising market, a strategy that simply remains invested without any pulling out can generate a strong return without using any market information. A learned agent that produces a positive return therefore has not necessarily demonstrated, by itself, that it has learned a useful trading strategy in such markets.

Two questions become relevant. First, can the learned agent outperform the simple buy-and-hold strategy with which it is being compared? Second, does the agent actually use the information available to it when deciding what action to take? These questions are important when evaluating on data not used during training.

This dissertation therefore focuses on both the financial performance of reinforcement learning agents and whether their decisions show evidence of using the information provided to them.

1.2 Problem Statement

The main problem this dissertation addresses is whether reinforcement learning algorithms can learn a trading policy for a portfolio — containing only one asset in this case — that beats a simple benchmark strategy on unseen data.

However, financial performance alone is not enough to determine that an agent has actually learned a useful trading strategy. For example, an agent that buys whenever it can, would perform well in a generally rising market without using any of the market information provided to it. A positive return from such a policy could be incorrectly interpreted as proof that the agent learned when to buy, sell, or remain out of the market, whereas that was not the case.

Evaluating a learned trading agent therefore requires a second perspective alongside financial performance: **behaviour**. The study must analyse whether the agent's actions are a result of changes in the observed market and portfolio state. This helps to differentiate a policy that genuinely responds to the information it receives from one that simply follows a fixed trading pattern.

This dissertation therefore addresses two interrelated problems: whether reinforcement learning can produce better financial performance than simple benchmarks, and whether that performance stems from a learned policy that actually uses the information available to it.

1.3 Research Questions and Objectives

The aim of this dissertation is to evaluate three reinforcement learning algorithms — REINFORCE [5], Deep Q-learning (DQN) [6], and Proximal Policy Optimization (PPO) [7] — and determine whether they can learn profitable and genuinely intelligent trading strategies for the SPY exchange-traded fund.

The study is structured around five research questions:

1. **Performance:** Do the learned agents outperform simple benchmarks on unseen data after transaction costs are taken into account?
2. **Behaviour:** Do the learned agents use the market and portfolio information provided in their state representation when making trading decisions?
3. **State representation:** Does adding a drawdown feature to the main nine-feature state change the behaviour learned or the financial performance?
4. **Reward:** Does a risk-aware reward improve the balance between return and drawdown, rather than simply encouraging the agent to trade less?
5. **Data:** Do the amount and balance of the training data explain the learned behaviour, and does training on large, balanced simulated data produce the state-dependent trading the real data could

not?

These research questions are touched upon through eight measurable objectives, each traceable to a specific part of the dissertation

- O1.** Build a reproducible daily trading environment for SPY using monthly episodes, transaction fees, and action masking, and verify the environment (Chapter 4).
- O2.** Implement REINFORCE, DQN, and PPO and verify their implementations before running the main trading experiments (Chapters 3 and 4).
- O3.** Calibrate the fixed trade size on validation data so the agents' exposure to the market is not limited, providing a fair comparison to the buy-and-hold benchmark (Section 5.3).
- O4.** Apply behavioural analysis to determine whether the learned policies change their actions in response to the observed state (Section 5.4).
- O5.** Measure the effect of adding drawdown to the state while keeping the other experimental conditions fixed (Section 5.5).
- O6.** Compare the final learned agents with the benchmark strategies over 24 held-out test months, using six independently trained random seeds for each configuration (Section 5.6).
- O7.** Examine whether results change when a different reward formula is used (Section 5.7).
- O8.** Build a calibrated market simulator with controlled up, down, and flat regimes, retrain all three algorithms on thousands of balanced simulated episodes, and evaluate them per regime and on the real held-out months (Chapters 3–5).

These objectives create a structured way to evaluate both the financial performance and the learned behaviour of the three reinforcement learning methods.

1.4 Scope

The study is deliberately narrowed to one asset to make the experiments more controlled and easier to analyse.

All experiments use the SPY asset, an exchange-traded fund that tracks the S&P 500 index, with daily data from 2018 to 2025. The data are divided chronologically into training (2018–2022), validation (2023), and held-out for testing periods (2024–2025). A second, simulated dataset is generated from a regime-switching model calibrated on the training years alone; it supplies thousands of balanced episodes for the controlled study, while the real 2024–2025 months serve as its transfer test.

Each monthly episode starts with \$10,000 of capital and no asset holdings, and each transaction to open a position incurs a 0.05% fee. The available actions are discrete: the agent can buy a fixed fraction of the initial capital, sell the same amount, or hold. The action masking prevents invalid actions, such as attempting to sell when no position is held.

Three reinforcement learning algorithms are evaluated. REINFORCE and Deep Q-learning are implemented from scratch, while PPO is implemented using a popular library.

Several areas are intentionally outside the scope of this study. These include multi-asset portfolios, continuous position sizing, intraday trading data, and live deployment. Keeping the study within these boundaries limits the number of factors that can affect the results, making it easier to identify which specific factor causes a change in performance or behaviour during testing rather than multiple changes at once.

1.5 Contributions

This dissertation makes five main contributions:

1. A framework for evaluating trading agents.

The study uses a controlled process in which experiment variables, such as trade size and transaction cost, are set using validation data and finalized before moving on to the held-out test data. The trading environment, state features, rewards, and benchmark strategies are also verified through tests to identify and correct problems before final evaluation,

2. A behavioural analysis for testing state utilization.

A separate analysis was developed to determine whether a policy learned actually responds to the information in its state, rather than simply following a fixed pattern of available actions.

In this evaluation, steps are grouped by their action mask. When a policy selects different actions at steps with the same mask pattern, the difference cannot be explained by the available actions alone, thereby demonstrating a change in its action based on the current observed state. The analysis does not identify which individual state feature drives this dependence, nor does it make claims about the learned network's internal representations.

This reveals a difference between algorithms that would not be seen from financial return figures alone: all six Deep Q-learning seeds were state-dependent, while none of the twelve policy-gradient seeds (REINFORCE and PPO) were (Section 5.4).

3. An explanation for the negative results.

None of the three learning algorithms outperformed buy-and-hold across the 24 held-out test months. The behavioural analysis helps explain this outcome. The policies that achieved returns closest to the benchmark, REINFORCE, generally did so through a fixed pattern of high exposure, while the algorithm that did use its state information, Deep Q-learning, reduced market participation rather than successfully understanding market conditions to perform better.

Although the overall result is negative, evidence from the six training seeds per configuration, behavioural analysis, and peripheral experiments on state representation, reward function, and transaction costs supports it as a finding rather than a failure.

4. A calibrated market simulator for controlled evaluation.

The study builds a regime-switching market generator whose drift, volatility, and intraday range are fitted on the real training years alone. The generator supplies thousands of balanced training episodes and, more importantly, labelled test conditions: every test episode is known to be a rising, falling, or flat market, so the evaluation can state how each agent behaves in each condition rather than in whatever mixture one historical period happens to contain.

5. Evidence that the data, not the algorithms, set the limit.

Retraining the same three algorithms on 3,000 balanced simulated episodes, with a smaller network, inverts the real-data ranking. Deep Q-learning, the worst method on real data, becomes profitable and state-dependent on all six seeds, earning \$126.94 per episode on a test set where no fixed trading pattern is profitable, and transferring profitably to the real held-out months. The policy-gradient methods collapse to fixed corner policies once always-buying stops paying. The comparison shows that the real-data failure was a property of 60 imbalanced training episodes, not of the learning methods.

1.6 Dissertation Structure

The dissertation is organized as follows. Chapter 2 reviews existing research on reinforcement learning for financial trading and portfolios and explains how this study's evaluation approach relates to previous work.

Chapter 3 presents the mathematical formulation of the portfolio trading problem, describing the trading environment as a Markov decision process; the state, action, and reward formulas; the three reinforcement learning algorithms; and their optimisation functions.

Chapter 4 describes the system implementation, data pipeline, verification process, trading environment, and the corrections made during development.

Chapter 5 presents the experiment results and analyses them in relation to the five research questions. It reports the trade-size calibration, behavioural analysis, benchmark comparison, state representation and reward experiment, and transaction-cost calibration on the real data, then diagnoses the data bottleneck behind those results and reports the simulated-data study with its controlled per-regime evaluation and real-data transfer test.

Finally, Chapter 6 concludes the dissertation by summarising the main findings, critically evaluating the work, discussing its limitations and possible directions for future research.

Chapter 2

Background and Literature Review

This chapter introduces the main concepts and prior work on which the rest of the dissertation is based. It begins with an overview of reinforcement learning and explains why the portfolio trading problem can be formulated using the reinforcement learning framework. It then presents the two algorithm families used in this study: value-based methods, from Q-learning extending to Deep Q-learning (DQN), and policy-gradient methods, from REINFORCE to Proximal Policy Optimisation.

The second half of the chapter focuses on reinforcement learning in financial trading. It reviews findings from previous studies, particularly how researchers have evaluated trading agents and the limitations that commonly appear in this area of research.

The chapter also considers the market-efficiency perspective, which provides an important baseline for measuring trading performance—judging whether a learning algorithm can consistently find profitable trading opportunities.

The chapter concludes by bringing these points together and identifying the small gap addressed by this dissertation—whether financial performance is sufficient to establish that a trained trading policy actually uses the information available to it.

2.1 Reinforcement Learning for Sequential Decisions

Most machine learning taught at the master's level is based on supervised learning, where a model is given inputs together with the correct answers and learns to reproduce the desired output. Reinforcement learning (RL) works differently. Instead of being given the correct action, the learner, known as the **agent**, interacts with an environment and learns from the results of its own decisions [1].

This interaction is a repeated process. At each step, the agent observes the current state, selects an action, and receives a reward, then moves to the next state. The reward indicates how useful the action was, but it does not directly tell the agent whether the action was correct. The agent must therefore learn through experience which actions are likely to produce higher rewards over time.

Trading fits naturally into this framework. A trader looks at the markets, decides whether to buy or sell more units of shares, and then waits for the effect of that decision on his portfolio account. A decision made early in a month affects the portfolio permanently from that point onwards. The agent must therefore learn not only whether a decision was profitable, but also how earlier actions contribute to later outcomes. This is known as the **credit-assignment problem** and is an important part of the mathematical formulation presented in Chapter 3.

The absence of explicit labelled answers also creates a practical issue. In supervised learning, an incorrect

prediction can be compared directly with a known target during training. In reinforcement learning, however, there is no equivalent answer to compare against during training. An experiment can therefore run without producing an obvious error in the environment, reward calculation, state representation, or action handling and still learn the wrong behaviour. This was an important consideration in this project and is one reason the verification tests were developed for use before the main experiments, as described in Chapter 4 and Appendix A.

The broader history of reinforcement learning systems shows this approach can learn complex sequential tasks. TD-Gammon, for example, learned to play backgammon strongly through self-play [8]. Later, Deep Q-learning achieved human-level performance across several Atari games using raw image inputs [6]. The analogy with trading is not that financial markets are equivalent to games but that in both cases, the agent is not given the correct action for each situation. Instead, they learn from a numerical reward or game score. In a trading environment, financial outcomes such as changes in portfolio wealth play this role.

The standard mathematical framework for representing sequential problems is the **Markov Decision Process (MDP)**. Its foundations are closely tied to Bellman's work on dynamic programming [9], and Sutton and Barto present a modern reinforcement learning formulation [1]. An MDP represents a problem through states, actions, transitions to the next state, rewards, and a discount factor, assuming that the current state contains all the information needed to make the next decision. In financial markets, however, it is difficult to show that any collection of market observations can fully capture the information relevant to future prices.

For this reason, Chapter 3 uses a practical state representation rather than claiming that the Markov assumption is perfectly satisfied: it captures market, portfolio, and timing information rather than claiming to reconstruct the complete market state.

The study also examines this design experimentally by testing whether adding drawdown information to the state changes the behaviour or performance of the learned agents. The next section discusses the two main families of reinforcement learning algorithms considered in this dissertation.

2.2 Two Families of Learning Algorithms

The algorithms used in this dissertation belong to two different families of reinforcement learning methods. This section discusses the main ideas behind each family and explains why comparing them on the same trading problem is useful.

Value-based methods estimate the value of taking an action and derive a policy from those estimates, whereas policy-gradient methods select the action with the highest probability to optimise the policy itself. These differences matter because they offer two different learning mechanisms for the same trading problem.

2.2.1 Value-based methods: from Q-learning to DQN

The Value-based family approaches the question of what action to take indirectly. Rather than learning the policy directly, the agent learns how valuable different actions are in a given state and then selects the action with the highest estimated value.

This approach started from **temporal-difference learning** [10], which updates value estimates using the difference between sequential predictions rather than waiting until the end of an episode to observe the final outcome. **Q-learning** turned this idea into an off-policy method that directly estimates the action-value function [11]. Under appropriate conditions, including continued exploration of state-action pairs and suitable learning-rate decay, tabular Q-learning converges to the optimal action values for finite state spaces. This convergence property does not depend on the particular actions used while learning.

In tabular Q-learning, a separate value is stored for every state-action pair. This works well when the number of possible states is small, such as in a simple grid-world, but becomes impractical when the state contains continuous financial information. Two very similar market states would also be treated as completely separate entries in the table. As a result, experience gained in one state would not automatically help the agent in a similar state.

Deep Q-learning (DQN) [6] addresses this problem by replacing the table with a neural network that estimates the action values. The network can therefore generalise across similar states, allowing information to be shared, but it introduces instability. However, combining neural networks with bootstrapping and off-policy learning can make training unstable [12, 1].

DQN then introduces two important mechanisms to improve this stability. **Experience replay** [13] stores previous transitions in a replay buffer and samples them later in random minibatches during training. This reduces the strong correlation present in consecutive experiences, allowing each transition to contribute to multiple updates. A **target network** is also used as a periodically updated copy of the value network to construct the learning target, preventing it from changing with every network update as the parameters are being optimised.

These mechanisms are highly useful for financial data, where consecutive daily observations are naturally correlated. Both were implemented from scratch in this project, as described in Chapter 4. As a result, the implementation shows the underlying learning mechanics rather than treating the algorithm as a library component.

Several Other improvements have since been proposed to address the limitations of the original DQN formulation. For example, **Double Q-learning** [14] was developed partly to reduce overestimation bias in action values that can arise in standard Q-learning when the same estimates are used to select and evaluate actions.

Such modifications are relevant to the broader DQN literature or discussed as possible future improvements but are not included in this study's present implementation, allowing the core DQN method to remain relatively simple and easier to verify.

2.2.2 Policy-gradient methods: from REINFORCE to PPO

Policy-gradient methods take a more direct approach to the decision problem. Instead of first estimating the value of each action to derive a policy from those estimates, the agent learns a policy that maps states to actions and adjusts its parameters to increase the expected return. The policy is represented by a differentiable model, typically a neural network, whose parameters are updated for the objective.

This study uses the foundational policy-gradient algorithm **REINFORCE** [5]. One of its main advantages is its simplicity, as it makes the relationship between the policy, observed states, selected actions, and expected return relatively transparent. It does not require bootstrapping, experience replay, or a separate target network. The policy is updated using returns from complete episodes, making the learning process straightforward to understand. For these reasons, REINFORCE is used as the main policy-gradient implementation in this dissertation.

The main weakness of REINFORCE is the high variance of its updates. The learning signal is based on Monte Carlo returns, which can vary considerably between episodes. This can make learning slow and unstable when rewards are noisy.

Later policy-gradient methods introduced techniques to address this problem. Actor-critic methods use a learned value estimate as a baseline to reduce the variance of the policy update [15]. The actor represents the policy while the critic estimates value, allowing the learning signal to use an advantage rather than the raw return. **A2C** is one example of this approach [16]. **Generalised Advantage Estimation (GAE)** [17] is another way of controlling the bias-variance in the advantage estimate.

A second problem is the size of policy updates. An update that changes the policy too aggressively can cause training to become unstable and degrade performance sharply. Trust-region methods [18] address this by restricting how far the policy can change between successive updates. **Proximal Policy Optimisation (PPO)** [7] uses a simpler approach by clipping the policy update when it moves too far from the previous policy.

PPO is widely used in reinforcement learning applications, including trading frameworks such as FinRL [19], and is therefore included in this study as a library-based policy gradient method for comparison while implementing REINFORCE from scratch. This allows for both a transparent baseline policy-gradient method and a more modern, stabilised method.

The trading environment also contains actions that are not always valid. For example, an agent should not be allowed to sell when it has no position or buy when there is insufficient cash. All three learning methods use **invalid-action masking** to prevent such actions from being selected while preserving only valid actions. The use of action masking in reinforcement learning, including its interaction with policy-gradient methods, has been examined by Huang and Ontañón [20]. In this dissertation, masking is part of the common environment rather than a difference between algorithms, ensuring that the comparison is made under the same set of trading constraints.

2.2.3 Why compare the two families?

The two algorithm families have different implementation trade-offs.

Value-based methods such as DQN can learn from individual transitions and reuse previously collected experience through replay. However, they rely on bootstrapping value estimates, meaning the agent learns partly from estimates produced by other learned estimates in the same learning process.

Policy-gradient methods instead optimise the policy directly with respect to the objective. In its basic form, REINFORCE provides a balanced policy-gradient estimate, but it requires new experiences for each update from fresh trajectories, and gradient estimates can suffer from high variance. PPO improves the stability of this family through constrained policy updates.

Neither family can therefore be assumed to be better for the portfolio trading problem considered here. Comparing them in the same environment and under the same evaluation procedure is therefore useful. If both families develop similar behaviour, this may suggest that the behaviour is mainly caused by the structure of the problem, reward, state representation, or environment. If their behaviour differs, the difference becomes evidence that the choice of learning method influences the resulting strategy.

This comparison becomes an important part of the experimental design in Chapter 5, as it is not limited to only asking which algorithm produces the highest financial return. It also asks whether the resulting policies learn fundamentally different ways of responding to the trading environment. The next section discusses how previous research applies these algorithms to financial markets.

2.3 Reinforcement Learning for Trading

The use of reinforcement learning (RL) for trading has been studied for many years, but the application comes with several well-known challenges. One of such early influential work is Moody and Saffell on direct reinforcement methods [2], which uses gradient-based methods to learn trading policies directly. Their work raised an important point: the reward should account for risk. If an agent is rewarded only for increasing profit, it may have a strong incentive to remain highly exposed to the market without explicitly rewarding control of downside risk.

This observation is highly relevant to this dissertation, giving us reason to run two experiments in Chapter 5: first, to examine what happens when the agents are trained using a simple wealth-change reward and the objective is tied only to portfolio growth; and second, whether adding a drawdown penalty to the reward changes their behaviour.

Drawdown is also introduced separately as an additional state feature. These two uses of drawdown act differently. Adding drawdown to the state changes the information available to the agent, whereas adding a drawdown penalty in the reward changes what the agent is encouraged to optimise. One changes what the agent can see, while the other changes what it is rewarded for. However, the specific drawdown penalty used in this dissertation is also an experimental design choice to be evaluated rather than a reward formulation from previous literature.

Keeping these changes separate allows their effects to be examined separately. The experiments determine whether either change affects the learned behaviour, rather than assuming what the resulting behaviour should be.

Later work extended these ideas using more complex models. Deng et al. [3], for example, combined direct reinforcement with deep recurrent representations. Jiang et al. [21] applied deep reinforcement learning to multi-asset portfolio management using cryptocurrency data, using a different formulation for the action, which represents portfolio allocation rather than a discrete buy, sell, or hold decision.

More recently, frameworks such as FinRL [19] have made it easier to build trading environments and apply these reinforcement learning algorithms to financial experiments, but it often defaults to PPO-based methods.

These frameworks are useful for experimentation, but they can also hide parts of the underlying learning process. This was one reason for implementing REINFORCE and DQN from scratch in this dissertation, allowing their behaviour and training mechanisms to be examined directly.

The study most closely related to the present work is Théate and Ernst [22]. They evaluate a DQN-based trading agent on a test portfolio set of 30 stocks, comparing it with buy-and-hold and traditional active strategies. Their results report that the DQN performs better on average than the benchmark strategies, although they include several active strategies rather than only passive buy-and-hold, which is more relevant to this dissertation. For SPY, which is also used in this dissertation, their reported Sharpe ratio for the DQN was 0.834, the same as buy-and-hold.

Across multiple stocks, they describe the agent's performance as identical or very close to passive strategies and suggest that the agent can tend towards a passive strategy when active trading becomes more uncertain. The variations between training runs showed that the agent traded less as transaction costs increased, eventually becoming passive. They also report non-negligible variance across identical training runs. These findings make Théate and Ernst study relevant to this study's evaluation.

There are, however, important differences between their experimental design and the one used here. Their study trains a separate agent for each stock, allows short positions, and uses a continuous test period, whereas this dissertation studies one ETF, does not permit shorting, and uses independent monthly episodes with a fixed trade size. Their benchmarks also differ from this study's buy-and-hold strategies, preventing a direct numerical comparison between the studies.

More importantly, although their results show that the learned agent can become increasingly passive, they do not directly test how the learned policy generates its decisions: whether the policy responds to its observations or simply follows a fixed trading pattern. Both behaviours could produce similar trading trajectories.

This difference is fundamental to this dissertation, as a policy can achieve a strong return without making meaningful use of market information, particularly when the asset always goes up over time. Chapter 5 therefore includes a specific behavioural analysis to determine whether the actions of each trained agent actually change when its observed state changes while the available actions remain the same.

Read together, the previous studies reviewed here vary widely in how RL-based trading is formulated and evaluated. Some focus on individual assets (Moody and Saffell [2], Deng et al. [3], and Théate and Ernst [22]), while others use multi-asset portfolios (Jiang et al. [21] and Guan and Liu [23]). Action representations range from discrete trading decisions [2, 3, 22] to continuous portfolio weights [21]. Reward functions also vary, from direct profit or return objectives [3, 21, 22] to risk-adjusted objectives such as Moody and Saffell's differential Sharpe formulation [2]. Most papers include transaction costs, but their assumptions and representations differ. Benchmark strategies also vary, from traditional trading strategies [22] to portfolio-selection methods [21].

There is also variation in the results across seed runs. Some studies acknowledge differences between training runs [22], but reported results are often based on individual runs. None of the reviewed studies also describe a procedure in which all experimental design decisions are selected using a validation period before being used on an unseen test period. Behavioural analysis is similarly limited, with previous work using methods such as trading trajectories [22] or feature attribution [23].

These differences make direct comparison between studies difficult and highlight an important issue: the evaluation procedure can be just as important as the choice of algorithm. The studies reviewed in this section were selected because they are particularly relevant to this dissertation's design, especially their use of RL algorithms, transaction costs, and benchmarks. Fischer's survey [4] provides a broader review of RL in financial markets.

2.3.1 Data requirements and simulated markets

There is one further issue the trading studies above rarely confront directly: how much data reinforcement learning needs. The landmark results of the field were all obtained with effectively unlimited experience. The Atari agents consumed tens of millions of frames per game [6], and AlphaGo sidestepped the shortage of human games entirely by generating its own training data through self-play [24]. A single asset's daily history offers nothing comparable. Eight years of daily bars is roughly 2,000 observations, and grouping them into monthly episodes leaves only dozens of episodes to learn from.

The common response in the literature is to train on generated data. Théate and Ernst, whose setting is the closest to this study's (a deep Q-network trading single stocks on daily bars), trained their agent "entirely based on the generation of artificial trajectories from a limited set of stock market historical data" [22]. The earlier studies solved the same problem differently: Moody and Saffell matched a tiny single-layer policy to scarce monthly data [2], while Deng et al. [3] and Jiang et al. [21] moved to minute and 30-minute bars, where a few years supply hundreds of thousands of observations. Open-source frameworks such as FinRL [19] and the Gym trading environments [25] structure the task so that any price series, real or synthetic, can drive the environment. Generators range from classical stochastic-process models, including regime-switching models built on hidden Markov chains [26], to learned generators such as time-series GANs [27]. A generator gives the experimenter two things a historical record cannot: as many training episodes as the learner needs, and control over which market conditions appear in training and in testing. The price is a domain gap, since no generator reproduces real markets exactly, so studies that train on synthetic data normally keep a real held-out period as a transfer test.

This dissertation follows that pattern. Its first study trains on real monthly episodes and shows what goes wrong at that data volume; its second trains on episodes from a calibrated regime-switching generator and tests under controlled conditions, with the real held-out months retained as a transfer test. Chapters 3 and 4 define the generator, and Chapter 5 reports both studies.

The next section therefore focuses more specifically on how to evaluate trading agents and on the weaknesses that can arise when financial performance is considered without sufficient behavioural and experimental controls.

2.4 Evaluating Trading Agents Honestly

A review of reinforcement learning in finance [4], together with research on financial machine learning [28, 29], identifies several common problems that can make trading results appear stronger than they are. Four of these issues are relevant to this dissertation, and each one influenced the experiment design. These are not claims that every reviewed study states as problems; they are risks that an experiment should account for when being designed.

The first issue is **weak benchmarks**. A learned trading strategy can appear successful if it is compared with a benchmark that invests less capital or does not represent a realistic alternative. For this reason, this study uses buy-and-hold as the main benchmark, with the same initial capital as the learned agents. The trade size was also calibrated in Chapter 5 so the agents could reach a comparable level of market exposure. This keeps the comparison fair, as the agents are not disadvantaged simply by being unable to invest enough capital.

The second issue is **look-ahead bias**. Financial data are ordered in time, so treating it like randomly ordered observations and shuffling the data before splitting could place future information in the training set and produce overly optimistic results. The data in this study are therefore divided chronologically: 2018–2022 for training, 2023 for validation, and 2024–2025 as the held-out test period. The test months were not used to select hyperparameters or make any design decisions, as described in Chapter 4.

The third issue is **ignoring transaction costs**. A strategy that trades frequently may look profitable when every trade is assumed to be free, but the same strategy could lose much of its return once real trading costs are included. In this dissertation, transaction fees are part of the trading environment and are therefore included while the agents are learning rather than only after training. Chapter 5 also varies the assumed fee to examine whether the learned strategies trade less when trading becomes more expensive.

The fourth issue is **backtest overfitting**. Repeatedly adjusting a strategy after looking at its test results gradually turns the test set into another training set. Bailey et al. [28] discussed this problem for investment backtests, and López de Prado [29] discussed methodology for controlling leakage and data reuse. The defence here is that all design choices were selected using the training and validation data, then evaluated once on the held-out test months, without further tuning.

A related concern comes from deep reinforcement learning itself. Henderson et al. [30] show that reported **results can vary substantially across random seeds**, making a single successful training run a weak

basis for comparison. Each configuration in this study is therefore trained with six different seeds, and the spread is reported alongside the mean rather than presenting one favourable run as representative.

Financial economics provides more context for understanding the results. Under the **efficient market hypothesis**, current prices should already reflect the information contained in past prices [31]. If this assumption holds, an agent whose market information consists only of price history should not be expected to consistently earn higher risk-adjusted returns than the market itself. In this study, buy-and-hold on the single fund provides the practical benchmark against which this expectation is tested.

Portfolio theory also emphasises that return cannot be considered separately from risk [32]. A strategy that earns less may still be useful if it achieves a large enough reduction in risk. The evaluation therefore reports mean return together with drawdown and the Sharpe ratio [33]. No single metric decides whether a strategy is successful; the results are read together so that the trade-off between return and risk remains visible.

This approach does not aim to guarantee a positive result for the learning algorithms. If an agent does not beat buy-and-hold, that outcome is treated as a finding to explain rather than a failure to hide. Understanding why a strategy does not improve on a simple benchmark is therefore just as important as reporting successful results. However, even a carefully controlled financial experiment only establishes what happened to the portfolio; it does not establish what the agent learned to do. That distinction leads to the gap considered in the next section.

2.5 The Gap Addressed In This Study

Financial performance alone does not tell us what a trading agent has actually learned. An agent can follow a simple rule, such as buying whenever possible, and still perform well during a rising market without using any market information. At the same time, an agent that genuinely responds to market conditions may earn less over the same period. For this reason, financial performance and agent behaviour need to be evaluated separately, and the second is needed to interpret the first.

This distinction matters because most of the trading studies reviewed in this chapter focus mainly on financial measures such as return, profit, Sharpe ratio, and drawdown. More recent work also examines how trained models make their decisions, including feature-attribution methods by Guan and Liu [23]. These methods can identify which inputs are related to a model's output. However, they do not establish whether the policy changes its actions as its inputs change. A policy that produces the same action regardless of its observations leaves no variation in action for an attribution method to explain.

Checking whether a learned agent responds to its observations is already used in reinforcement-learning research and in the wider evaluation of machine-learning systems. This dissertation therefore does not claim to introduce sensitivity testing or behavioural analysis as new concepts, as there are related ideas that exist in input sensitivity.

The more specific issue addressed here is how to test dependence on state information in a trading environment with constraints on actions. The agent's cash and holdings determine which actions are legal:

it cannot buy when it has insufficient cash and cannot sell when it has no position. The set of available actions can therefore change during an episode, and the selected action may change simply because an action became unavailable rather than because the agent responded to market information, and this would risk counting these forced changes as evidence of state use.

To account for this, this dissertation defines state dependence intentionally: **trained policies must select different actions in situations where the same actions are legally available, with the difference occurring in the observed state.** Conditioning the analysis on action masking reduces the chance of misinterpreting changes caused by the constraints as evidence that the agent is using its observations.

The diagnostic, however, answers a limited question — does the policy show any evidence of changing its decisions in response to its observed state? It does not identify which feature caused the change, whether the response was economically sensible, or whether it improved financial performance.

Among the trading studies reviewed here, policy responsiveness is not reported as a quantitative evaluation result alongside financial performance. This is a narrower claim than saying previous research has never used sensitivity testing or behavioural analysis. Established practices such as chronological splits, transaction-cost modelling, benchmark comparisons, and multi-seed reporting are also not presented as new contributions here.

Théate and Ernst [22] show why this distinction matters. Their DQN moved towards passive trading under some conditions, but passive behaviour could come from a policy that responds to market information or from one that follows a fixed trading rule. Without directly testing the relationship between the observed state and the chosen action, both cannot be explained separately.

This dissertation addresses this gap through a controlled comparison of three different learning algorithms. REINFORCE, Deep Q-learning, and PPO are trained using the same observed state, the same masked action space, the same fee-inclusive reward, and the same chronological splits. Their financial performance is compared against benchmark strategies with the same capital, and their behaviour is examined using state-dependence analysis. The diagnostic is reported for each seed alongside the financial results rather than being treated as a technical implementation check.

Two further experiments examine whether changing the information available to the agent or changing its incentive affects the learned behaviour. The first adds drawdown to the state, allowing the agent to observe its recent portfolio losses. The second adds a drawdown penalty to the reward, changing the incentive to take risks. One changes what the agent can see, while the other changes what the agent is encouraged to optimise.

The contribution of this dissertation is therefore not the invention of a new evaluation method, but the controlled combination of experimental practices established from prior works to measure the agent's state-dependent behaviour along with financial performance. This allows the evaluation to distinguish between an agent that earns a strong return by following a simple fixed rule and one whose decisions actually respond to the information provided.

A second gap follows from the data-requirements discussion in Section 2.3.1. The single-asset trading

studies reviewed here train on historical data alone and rarely ask whether that history is large enough for the networks being trained on it. This dissertation makes the question explicit: it measures what the algorithms learn from 60 real episodes, diagnoses the failure, and then repeats the comparison on thousands of balanced generated episodes where the data volume is no longer the binding constraint.

These points lead directly to the five research questions introduced in Chapter 1. The first asks whether the learned agents can outperform meaningful benchmarks after transaction costs. The second asks whether their decisions depend on the information in their states. The third asks whether adding drawdown to the state changes behaviour or performance. The fourth considers whether penalising drawdown through the reward produces a different risk-sensitive outcome. The fifth asks whether the amount and balance of training data explain the learned behaviour, and whether fixing the data fixes the behaviour. The next chapter formalises these questions through the MDP formulation, state representation, trading rules, reward functions, the simulated market model, and the algorithms used in the experiments.

Chapter 3

Methodology & Mathematical Formulation

This chapter explains how the portfolio problem can be treated as a Markov Decision Process (MDP) [9, 1], i.e the trading process is treated as a sequence of step-by-step decision-making problems. The objective is to develop an environment in which an autonomous trading agent can learn an optimal decision-making behaviour by interacting with a simulated market. At each step, the portfolio system observes current market and portfolio conditions, takes an action such as buying, selling, or holding, receives a result based on how the portfolio's wealth changed, and then moves to the next step.

The chapter then explains the two main learning methods used in the project: **REINFORCE** [5]—the Monte-Carlo policy gradient, which learns by improving its strategy based on the rewards it receives—and **Deep Q-learning** [11, 6], which learns which actions are likely to give the best results. Finally, **PPO** is introduced as an additional method in the ablation row to compare against the other approaches, as it is designed to make the policy learning process more stable.

Table 3.1 collects every symbol used for reference.

Table 3.1: Notation used throughout Chapters 3–5.

Symbol	Meaning
$s_t \in \mathbb{R}^9$	state observed by the agent at trading day t , containing the nine features used by the trading system (Table 3.3)
$a_t \in \{0, 1, 2\}$	action selected by the agent: Hold, Buy, or Sell.
P_t	asset price at day trading t ;
r_t	reward received after an action, calculated from the change in portfolio value after transaction costs.
C_t	amount of cash available to the agent at trading day t
h_t	number of (fractional) units of the asset held at trading day t .
$W_t = C_t + h_t P_t$	Total portfolio value, calculated as untouched cash plus the current value of the asset held;
C_0	initial cash available at the beginning of an episode.
c	proportional transaction fee for taking a trade, set to 0.05% per transaction
$A = f C_0$	fixed monetary size of each trade; the fraction f is selected on validation data (Section 5.3)
γ	discount factor used to determine the importance of future rewards

Symbol	Meaning
τ	a trajectory (one complete trading episode), corresponding to one month of daily trading; related to t & T
T	episode length in days (with $T \approx 18-23$)
t	1 trading day
$G(\tau)$	The total return obtained over a complete episode
G_t	The accumulated future reward from time t onwards.
$\pi_\theta(a s)$	policy network that determines the likelihood of choosing action a given the current state s .
$Q_\phi(s, a)$	The Q-network, which estimates the value of taking a particular action in a given state; frozen target copy
Q_ϕ^-	A fixed copy of the Q-network used as a target during Deep Q-learning.
$J(\theta)$	objective function representing the expected return of the policy
ε	exploration rate used by the (ε -greedy) action-selection strategy

3.1 Overview: Trading as a Markov Decision Process

Reinforcement Learning (RL) trains an agent to make sequential decisions by interacting with an environment. The environment is modelled as an MDP where decisions are made at regular time steps, represented by the tuple $(\mathcal{S}, \mathcal{A}, \mathcal{P}, \mathcal{R}, \gamma)$: at step t the agent observes the current state $s_t \in \mathcal{S}$, selects an action $a_t \in \mathcal{A}$, receives a numerical reward based on the outcome r_t , and the environment moves to the next state s_{t+1} according to the current state and the action selected by the agent $P(s_{t+1}|s_t, a_t)$.

In this project, the trading environment is modelled in the same way, providing a formal description of how the trading agent interacts with the market over time; one time step t corresponds to one trading day.

The state s_t represents all information available to the agent on day t to fully capture its position within the trading episode. This includes the selected market features, the agent's current portfolio position, and its available cash.

Based on this state, the agent selects one of three possible actions $a_t \in \{0, 1, 2\}$: *Hold the current position, Buy a fixed amount, or Sell a fixed amount*. The selected trade is then executed, with transaction fees c applied where appropriate.

After taking an action, the environment calculates the change in the agent's total portfolio value and returns it as the reward r_t .

The environment advances to the next trading day s_{t+1} , and the process repeats throughout the episode τ , allowing the agent to update its policy π_θ or value estimates Q_ϕ based on the experience collected.

The interaction can therefore be represented as: $s_t \rightarrow a_t \rightarrow r_t, s_{t+1}$

Figure 3.1 illustrates the interaction between the trading agent and the environment.




Figure 3.1: The agent–environment interaction loop for the trading MDP. One cycle corresponds to one trading day.

To make the MDP formulation easier to understand, Table 3.2 compares the main components of the trading task with those of a familiar game-playing task, Flappy Bird, which was used initially to understand Reinforcement Learning.

Table 3.2: Mapping a popular game MDP to the portfolio trading MDP.

MDP component	Flappy Bird	Trading (this work)
State s_t	Bird's position and speed, together with the position of the next pipe	Market information and the agent's current portfolio position
Action a_t	Flap or do not flap	Hold, Buy, or Sell
Reward r_t	Survival score based on how long the bird remains in play	Change in portfolio wealth after trading fees
Episode τ	One complete game session	One calendar month of daily trading
Terminal rule	The game ends when the bird crashes or hits the ground	The episode ends at the final trading day, with any remaining position sold

This analogy also highlights a design principle used throughout the project: as the Flappy Bird state contains the information relevant to the next decision, the trading state includes the relevant information available to the agent at each trading step, without being given future information. Section 3.3 discusses the design of the trading state in more detail.

3.2 Episode Definition and Data Structure

3.2.1 Episode definition (τ)

Each episode (τ) represents one calendar month of daily closing prices, from P_0 to P_T ($P_0, P_1, \dots, P_T$) where T is typically between 18 and 23 trading days per month, excluding public holidays when markets aren't open and weekends of market inactivity ($T \approx 18-23$). At the beginning of each episode, the agent starts with cash and no asset holdings. It can then trade throughout the month based on the available information. At the end of the final trading day t , any remaining position is automatically sold. This ensures that no position is carried over from one episode to the next.

3.2.2 Independence of Trading Episodes

For the purposes of the experiment, each monthly episode is treated as a separate and independent trading period. This is similar to treating each game session in Flappy Bird as a separate episode.

However, this is a simplification because market conditions in consecutive months are not truly independent. For instance, events and trends in one month do influence the market in the following month. The resulting data are therefore not independent over time.

Modelling these temporal relationships in greater detail would require time-series or partially observable approaches, which are outside the scope of this project. The implications of this assumption are discussed further in the conclusion.

3.2.3 Choice of Episode Length

Monthly episodes were selected to provide a balance between the length of each trading period and the total number of training examples available.

A single year of daily data provides approximately 12 monthly episodes and around 250 daily trading decisions. Using eight years of data from 2018 to 2025 gives 96 monthly episodes and approximately 2,000 daily transitions. These episodes are divided chronologically into 60 months for training, 12 months for validation, and 24 months for the held-out test period, as detailed in Chapter 4. The data are not shuffled, which prevents information from later periods from being used during training and helps avoid look-ahead leakage.

The same environment could also be adapted to shorter intervals, such as minute-level data, without a fundamental change to its design, but intraday trading is outside the scope of this study.

3.2.4 Selection of the Historical Period

The experiments use daily market data from 2018 to 2025 as a fixed period. Earlier data, including data from 2016 and before, were not included in the main experimental period because the study aims to evaluate the learning algorithms under relatively recent and varied market conditions. These selected periods

contain a range of market conditions, including relatively stable periods in 2018, the market decline associated with the COVID-19 pandemic and its subsequent recovery, the later period of increased inflation and interest rates, and the war in Ukraine in 2022, which caused weaker market conditions. As a result, the dataset is not limited to a single type of market condition, providing a useful range for evaluating the behaviour of the portfolios trading agents without extending the dataset into substantially older market years.

Including older data, such as 2016 or earlier, would increase the number of observations but could also introduce market conditions that differ considerably from those in the selected period, such as older microstructure and fee regimes. Earlier historical data are therefore left to future work as an additional robustness test of how well the trained agents generalise across market regimes.

3.2.5 The Simulated Market Model

Even the full 2018–2025 period yields only 60 training episodes, and Chapter 5 shows this is far too few for the networks to learn rather than memorise. The study therefore also defines a simulated market, and uses it as the primary source of training data. The idea follows a standard pattern in reinforcement learning: when real experience is scarce, generate experience from a model of the task and keep the real data for a final test [24].

The simulated market has three **regimes**: up, down, and flat. An episode is a sequence of 21 daily bars drawn entirely from one regime m . Within the regime, daily log-returns are independent draws from a normal distribution,

$$r_t = \ln \frac{P_t}{P_{t-1}} \sim \mathcal{N}(\mu_m, \sigma_m^2),$$

so the price follows a geometric random walk whose drift μ_m and volatility σ_m belong to the regime. High and low prices extend each day's move by random wicks sized to match the real average true range, and volume is drawn from a lognormal distribution. The regime parameters are estimated from the real training months only; Chapter 4 reports the fitted values.

The model is deliberately the simplest one that supports a controlled experiment. Because the generator's regime is known for every episode, the evaluation can report how a policy behaves in rising, falling, and flat markets separately, which no historical test period can guarantee to provide. Because episodes are generated, the training set can be as large and as balanced as the experiment requires. And because consecutive generated episodes share nothing, the independence assumption of Section 3.2.2, which is a simplification for real months, holds exactly in simulation.

The cost of this simplicity is realism. Real returns are not independent across days, volatility clusters within months, and regimes change mid-month. A policy trained on the simulated market therefore faces a domain gap when applied to real data, and Chapter 5 measures that gap directly by evaluating the simulation-trained agents on the real held-out months.

It should also be noted that real financial markets are more complex than any MDP and cannot be assumed to be fully observable. The MDP used here is a simplified representation of the portfolio trading problem; the information the agent cannot observe, and the consequences of this, are discussed with the state representation in Section 3.3.

3.3 State Representation Design

In reinforcement learning, how information is presented to the agent plays an important role in the quality of decisions it can learn. The state representation must provide enough information for the agent to understand the current situation and choose an appropriate action.

In this portfolio trading problem, this requires capturing information about both the market and the agent's own portfolio. The agent needs to know how the asset has been moving, but it also needs to know how much cash is available, whether it currently holds the asset, whether that position is profitable, and how much time remains before the trading period ends.

A state that leaves out important information can make it difficult for the agent to learn effective strategies, even if the learning algorithm is suitable for the task. This is because the agent can make decisions only using the information provided. Therefore, the state representation in this study was developed gradually by first examining simpler representations and then adding information as needed to address specific limitations of the previous representation.

The objective was not to create an unnecessarily complex state representation with many variables, but to design a compact representation containing all relevant information needed for decision-making within the defined portfolio trading environment, while avoiding features that simply duplicate information already available from other variables

3.3.1 Information Available to the Agent

As discussed in Section 3.1, the state should see the whole by providing the agent with information relevant to its decision, as the Flappy Bird state does. In a trading environment, this information can be divided into two halves: **the agent's own portfolio** and **the wider market**. These two areas behave very differently and should therefore be considered separately when designing the state.

The first area is the agent's own portfolio. This includes information such as the amount of cash available, how much of the asset it owns, what it paid to enter a trade position, and the time remaining in the trading period. The agent controls this half; it is not hidden or estimated, and it can be completely represented. Given the sequence of the asset prices and the agent's actions, the portfolio can be updated directly using the trading environment's rules.

The second area is the market. Here, complete information about the market is impossible in principle and cannot be provided to the agent because many factors affecting an asset's price are not directly observable from the available price data. For example, the agent cannot observe other market participants' decisions, expectations, or intentions. Similarly, information such as news, market sentiment, and macroeconomic

events may influence prices without appearing directly in the closing price at the end of the episode.

Adding more features calculated from past prices can give the agent additional information about recent market behaviour, but it cannot reveal everything happening in the market. The attainable goal is therefore a state that is complete on the portfolio side and rich enough on the market side: **the portfolio information can be represented completely because the environment directly controls and records it, whereas market information can only approximate the wider market.** Whether the additional market features provide useful information is ultimately an empirical question, which is considered when evaluating the state representation later in the dissertation.

3.3.2 Criteria for a Suitable Trading State

It is tempting to design a state by asking which features look useful, but this is a weak approach as it says nothing about whether the object being trained is an MDP at all.

Before deciding on the final state representation, three conditions were considered. These provide a basis for deciding whether the proposed state contains the information required for the portfolio trading problem.

1. Reward measurability

The reward should be calculable from information available to the agent and the resulting transition. If the reward depends on information not contained in the state or the transition, the agent is effectively being evaluated using information it cannot observe.

The primary reward used in this study is the change in portfolio wealth, $r_t = W_{t+1} - W_t$. The final state includes cash and asset exposure as fractions of the initial capital, and these two features combine to give the current wealth, so the reward can be calculated from the observed transition (s_t, a_t, s_{t+1}) with no hidden portfolio variable. Section 3.5 defines the reward and shows this derivation in full.

2. Transition consistency

The state should also contain enough information for the next state to be determined from the current state, the selected action, the next observed market price, and the trading environment rules.

This matters because the agent should not need to access hidden historical variables to describe how its portfolio changes after an action.

For example, the current portfolio exposure can be used to determine the position's value, while the current unrealised profit/loss contains information about the position's entry price. Together with the current price and the environment's trading rules, these quantities provide the information required to update the portfolio after an action.

This keeps the state representation consistent with the defined trading environment rather than requiring the environment to retrieve additional portfolio history that is not represented in the state.

3. Decision sufficiency

The state should contain enough information for the agent to make a decision based on the current trading situation without ambiguity. If two different situations look identical to the agent but require different actions, then the state is missing important information, making the representation incomplete.

For example, consider the following two situations:

- the asset price has increased by 2% today, and the agent does not hold the asset;
- the asset price has increased by 2% today, but the agent holds a position that is currently making a loss.

In both cases, the observed price movement is the same. However, the most appropriate action may be different. In the first situation, the agent may consider opening a position, whereas in the second, it may consider reducing or closing its existing position to minimize its losses.

The state representation must therefore include information about both the market and the agent's own portfolio.

3.3.3 Development of the State Representation

The state was developed in stages, starting from the smallest representation that could describe a trading decision and adding information only where the previous stage left the agent unable to distinguish situations that require different actions. Table 3.4 at the end of this section summarises the full progression.

The starting point was a price-only state, $s_t^{(0)} = [\Delta P_t]$, where $\Delta P_t = (P_t - P_{t-1})/P_{t-1}$ is the one-day return. This is insufficient because the same market movement can occur in very different trading situations. A positive daily return means one thing when the agent holds nothing and may consider entering a position, and another when it already holds a large position and may consider reducing it. The agent cannot see its own portfolio.

Portfolio, position, and time information were therefore added in turn, producing the following features:

1. Normalised cash (c_t), defined as $c_t = C_t/C_0$, is the proportion of the initial capital that remains available for future decisions. Using a normalised value keeps the information consistent regardless of the starting account size.

2. Normalised exposure (v_t), defined as $v_t = h_t P_t / C_0$, is the current market value of the open position relative to the initial capital. This replaces the raw unit count h_t because the same number of units can represent a small or large investment depending on the asset price. Together, the two portfolio features recover the agent's current wealth as $c_t + v_t = W_t / C_0$.

3. Unrealised profit/loss (PnL_t):

$$PnL_t = \frac{P_t - \bar{P}_t^{entry}}{\bar{P}_t^{entry}}$$

where $\bar{P}_t^{\text{entry}}$ is the average price at which the current position was opened. This tells the agent whether its position is in profit or loss — two situations that look identical from the market side but may call for different actions. Holding an asset after a price increase may be profitable if it was bought at a lower price, or losing money if the position was entered at a higher price before the price declined.

4. Episode clock ($\square_t$), defined as $\tau_t = t/T$, where T is the total length of the episode. A value close to 0 means the episode has just started, while a value close to 1 means the forced close at the final trading day is approaching. This matters because the same action can have different consequences depending on when it is taken: buying early leaves the position time to grow, whereas buying shortly before the episode ends does not.

3.3.4 Final State Representation

Although the previous stages successfully introduced information about the agent's portfolio and position, they provide limited understanding of the broader market conditions. A trading decision should not depend only on the most recent price change, as different market situations exist, producing similar short-term movements.

For this reason, the agent also needs information that helps it distinguish between market conditions: whether the asset is in a short-term upward or downward trend, whether the market is stable or highly volatile, and whether the current price is above or below its recent average level. Additional market-related features are therefore included to give the agent a more complete view of the trading environment.

The final state representation is defined as:

$$s_t^{(4)} = [\Delta P_t, mom_t, vol_t, gap_t, rng_t, \tau_t, c_t, v_t, PnL_t] \in \mathbb{R}^9$$

The following variables were added:

1. Momentum

$$mom_t = \frac{P_t - P_{t-k}}{P_{t-k}}$$

where $k = 5$ previous days (about one trading week). Momentum gives the agent a longer view of the recent price direction, helping it differentiate a temporary movement from a more consistent trend developing over several timesteps.

2. Volatility

$$vol_t = \sigma(\Delta P_{t-k+1}, \dots, \Delta P_t)$$

where σ is the standard deviation of the last $k = 5$ daily returns. This helps the agent recognise differences

in market uncertainty: a 2% price increase during a relatively stable period is different from a 2% increase during a period of large and frequent price movements, which indicates a riskier situation.

3. Price trend position

$$gap_t = \frac{P_t}{MA_{m,t}(P)_t} - 1$$

where $MA_{m,t}(P) = \frac{1}{m} \sum_{i=0}^{m-1} P_{t-i}$ is the simple moving average of the closing price over the previous ($m = 20$) days (about one trading month). A positive value means the current price is above its recent average level and a negative value means it is below, helping the agent identify potential trends or deviations from normal price behaviour.

4. Normalised Trading Range

The normalised trading range is based on the Average True Range (ATR):

$$rng_t = \frac{ATR_{n,t}}{P_t}$$

The True Range at timestep t is $TR_t = \max(H_t - L_t, |H_t - C_{t-1}|, |L_t - C_{t-1}|)$, where H_t and L_t are the current high and low prices and C_{t-1} is the previous closing price. The ATR is the rolling arithmetic mean of the last ($n = 14$) True Range values (about two trading weeks):

$$ATR_{n,t} = \frac{1}{n} \sum_{i=0}^{n-1} TR_{t-i}$$

The ATR therefore measures the size of recent price movement using the asset's high, low, and closing prices, and dividing it by the current price creates a normalised measure that can be compared across different price levels.

The nine features and what each tells the agent are summarised in Table 3.3.

Table 3.3: Full State Definition

Table 3.3: The nine-feature state representation ($s_t \in \mathbb{R}^9$) used in the main experiments. C_0 denotes the initial cash, h_t the number of asset units held, and $W_t = C_t + h_t P_t$ the portfolio wealth at time t . the mark-to-market wealth. The feature timestep lengths are set to ($k=5$) for momentum and volatility, ($m=20$) for the moving-average reference, and ($n=14$) for the Average True Range.

#	Feature	Definition	What it tells the agent
Market			
1	ΔP_t	$P_t - P_{t-1} / P_{t-1}$	The most recent price movement

#	Feature	Definition	What it tells the agent
2	mom_t	$P_t - P_{t-k} / P_{t-k}$	whether that recent price movement is part of a short-term trend
3	vol_t	$\sigma(\Delta P_{t-k+1}, \dots, \Delta P_t)$	how much the asset price has been varying recently
4	gap_t	$P_t / MA_m(P) - 1$	how far the current price is from its recent average
5	rng_t	$ATR_n(P) / P_t$	The size of the recent price ranges, which closing prices alone cannot show
Portfolio			
6	c_t	C_t / C_0	how much of the initial cash remains available as cash
7	v_t	$h_t P_t / C_0$	The current value of the open position relative to the initial capital
Position			
8	PnL_t	$(P_t - \bar{P}_{tentry}) / \bar{P}_{tentry}$	whether the open positions are in profit or loss
Time			
9	τ_t	$t / T \in [0, 1]$	how much time remains before the trading episode ends

3.3.5 Feature Redundancy and Excluded Information

The state was not designed by adding every potentially useful market indicator. Features were included only where they provide information that cannot already be recovered from the other variables in the state.

This is important because adding a feature that simply repeats existing information increases the size of the input without providing the agent with anything new.

The following features were excluded:

1. Running portfolio return

For example, the portfolio's current return relative to its starting capital does not need to be included as a separate feature because it can already be calculated from the normalised cash and asset exposure as $c_t + v_t$. Adding it would duplicate information already available in the state.

2. Raw position size

The raw number of units held, h_t , is also not included in the final state. The financial value of the position is more directly represented by:

$$v_t = \frac{h_t P_t}{C_0}$$

This provides information about the actual exposure rather than only the number of units.

3. Raw asset price

The raw asset price is also not required as a separate feature. The state instead uses returns and normalised values, which describe price movements without depending on the standard price level.

The general principle is therefore to include features that provide useful information while avoiding variables that can already be reconstructed from information available elsewhere in the state.

3.3.6 The Risk-Aware Extension: Drawdown State

The main state representation contains the information needed for the change in wealth reward. However, it does not describe the history of the portfolio's losses: current wealth alone does not show how the portfolio reached its current value. For example, two portfolios can have exactly the same wealth at the current timestep but may have followed very different paths:

- one portfolio may have increased steadily over time;
- another may have suffered a large loss before recovering to the same level.

To address this, an additional feature representing portfolio drawdown is introduced as a separate risk-aware extension:

$$dd_t = \frac{\max_{u \leq t} W_u - W_t}{\max_{u \leq t} W_u}$$

where dd_t represents the current decline in portfolio wealth from the highest wealth level reached so far during the episode.

A value of zero means that the portfolio is currently at its highest wealth level reached so far in the episode. A larger value indicates that the portfolio has fallen further below its previous peak.

For this reason, drawdown is treated separately from the main state representation. This makes it possible to evaluate whether explicitly providing information about portfolio loss history affects the agent's behaviour and performance.

The resulting risk-aware state is therefore:

$$s_t^{risk} \in R^{10}$$

where:

$$s_t^{risk} = [s_t, dd_t]$$

Since the original state s_t contains nine features, adding drawdown produces a **ten-dimensional** risk-aware state.

3.3.7 Feature Calculation and Information Timing

The timing of feature calculation matters because the agent should receive only the information available when making its trading decision. Using future information would give the agent an unfair advantage and could lead to overly optimistic results.

For this reason, the market features are calculated using information available at or before timestep (t). The current price and past prices can be used to create the state, but the agent must not use future prices when it decides its action (a_t).

The same principle applies to the features that use historical windows: momentum uses the previous ($k = 5$) timesteps, the moving average the previous ($m = 20$) prices, and the normalised trading range the ($n = 14$)-day ATR, as reported in Table 3.3.

Therefore, each state is based only on information that is available at the time the decision is made. This prevents future market information from entering the agent's observation and helps avoid **look-ahead bias**, where a model appears to perform well because it has indirectly been given information that would not have been available during actual trading.

3.3.8 State Representation Limitations

Although the final state representation gives the agent a wider range of information, the trading environment is still a simplified representation of a real financial market.

The agent does not have access to several types of information that can influence real-world trading decisions, including:

- order flows;
- market sentiment;
- financial news;
- macroeconomic conditions;
- relationships between the asset and other financial assets.
- The actions and intentions of large financial institutions.

These limitations mean that the state should not be interpreted as a complete description of the real financial market. Because this relevant information cannot be represented in the environment, it is better to describe this as a *partially observable Markov decision process (POMDP)* rather than a fully observable one.

However, this dissertation does not aim to replicate every aspect of real-world financial markets. The purpose is to compare and evaluate reinforcement learning algorithms within a controlled trading environment for an asset.

The proposed state representation therefore provides sufficient information for the defined task and can be **considered complete with respect to the information used by the defined trading environment**, which includes recent market behaviour, the current portfolio position, the performance of its current position, and the amount of time remaining in the trading period, while acknowledging that the model does not capture all factors present in real-world trading.

3.3.9 Summary

The final state was not selected because it contained a particular number of variables, but because its components collectively describe the market conditions, portfolio exposure, position performance, and remaining trading episode available within the defined environment. Table 3.4 summarises the design progression and the limitation each stage resolved.

Table 3.4: The staged development of the state representation. Each stage resolves the limitation named in the row above it. ΔP_t is the one-step return, c_t cash and v_t exposure as fractions of initial capital, h_t the units held, τ_t the episode clock and dd_t the fall from peak wealth.

Stage	State	Limitation addressed by the next stage
S_0	$[\Delta P_t]$	The agent can observe recent price movement but has no information about its own portfolio. It cannot tell whether it already holds the asset, or whether it has enough cash to open a position.
S_1	$[\Delta P_t, c_t, h_t]$	The agent now knows its available cash and the number of units it holds but not what they are worth. For example, ten units could represent a small or large investment depending on the asset price.
S_2	$[\Delta P_t, c_t, v_t, PnL_t]$	The agent can now determine the value of its position and tell whether its in profit or loss, but it does not know where it stands in the month. A position opened at the beginning of the month can therefore appear identical to one made immediately before the required closing of the position at the end of the month.
S_3	$[\Delta P_t, c_t, v_t, PnL_t, \tau_t]$	The agent's own circumstances are now fully represented, but its view of the market is still limited to the most recent price return. It cannot tell a stable trend from a reversal, or between relatively stable and highly volatile conditions.

Stage	State	Limitation addressed by the next stage
S_4	adds momentum, volatility, distance from trend, and daily range (Table 3.3)	The state provides a broader representation of the market while retaining the portfolios information. It is treated as the main state used in the experiments.
S_{risk}	adds dd_t	Adds information about the portfolio's previous losses from its highest wealth level. Required when the experiment considers risk history

The separate tenth feature, drawdown, is used only for the risk-aware experiments described in Section 3.3.6.

The main state representation provides the basis for defining the available actions and applying the reinforcement learning algorithms described in the following sections.

3.4 Action Space Design

3.4.1 Overview

The action space defines the choices available to the trading agent at each timestep. Its design matters because it determines how the agent interacts with the market and what its learned policy means.

A poorly designed action space can introduce unnecessary complexity or create unrealistic trading behaviour. For example, allowing the agent to choose any number of shares to buy or sell adds learning challenges: the agent must simultaneously learn when to trade, which direction to take, and how much to trade. This can make it difficult to determine whether poor performance results from an ineffective trading strategy or from difficulty learning appropriate position sizes.

To keep the focus on the trading decisions made by the reinforcement learning algorithms, this study uses a **simplified discrete action space**. The agent decides on the action (trading direction), while the environment determines the amount traded.

This separates the trading problem into two decisions:

1. **When and in which direction to trade** — this is the trading policy being learned by the reinforcement learning algorithm.
2. **How much capital to trade** — this is fixed by the environment and treated as an experimental setting.

This approach keeps the action space simple and makes it easier to compare the learning algorithms on the same portfolio trading task. It also means that differences in performance can be interpreted mainly in terms of the agents' ability to learn trading decisions, rather than their ability to learn position sizing.

3.4.2 Trading Actions and Trade Size

The action space is defined as $A = \{0, 1, 2\}$, where action 0 **holds** the current position, action 1 **buys** a fixed amount, and action 2 **sells** a fixed amount. At each timestep, the agent selects one of these three actions using the information available in the current state.

The **hold** action leaves the portfolio unchanged. No transaction takes place: the available cash and the number of units held remain the same, and no transaction fee is charged. This allows the agent to remain in its current position when it does not consider buying or selling to be appropriate.

The **buy** and **sell** actions change the agent's exposure to the asset by a fixed monetary amount:

$$\Delta = f C_0$$

where C_0 is the initial capital and f is the fraction of that capital allocated to one transaction. A fixed monetary amount is used rather than a fixed number of shares because the cost of one share depends heavily on the price level: buying one share of an asset priced at £10 is a small investment, while buying one share priced at £1,000 is a large one, even though both would be described as buying one share. Fixing the monetary size keeps every trade at the same proportion of the initial capital throughout the dataset, regardless of the asset price.

The fraction f is treated as an experimental setting rather than a fixed design constant. It is selected using the validation data before the final experiments, as reported in Section 5.3.

The number of units bought or sold at price P_t is:

$$\Delta h_t = \pm \frac{\Delta}{P_t}$$

with the positive sign for a buy and the negative sign for a sell, and with fractional ownership allowed, similar to a real-world portfolio. A sell is valid only when the agent holds enough units for the transaction. Because the trade size is fixed, the agent can build or reduce its position gradually over several decisions while still operating within a simple three-action framework.

3.4.3 Action Constraints

Although the agent has three possible actions, some actions may not be possible at a particular timestep because of the current portfolio position.

For example:

- the agent cannot buy if it does not have enough available cash;
- the agent cannot sell if it does not hold enough of the asset to withdraw.

To handle these situations, the environment uses a **legal action mask**. The mask does not provide any additional market information to the agent. It simply indicates which actions can be carried out given the current portfolio status.

The mask is represented as $M_t = [m_{\text{hold}}, m_{\text{buy}}, m_{\text{sell}}] \in \{0, 1\}^3$, where $M_t(a) = 1$ indicates that action a is available and $M_t(a) = 0$ indicates that it is not.

The availability of each action is determined from the portfolio state:

- available cash determines whether a **buy** action can be carried out;
- the current asset holding determines whether a **sell** action can be carried out;
- **hold** remains available throughout the episode.

Because the mask is derived directly from information already in the portfolio state, it is not treated as an additional feature of the state representation. Instead, it acts as a constraint on the agent, preventing it from selecting actions the environment cannot execute.

3.4.4 Transaction Costs

A transaction fee of $c = 0.05\%$ is charged whenever the agent carries out a buy or sell action, so a transaction of value Δ incurs a fee of Δc . The same fee calculation is applied to both purchases and sales.

Including transaction costs matters because it prevents the agent from overly buying and selling repeatedly. Without these fees, the agent could benefit from making many trades even when price movements are very small, making frequent trading appear more profitable than it would be in practice.

The inclusion of fees therefore introduces a basic cost of trading into the environment and provides a more realistic basis for evaluating learned trading strategies.

3.4.5 Relationship Between the State and Actions

The action space was designed to match the state representation of Section 3.3. Deciding whether to buy or sell depends on the agent's current exposure, its available cash, whether an existing position is in profit or loss, the time remaining in the episode, and the current market conditions — exactly the information the state provides. The agent decides whether to hold, buy, or sell, while the environment determines the amount traded and prevents invalid actions from being carried out.

3.5 Reward Function Design

3.5.1 Purpose of the Reward Function

The reward function tells the reinforcement learning algorithm how well it performed after taking an action and guides the agent toward more desirable actions. Unlike supervised learning, where the model is given the correct answer, the agent learns by taking actions and receiving feedback on the results of those actions.

In a trading environment, the reward function needs to accurately reflect the actual objective of the investment strategy. If it is poorly designed, the agent may learn behaviours that are not desirable, such as trading too frequently, taking unnecessary risks, or performing well during training but poorly under realistic conditions.

The main goal of this dissertation is to investigate whether reinforcement learning algorithms can learn profitable trading strategies for the selected asset while accounting for transaction costs. For this reason, the primary reward is based on the **change in portfolio wealth** from one timestep (trading day) to the next.

3.5.2 Wealth-Based Reward

The main reward used in this study is based on the change in the portfolio's total value:

$$r_t = W_{t+1} - W_t$$

where portfolio wealth at timestep t is given by:

$$W_t = C_t + h_t P_t$$

Here:

- C_t represents the cash available in the portfolio;
- h_t represents the number of asset units held;
- P_t represents the current asset price.

The reward therefore measures how much the portfolio's total value changes from one timestep (trading day) to the next after the agent's action.

A **positive reward** means the portfolio increased in value, while a **negative reward** means it decreased. This gives the agent a direct measure of the financial outcome of its decisions.

3.5.3 Including Transaction Costs in the Reward

Transaction fees are included when calculating the reward so that the result reflects the actual cost of carrying out a trade.

When a buy or sell action is taken at timestep t , the portfolio value is updated after the transaction fee is deducted:

$$W_{t+1} = C_{t+1} + h_{t+1} P_{t+1} - Fees$$

The reward is then calculated from the resulting change in portfolio wealth. This means that the agent is rewarded or penalised based on the value of the portfolio **after trading costs have been taken into account**.

3.5.4 Wealth Change As the Primary Reward

The change in portfolio wealth was chosen as the main reward for several reasons.

First, it directly reflects the main objective of the trading strategy: **increasing the value of the portfolio**. This also makes the reward closely related to the metric used to evaluate investment performance in the real world.

Second, it satisfies the requirements of the MDP defined in Section 3.2 because the reward can be obtained from portfolio information in the state representation.

The current portfolio wealth can be recovered from the cash and exposure features in the state representation:

$$c_t + v_t = \frac{W_t}{C_0}$$

Therefore, the change in wealth can be written as:

$$W_t = C_0 [(c_t + v_t)]$$

$$r_t = C_0 [(c_{t+1} + v_{t+1}) - (c_t + v_t)]$$

This shows that the reward depends only on the observed transition (s_t, a_t, s_{t+1}) , without requiring additional information unavailable to the agent.

3.5.5 Risk-Aware Reward Extension

Although the main objective is to increase portfolio wealth, focusing only on returns may encourage the agent to take more risk than desired. For example, a strategy may achieve a high final return but experience large losses along the way, which is not a desirable trading behaviour.

Therefore, a risk-adjusted reward that accounts for changes in portfolio drawdown is considered as an additional experimental condition:

$$r_t^{\text{risk}} = \Delta W_t - \lambda C_0 \Delta dd_t$$

where:

- ΔW_t is the change in portfolio wealth;
- Δdd_t is the increase in portfolio drawdown at step t (and is zero when drawdown does not rise);
- C_0 is the initial capital;
- λ controls the penalty applied to increases in drawdown.

Drawdown measures how far the current portfolio wealth has fallen from the highest level reached so far:

$$dd_t = \frac{\max_{u \leq t} W_u - W_t}{\max_{u \leq t} W_u}$$

where W_t is the portfolio wealth at time t , and W_u is the wealth at any earlier time u up to t . $\max_{u \leq t} W_u$ is therefore the highest wealth reached so far in the episode.

Δdd_t is therefore the increase in drawdown from one trading day to the next. Multiplying it by the initial capital C_0 keeps the reward in dollars, so a penalty of $\lambda = 0$ lets the wealth-change reward remain the same, while a larger λ discourages actions that worsen the drawdown.

However, this reward requires access to the portfolio's previous performance because the previous highest wealth reached up to time t must be known in order to calculate drawdown.

For this reason, the risk-aware reward is used only with the **risk-aware state extension** described in Section 3.3.6. This ensures that the information needed to calculate the reward is available to the agent.

The wealth-change reward remains the main reward used in the baseline experiments. The risk-aware version is treated separately so that the effect of explicitly including information about portfolio loss history can be examined.

3.5.6 Limitations of the Reward Function

Although the wealth-change reward provides a clear and direct measure of trading performance, it does not capture every aspect of a desirable trading strategy.

For Example, it does not directly penalise:

- large falls in portfolio value;
- unstable performance in returns;
- high levels of volatility.

A strategy could therefore achieve a strong overall return while experiencing substantial losses during the trading period. The limitations are addressed by the risk-aware reward in Section 3.5.5 by adding a penalty when drawdown increases.

However, adding more objectives changes the learning problem and what the agent is being trained to optimise. This can make comparisons between algorithms less straightforward because improvements may result from the different reward definition rather than from the learning algorithm itself.

Therefore, the **wealth-change reward is used as the common baseline**, while the risk-aware reward is considered separately as an additional experiment.

3.6 Objective Function

The objective function describes what the reinforcement learning algorithms are trying to achieve during training. In this study, the aim is to learn a trading policy that produces the highest possible expected return over each trading episode.

The reward function introduced in Section 3.5 measures the result of an individual decision, while the objective function extends this idea by combining the rewards received throughout the episode.

3.6.1 Episode Return

An episode/trajectory can be represented by the sequence:

$$\tau = (s_0, a_0, r_0, s_1, a_1, r_1, \dots, s_T)$$

where s_t is the state observed at timestep t ,

a_t is the action selected by the agent, and

r_t is the reward resulting from that action.

The episode contains T trading decisions and ends when it reaches the final market observation. Any position that remains open at this point is then closed according to the rules of the trading environment.

The total discounted return from the beginning of the episode is:

$$G_0 = \sum_{t=0}^{T-1} \gamma^t r_t,$$

Where $\gamma \in (0, 1]$ is the discount factor.

The same calculation can be made from any point within the episode. The return from timestep t , commonly referred to as the **reward-to-go**, is:

$$G_t = \sum_{t'=t}^{T-1} \gamma^{t'-t} r_{t'}.$$

This represents the rewards expected from the current timestep onwards. It is particularly relevant to the REINFORCE implementation because it associates an action with the rewards that follow it, rather than assigning the episode's total return to every action taken in that episode.

A discount factor of $\gamma = 0.99$ is used throughout the experiments. The trading episodes contain approximately 18–23 daily observations, so this value results in relatively mild discounting over the length of an episode: $0.99^{20} \approx 0.82$, so a reward received near the end of the monthly episode still contributes around 82% of its original value to the discounted return.

Using $\gamma = 0.99$, therefore, maintains the standard reinforcement-learning formulation for discounted returns while still ensuring the objective is maximising end-of-month wealth.

3.6.2 Expected Return

The outcome of a single trading episode is not enough to judge the quality of a policy because each episode can reflect different market conditions experienced during that month. For instance, a policy that performs well in one month may not perform equally well in another. Therefore, the objective is defined using the expected returns the policy earns across different trading episodes.

The policy is represented by

$$\pi_{\theta}(a|s),$$

where θ represents the trainable parameters of the policy network. For a given state, the policy assigns probabilities to the available actions, allowing the agent to choose between Buy, Hold, and Sell.

$$\pi_{\theta}(a_t | s_t) = \begin{bmatrix} P(a_t = \text{Hold} | s_t) \\ P(a_t = \text{Buy} | s_t) \\ P(a_t = \text{Sell} | s_t) \end{bmatrix}$$

The expected return is defined as:

$$J(\theta) = \mathbb{E}_{\tau \sim p_{\theta}(\tau)} [G_0(\tau)]$$

The learning objective is therefore:

$$\theta^* = \arg \max_{\theta} J(\theta).$$

In simple terms, the aim is not for the agent to find a policy that performs well in only one trading period. Instead, the agent should learn policy parameters that produce good, expected performance across the different monthly market episodes used during training.

The policy parameters affect the objective only through the actions the agent selects: changing θ changes the probabilities of choosing Buy, Hold, or Sell in each state, and these decisions affect the portfolio and, in turn, the rewards collected over the episode.

3.6.3 Objective and the Trading Environment

The formulation above also separates the trading environment's role from the learning algorithm's. The environment determines the market data, the current portfolio state, the transaction costs, which actions are currently valid, the resulting portfolio wealth, the reward, and whether the episode has ended. The learning algorithm determines only how the agent updates its policy or value function using that experience.

This separation is important for comparing the different learning approaches in this dissertation. **REINFORCE**, **DQN**, and **PPO** use the same state representation, action space and masking, market data, and reward function, so differences in their results stem from how they learn rather than from the conditions they operate under.

3.7 REINFORCE: Monte-Carlo Policy Gradient

REINFORCE is the first policy-gradient method used in this study. Unlike value-based methods, which estimate how useful different actions are in a given state, **REINFORCE directly learns a policy** that gives probabilities to the available actions.

It provides a useful baseline because its main idea is straightforward. The agent selects actions according to its current policy, observes the resulting rewards, and then adjusts the policy so actions can lead to better returns in the future.

3.7.1 Policy Representation

The policy is represented by a neural network with trainable parameters θ . Given the current state s_t , the network produces one output, or logit, for each of the three possible actions {Buy, Hold, Sell}. Actions that are not permitted at the current timestep using the action mask described in Section 3.4 are removed. The remaining logits are then passed through a SoftMax function to produce the policy probabilities $\pi_{\theta}(a_t|s_t)$.

During training, the agent samples an action from this probability distribution. This lets it try different actions and explore alternative trading decisions rather than always choosing the action with the highest probability.

During evaluation, the agent selects the valid action with the highest probability without exploring.

The policy network used in the real-data study has **two hidden layers, each with 128 units**, and uses the hyperbolic tangent tanh as the activation function. The final layer produces one output for each of the three available actions. The simulated-data study keeps the same structure but narrows the layers to 64 and 32 units, for reasons of model capacity discussed in Chapter 4.

3.7.2 Deriving the Policy-Gradient Objective

The objective introduced in Section 3.6 is an expected return over trading episodes:

$$J(\theta) = \mathbb{E}_{\tau \sim p_\theta(\tau)} [G_0(\tau)].$$

To improve the policy, the learning algorithm needs to determine how a small change to the policy parameters (θ) affects this expected return. This requires calculating the gradient:

$$\nabla_\theta J(\theta).$$

The difficulty is that the probability of a complete trading episode $p_\theta(\tau)$ depends on the initial state, the actions selected by the policy, and the resulting environment transitions $(s_0, a_0, r_1, s_1, a_1, r_2, \dots, s_T)$. It can be written as:

$$p_\theta(\tau) = \rho(s_0) \prod_{t=0}^{T-1} \pi_\theta(a_t | s_t) P(s_{t+1} | s_t, a_t),$$

where $\rho(s_0)$ represents the probability distribution of initial states and

$P(s_{t+1} | s_t, a_t)$, represents the environment's transition to the next state s_{t+1} given the current state and action (s_t, a_t) .

The important point is that the policy parameters (θ) affect the action probabilities $\pi_\theta(a_t | s_t)$ but not the market transition process. REINFORCE therefore does not need to model this market transition explicitly, just the action probabilities.

So, we use the likelihood-ratio identity, which provides a way of expressing the gradient $\nabla_\theta J(\theta)$ using the probability of the sampled episodes $p_\theta(\tau)$:

$$\nabla_\theta p_\theta(\tau) = p_\theta(\tau) \nabla_\theta \log p_\theta(\tau).$$

The policy gradient can then be expressed as:

$$\nabla_\theta J(\theta) = \mathbb{E}_{\tau \sim p_\theta} [\nabla_\theta \log p_\theta(\tau) G_0(\tau)].$$

Since only the policy terms depend on (θ), this simplifies to

$$\nabla_\theta J(\theta) = \mathbb{E}_{\tau \sim p_\theta} \left[\sum_{t=0}^{T-1} \nabla_\theta \log \pi_\theta(a_t | s_t) G(\tau) \right].$$

This is the basic policy-gradient expression REINFORCE uses in this study.

In simple terms, the equation adjusts the probability of an action based on the outcome that follows it. Actions resulting in higher returns are encouraged by increasing their probability, while actions associated with lower returns are made less likely.

3.7.3 Reward-to-Go

Using the complete episode return $G(\tau)$ for every action is inefficient because an action taken at time t cannot influence rewards already received before that time. REINFORCE therefore weights each action by the reward-to-go G_t defined in Section 3.6, so only rewards from time t onwards are used when evaluating the outcome associated with that decision.

The policy gradient can therefore be written as:

$$\nabla_{\theta} J(\theta) = \mathbb{E} \left[\sum_{t=0}^{T-1} \nabla_{\theta} \log \pi_{\theta}(a_t | s_t) G_t \right].$$

where $\nabla_{\theta} \log \pi_{\theta}(a_t | s_t)$ is the direction in which the policy's action probability changes, and G_t is the return obtained after taking action a_t . This links each decision directly to the rewards that followed it.

3.7.4 Monte-Carlo Estimation

The expectation $\mathbb{E}$ in the policy-gradient equation cannot normally be calculated because the full distribution of possible market trajectories is unknown. REINFORCE therefore uses a Monte Carlo approach: the agent interacts with the trading environment and uses the resulting episodes to estimate the gradient.

For N sampled episodes, the gradient can be approximated as:

$$\nabla_{\theta} J(\theta) \approx \frac{1}{N} \sum_{n=1}^N \sum_{t=0}^{T_n-1} \nabla_{\theta} \log \pi_{\theta}(a_t^{(n)} | s_t^{(n)}) G_t^{(n)}.$$

The superscript n represents the number of the sampled episode.

The gradient is therefore estimated from the average experience collected from the sampled episodes.

This is why REINFORCE is described as a **Monte-Carlo policy-gradient method**: The method does not require the agent to know in advance how the market will move or to build a separate model of future price changes. Instead, it learns how to estimate the expected policy gradient from complete sampled episodes produced by its interaction with the environment.

3.7.5 Loss Function Used for Training

Most neural-network frameworks are designed to minimise a loss function, while reinforcement learning is to maximise expected rewards $J(\theta)$. The policy-gradient objective can therefore be written as a loss

by introducing a negative sign:

$$L(\theta) = -\frac{1}{N} \sum_{n=1}^N \sum_{t=0}^{T_n-1} \log \pi_{\theta} \left(a_t^{(n)} | s_t^{(n)} \right) G_t^{(n)}.$$

The negative sign changes the original maximisation problem into a minimisation problem that can be handled by the PyTorch optimiser.

The logarithm appears as a direct result of the likelihood-ratio derivation above; it is not an additional trading assumption.

The implementation uses the **Adam optimiser** [34] with a learning rate of 10^{-3} .

During implementation, the sampled returns may also be normalised before being used in the update. This helps control the gradient magnitude across different training batches and can improve numerical stability.

3.7.6 Strengths and Limitations

REINFORCE is a suitable baseline for this project because its formulation is relatively straightforward: it directly learns the policy and does not require a separate model of the market, so performance differences against more advanced algorithms are easier to interpret.

Its main limitation is that **its gradient estimates can have high variance** because they depend completely on episode returns, and different episodes can produce very different outcomes even in similar trading situations. This can make learning less stable, particularly when training data is limited. Reward-to-go and standardised returns reduce some of this variation, but they do not remove the algorithm's underlying high-variance nature.

For this reason, this provides motivation for comparing REINFORCE with the other methods, **DQN and PPO**, to see whether they can achieve a more stable or effective learning in the same trading environment.

3.8 Deep Q-Learning

Deep Q-learning provides a **value-based alternative** to REINFORCE. Rather than directly learning the probability of selecting each action, it estimates the value of each possible action in the current state. This comparison is useful because both methods address the same portfolio trading problem but learn in different ways.

3.8.1 Action-Value Function

The action-value function, commonly referred to as the **Q-function**, is defined as $Q^{\pi}(s, a)$:

$$Q^{\pi}(s, a) = \mathbb{E}_{\pi} \left[\sum_{k=0}^{\infty} \gamma^k r_{t+k} \right].$$

This represents the expected future return from taking action a in state s and then continuing to follow policy π .

The optimal action-value function is written as $Q^*(s_t, a_t)$. It represents the highest expected return that can be achieved from a given state when the best available actions are selected

For the optimal action-value function, the Bellman optimality equation [9] is:

$$Q^*(s_t, a_t) = \mathbb{E} \left[r_t + \gamma \max_{a'} Q^*(s_{t+1}, a') \right].$$

In simple terms, the value of an action depends on two things:

1. the reward received immediately after taking the action; and
2. the value of the opportunities available afterwards.

This makes the Q-function a useful representation for the portfolio trading problem because the agent has a small action space.

In practice, the optimal Q-function is approximated by a neural network. Once the values of the available actions have been estimated, the agent selects the action with the highest Q-value:

$$a_t^* = \arg \max_{a \in \mathcal{A}_{\text{legal}}(s_t)} Q^*(s_t, a).$$

Only legal actions are considered when making this choice. This is important in the trading environment because, for example, the agent cannot buy when it does not have enough cash or sell when it does not hold enough of the asset.

3.8.2 From Q-Learning to Deep Q-Learning

Traditional tabular Q-learning stores a separate value for every state-action pair in a table. This approach is not suitable for the present portfolio trading problem here because the state contains continuous variables and several different market and portfolio features.

For example, volatility and portfolio exposure can take many different values. Converting all nine state features into discrete categories would first require discretising the continuous variables and then creating a very large state space that would also lose information by grouping different values together.

Deep Q-learning avoids this problem by replacing the table with a neural network $Q_\phi(s, a)$, where ϕ represents the trainable parameters of the network.

The network takes the nine-dimensional state $s_t \in \mathbb{R}^9$ as input and estimates the Q-value for each available action. These values are then used to select the most valuable valid action, as in Section 3.8.1.

The DQN used in the real-data experiments has **two hidden layers with 128 units each** and uses ReLU activation functions. As with the policy network, the simulated-data study narrows the layers to 64 and 32 units.

3.8.3 Experience Replay

Successive decisions within a trading episode are naturally related because they come from consecutive trading days. If the network were trained directly on these transitions in their original order, it could produce highly correlated consecutive updates.

DQN addresses this issue using an **experience replay buffer** [13, 6]. Each interaction with the environment is stored as a transition in a memory buffer:

$$(s_t, a_t, r_t, s_{t+1}, M_{t+1}, d_t),$$

where:

- s_t is the current state;
- a_t is the action selected by the agent;
- r_t is the reward received;
- s_{t+1} is the next state;
- M_{t+1} is the action mask for the next state;
- d_t indicates whether the episode has ended.

During training, the agent samples random minibatches of previously stored transitions from the replay buffer so the network does not rely only on the most recent transition.

The replay buffer has a capacity of **50,000 transitions**, with **64 transitions sampled per minibatch**.

This lets the network learn from reusable experiences collected at different points during training and reduces the dependence between consecutive training samples.

3.8.4 Target Network

Using the same network to calculate both the current Q-values and the target values can make training unstable. This is because the network parameters change during training, which also causes the target used for learning to change.

To reduce this problem, DQN uses a second network known as the **target network** [6], $Q_{\phi^-}(s, a)$.

The target network is kept separate from the main, or online, Q-network. Rather than updating it after every gradient optimisation step, its parameters are periodically copied from the main Q-network.

Keeping a fixed number of updates for the target provides a more stable reference for calculating the temporal-difference target and reduces how rapidly the learning target changes while the main network is being updated.

In this implementation, the target network is synchronised with the online network every **500 training steps**.

3.8.5 Temporal-Difference Learning

Unlike REINFORCE, which waits until the end of an episode to calculate the return, DQN updates its estimates using a one-step temporal-difference target.

For a transition, (s_t, a_t, r_t, s_{t+1}) , the target Q-value y_t is calculated from the immediate reward and the highest estimated value of the best valid action in the next state:

$$y_t = r_t + \gamma \max_{a' \in \mathcal{A}_{\text{legal}}(s_{t+1})} Q_{\phi^-}(s_{t+1}, a').$$

Action masking is applied here because the agent should not assign value to actions it cannot take in the next state.

When the episode ends, there are no future trading decisions or rewards to consider. The target must therefore account for this condition:

$$y_t = r_t + \gamma(1 - d_t) \max_{a' \in \mathcal{A}_{\text{legal}}(s_{t+1})} Q_{\phi^-}(s_{t+1}, a').$$

Where d_t indicates whether the episode has ended.

When $d_t = 1$, the second term $(1 - d_t)$ becomes zero, leaving $y_t = r_t$.

This is particularly important in the trading environment because any remaining position is closed at the end of each monthly episode. The agent cannot take another trading action after this point. Including a future Q-value beyond this state would incorrectly treat the episode as if further decisions were possible.

3.8.6 DQN Loss

The online Q-network is trained by reducing the difference between its predicted Q-value and the temporal-difference target.

For a minibatch of B transitions, the loss is defined as:

$$L(\phi) = \frac{1}{B} \sum_{b=1}^B [Q_{\phi}(s_b, a_b) - y_b]^2$$

This is the **mean squared error (MSE)** between the network's predicted value and the target value.

The main difference from REINFORCE is how it obtains the learning signal. REINFORCE updates the policy using returns calculated from sampled episodes, whereas DQN updates its Q-values using one-step targets that include an estimate of the future value.

3.8.7 Exploration

If the agent always selected the action with the highest current Q-value, it would follow a purely greedy strategy. This can be problematic early in training because the Q-values are initially poor estimates, and the agent may settle on actions before exploring other available alternatives.

To encourage exploration, the DQN uses an ϵ – *greedy* strategy

With probability $1 - \epsilon$, the agent selects the valid action with the highest estimated value, and with probability ϵ it instead selects an action randomly from the set of valid actions.

The exploration rate starts at $\epsilon = 1.0$ and is gradually reduced to $\epsilon = 0.05$ over the first half of training.

This means that early in training, the agent explores the available actions frequently, while later, as training progresses, it increasingly relies on the most valuable Q-values it has learned.

Also, the agent considers only valid actions when exploring. It cannot select an invalid action simply because it is in its exploratory phase.

3.8.8 Strengths and Limitations

DQN has several advantages for this problem. Experience replay lets the agent reuse transitions collected at different stages of training, the target network stabilises learning, and the separate value estimates make it possible to directly compare the expected value of Buy, Hold, and Sell.

Its limitations are the mirror of these strengths. Training can be affected by the scale of the reward and target values, unlike the REINFORCE implementation, where returns are standardised within each episode. DQN also uses bootstrapping: the target partly relies on Q-value estimates given by the network rather than observed rewards, so inaccurate estimates can carry errors into later updates.

These differences make DQN a useful comparison to the policy-gradient algorithms used in this study rather than using another implementation of the same learning architecture.

3.9 PPO: Proximal Policy Optimisation

Proximal Policy Optimisation (PPO) [7] is used as the third learning method in this study to provide a more stable policy-gradient algorithm that can be compared with REINFORCE.

REINFORCE updates the policy using the returns obtained from sampled episodes. Although this approach is straightforward, the policy can change considerably between updates, specifically when the sampled returns are large, making training unstable.

PPO reduces this problem by limiting how much the policy changes during each update. It does this by comparing the new policy with the policy from the training data.

3.9.1 Actor-Critic Structure & Advantage Function

PPO uses two neural networks: an **actor** and a **critic**. The actor represents the portfolio trading policy $\pi_{\theta}(a|s)$, where θ represents the parameters of the policy network, and the critic estimates the value of the current state, $V_{\phi}(s)$, where ϕ represents the parameters of the value network.

The actor selects actions, while the critic estimates the expected future return from a given state. The critic therefore provides a reference to compare the outcome of a policy's action against.

This lets PPO use an **advantage** value rather than relying directly on the return of a full episode, like in REINFORCE.

The advantage at time t can be represented as

$$A_t = G_t - V_{\phi}(s_t),$$

The advantage shows whether the selected action performed better or worse than expected for the current state. A positive advantage means the outcome was better than expected for the current state, while a negative advantage means it was worse than expected.

For instance, if the expected return from a state $V_{\phi}(s_t)$ is relatively low but the selected action a_t produces a much higher return G_t , the advantage will be positive. The policy $\pi_{\theta}(a|s)$ should therefore increase the likelihood of taking that action in similar situations.

The implementation uses **Generalised Advantage Estimation (GAE)** to calculate the advantages. Using an advantage estimate helps reduce the variance of policy-gradient updates.

3.9.2 Policy Ratio

PPO measures how much the policy has changed by comparing the new policy with the policy that generated the sampled data.

The probability ratio ρ_t is defined as:

$$\rho_t(\theta) = \frac{\pi_{\theta}(a_t|s_t)}{\pi_{\theta_{old}}(a_t|s_t)}.$$

ρ_t is used here to differentiate the policy ratio from the trading reward function r_t .

If $\rho_t(\theta) = 1$, the probability of the selected action has not changed from the previous policy.

A ratio greater than 1 means the new policy gives the selected action a higher probability than the previous policy, while a ratio below 1 means its probability has decreased.

Without restrictions, a single update could change the policy substantially. PPO uses the ratio to decide whether a policy update has moved too far from the policy that generated the sampled experience, and it limits that ratio from being used in the policy objective.

3.9.3 Clipped Objective

PPO controls policy updates using a clipped objective:

$$L^{CLIP}(\theta) = \mathbb{E}_t [\min(\rho_t(\theta)A_t, \text{clip}(\rho_t(\theta), 1 - \epsilon, 1 + \epsilon)A_t)].$$

This project uses the standard implementation, $\epsilon = 0.2$, which gives a clipping range that limits the policy ratio to approximately $[0.8, 1.2]$.

Clipping does not prevent the policy from moving outside this range or changing by more than 20%. Instead, once the probability ratio moves beyond the specified range, the objective limits improvement in that direction to prevent the update from becoming even larger.

For example, if an action has a positive advantage, increasing its probability can improve the objective. However, once the change becomes sufficiently large, the clipped objective prevents the update from gaining additional benefit that increases that probability further.

In this way, PPO improves the policy while avoiding large changes between continual updates. This provides a smaller, more controlled learning process than directly updating the policy from sampled returns, like in REINFORCE.

3.9.4 Action Masking in PPO

The action masking described in Section 3.4 also applies to PPO [20]. The agent should not be able to select an action that the trading environment cannot execute.

For this reason, the implementation uses **MaskablePPO** instead of the standard PPO implementation. The action mask is applied after the policy determines the probability of each action, so it excludes actions that are not currently possible.

For example, when the available cash is less than the amount required to open an asset position, the Buy action is unavailable. Likewise, the Sell action is unavailable when the agent does not hold enough of the asset. All three algorithms therefore face the same mask described in Section 3.4.

3.9.5 PPO in the Experiment

PPO is included to provide a third comparison with REINFORCE and DQN. It is implemented using **MaskablePPO** from sb3-contrib [35]. The policy uses the same two-hidden-layer architecture as the other neural-network agents, 128 units per layer in the real-data study and 64 and 32 units in the simulated-data study, and the remaining PPO settings use the library defaults. PPO also uses the same trading environment, state representation, transaction costs, initial capital, data split, and evaluation procedure as

the other methods (Section 3.6.3), such that differences in performance mainly reflect how each algorithm learns.

3.10 Algorithms

The three learning procedures, as implemented, are listed in full as Algorithms 1–3 in Appendix B. REINFORCE and Deep Q-learning are built from scratch following the formulations in Sections 3.7 and 3.8; MaskablePPO is the library comparator of Section 3.9 and applies the same action mask. The next section states why all three methods are retained in the experimental comparison.

3.11 Comparison of the Three Learning Approaches

REINFORCE and DQN are the primary learning methods because they directly represent the two main reinforcement learning algorithms being investigated: policy-gradient learning and value-based learning.

PPO is included as an additional policy-gradient comparison to compare the performance of the basic Monte Carlo policy-gradient algorithm with a more stable policy-gradient algorithm.

Using all three algorithms provides a stronger comparison than relying on a single reinforcement learning algorithm. If similar trading behaviour is noticed across the three algorithms, it gives greater confidence that the behaviour is related to the trading environment or state representation rather than a specific algorithm. On the other hand, if the algorithms produce different results, these differences can help show how the choice of learning method affects the resulting trading strategy on the portfolio.

Chapter 4

System Design and Experiment Framework

This chapter explains how the mathematical formulations and design choices from Chapter 3 were implemented in the portfolio trading system using reinforcement learning. The state representation, action space, reward function, and learning algorithms described previously were converted into a working experimental system for evaluation.

4.1 Introduction

The system is a Python-based system, divided into separate components for market data preparation, the trading environment, the learning algorithms, experimental configurations, and the evaluation process. This structure allows easy modification of each component without affecting the rest of the system. More importantly, it allows the same trading environment to be used by different algorithms, REINFORCE, DQN, and PPO, with little difference in their experimental conditions.

REINFORCE and Deep Q-Learning (DQN) were implemented from scratch using PyTorch [36]. PPO was implemented using the MaskablePPO implementation from sb3-contrib [35]. The Gymnasium interface [37] simulates the trading environment and is used across all the learning algorithms.

The implementation provides two data pipelines. The first prepares real SPY market data. The second generates simulated market episodes under controlled conditions, and exists because the real data turned out to be too small to train the networks without memorisation. Both pipelines produce episodes in the same format, so the environment, the agents, and the evaluation code do not know which kind of data they are running on.

The implementation also addressed several practical issues stated in the methodology that the mathematical formulation could not capture. These include limiting the agent to only valid actions, ensuring proper separation between training and testing data, closing any remaining position at the end of each monthly episode, and ensuring the risk-aware states and rewards are used separately. This helps ensure that the experiment setup follows the constraints assumed in Chapter 3 to the letter. These details matter heavily because even correct mathematical formulations can produce unreliable results if implemented improperly.

4.2 Data Pipeline

4.2.1 Market Data Preparation

The first stage of the implementation converts historical market data into the observations the simulation trading environment needs. The experiment uses daily **OHLCV market data (open, high, low, close,**

and volume) of the SPDR S&P 500 ETF Trust (SPY) to construct features.

SPY was selected because it is a highly traded ETF that represents the broader US equity market. Using an index ETF rather than an individual company reduces the effect of company-specific events and provides a universal asset for the experiments.

Although most formulations use closing prices, the implementation keeps the open, high, low, close, and volume data. The high and low prices are necessary to calculate the range for the ATR feature used in the state representation.

The main dataset period covers **January 2018 to December 2025**. This period was stated in Chapter 3 to capture a mix of market conditions while keeping the main experiment within a significant historical period.

The data are obtained using the **yfinance interface** and are cached locally after processing. This ensures that later experiments use the same prepared dataset rather than retrieving and processing the data again.

4.2.2 Feature Construction & Continuous Calculation

The state features are calculated before separating the data into monthly episodes. This is necessary because some of the features depend on previous data.

For example, momentum uses a five-day window, while the moving average and ATR features use longer windows. Their calculations at the beginning of a month should still use data from previous months and should not restart simply because the environment starts a new monthly episode.

Calculating these features only in each monthly episode would mean that the first few days of every month can't access or use data from the previous month. This would artificially shorten that feature's historical window at the beginning of the experiment.

To avoid this, the pipeline first creates a continuous time series dataset and calculates all features that roll from the past months across the full dataset. It then creates the monthly episodes from the resulting feature data.

More observations are used before the start of the main experiment period to provide the historical data needed to calculate rolling features. For instance, the longest feature window, ATR, uses 20-day periods, so the pipeline includes 21 additional previous day bars before the first actual observation. These past data are used only for calculating the initial features and are not included in the main experiments.

For this reason, the implementation does not begin feature calculation exactly at the first observation of January 2018. Instead, additional market data is loaded before the experimental period begins. This warm-up data provides the previous observations necessary to calculate the first valid features with rolling window calculations.

The final pipeline uses a warm-up period beginning in September 2017. The data from this period are used to initialise the rolling calculations and are then removed before the experimental episodes are created.

This approach avoids using the dataset's first observations with incomplete rolling statistics and treats the first experiment's observation the same as later observations, with the required historical context already available.

The pipeline also checks the processed data for missing values before creating monthly episodes. If the data is not available, the data preparation process raises an error rather than silently producing incomplete features.

Overall, the process ensures that each feature at time t is calculated using information available at or before that time, while still allowing rolling features to use historical observations.

4.2.3 Dataset Split & Monthly Episodes

After feature construction, the continuous dataset is divided into monthly episodes by calendar month. Months containing fewer than 15 trading days are excluded from the dataset. After this filtering, the dataset is left with 96 monthly episodes.

The experimental data are separated into training, validation, and testing periods. The split is chronological rather than random because the observations form a time series and the aim is to apply the learned policies to current markets.

Figure 4.1 shows the split.




Figure 4.1: Chronological split of the 96 monthly episodes: 60 training months, 12 validation months, and 24 held-out test months.

The episodes are kept in chronological order within the splits and are not randomly shuffled between the three periods. The training set covers January 2018 to December 2022 (60 monthly episodes) and is used to learn the algorithms' parameters. The validation set covers 2023 (12 months) and is used for decisions made during development, such as the fixed trading amount. The test set contains the final 24 months, January 2024 to December 2025, and is kept separate from these decisions for only the final evaluation.

Maintaining a chronological split is important for financial data [29]. This prevents observations from later periods from entering the training set and influencing the model's development. This would not reflect how a trading strategy would be developed in practice.

The training period includes most market conditions, including the volatility experienced in 2018, the COVID-19 market crash and recovery in 2020, and the 2022 downtrend. The 2023 period is used for validation, while 2024 and 2025 form the final test period.

The separation between validation and test data is especially important for the experiments in Chapter 5. A configuration that performs well on the validation period may therefore be selected to examine its performance on the held-out test period. The test results are then used to compare the learned strategies.

The processed episodes and their metadata are cached locally so that the same dataset can be used consistently across all experiments.

4.3 Simulated Market Data

This section describes the second data pipeline. It generates simulated market episodes for the controlled study reported in Chapter 5, and it exists because the real dataset cannot supply enough training episodes for the learning algorithms to do more than memorise.

4.3.1 Why the Real Dataset Is Not Enough

The real training split contains **60 monthly episodes**. The network described in Chapter 3 has two hidden layers of 128 units, which is more than 18,000 trainable parameters for the nine-feature state. A model with 18,000 free parameters fitted to 60 examples can match every training month exactly without learning any rule that transfers to new months. The pilot study in Chapter 5 shows exactly this behaviour.

Sixty episodes is also far below what reinforcement learning normally uses. The Atari agents of Mnih et al. trained on millions of frames [6], and AlphaGo generated its own training games through self-play precisely because no collection of human games was large enough [24]. When the real data is too small, the standard remedy is to train on generated data and keep the real data for a final transfer test.

The real training years also lean upward. Of the 60 training months, 29 rose by more than 2% and only 16 fell by more than 2%. The held-out test years 2024–2025 rise almost throughout. An agent that always buys therefore scores well on both periods without reading its state once. This is a class-imbalance problem: the data rewards one constant answer, so the constant answer is what gets learned.

4.3.2 Generator Design

The generator produces monthly episodes under three **market regimes**: up, down, and flat. One episode is 21 daily bars drawn from a single regime, so every test result can be attributed to a known market condition.

Within a regime, daily log-returns are drawn independently from a normal distribution with that regime's drift and volatility. High and low prices extend each day's close-to-close move by random wicks whose size is fitted to the real average true range. Volume is drawn from a lognormal distribution; only relative volume enters the feature set, so its absolute scale does not matter.

Each generated episode passes through the same feature code as the real pipeline. A warm-up lead of 21 bars is generated before each episode and discarded after the rolling features are computed, exactly as the real pipeline discards its 2017 warm-up data.

The generator is deliberately simple. Richer market generators exist, from hidden Markov models [26] to adversarial time-series generators [27], and they produce more realistic price paths. Realism is not the goal here. The goal is control: the study needs to state how each agent behaves when the market rises, falls, or drifts sideways, and that requires knowing which condition each test episode belongs to.

4.3.3 Calibration

The generator's parameters come from the **real training months only**. Each of the 60 real training months is labelled by its return: up if it gained more than 2%, down if it lost more than 2%, and flat otherwise. The daily log-returns of the months in each class are pooled, and their mean and standard deviation become that regime's drift and volatility. The wick size is fitted so the generated average true range matches the real per-regime mean. Table 4.1 shows the fitted values.

Table 4.1: Generator parameters fitted on the 60 real training months.

Regime	Months	Daily drift	Daily volatility	Mean ATR/close
Up	29	+0.244%	0.967%	1.38%
Down	16	−0.358%	2.000%	1.97%
Flat	15	+0.012%	1.095%	1.25%

The fitted values behave as market data should. Down months carry twice the volatility of up months, which is the volatility clustering seen in real crashes. As a calibration check, the generated episodes reproduce the target average true range to within 4% in every regime. The real validation and test months are never read by the generator, so the transfer test in Chapter 5 stays untouched by any generation decision.

4.3.4 Simulated Splits

The simulated dataset contains **3,000 training episodes, 300 validation episodes, and 600 test episodes**, each split holding equal numbers of up, down, and flat episodes. The three splits are independent draws from the same generator under different random seeds.

This design fixes two problems of the real split at once. First, the volume problem: 3,000 episodes of 20 decisions each give 60,000 transitions per pass over the data, against 1,200 for the real training split. Second, the distribution problem: train and test are now drawn from the same distribution by construction, so a policy that fails on the test episodes cannot blame a regime shift between the periods.

The balance also removes the constant-answer shortcut. An agent that always buys earns well in the up third, loses more in the down third, and roughly breaks even in the flat third. Buying regardless of state now loses money on average, so a profitable policy must read its observations. The training procedure for this study is described in the next section alongside the real-data procedure.

4.4 Training, Validation and Testing Procedure

The experiments use a fixed sequence that separates training, validation, and test periods. This separation ensures that the test data are not used when making decisions about the model or experimental setup.

4.4.1 Training

During training, each learning algorithm interacts with the monthly episodes in the training period. Each algorithm is given the same market data, state representation, portfolio conditions, and trading constraints. The main difference between the algorithms is how they use the collected experience to update their parameters.

REINFORCE updates the policy using the reward-to-go calculated from completed monthly episodes. DQN updates its Q-network using transitions sampled from the replay buffer. PPO collects fixed-length rollouts and uses the resulting batches for several optimisation epochs.

Each configuration is trained using six random seeds. Using multiple seeds reduces the dependence on using one particular network initialisation or a sequence of sampled experiences and allows the variation between runs to be reported [30].

The trained model from each seed is saved with its configuration so that it can be evaluated later using the same settings.

4.4.2 Validation

The validation set contains the 12 monthly episodes from 2023.

This period is used for decisions that need to be made before the final test. For example, it determines the fixed transaction amount traded for each Buy or Sell action.

The candidate transaction sizes are 0.05, 0.10, 0.25, and 0.50 of the initial capital.

The selection is based on the maximum portfolio exposure that can be achieved from each action rather than on trading performance during the test period. This prevents the action from being chosen simply because it happens to perform well on the final evaluation data.

Once the required configuration decisions have been made, the selected settings are fixed for testing.

4.4.3 Testing

The final evaluation uses the held-out data from January 2024 to December 2025.

The test data are not used to make configuration decisions. This ensures that results provide an independent assessment of the methods.

Each trained model is evaluated in the same trading environment and using the same evaluation procedure, which includes the initial capital, transaction costs, action constraints, and monthly episode struc-

ture.

For the learned agents, actions are selected intentionally by choosing the highest-probability action for the policy-based methods or the highest-valued action for DQN.

The portfolio outcomes are then recorded for each monthly test episode and each random seed. This produces a consistent set of results that can be used to compare the different learning algorithms.

The stored results therefore contain sufficient information to compare algorithms, state representations, reward functions, and benchmark strategies, which are presented in Chapter 5.

4.5 Experiment Configuration

The implementation was designed so that different experiments could be run by changing their settings without changing the overall trading environment or learning algorithm settings.

For this reason, experiments are defined through configuration files or parameters rather than by manually changing the source code for each run. Each experiment is therefore linked to a specific configuration that records the conditions used to produce it.

This separates experiment settings from the environment's implementation and learning algorithms, allowing the same code to run different state, reward, and action configurations.

4.5.1 Configuration Parameters

The configuration records the main settings that can affect an experiment's outcome. These include:

- learning algorithm;
- state representation and state dimensionality;
- reward formulation;
- transaction fee;
- fixed transaction amounts;
- risk-penalty coefficient, where applicable;
- action model;
- random seed;
- training settings;
- evaluation settings.

Using this approach reduces the need to manually change the source code between experiments. It also makes the experimental process easier to track, since each result can be linked directly to the configuration that produced it.

This is particularly important for the experiments reported in Chapter 5, where several state, reward, and action configurations are compared.

4.5.2 Algorithm Configuration

The experiment uses the same starting capital and transaction cost across all three learning algorithms. The initial capital is set to $C_0 = \$10,000$, and a transaction fee of 0.05% is applied to each trade opened.

The main settings used for REINFORCE, DQN, and PPO are shown below.

Parameter	REINFORCE	DQN	PPO
Network	MLP with 2 hidden layers of 128 $d \rightarrow 128 \rightarrow 128 \rightarrow 3$	MLP with 2 hidden layers of 128 $d \rightarrow 128 \rightarrow 128 \rightarrow 3$	<i>MlpPolicy</i> $d \rightarrow 128 \rightarrow 128$
Activation	tanh	ReLU	
Optimiser	Adam	Adam	Adam
Learning rate	10^{-3}	10^{-3}	3×10^{-4}
Discount factor γ	0.99	0.99	0.99
Replay buffer	—	50,000	—
Batch size	Episode	64	64
Target network update	—	Every 500 steps	—
Exploration	Policy sampling	$\epsilon - greedy$	Policy sampling
ϵ schedule	—	$1.0 \rightarrow 0.05$	—
PPO clipping	—	—	0.2
GAE	—	—	0.95
PPO epochs	—	—	10
Entropy coefficient	—	—	0
Training budget	80,000 steps	80,000 steps	80,000 steps
Random seeds	42–47	42–47	42–47

All three methods use the same training budget of 80,000 steps and the same six random seeds to keep the computational budget similar throughout when comparing performance.

The simulated-data study changes two of these settings, and only these two. The network shrinks to a pyramid of two hidden layers, $d \rightarrow 64 \rightarrow 32 \rightarrow 3$, which is about 2,800 parameters for the nine-feature state. The training budget rises to 300,000 steps, which is roughly 15,000 episode draws over the 3,000

simulated training episodes. The old configuration had 18,000 parameters and 60 episodes; the new one has 2,800 parameters and 3,000 episodes. The data now outnumber the parameters instead of the other way around. Everything else, including the activation functions, optimiser, learning rates, discount factor, and the six seeds, is unchanged, so the two studies differ only in the data and in the network sized to match it.

The experiments consider the different state and reward configurations described in Chapter 5. These include the standard nine-feature state, the ten-feature state with drawdown, the primary wealth-change reward, the risk-aware reward, and the alternate action configurations.

The implementation also treats each trading position size as a configurable parameter. Its final value is selected using the predefined validation procedure. Chapter 5 reports the different values and their effect on trading results as part of the experimental findings.

4.5.3 Benchmark Implementation

The learned agents are compared with benchmark strategies to provide clear grounds for assessing their performance. These benchmarks are simple trading behaviours that do not require reinforcement learning.

The main benchmark is a buy-and-hold strategy where the initial capital is used to purchase the asset at the start of the episode. The position is then held until the final trading step, when it is closed. The same transaction fee and position closure used for the learned agents apply to this benchmark.

This approach is different from gradually purchasing the asset using the same fixed trade amount available to the learning agents. A gradual purchasing strategy may take several trading days to build a position and, if the fixed trade size is small, it may not become fully invested before the monthly episode ends.

For this reason, the gradual purchase strategy is treated as a separate reference strategy rather than as buy-and-hold. The conventional buy-and-hold strategy gives the learned agents a straight performance target.

The benchmark strategies use the same portfolio conditions, transaction costs, and evaluation procedures as the learned agents. This keeps the comparison consistent and reduces the chance that differences in results stem from inaccuracies in calculating portfolio performance rather than from the strategies themselves.

4.6 Implementation Verifications

Before using the experimental results for analysis, the implementation was tested to ensure that it followed the definitions and constraints established in Chapter 3.

These tests check that the environment, data pipeline, reward calculation, and action constraints behave as intended.

The final implementation includes 21 automated tests covering all major components, all of which pass. Appendix A lists every test and the condition it verifies; this section summarises what each group of tests protects.

4.6.1 Environment and Portfolio Tests

The environment tests focus on the correctness of the portfolio calculations during trading.

The tests verify that portfolio wealth matches the cash balance and the market value of current holdings. They also check that Buy and Sell actions update the portfolio correctly and that transaction fees are included in all transactions.

The tests further verify that invalid actions are handled correctly and that each monthly episode ends without an open position.

These checks are important because an error in portfolio accounting would affect the reported rewards and performance measures for all three learning algorithms. Such an error could therefore lead to incorrect conclusions even if the learning algorithms were implemented correctly.

4.6.2 Feature and Data Tests

The feature tests check that the state representation follows the intended information structure defined in Chapter 3.

The tests verify that features do not use future information and that rolling features access the necessary historical information when moving from one monthly episode to the next. They also check that the state dimensions and feature ordering match the reported configurations.

The chronological split is tested to ensure that the training, validation, and test periods remain separate.

This testing process identified an important implementation issue. In an earlier version of the pipeline, rolling features were recalculated separately for each monthly episode, so the rolling windows restarted at the beginning of every month instead of continuing from the previous observations. This caused incomplete feature windows in approximately **43.1% of episode steps**. The correction — calculating rolling market features on the continuous dataset before dividing it into monthly episodes — is described with the other corrections in Section 4.6, and the corrected pipeline was tested again to confirm that historical information is retained across episode boundaries.

4.6.3 Reward and Risk Tests

The reward tests verify that the primary wealth-change reward corresponds to the change in portfolio wealth with no risk penalty applied.

The risk-aware reward is tested separately to see whether increasing drawdown does not lead to an increase in the reward due to the risk penalty applied.

An additional test checks the relationship between the risk-aware reward and the state representation. A risk-aware reward cannot be paired with a state that does not have the drawdown feature.

The implementation uses this restriction so an inconsistent combination of state and reward settings cannot run accidentally.

These tests therefore ensure that the relationship between state representation and reward design described in Chapter 3 is also maintained in the implemented system.

4.6.4 Action-Masking Tests

The action-masking tests verify that the environment correctly identifies invalid actions that cannot be executed under the current portfolio conditions.

The tests check that the Buy action is unavailable when the portfolio has insufficient cash and that the Sell action is unavailable when the agent does not hold enough of the asset. They also verify the different configurable fixed trading amounts.

A defensive check is also included in the environment so that an invalid action supplied externally is converted to Hold rather than being executed.

Logs monitor these checks during the experiments. After the final experiment, the recorded rate of invalid actions was zero during evaluation, confirming that the action masks accurately prevented the learning algorithms from executing any invalid trades during the testing period.

4.7 Implementation Corrections

4.7.1 Problems Identified During Development

The verification process identified several implementation issues before the final experiments were run.

The main issues involved the feature calculation, benchmark implementation, and the risk-aware reward configuration.

First, the rolling market features using past observations did not initially contain the needed historical context when a new monthly episode began. These meant that some features were calculated using shorter windows than intended.

Second, an earlier version of the buy-and-hold benchmark used the same transaction size as the learning algorithms. This did not represent a fully invested buy-and-hold strategy, where an asset is bought at the beginning of an episode and held untouched till the end of the experiment, and it caused understated exposure and benchmark performance.

Third, the risk-aware reward could initially be used without requiring the drawdown feature in the state. This created a mismatch between the information needed by the reward and the information supplied to the agent because drawdown information was unavailable.

These issues were treated as implementation errors that needed correcting because leaving any of them in place would affect the validity of comparisons between the learning algorithms and benchmark strategies.

4.7.2 Correction and Re-verification

Each identified issue was addressed through a general correction process:

1. the problem was identified through testing or auditing the implementation;
2. its potential effect on the experiment results;
3. the relevant part of the implementation that needed correcting;
4. adding or updating verification tests where necessary after the correction;
5. completely running the verification tests again.

The buy-and-hold issue, for example, shows why this process was necessary. The affected implementation understated the benchmark by approximately **\$38.95 per month**, or **21.6%** before correction.

These values describe only the incorrect implementation, which has been discarded, and are reported here only to show the potential impact of the error. They are not presented as findings about the asset's or the market's performance.

After the corrections, the complete automated batch tests were run again, and all 21 tests passed.

The final experiment results in Chapter 5 come only from the corrected implementation, not the earlier incorrect versions.

4.8 Complete Implementation Pipeline

The complete implementation pipeline is a sequence of stages, from collecting historical market data to storing the final evaluation test results.

The first stage loads the SPY OHLCV data, including an initial warm-up period needed to calculate the rolling features. The market features are then calculated on the continuous dataset using only information available up to said observation. Once feature construction is complete, the data are divided chronologically into the training, validation, and test datasets.

The observations are then grouped into monthly episodes. At the start of each episode, the portfolio resets to the initial capital with no asset holdings. At each timestep, the learning algorithm receives the current state and selects one of the valid available actions. The trading environment then executes the action with the transaction fee, updates the portfolio, returns the reward, and moves to the next state.

REINFORCE, DQN, and PPO all interact with the same trading environment. Their main difference is how each algorithm processes the experience to update their model parameters.

After training and validation, the final experiment settings are fixed, and the trained models are evaluated on the held-out 2024-2025 test datasets. The resulting metrics, monthly results, training information, configuration settings, and performance measures are stored together with information about the software environment used to produce them.

This pipeline clearly links the methodological decisions defined in Chapter 3, their implementation in the system, and the experiment results presented in Chapter 5.

Figure 4.2 summarises the flow.




Figure 4.2: System architecture showing the data flow described in Section 4.7. All learned algorithms and benchmark strategies are evaluated using the same episode runner, ensuring that the comparison uses the same evaluation procedure.

Chapter 5

Experimental Results and Analysis

5.1 Introduction

This chapter reports and analyses the experiment results from the trading environment and learning algorithms described in Chapter 4. The experiments examine the financial performance of the learned agents, the effects of two state representations, and their respective reward functions.

The experiments are structured around four questions that directly relate to the objectives of the dissertation to increase portfolio wealth:

1. **Performance:** Do the learned agents outperform appropriate simple benchmarks on unseen data after accounting for transaction costs?
2. **Behaviour:** Do the learned agents make use of the market and portfolio information provided in the state representations?
3. **State representation:** Does adding the drawdown feature to the main nine-feature state affect the learned behaviour or financial performance?
4. **Reward:** Does the risk-aware reward improve the balance between return and drawdown, rather than simply encouraging the agent to trade less?
5. **Data:** Do the size and balance of the training data explain the behaviour of the learned agents, and does fixing the data fix the behaviour?

The chapter contains two studies. The first is the **real-data study**: the agents are trained on the 60 real training months and evaluated on the held-out 2024–2025 months. Its results answer questions 1 to 4 and expose a problem that motivates the fifth. The second is the **simulated-data study**: the same agents, trained on 3,000 balanced simulated episodes from the generator of Chapter 4 and evaluated under controlled market regimes, with the real test months kept as a transfer test. The second study is the corrected version of the first, and the chapter reports both because the failure of the first is what the second one explains.

A key part of the analysis is determining whether the learned strategies outperform simple strategies. Buy-and-hold, for example, is an important benchmark here. A learned agent producing a positive return does not necessarily indicate that it has learned a useful trading strategy if a simple investment strategy achieves a similar or better result.

For this reason, the analysis considers both **financial performance** and **agent behaviour**. Financial performance is evaluated using measures including wealth change, exposure, trading activity, drawdown,

and Sharpe ratio. Behavioural analysis examines whether an agent's actions change as the state changes. This helps to identify a policy that responds to state information rather than one that follows a fixed trading pattern.

5.2 Experiment Protocols and Evaluation Criteria

This section describes how the experiments in this chapter were conducted and how the reported results should be interpreted.

5.2.1 Experimental Configurations

Each configuration trained REINFORCE, Deep Q-learning (DQN), and Proximal Policy Optimisation (PPO) separately for 80,000 steps per seed in the portfolio environment,

To account for randomness in reinforcement learning [30], six random seeds were used, with an initial capital of $C_0 = \$10,000$ at the start of every monthly episode. This means six independent training runs all with the same configuration, allowing us to analyse the consistency of each algorithm's results rather than relying on a single training run.

During evaluation, all policies acted deterministically: policy-gradient algorithms select the action with the highest probability, while value-based algorithms select the action with the highest estimated value.

The 2023 validation period was used to make every experimental decision reported in this chapter. This included selecting the trade size in Section 5.3, the risk-penalty weight in Section 5.7, and the transaction fee in Section 5.8. These choices were fixed before using them for the main experiments in the 24 held-out test months from 2024–2025.

There are three main learned experiments in this chapter:

- The primary experiment evaluating the nine-feature state with the wealth-change reward.
- The ten-feature state that adds drawdown while keeping the same wealth-change reward.
- The third experiment using the ten-feature state with the risk-aware reward, which actually requires drawdown.

5.2.2 Evaluation Metrics

There are five financial measures used in each experiment report. They are averaged over the episodes and, for learned agents, across the six training seeds:

- Mean wealth change (ΔW) reported in dollars per month;
- Mean exposure (Δv_t), which represents the proportion of initial capital invested in the asset and is averaged across trading days;
- Trades per month;

- Mean episode drawdown (Δd_t);
- and the monthly Sharpe ratio [33], a financial metric that calculates the average return relative to the variation of those daily returns during the month.

The tables also report the **seed spread**, which is the difference between the best and worst result across the six seeds. This is included because there might be large differences between training runs which the average alone may hide [30]. A large spread indicates that the algorithm did not produce a consistent trading behaviour across the different runs.

Financial performance is then followed by a behavioural analysis, as returns alone cannot show whether a policy actually used the information in its state. For example, a policy that buys simply because buying is allowed can perform well during a rising market without responding to any market information.

This led to **state-dependence analysis**, which was used to test whether a policy's actions changed when its observed state changed while the valid actions remained the same. For instance, if a policy selected different actions at different steps with the same actions available, the difference could be attributed to the observed state. The tables report how many of the six seeds met this condition.

5.2.3 Benchmark Strategies

Two non-learning strategies are used as benchmarks. They use the same trading environment, portfolio conditions, and transaction fees as the learned agents.

Buy-and-hold invests the full initial capital on the first trading day and holds the position until the end of the month, when the position is closed. This is the main benchmark as it is a simple passive investment strategy.

Always-buy ($0.25C_0$) purchases a fixed trade amount equal to 25% of the initial capital whenever buying is available, continuing until it has fully invested in that asset. This strategy is useful for comparing against policies that achieve their returns through repeated buying rather than by using information from its observed state, as it also does not use any market or portfolio information when deciding whether to buy.

5.3 Fixed Trade Size Tuning

Before comparing the learned agents with the benchmarks, the fixed transaction size used by the learning agents must be finalized. Determining this trading-action size matters because the amount traded at each decision point determines how quickly the agent can build or reduce its position during a monthly episode.

In the trading environment, each Buy or Sell action changes the trade by a fixed fraction of the initial capital. A smaller value gives the agent finer control over its position but also means that more trading decisions are needed to reach a high level of asset exposure. A larger value lets the agent build its position more quickly but provides less control over the size of each change.

This creates a problem between the learned agents and the buy-and-hold benchmark. Buy-and-hold invests the available capital at the beginning of the episode, while a trading agent must build its position through repeated actions. If the trade size uses a small, fixed transaction size, the agent will need several trading days to reach a similar exposure as buy-and-hold and may not overall reach the same exposure within a typical month. Poor results could therefore arise from the fixed trade amounts rather than from the learned policy.

The trade size was therefore calibrated using the validation data to select the most stable size before running the main comparison between learning algorithms on the held-out data.

5.3.1 Exposure available under different transaction sizes

The validation experiment measures the maximum exposure reachable during a month using a policy that always buys whenever the action is available, for each fixed transaction size. The validation set contains 12 monthly episodes, with an average length of approximately 20.8 trading days. Table 5.1 and Figure 5.1 show the results.

Table 5.1: Attainable exposure for each fixed trade amount on the validation months. The policy buys whenever a purchase is valid.

Trade amount	Buy Days required to purchase a full asset	Mean maximum exposure	Mean wealth change
$0.05C_0$	20	0.450	+\$95.49
$0.10C_0$	10	0.690	+\$121.79
$0.25C_0$	4	0.838	+\$167.39
$0.50C_0$	2	0.886	+\$162.77
Buy-and-hold	—	0.910	+\$166.92




Figure 5.1: Exposure ceilings from Table 5.1. The $0.25 C_0$ trade amount (highlighted) is the smallest that clears the 90% target of buy-and-hold's exposure; the $0.10 C_0$ amount cannot reach it even in principle.

The results show that the original $0.10C_0$ trade size restricts the agent's asset exposure. This trade size requires ten buy days to fully invest in that asset, but the average episode is only 20.8 trading days long, so half a typical month passes before full investment is even possible. Its mean maximum exposure is 0.690, which is also significantly below the 0.910 asset exposure for the simple buy-and-hold strategy.

Hence, even a policy that buys whenever possible cannot reach the same level of exposure as a simple benchmark during a typical monthly episode. This matters when analysing the agent's performance. A low return may result from poor decision-making, but it may also result from the trading size preventing the agent from investing enough in the asset. The two outcomes should be treated distinctly.

The $0.25C_0$ trade size was therefore selected as the primary configuration, reaching 0.838 exposure, which corresponds to approximately 92% of the buy-and-hold exposure. This selection was based on the criteria that exposure must reach at least 90% of the buy-and-hold, since no trade size can reach that level completely.

The buy-and-hold exposure of 0.910 gives an average exposure of approximately 0.819 under this criterion. The $0.25C_0$ trade size therefore satisfies the requirement as the smallest tested trade size whose exposure is close to the benchmark.

Increasing the trade size setting to $0.50C_0$ raises exposure slightly, from 0.838 to 0.886, but makes each individual action much larger, giving the agent less control over gradual exposure changes.

Based on this calibration, the $0.25C_0$ size balances maximum asset exposure and control over position changes. It gives the learned agents a reasonable opportunity to reach the benchmark level without using a much coarser $0.50C_0$ trade size.

5.4 Do the Learned Policies Use Their Observations?

The Financial performance of the portfolio trading agent on its own does not show whether a learned agent is actually using the information contained in its state representation. For example, an agent that buys whenever buying is possible can perform well in a rising market without using any of the market information its state provides.

This requires examining the behavioural patterns of the agents' learned strategies to identify whether a policy selects a different action when the state changes, provided the available valid actions remain the same. This shows that the policy responds to its observed state rather than simply following action constraints and can be classified as **state-dependent**.

The experiment is performed on the validation data using the calibrated $0.25C_0$ trade size decided in Section 5.3, with the nine-feature state and the primary wealth-change reward. This test aims only to analyse behavioural differences between the algorithms by determining whether changes in the information provided to the agent are reflected in its decisions. It is not used to determine the final financial performance of the agents.

Every cell was trained from scratch for this study: six seeds per method, 80,000 steps each.

5.4.1 Behaviour of the learned policies

Table 5.2: Behavioural Results on the validation months (0.25 C_0 trade size, nine-feature state, wealth-change reward, six freshly trained seeds). The always-buy strategy shows the maximum asset exposure this trade size can reach.

Policy	Mean ΔW	Seed spread	Exposure	Trades	State-dependent
Always-buy	+\$167.39	—	0.838	—	—
REINFORCE	+\$154.96	\$37.29	0.807	10.3	0/6
Deep Q-learning	+\$86.16	\$64.20	0.325	8.5	6/6
PPO	+\$118.55	\$167.39	0.596	3.7	0/6




Figure 5.2: Results for each seed from Table 5.2. Hollow-marker seeds show actions that don't depend on their state information; filled-marker seeds are state-dependent ones. Policy-gradient methods (REINFORCE, PPO) produce policies that cluster to either high exposure or collapse towards zero, while all six Deep Q-learning seeds fall between those corners.

Table 5.2 summarises the outcome and Figure 5.2 shows every seed individually. The columns read as follows. Mean ΔW is the average wealth change per month, calculated over the 12 validation months and then over the six seeds. Seed spread is the difference between the best and worst seeds; a spread larger than the mean produced noticeably different results and did not agree on a strategy. State-dependent counts how many of the six seeds act according to the experimental behaviour being analysed.

The REINFORCE results are clearer when the individual seeds are read. Four of its six seeds finished at exactly +\$167.39 with an exposure of 0.838. This matches the always-buy strategy, which was set as the

upper limit. These policies buy on the first four days when buying is valid, hold the position, and close at the end of the month. The other two seeds earned +\$130.10 while trading 20.8 times per month. Their actions followed a fixed pattern of trading on each bar rather than changing in response to the market state.

PPO produced three different behaviours across its six seeds. Four seeds reached the always-buy result of +\$167.39. One seed never traded at all, producing a wealth change of \$0.00 and an exposure of 0.000. The remaining seed earned +\$41.73, relying on only two trading styles. This produced a seed spread of \$167.39, showing the large difference between the strongest and weakest PPO runs.

Deep Q-learning is the only method where none of its seeds reached either extreme behaviours, unlike the policy-gradient methods. Its results ranged from +\$53.23 to +\$117.43, with a mean wealth change of +\$86.16 and an average exposure of 0.325, less than half the 0.838 standard exposure of the always-buy strategy.

5.4.2 State-dependence analysis

The final column in Table 5.2 summarises the state-dependence test, with the count representing the difference in behaviour across the six seeds

For REINFORCE, the result of 0/6 means that none of the six independently trained policies changed their action when the state changed while the same actions remained valid. The four seeds that reached the maximum exposure followed the same basic rule of buying whenever possible, while the other two followed a fixed trading pattern. The diagnostic therefore found no evidence that these policies used market and portfolio features to change their decisions.

Deep Q-learning produced the opposite result. All six seeds were classified as state-dependent. Each seed selected different actions in states where the same valid actions were available. This indicates that the observed current state influences the decisions made by the DQN policies.

PPO produced the same result as REINFORCE, with 0/6 state-dependent seeds. Taken together, the splits are categorical by algorithm family across the three methods: all six DQN seeds were state-dependent, compared with none of the 12 policy-gradient seeds from REINFORCE and PPO.

5.4.3 Interpretation of the learned trading behaviour

The different results across the three methods suggest that the learning approach had a strong influence on the type of trading behaviour that developed. REINFORCE and PPO frequently moved towards simple policies that did not change their actions in response to the observed state, whereas Deep Q-learning showed state-dependent behaviour across all six seeds.

One possible explanation is the difference between policy-gradient and value-based learning. REINFORCE updates the probability of an action according to the return that follows it. When the market generally favours holding an asset, repeatedly buying can produce a strong return without needing to

sample other policies. The learning signal may therefore settle on a simple high-exposure strategy for higher rewards rather than exploring.

PPO uses a different update rule from REINFORCE, but it still learns a policy directly. Its clipped objective prevents large policy changes, which can make learning more stable, but it does not need the policy to become state-dependent. If a simple trading rule produces consistently positive returns, the policy may face little pressure to learn a more complex relationship between state and action. The extreme seed outcomes are what hint at this. Seeds from both policy-gradient methods finish at the always-buy upper limit, since that is the most rewarding.

Deep Q-learning learns differently. Rather than adjusting action probabilities directly, it estimates the value of each available action in the current state and picks the action with the highest estimate at each step. When future estimated values differ, the chosen action changes accordingly, and the agent has to distinguish between these actions. The state-dependent behaviour therefore becomes a structural property of this algorithm rather than an achievement, as it necessarily depends on the changes in the market and portfolio state.

These explanations should be treated as interpretations of observed behaviours rather than as direct evidence of the internal learning process. This experiment was not designed to find out the precise reason the algorithms moved to these different strategies. They are, however, consistent, as different seeds show that learning algorithms affect how available state information is used to develop a trading strategy.

5.5 Effect of the State Representation

Chapter 3 defines the nine-feature state as the main state representation used in this study. It includes market, portfolio, position, and time information. Drawdown is intentionally treated separately as an additional risk-related feature rather than as the last stage of the state design.

This experiment therefore focuses on a specific question: does adding drawdown to the nine-feature state produce a meaningful change in performance or behaviour when the reward function remains unchanged?

The comparison keeps the following conditions fixed:

- The same wealth-change reward;
- the same $0.25C_0$ trade size;
- the same held-out test period;
- the same six random seeds.

The intended difference between the two configurations is therefore only the inclusion of drawdown in the state.

5.5.1 Nine-Feature vs Ten-Feature State

Method	State	ΔW (\$)	Spread	Exposure	Trades	DD	Sharpe	State-dep.
REINFORCE	9-feature	+171.51	24.7	0.810	10.3	0.0266	0.675	0/6
REINFORCE	10- feature	+167.39	24.7	0.794	13.0	0.0263	0.670	0/6
Deep Q	9-feature	+74.68	71.4	0.312	8.2	0.0124	0.507	6/6
Deep Q	10- feature	+77.81	89.5	0.354	7.4	0.0140	0.505	6/6
PPO	9-feature	+127.35	179.8	0.598	3.7	0.0195	0.562	0/6
PPO	10- feature	+159.97	94.0	0.759	7.3	0.0250	0.671	1/6

Table 5.3: Results of the nine- and ten-feature states on the test data (wealth-change reward, 0.25 C_0 trade size, six seeds). DD shows the mean drawdown within each episode, while State-dep. shows the number of seeds where actions changed with the observed state.

The effect of adding drawdown and the changes in mean wealth performance for REINFORCE and DQN are small.

For REINFORCE, wealth change falls from \$171.5 with the nine-feature state to +\$167.39 with the ten-feature state. The difference is only $-\$4.12$ per month against a seed spread of \$24.70 for both configurations, which is considerably larger than the mean wealth change.

The behavioural result also remains unchanged. None of the six seeds are classified as state-dependent in either configuration. Adding drawdown therefore did not produce a noticeable change in the behaviour of the REINFORCE policies.

DQN shows a small change in the opposite direction. Wealth increases from +\$74.68 to +\$77.81, a change of +\$3.13 per month. The seed spreads are \$71.40 and \$89.50, respectively, so the change between the two state configurations is again much smaller than the variation across seeds.

The behavioural result is unchanged for DQN. All six seeds remain state-dependent with both state representations.

PPO appears to improve, from +\$127.35 to +\$159.97, gaining \$32.62. However, the result needs to be considered alongside the variation between seeds. The nine-feature configuration has a spread of \$179.80—larger than the mean itself—but it narrows to \$94.00 in the ten-feature configuration, with only one seed becoming state-dependent.

The change in the mean is therefore not enough, on its own, to show that adding drawdown caused the improvement.

5.5.2 Analysis When Drawdown Was Added

The other performance measures columns tell the same story in different parts. For DQN, average exposure rose from 0.312 to 0.354, and mean drawdown from 0.0124 to 0.0140 after adding that feature to the state — return, exposure, and risk all moved up together by similar proportions. Its Sharpe ratio, however, changes very little, from 0.507 to 0.505.

The patterns noticed here suggest the additional drawdown feature did not lead to more effective risk-aware decisions, but to slightly larger asset exposure

The changes for REINFORCE barely move. Exposure falls from 0.810 to 0.794, while drawdown decreases slightly from 0.0266 to 0.0263, as expected of a policy algorithm that does not read the state it is given.

The small changes in the financial measures therefore do not indicate a change in the basic policy behaviour.

Overall, the results do not show a consistent improvement in either performance or behaviour from adding drawdown to the state while using the default wealth-change reward.

5.5.3 Interpretation

These results do not mean that drawdown information is unimportant. They show that, in this environment, without changing the reward objective, the additional information did not provide a clear advantage— an agent maximising wealth change in general rising markets has little use for its own loss history.

This matters because Chapter 3 introduced drawdown to provide information that could support risk-sensitive decisions. However, this feature becomes relevant only when the agent's learning objective makes larger drawdowns costly while increasing portfolio wealth, which is precisely what the risk-aware reward experiment in Section 5.7 tests.

5.6 Held-Out Test Results

The main performance question of how each learned strategy, in its final configuration, compares with the simple benchmarks on unseen data is evaluated on the 24 held-out months from January 2024 to December 2025.

At this point, the main configuration has been fixed:

- nine-feature state;
- wealth-change reward;
- $0.25C_0$ trade size;
- 0.05% transaction fee;

- six random seeds.

5.6.1 Overall Comparison with the Benchmark Strategies

Policy	ΔW (\$)	Spread	Exposure	Trades	DD	Sharpe	State-dep.
Buy-and-hold	+180.25	—	0.913	2.0	0.0310	0.638	—
Always-buy (0.25C₀)	+179.75	—	0.841	5.0	0.0274	0.684	—
REINFORCE	+171.51	24.7	0.810	10.3	0.0266	0.675	0/6
Deep Q-learning	+74.68	71.4	0.312	8.2	0.0124	0.507	6/6
PPO	+127.35	179.8	0.598	3.7	0.0195	0.562	0/6

Table 5.4: test results over the 24 held-out months (2024–2025): benchmarks and the three learned agents in their final configuration, six seeds each.

Table 5.4 is the central result of this dissertation.

Two simple strategies are used as benchmarks. Buy-and-hold invests the full capital on the first day and is the benchmark an investor would actually compare against. Always-buy at $0.25C_0$ buys an asset at a quarter fraction of the initial capital. This takes trades with every valid action until it has fully invested in the asset. It represents a strategy that buys unconditionally at the fixed trade size, without requiring observations.

The central result is straightforward: none of the three learned agents outperformed buy-and-hold during the held-out test period.

Buy-and-hold earned +\$180.25 per month, equivalent to a 1.8025% increase of the initial capital per monthly episode, and no learned agent reached this. The closest was REINFORCE at +\$171.51, which was \$8.74 short of buy-and-hold per month; PPO earned +\$127.35 and Deep Q-learning +\$74.68.

Always buying at $0.25C_0$ earned +\$179.75, within \$0.50 of buy-and-hold, and it sets the standard a learned agent must reach/beat before it can claim any skill, as matching it needs no market information at all — only buying whenever buying is available.

Over the full two years, the gap compounds visibly in Figure 5.3: buy-and-hold accumulates to +\$4,326 in monthly wealth change at the end of 2025, REINFORCE to +\$4,116 closely following buy-and-hold, PPO to +\$3,056, and Deep Q-learning at the lowest value of +\$1,792.




Figure 5.3: Cumulative sum of average monthly wealth changes across the test months. Each monthly episode resets to C_0 , so the lines show the sum of monthly gains and losses rather than compound; the right axis expresses percentages of initial capital C_0 . REINFORCE follows Always-buy ($0.25 C_0$) and buy-and-hold almost exactly, while Deep Q-learning shows a flatter line due to low exposure, including through the March–April 2025 decline.

5.6.2 Return and market exposure

The exposure column explains most of the return column. The REINFORCE result is particularly interesting; its return was close to buy-and-hold at +\$171.51. Its average exposure was 0.810, compared with 0.913 for buy-and-hold, and none of its six seeds were state-dependent, suggesting the high return came mainly from staying invested rather than learning when to enter or leave the market. Its behaviour was therefore closer to the Always-buy at $0.25C$ accumulation strategy. Hence why it sits within a few percent of Always-buy results (+\$179.75, exposure 0.841).

DQN produced a different result. All six seeds were state-dependent, showing that the agent changed its actions in response to state information, and chose to participate the least, with the lowest exposure of 0.312 and a wealth change of \$74.68. These results are roughly a third of the buy-and-hold benchmarks, so the lower return may therefore be correlated with reduced market participation.

PPO produced a return between REINFORCE and DQN, with an average wealth change of \$127.35 and exposure of 0.598. However, its seed spread was 179.8, which is larger than its mean return. This shows that the six PPO runs produced different outcomes and should not be treated as a single stable trading strategy.

5.6.3 Risk and risk-adjusted performance

Lower returns can still be desirable if they follow a sufficient reduction in risk. For this reason, drawdown and Sharpe ratio are also used alongside return. Figure 5.4 puts every learned configuration in this chapter on that scale.

Buy-and-hold earned \$197 per unit of exposure. An agent that merely holds less of the same asset lands on the dashed proportional line; genuine downside management would appear above it.




Figure 5.4: Average wealth change compared with exposure (left) and daily drawdown (right) for every learned configuration in this chapter. The dashed lines show the proportional relationship with buy-and-hold. None of the learned configurations show a clear improvement over either line.

Deep Q-learning is an informative case. It earned 41.4% of the buy-and-hold benchmark's return at 34.2% of its exposure (0.312) and 40.0% of its drawdown (0.0124) — return and risk fell equally, which is more like scaling down. Its Sharpe ratio of 0.507 versus buy-and-hold's 0.638 confirms it earned less return per unit of risk exposed to, not more.

The results therefore do not indicate that DQN achieved better risk-adjusted performance than the benchmark. Its lower drawdown largely came from lower exposure, not from a strategy that intentionally avoided the market's decline periods.

5.6.4 Behaviour During the March 2025 Decline

March 2025 provides a useful example because it was the worst test month for the buy-and-hold benchmark. During this month, buy-and-hold lost \$398.41, i.e. 3.9841% of the initial capital, with a drawdown of approximately 5.6%. If the state-dependent agent, Deep Q-learning, had recognised the conditions, this would have been the month to step aside. Instead, DQN had a mean exposure of 0.490 in March—higher than its overall average exposure of 0.312 for the entire test period—and lost \$184.04, 46% of the benchmark's loss.

These results are important when interpreting the state-dependence finding. Even in a month when reducing exposure before the decline could have helped, it did not significantly reduce exposure, and these March 2025 losses match its equivalent market participation. Hence, its lower drawdown cannot be attributed to knowing when to avoid bad market periods, but to holding less overall.

Figure 5.5 extends to all 24 months. Each pair of bars is one test month: the dark bar represents buy-and-hold's wealth change, and the orange bar represents the Deep Q-learning for the same month. The two series appear to rise and fall together — the agent's best month is the benchmark's best month (May 2025, +\$390.80 compared to +\$543.11) and its worst is the benchmark's worst (the discussed March 2025).




Figure 5.5: Monthly wealth changes of buy-and-hold and Deep Q-learning over the 24 test months. The agent’s bars match a 1/3 scaling of the benchmark’s bars in almost every month: In March 2025, the worst month, it lost \$184.04 compared to the benchmark’s \$398.41 loss.

Buy-and-hold lost money in seven of the 24 months, and the agent ended positive in three of those seven. Those three months, however, produced little profit compared to the benchmark losses

In the two heavy down months, April 2024 (−\$396.08) and March 2025 (−\$398.41), the agent was also at a loss—approximately 24% and 46% of the benchmark’s loss, respectively—correlating to its 34% average market participation and one-third scaling to the benchmark. Where drawdown management mattered most, the agent still left the position open to losses.

5.6.5 Trading behaviour and Seed consistency

The results also show considerable differences in consistency between the algorithms. To examine this, we analyse seed spreads, which measure differences between runs.

REINFORCE had the smallest seed spread (\$24.70), showing seeds agree with each other. However, this consistency should be interpreted with its behavioural results. All six seeds were classified as state-independent and exhibited similar behaviour, precisely because they agree on ignoring market information and generally being highly exposed, as that is more rewarding. The small spread therefore does not necessarily indicate that REINFORCE learned a more reliable trading strategy.

DQN had a larger spread of \$71.40. This is consistent with the six seeds learning different state-dependent policies, resulting in greater variations in their returns.

PPO had the largest spread of \$179.80, which was greater than the mean return, also showing that individual seeds ranged from near-inactivity to near-full investment in Fig. 5.2, and none were state-dependent.

These results show why evaluating a single random seed would not provide enough evidence for a reliable conclusion about the learning algorithms. Behaviour and performance can vary considerably from one training run to the next.

5.7 Effect of the Risk-Aware Reward

Earlier experiments showed that the main wealth-change reward did not always encourage agents to respond differently to changes in the observed state. It simply rewards changes in portfolio wealth.

Chapter 3.5 identifies the limitation of this objective: it does not penalize large portfolio drawdowns, raising the question of whether explicitly including portfolio drawdown in the reward would lead to more risk-aware behaviour.

To examine this, the risk-aware reward was introduced in which it was modified to include a penalty for increases in drawdown:

$$r_t^{risk} = \Delta W_t - \lambda C_0 \Delta dd_t.$$

where

Δdd_t is the increase in portfolio drawdown at trading day t (and is zero when drawdown does not rise), and coefficient λ controls the penalty effects. The initial capital C_0 factor puts Δdd_t in dollars so that λ is dimensionless.

The risk-aware reward is paired with the ten-feature state using the drawdown feature because it provides the agent with information about its current position relative to its previous maximum wealth.

5.7.1 Selection of the Penalty size

The penalty size was selected using the validation set and Deep Q-learning as the selector because it is the only algorithm whose policies consistently respond to the state (Section 5.4). The held-out test period was not used for this decision.

A selection rule was created. The penalised configuration must keep at least half (50%) of both the unpenalised return—to still earn a decent amount — and the unpenalised exposure—to still participate in the market.

The unpenalised DQN configuration had a mean wealth change of \$78.35 and exposure of 0.445. The resulting minimum requirements were therefore:

$$\Delta W \geq \$39.17$$

and

$$Exposure \geq 0.223.$$

The results of the penalty sweep are shown in Table 5.5.

λ	Mean ΔW	Exposure	Drawdown	DD change	Passes
0 (unpenalised)	+\$78.35	0.445	0.0163	—	—
0.05	+\$88.78	0.401	0.0128	−21.5%	yes
0.10	+\$67.30	0.344	0.0124	−24.0%	yes
0.25	+\$19.90	0.074	0.0035	−78.4%	no
0.50	+\$1.73	0.002	0.0001	−99.4%	no
1.00	\$0.00	0.000	0.0000	−100%	no
2.00	\$0.00	0.000	0.0000	−100%	no

Table 5.5: Risk-penalty size on the validation months (Deep Q-learning, six seeds). The rule requires $\Delta W \geq +\$39.17$ and exposure ≥ 0.223 ; the largest passing value that meets both requirements is $\lambda = 0.10$.

Deep Q-learning on the validation months




Figure 5.6: Results from the penalty size of Table 5.5. Left: mean wealth change against λ with the requirement floor of +\$39.17. Right: exposure and drawdown relative to their values without penalties. Between $\lambda = 0.10$ and 0.25, exposure drops sharply rather than decreasing gradually.

The results show a clear change in behaviour of the penalty λ as the size increases. At $\lambda = 0.10$, the agent still maintains a meaningful level of exposure and a return of \$67.30.

Increasing the penalty to 0.25 causes exposure to fall sharply to 0.074. The agent loses about 78% of its exposure, indicating reluctance to trade. At $\lambda = 0.50$ and above, the agent stops trading completely.

This result depicts an important limitation of the reward formulation. The simplest way for an agent to reduce drawdown is to hold less of the asset or simply not hold it at all. In theory, a portfolio that remains in cash is not exposed to asset movements and therefore does not experience market drawdowns. Once the drawdown penalty becomes too large, market participation becomes unattractive.

Based on these criteria, the largest coefficient that satisfied both the return and exposure requirements was therefore:

$$\lambda = 0.10.$$

This became the fixed penalty value for evaluating the main risk-aware experiment.

5.7.2 Effect on the Test Set

The selected risk-aware reward was then evaluated on the held-out test set.

Algorithm	Reward	ΔW (\$)	Spread	Exposure	Trades	DD	Sharpe
REINFORCE	wealth	+167.39	24.7	0.794	13.0	0.0263	0.670
REINFORCE	risk	+167.39	24.7	0.794	13.0	0.0263	0.670
Deep Q	wealth	+77.81	89.5	0.354	7.4	0.0140	0.505
Deep Q	risk	+52.90	43.1	0.178	5.4	0.0082	0.416
PPO	wealth	+159.97	94.0	0.759	7.3	0.0250	0.671
PPO	risk	+21.48	128.9	0.108	0.7	0.0037	0.107

Table 5.6: *Wealth-change versus risk-aware reward ($\lambda = 0.10$) on the held-out months, ten-feature state, six seeds.*

The DQN result gives the clearest evidence of the effect of the risk-aware reward. For Deep Q-learning, the mean drawdown falls from 0.0140 under the standard reward to 0.0082 with the risk-aware reward, a reduction of approximately 41%. However, its return also falls from +\$77.81 to +\$52.90, a reduction of about 32%. Exposure falls from 0.354 to 0.178, meaning that the agent decides to hold only half as much of the asset.

The results do not suggest improved risk-adjusted performance. Instead, the reward mainly encouraged the agent to reduce its exposure to the asset, and the Sharpe ratio also decreased from 0.505 to 0.416.

The PPO results show the same effect more strongly. With the risk-aware reward, PPO's exposure fell to 0.108, and its average trading activity dropped below one trade per month, from 7.3 to 0.7. As a result, its mean wealth change decreased from \$159.97 to +\$21.48.

REINFORCE shows no noticeable change. Its results are identical for both rewards. This is consistent with the earlier behavioural findings: REINFORCE policies do not change their trading behaviour simply by changing the reward.

5.7.3 Interpretation of the risk-aware return

The risk-aware reward changed the behaviour of the state-dependent agents, but there is no clear evidence that it led to better risk management. Instead, the results suggest that the penalty mainly encouraged the agents to reduce their exposure to the asset.

These results do not show that risk-aware reinforcement learning is ineffective in general. However, they do show a limitation of the reward formulation used in this study. An agent’s policy could reduce drawdown to zero by simply holding less of that asset or by never entering the market, but this would not satisfy the main objective of generating useful investment returns.

A more effective objective would therefore need to consider both return and risk, rather than mainly penalising exposure to drawdown. The relevant question becomes whether the agent can **reduce drawdown while still leading to meaningful returns and exposure**. For example, the objective could reward the agent for achieving higher returns for a given level of risk. Chapter 6 discusses other possible alternatives in future work.

5.8 Transaction-Fee Sizing

Transaction costs are an important part of the trading environment because frequent trading can reduce the returns produced by a strategy, even when individual decisions appear reasonable. A strategy may appear profitable before fees are included but become less attractive once the cost of each trade is taken into account.

The main experiments use a transaction fee of 0.05%. To test whether the observed behaviour changes as this fee varies, the fee was swept from 0 basis points (0.00%) to 50basis points (0.50%) on the validation months with the nine-feature state, and everything else held fixed.

The sweep was run on all algorithms, each with six freshly trained seeds per fee level. Table 5.7 and Figure 5.7 report the outcome.

Fee	B&H ΔW	REINF. ΔW	REINF. trades	Deep Q ΔW	Deep Q trades	PPO ΔW	PPO trades
0bps(0.00%)	+\$177.09	+\$177.56	5.0	+\$108.00	10.9	+\$155.35	4.5
5bps(0.05%)	+\$166.92	+\$154.96	10.3	+\$86.16	8.5	+\$118.55	3.7
10bps(0.10%)	+\$156.76	+\$143.74	4.7	+\$77.60	8.9	+\$78.50	2.7
25bps(0.25%)	+\$126.33	+\$88.76	10.3	+\$32.35	3.9	+\$26.40	1.2
50bps(0.50%)	+\$75.83	+\$41.31	2.8	+\$16.32	1.9	+\$12.60	1.0

Table 5.7: Transaction Fee sizing on the validation months; six seeds. Trades are per month. Fees are shown in basis points and as a percentage of trade value.

5.8.1 Effect of transaction costs on returns & trading behaviour

As expected, higher transaction fees reduce the returns of the different algorithm strategies. The effect is important for strategies that trade frequently because more of their portfolio gains go to transaction costs.




Figure 5.7: Transaction Fee sizing from Table 5.7. Left: average wealth change falls for every strategy as fees rise. Right: average trades per month as fees rise.

Deep Q-learning shows the clearest reduction in trading activity as the fee of taking a trade increases. Its average number of trades falls from 10.9 per month at 0bps (0.00%) to 1.9 at 50bps (0.50%), an 83% reduction. As each action gets more expensive, the agent acts less, reducing exposure. This is consistent with the algorithm correctly incorporating transaction costs into its environment.

REINFORCE does not show the same clear pattern. Its trade counts vary, running from 5.0, 10.3, 4.7, 10.3, and 2.8, respectively, with no consistent decrease in trades taken as the fee rises. This aligns with the earlier state-independent results: an agent that does not rely on information from its current state to make decisions is less likely to adjust its behaviour when trading costs change.

PPO's average trade count also falls from 4.5 to 1.0, showing that many seeds are pushed to stop trading due to higher fees. This is also why the average return collapses by 77%, from +\$177.56 to +\$41.31, a difference of \$136.25 per month, even though the average trade count moves only gently.

5.8.2 Interpretation

The transaction-fee experiment therefore reflects another difference between the learning algorithms in terms of their change in behaviour. Deep Q-learning changes its trading behaviour as trading costs increase, even though its overall returns are below the buy-and-hold benchmark. This mirrors the fee response reported by Théate and Ernst, whose agent also traded less as the assumed costs rose [22]. REINFORCE produces returns generally closer to the buy-and-hold strategy, and it also shows higher exposure that does not respond consistently to changes in transaction costs. PPO, on the other hand, shows a major collapse in returns as fees increase, while its trade counts decline more gently. Overall, the learned agents lose proportionally more because they trade more often in this experiment.

This indicates that the learning system is capable of responding to different changes when the learned policy is sufficiently state-dependent.

5.9 The Data Bottleneck

This section explains why the real-data results of the previous sections look the way they do. The explanation is not about any single algorithm. It is about how little data the agents were given.

The training split contains 60 monthly episodes. The networks used in the real-data study have more than 18,000 trainable parameters. A model with that much capacity can reproduce 60 training months exactly, the way a student who memorises 60 practice problems scores perfectly on those problems and then fails a fresh exam. Nothing forces the network to learn a rule that transfers, because memorising the answers is cheaper.

The training months are also imbalanced. Twenty-nine of the 60 months rose by more than 2% and only 16 fell by more than 2%, and the test years 2024–2025 rise almost throughout. Under that distribution, “always buy” is a nearly optimal answer that requires no state information at all. This is the same trap as training a classifier on a dataset that is 99% one class: predicting the majority class scores well and learns nothing. The always-buy policies of the policy-gradient agents in Section 5.4 are exactly this artefact.

The same reasoning explains why Deep Q-learning performed worst on the real test months. A value-based method must estimate the value of each state it visits, which needs many visits to similar states under a stable state distribution. Sixty episodes from a shifting market provide neither. The earlier behavioural finding stands, six of six DQN seeds did condition on their state, but the values behind those decisions were estimated from too little data, so conditioning on them did not pay.

None of this can be tested on the real data, because the real data is the problem. Testing it needs a setting where the number of episodes, the balance of market conditions, and the match between training and test distributions are all controlled. That is what the simulated-data study provides. The next section reports it.

5.10 Simulated-Data Study: Controlled Regime Results

This section reports the main study of the dissertation. The three algorithms are retrained on the 3,000 balanced simulated episodes described in Chapter 4, using the smaller pyramid network of about 2,800 parameters and a budget of 300,000 steps. Evaluation uses 600 fresh simulated episodes, 200 in each regime, plus the 24 real test months as a transfer test. Six seeds per algorithm, as before.

5.10.1 Benchmarks Under Balanced Regimes

The balanced test set changes what the two benchmarks earn, and this change is the point of the design. Buy-and-hold now loses \$60.29 per episode on average. It gains \$510.17 in the average up episode, loses \$710.06 in the average down episode, and roughly breaks even when the market is flat. Always-buy at $0.25C_0$ behaves the same way at slightly smaller magnitudes, losing \$58.70 per episode. A policy that never trades ends every episode with exactly its starting capital, a wealth change of \$0.00.

On the real test months, buying and holding earned \$180.25 because the period only rose. On the balanced

episodes, no fixed pattern is profitable: constant buying loses money and abstaining earns nothing. Any agent that earns money here must be reading its observations and behaving differently in different regimes. The benchmarks are no longer a wall to climb; they are a test of whether the agent conditions on its state.

5.10.2 Main Results

Deep Q-learning earned **+\$126.94 per episode** on average across the six seeds, with a seed standard deviation of \$7.30. Every other strategy, learned or fixed, lost money or broke even on the balanced test set. All six DQN seeds were classified as state-dependent, and all six were profitable. Table 5.8 summarises the study.

Table 5.8: Simulated-data study: mean wealth change (\$) per episode on the 600 balanced test episodes, per-regime means, and the real-data transfer test. Mean and standard deviation are across six seeds.

	Sim test	Std	Up	Down	Flat	Real
DQN	+126.94	7.3	+473.72	−103.35	+10.45	+130.61
REINFORCE	−5.71	52.2	+228.58	−248.32	+2.60	+73.04
PPO	−11.16	27.3	+67.69	−99.04	−2.14	+25.84
Buy-and-hold	−60.29	—	+510.17	−710.06	+19.00	+180.25
Always-buy ($0.25C_0$)	−58.70	—	+477.22	−663.91	+10.58	+179.75

The per-regime columns show what the Deep Q-learning agents actually learned. In up episodes they earn \$473.72 of the \$477.22 that always-buying earns, so they capture 99% of the available upside. In down episodes they lose \$103.35 where always-buying loses \$663.91, an 84% reduction in the loss. In flat episodes they roughly match the market. That is the state-dependent behaviour the real-data study looked for and could not find: stay invested while the market rises, step aside while it falls. Figure 5.8 shows both views: the per-regime means on the left, and every seed’s simulated result against its real transfer result on the right.

5.10.3 The Policy-Gradient Methods Collapse to Corners

REINFORCE and PPO did not reach that policy reliably. Across the six REINFORCE seeds, three converged to never trading (a wealth change of exactly \$0.00), two converged to always buying (−\$58.70, matching the buy-and-hold row to the cent), and one learned a genuinely conditioned policy. PPO was similar: five seeds never traded and one adopted a fixed buying pattern. No PPO seed was state-dependent.

The corners are not senseless. On balanced data, unconditional buying loses money, so never trading is a locally safe answer: it earns nothing, but it cannot lose. A policy-gradient method that drifts toward either deterministic corner stops exploring, because the gradient of a near-deterministic policy carries almost no signal. The value-based method does not have this failure mode, since epsilon-greedy exploration keeps visiting both actions regardless of what the current policy prefers.




Figure 5.8: Simulated-data study. Left: mean wealth change per episode by test regime for each algorithm and for always-buy. Deep Q-learning keeps almost all of the up-regime profit while cutting the down-regime loss by 84%. Right: each seed’s simulated test result against its real transfer result; filled markers are state-dependent seeds, hollow markers are mask-only seeds. All six Deep Q-learning seeds cluster in the profitable upper-right; the policy-gradient seeds sit at the never-trade origin or the always-buy corner.

The single conditioned REINFORCE seed matters, though. It earned \$83.13 on the simulated test, keeping \$417.02 of the up-regime profit while cutting the down-regime loss to \$162.11. The policy the study wants is therefore inside REINFORCE’s reach; the training process just does not find it dependably. One seed in six is an existence proof, not a method.

5.10.4 Transfer to the Real Test Months

Each trained policy was also evaluated, unchanged, on the 24 real test months. The Deep Q-learning agents earned \$130.61 per month on average (seed range \$115.24 to \$140.92), against \$180.25 for buy-and-hold.

The gap has a clean explanation. The real 2024–2025 period rises almost throughout, and a regime-aware agent pays for insurance it never needs: whenever its features suggest a possible downturn, it reduces exposure, and in a period with no sustained downturn every such reduction costs return. The simulated down episodes show what that insurance buys, an 84% smaller loss when the market does fall. The real test years simply never presented the conditions the agent learned to defend against.

This is also a domain-gap measurement. The generator produces simpler price paths than the real market, so a policy trained on it cannot be expected to transfer at full strength [24]. That the transferred agents remain profitable in every seed, and that their ordering matches the simulated ordering, suggests the learned regime response survives the gap even though its edge shrinks.

5.10.5 What the Controlled Study Establishes

Three findings, each impossible to obtain from the real data alone. First, the task is learnable: given enough balanced data, a value-based agent learns a profitable, state-dependent policy on every seed.

Second, the real-data failure was a data failure: the same algorithms, environment, and reward that failed on 60 imbalanced months succeed on 3,000 balanced episodes, with the network being smaller rather than larger. Third, the algorithm ranking inverts and the inversion is informative: the method that was worst on memorisable data is the only method that wins when the data is large enough to reward conditioning, while the methods that looked strong on real data collapse to fixed corners once always-buying stops paying. The next section brings both studies' findings together.

5.11 Discussion of the Main Findings

The experiments produced several findings about both the performance of the learning agents and the way they behaved within the trading environment. They also answer the research questions of Section 5.1. The first findings below come from the real-data study; the section then turns to what the simulated-data study adds and corrects.

5.11.1 The Benchmark Remains Difficult to Beat on Real Data

The main real-data finding is that none of the learned agents could outperform the simple buy-and-hold benchmark on the held-out test period. Buy-and-hold achieved a wealth change of \$180.25, compared with \$171.51 for REINFORCE as the best learned agent, while every other agent fell further behind with \$127.35 for PPO, and \$74.68 for DQN.

However, it is important to note that the test period consists of generally favourable market conditions. Therefore, the experiments cannot simply be viewed as evidence that reinforcement learning failed, as staying invested in generally rising markets is already a strong strategy

Buy-and-hold is therefore a difficult benchmark to improve upon, which is consistent with the market-efficiency expectation set out in Chapter 2 [31]. A learned strategy that reduces exposure must therefore focus on periods when leaving the market is advantageous enough to compensate for the return not gained while being out of the market.

The experiments provide limited evidence that the tested agents could make this decision and overall did not demonstrate such behaviour.

5.11.2 Additional State Features Did Not Lead to Better Performance

The state comparison did not show a consistent advantage from adding drawdown to the nine-feature state. This is important because the state representation was designed to include information about market movements, portfolio position, wealth, and risk.

For REINFORCE and DQN, the changes in return were -\$4.12 and \$3.13, respectively, in the nine-to-ten-feature comparison. This does not necessarily mean the features were poorly selected, nor that the nine-feature state is universally sufficient. It means that, in this particular environment and with this wealth-change objective, additional market information produced no measurable improvement. The later

risk-aware experiments show the behaviour changing only when the objective itself included drawdown, and mainly through reduced market participation.

5.11.3 Why the Algorithms Ranked the Way They Did

One of the clearest differences between the algorithms is behavioural rather than financial. REINFORCE and PPO often produced policies that were classified as state-independent, so their relatively strong real-data returns need careful interpretation: a policy that buys whenever possible can perform well in rising markets without using the information provided to it. Deep Q-learning was state-dependent in all six seeds, yet earned the least on the real test months.

All three methods condition on the same state, maximise the same reward, and see the same data, so the difference between them cannot be that one “uses the state” and the others do not. The difference lies in what each method must estimate. A value-based method has to estimate the value of every state it acts in, which takes many visits to similar states under a stable distribution. Sixty episodes from a shifting market cannot support those estimates, so Deep Q-learning conditioned on values that were largely noise. The policy-gradient methods do not estimate values at all; they only need a policy that scores well on the training months, and on upward-leaning data the always-buy policy is that answer at zero cost.

The simulated-data study confirms this account by inverting the ranking. Given 3,000 balanced episodes and a stable generating distribution, the value estimates become reliable and Deep Q-learning wins on every seed, while the policy-gradient methods lose their free always-buy answer and collapse to corner policies. The ranking on either dataset alone would mislead; together they show that the ranking measures the data, not the algorithms.

5.11.4 Lower Risk Was Mainly Associated with Lower Exposure

The results show a consistent relationship between market participation and risk; specifically, risk was reduced mainly by participating less in the market.

Deep Q-learning achieved lower drawdown than buy-and-hold, but this was due to lower exposure, and overall, it led to lower returns in the same proportion.

This means lower drawdown should not be interpreted as evidence that DQN avoided the worst market movements, as this differs from identifying and avoiding specific periods of downside markets.

A proper risk-management strategy would ideally remain invested during rising market periods while reducing exposure during market downsides once detected. The current results do not provide enough evidence to conclude that the agent learned this type of selective risk management.

5.11.5 The Trade Size Affects Performance

The trade size experiment shows that the transaction size is part of the problem definition rather than merely an implementation design. A small trading size limits how much exposure an agent can build during a monthly episode, so poor performance may stem from the trade size rather than the quality of the

learned policy. The calibration in Section 5.3 addressed this by selecting the $0.25C_0$ size, which brings the attainable exposure close to the benchmark's. Notably, changing the trade size had a larger effect on performance than adding features to the state, which supports calibrating it before working on the main learned strategies.

5.11.6 Transaction Costs Provided a Useful Behavioural Test

The transaction-fee experiment showed that Deep Q-learning reduced its trading activity by approximately 83% as the fee rose from zero to 0.50%, while REINFORCE did not show the same response. This provides further evidence that Deep Q-learning was making greater use of the information available through the environment, although this additional responsiveness did not result in better financial performance.

5.11.7 Variation Between Seeds Is Important

This experiment showed considerable variation between random seeds, especially for the policy-gradient methods.

In some configurations, different seeds produced noticeably different behaviours. Some policies turned towards high exposure, while others traded less or became more conservative.

A seed producing a strong return would not be enough to conclude that the algorithm had learned a reliable trading strategy. Consistent behaviour across multiple seeds proves that, and becomes stronger evidence that, the observed result is derived from that algorithm rather than a product of one training run.

5.11.8 Implications for the Research Objectives

Overall, the two studies together answer the main research question of whether reinforcement learning algorithms can learn profitable, state-dependent trading strategies for a portfolio (in this case, the SPY asset).

On real data at the volumes a single asset provides, they cannot. Sixty training months support memorisation, not learning, and the agents either memorised the majority answer (always buy) or conditioned on value estimates too noisy to pay. No agent beat buy-and-hold on the held-out months.

Given enough balanced data, they can. On 3,000 simulated episodes, Deep Q-learning earned \$126.94 per episode on a test set where every fixed strategy loses, it did so on all six seeds, and the learned behaviour is the one the objectives asked for: near-full participation in rising markets and an 84% smaller loss in falling ones. The capability was never absent from the method. It was absent from the data.

The algorithm comparison also resolves differently than the real-data study suggested. The policy-gradient methods, REINFORCE and PPO, proved fragile in both settings, but for opposite surface reasons: on imbalanced real data they looked strong by adopting the majority action, and on balanced simulated data they collapsed into abstention or fixed buying. One REINFORCE seed in six found the

conditioned policy, which shows the optimum is reachable but not reliably reached. Value-based learning with sustained exploration was the only approach that converted more data into a better policy on every seed.

The risk-aware reward findings from the real-data study are unchanged: the penalty reduced drawdown mainly by reducing participation, and a large penalty stopped trading altogether.

These findings replace the earlier working conclusion. The binding constraint in this problem was not the reward function, the state design, or the trade size, although each mattered at the margin. It was the amount and balance of training data, and once that constraint is lifted, the remaining differences between algorithms are differences in how they explore and what they must estimate.

5.12 Chapter Summary

The chapter first determines the final trade size because this affects how quickly an agent can build or reduce its position. It then examines whether the learned strategies rely on the information from the state representations, followed by a comparison between the main nine-feature state and the additional risk-aware state. It then presents the real-data test results, the risk-aware reward experiments, and the transaction-cost experiments. The chapter then diagnoses the data bottleneck behind those results, reports the simulated-data study built to remove it, and closes by relating both studies to the research objectives.

The real-data results show that none of the learned agents consistently outperforms the buy-and-hold benchmark over the held-out dataset, and the diagnosis attributes this to 60 training episodes feeding networks of 18,000 parameters under a strong upward class imbalance. The simulated-data study removes those conditions and the picture changes: Deep Q-learning earns \$126.94 per episode on a balanced test set where no fixed trading pattern is profitable, on all six seeds, and transfers profitably (\$130.61 per month) to the real test months. The policy-gradient methods collapse to fixed corner policies once always-buying stops paying. The conclusion that survives both studies is that the data, not the algorithms, set the limit on what was learnable, and that value-based learning was the only method able to exploit more data once it was available.

Chapter 6

Conclusion and Future Work

6.1 Summary of the Work

This dissertation examined whether reinforcement learning algorithms can learn profitable trading policies that utilize market and portfolio information from the observed states when trading assets in a portfolio.

The study designed a complete experimental system for trading the SPY exchange-traded fund. The system uses daily market data organised into monthly episodes, includes transaction costs and action masking, and uses a nine-feature state representation with an optional drawdown feature. Two reward formulations were tested: the primary wealth-change reward and an optional risk-aware reward that penalised drawdown, along with three learning algorithms: REINFORCE, Deep Q-learning (DQN), and Proximal Policy Optimisation (PPO).

Experimental decisions were made on the 12-month validation dataset (2023), while keeping the 24-month test set (2024-2025) separate and using it only for final evaluations.

The experiments followed a fixed sequence. First, the trade size was adjusted to ensure the agents could reach a reasonable level of exposure relative to the buy-and-hold benchmark. The behavioural experiment then determined whether the learned policies actually respond to the information in their states. Once these environment, calibration and behavioural choices had been decided, the main performance of the learned agents was evaluated and compared against the benchmark strategy on the unseen test period.

Additional experiments then tested whether adding drawdown to the state changed the results, whether a risk-aware reward improved loss management in the portfolio trade, and how the different learning methods responded to changes in transaction costs.

The real-data results then motivated a second study. The failures of the first were diagnosed as a data problem: 60 training episodes cannot train networks of 18,000 parameters to do anything except memorise. A regime-switching market simulator was calibrated on the training years, 3,000 balanced episodes were generated, the network was reduced to about 2,800 parameters, and all three algorithms were re-trained and evaluated under controlled rising, falling, and flat conditions, with the real test months kept as a transfer test.

6.2 Conclusions Against the Research Questions

The experiments in Chapter 5 address the five questions asked in 5.1.

6.2.1 Performance

The first question considered whether the reinforcement learning agents could outperform the buy-and-hold benchmark on unseen data.

The results do not provide any evidence that they could. No learned agent outperformed the simple buy-and-hold strategy on the held-out test period. Buy-and-hold produced a mean wealth change of +\$180.25 per month, while the best learned agent, REINFORCE, achieved +\$171.51. PPO achieved +\$127.35 and DQN +\$74.68.

The result should not be interpreted as reinforcement learning algorithms being incapable of learning profitable trading strategies. All agents produced positive average wealth changes. Instead, the more specific conclusion is that none of the tested algorithms learned a level of trading skill that could outperform the simple investment strategy on unseen data.

Transaction costs and the trade size also cannot explain the difference, since they were the same throughout the experiment, providing a basis for fair comparison. The test period, however, favoured remaining invested, with positive market performance in 17 of the 24 months. Under these conditions, buy-and-hold was difficult to beat as it benefits from the rising market direction.

6.2.2 Behaviour

The second question examined whether the learned policies genuinely use the state information provided to them.

The algorithms showed a clear difference in how they used the state. All six DQN seeds changed their actions in response to the observed state, while none of the twelve policy-gradient seeds from REINFORCE and PPO did so. This changes how the performance results should be interpreted. REINFORCE came close to the benchmark in terms of return, at +\$171.51 per month, but its policies largely shifted to buying rather than using the market information provided.

DQN however produced policies that consistently responded to changes in the observed state but produced a lower average return than REINFORCE. This suggests that state dependence and financial performance are not to be treated as equal objectives in the experimental setting. An agent can learn to use its observations without that behaviour necessarily yielding higher returns.

This differentiation is important because an agent can perform well for the wrong reason. During a generally rising market, a policy that remains invested can achieve a strong return without needing to learn when market conditions require entering, reducing, or exiting a position.

Looking only at return would have completely hidden the distinction between these behaviours. The behavioural analysis is therefore an important part of evaluating learned trading policies, as it revealed a clear difference between the value-based and policy-gradient methods.

6.2.3 State Representation

The third question examined whether adding a drawdown feature to the state representation improved agent performance or behaviour.

Adding drawdown to the nine-feature state produced only limited changes using the wealth-change reward. REINFORCE's mean monthly wealth change decreased by \$4.12, while DQN's increased by \$3.13. Both changes were small compared with the variation between their random seeds and cannot be reasonably treated as reliable improvements.

PPO's mean return increased by \$32.62, but its seed spread also changed heavily, from \$179.80 to \$94.00. This improvement cannot confidently be attributed to the drawdown feature itself.

Overall, the results suggest that providing additional information does not automatically produce a better policy or improve learning unless the objective specifies it. This is consistent with an important principle noticed from the experiments: **information in the state is only useful when the learning objective and optimisation process give the agent a reason to use it**. Hence, the agent must be able to learn a relationship between the additional information and actions that improve the objective being optimised.

6.2.4 Reward Formulation

The fourth question considered whether adding drawdown into the reward could encourage the agents to reduce losses and generate a more risk-aware behaviour.

The risk-aware reward reduced drawdown mainly by reducing market participation, not by improving understanding of market conditions. With the selected penalty of $\lambda = 0.10$, DQN's mean drawdown fell by 41%, from 0.0140 to 0.0082. However, its mean return also fell by 32%, from +\$77.81 to +\$52.90, while its exposure was approximately halved. Its Sharpe ratio fell from 0.505 to 0.416, so the reduced drawdown did not result in better risk-adjusted performance. The reward therefore encouraged the agent to hold less of the asset rather than to identify and avoid declining markets. This is also reflected in March 2025, when reducing exposure would have been particularly useful: DQN's exposure during the month was higher than its overall average, and it still experienced 46% of the benchmark's loss.

The result illustrates an important limitation of using drawdown as a standalone penalty. The easiest way for an agent to reduce exposure to drawdown is to reduce its participation in the market. This behaviour can theoretically satisfy the risk conditions of the reward without necessarily producing a better ability to identify when the market is likely to decline.

The penalty-weight analysis proves this interpretation. A relatively small increase in the penalty produced a major behavioural change, with exposure falling from 0.344 at $\lambda = 0.10$ to 0.074 at $\lambda = 0.25$. This indicates that the reward formulation created a strong correlation between market participation and risk avoidance.

6.2.5 Data

The fifth question asked whether the amount and balance of the training data explain the learned behaviour, and whether fixing the data fixes the behaviour. The simulated-data study answers yes to both.

On 3,000 balanced simulated episodes, Deep Q-learning earned +\$126.94 per episode across six seeds (standard deviation \$7.30) on a test set where buy-and-hold loses \$60.29 and doing nothing earns \$0.00. All six seeds were profitable and all six were state-dependent. The per-regime results show the behaviour the real-data study looked for: the agents captured 99% of the profit available in rising episodes and cut the falling-episode loss by 84% relative to always buying. The same policies remained profitable when transferred unchanged to the real 2024–2025 months, earning +\$130.61 per month against buy-and-hold's +\$180.25; the shortfall is the cost of defensive behaviour in a period that never fell.

The study also explains the real-data algorithm ranking. Deep Q-learning was worst on real data not because it is unsuited to trading but because value estimation needs many visits to similar states, which 60 shifting months cannot provide. Once the data supported the estimates, the ranking inverted. REINFORCE and PPO, which looked strong on real data by adopting the majority always-buy action, collapsed to fixed corner policies on balanced data, where no fixed action is profitable. One REINFORCE seed in six did learn the conditioned policy, so the optimum is within reach of policy-gradient methods, but the training process does not find it reliably.

The overall conclusion of the dissertation therefore has two halves. At real single-asset data volumes, none of the tested algorithms learns a genuine trading policy. With enough balanced data, value-based learning does, on every seed, and behaves sensibly in every market condition the simulator can produce.

6.3 Critical Appraisal of the Project

The most important lessons from this project were from the experimental methodology rather than the choice of reinforcement learning algorithm alone. Several of the most useful findings came from implementation problems identified during development and testing.

One of the strongest aspects of the project was the verification tests. The automated tests identified several implementation problems during development, **three of which could have affected the experiments' validity**. These included issues with rolling feature calculation at the end of monthly episodes and an implementation error in the buy-and-hold benchmark that reduced its performance by \$38.95 per month, about 21.6%. If this error had remained, the dissertation's main conclusion could have been misleading, as the learned agents would have appeared to outperform the benchmark.

The verification tests also noticed the experimental design problem in the trade size. Using the original $0.10C_0$ trade size in Chapter 3 would make the maximum attainable exposure only 0.690, compared with 0.910 for buy-and-hold. This meant that the learning agents were at a disadvantage before training even began. Selecting $0.25C_0$ allowed a fairer comparison while still letting the agents control their positions in smaller steps than the $0.50C_0$ alternative.

The complete set of 21 verification tests passed before the final experiments were accepted, giving a

higher confidence that the reported results reflect an appropriate experimental design rather than implementation errors.

The behavioural analysis was another important part of the study. Without it, the REINFORCE result of +\$171.51 could have been interpreted as evidence that the agent had learned an intelligent, effective trading strategy. The analysis showed this interpretation was not justified, as the policies acted like a consistent buyer because the market was consistently rising, completely changing the interpretation of the result.

The project also highlighted the importance of random seed variation. The results, particularly for PPO, varied considerably between independent training runs. PPO had a seed spread of \$179.80 in one configuration, which was greater than its mean return. This made it difficult to draw conclusions based on individual training runs, as that would have been unreliable. Reporting multiple seeds widened the evaluation, providing a more consistent view of how each algorithm behaved, although it increased the computational and experimental workload; however, six seeds are still lacking for more definite statistical inference.

The largest single lesson, however, was about data. The project spent most of its effort on the state representation, the reward, the trade size, and the verification of the environment, and all of that effort was necessary; none of it was sufficient. The binding constraint was 60 training episodes feeding 18,000 parameters, and no amount of feature or reward engineering could compensate for it. Diagnosing that constraint and rebuilding the study around simulated data changed the results more than every other design decision combined. The lesson for similar projects is to budget the data before budgeting the model: count the episodes, count the parameters, and if the second number is larger, change the data or the model before running anything else.

Overall, the project demonstrated that careful experimental design is as important as the choice of reinforcement learning algorithm. The main lessons came from testing the assumptions behind the experiments, rather than from simply comparing model returns. State dependence, action constraints, reward design, seed variation, and above all data volume affect how the final results could be interpreted. This made the experimental process more demanding but also produced a more reliable basis for evaluating the learned trading policies.

6.4 Limitations

The findings of this study should be considered within the scope of the experimental design, but there are several limitations.

6.4.1 One asset and a limited market period

The experiments use SPY as the only asset, and the real held-out period from 2024 to 2025 was generally favourable to holding the asset due to rising markets. The simulated-data study addresses the market-condition half of this limitation by testing under controlled falling and flat regimes, but only for simulated

prices; no real declining period was available for testing. The findings cannot therefore be automatically generalised to other assets or to real bear markets without further testing.

6.4.2 Simulator realism and the domain gap

The market generator draws each day's return independently within a regime. Real returns are not independent: volatility clusters, trends persist, and regimes change mid-month. The generator also holds one regime for a whole episode, so the agents never face a mid-month crash. These simplifications are the price of a controlled experiment, and the transfer test measures their cost directly: the Deep Q-learning agents kept roughly three-quarters of buy-and-hold's real-data return while being trained on simulated data alone. Conclusions about absolute profitability on real markets should therefore rest on the transfer numbers, not the simulated ones.

6.4.3 Policy-gradient training was not separately tuned

REINFORCE and PPO collapsed to corner policies on the balanced simulated data. Both were trained with the same hyperparameters as the real-data study, including PPO's default entropy coefficient of zero. Entropy regularisation or other exploration incentives might let the policy-gradient methods find the conditioned policy more reliably, as one REINFORCE seed shows it exists. The study reports the collapse as observed behaviour under shared settings rather than as a bound on what tuned policy-gradient methods could achieve.

6.4.4 Monthly Episode Structure

The trading environment treated each calendar month as a separate episode. Every episode started with the same initial capital of \$10,000, and any remaining position was closed at the end of the month. As a result, the results do not represent a single portfolio that continuously compounds over the full evaluation period as a real portfolio would.

This design made the experiments easier to control by keeping the episode length and starting conditions consistent. However, it also means that gains or losses from one month did not affect the starting capital or portfolio position in the following month, which does not represent a real-world portfolio. A longer continuous portfolio environment could therefore lead to a different behaviour, especially for policies that require accumulated wealth or changes in portfolio value over time.

6.4.5 Fixed Trade Size

The agents trade fixed fractions of the initial capital rather than choosing their position size continuously. This was calibrated to create a controlled, reasonable environment for comparison with the benchmark, but it restricts how much the agents can express their decisions. A continuous position-sizing model could potentially respond differently based on its exposure in a particular market condition. Rather than repeatedly applying fixed buy or sell actions, such an agent could select its desired portfolio exposure, producing different results in returns and trading behaviour.

6.4.6 Limited Number of Seeds and Lack of Significance Testing

Six seeds were used for each configuration, and the seed-level results and variation between them are reported throughout the results. This is stronger than relying on a single training run, but it does not provide sufficient statistics for formal claims about small differences. For example, the \$3.13 change observed for DQN and **-\$4.12** for REINFORCE when the state representation was changed is not large enough to be treated as a meaningful improvement.

6.4.7 Risk Penalty Selection was based on DQN

The reward's penalty weight λ was selected using DQN because it was the algorithm that consistently showed state-dependent behaviour. This does not guarantee that the same penalty value is appropriate for REINFORCE or PPO, as it was not independently optimised for every algorithm. A separate calibration for each algorithm could produce different results. This limitation is particularly relevant because the algorithms responded differently to the state and reward formulation.

6.4.8 The risk penalty only targets drawdown

The risk-aware reward penalises drawdown but does not directly measure whether the return achieved is sufficient for the level of risk taken. While drawdown is directly relevant to losses, it does not fully represent the portfolio risk. A reward that relates return to risk, rather than penalizing risk alone, was not tested.

Additionally, as shown in Section 5.7, the agents could reduce the penalty simply by reducing their exposure to the asset or not trading at all. The experiments showed that a lower drawdown value should not automatically be associated with evidence of better risk management.

6.5 Future Work

The limitations above suggest several directions for furthering the research. Each stems directly from a limitation or an unresolved result in the experiments. The behavioural analysis should also be applied to these future works to determine the agent's behaviour to changes alongside performance improvements.

6.5.1 Richer Market Generators

The simulated-data study used the simplest generator that supports a controlled experiment. The natural next step is a generator that keeps the control but adds the structure the current one omits: regime switches within an episode, driven by a hidden Markov chain, and volatility clustering within a regime. Learned generators such as time-series GANs offer more realism again, at the cost of transparency. Each step up in realism should shrink the domain gap that the transfer test measures, and the transfer test itself provides the yardstick for whether it does.

6.5.2 Longer and More Challenging Real Market Periods

The real-data transfer test could be extended over longer periods containing genuine declines, such as 2022 or the COVID-19 crash, using a walk-forward approach in which the simulator is recalibrated and the agents retrained as conditions change. The controlled study shows the agents defend well in simulated declines; a real decline is the missing test of that behaviour.

6.5.3 Predictive and Uncertainty-Aware State Information

The drawdown feature describes the portfolio's previous losses but does not provide information about future market movements. Future work could therefore include predictive information, such as probabilistic forecasts containing estimates of forecast uncertainty.

This would give a state-dependent agent more useful information for making decisions about future market conditions rather than relying only on historical data. It would also directly test whether DQN's ability to use the state leads to more useful market understanding and trading decisions.

6.5.4 Risk-Efficiency Objectives

The most direct correction to the drawdown penalty could be to replace it with a reward that considers return relative to risk, such as a Sharpe-based or drawdown-relative objective [2]. This would make simply reducing market exposure a less effective way of improving the reward, as seen in the current experiments.

The results in Section 5.7 make this future rework highly relevant because the tested penalty mainly caused agents to reduce exposure rather than learn when downside risk is increasing.

6.5.5 More Seeds for Better Statistical Analysis

More random training seeds would improve the reliability of comparisons [30] between algorithms and configurations. Future work should combine this with appropriate statistical analysis to determine whether observed differences between configurations are likely to represent genuine results rather than random noise.

6.5.6 More flexible trade sizes

Future experiments could allow the agent to choose its position size continuously rather than using fixed trade amounts. This would remove the exposure limit created by the current environment and allow the agent to express different levels of confidence by adjusting its position size, potentially enabling smoother portfolio adjustments.

6.6 Closing Remarks

This dissertation examined whether standard reinforcement learning could learn a genuine trading policy for a single asset. On real data it could not: the best-performing learned agent finished \$8.74 per

month below buy-and-hold, and it did so by ignoring its market observations to follow a simple trading behaviour.

The study's contribution is what it did with that failure. The behavioural analysis showed the strong-looking returns were not evidence of learning; the diagnosis traced the failure to 60 training episodes feeding 18,000 parameters; and the simulated-data study removed that constraint and showed the same algorithms succeeding, with Deep Q-learning profitable and state-dependent on every seed and its policies transferring, still profitable, to real unseen months. The negative result and the positive one are the same finding viewed from two sides: what these methods learn is decided by the data they are given.

The implementation, the market simulator, the experimental framework and the behavioural analysis developed in this study provide a useful basis for future work. More advanced models and richer generators can build on this framework while being evaluated against the same fair standards.

References

- [1] R. S. Sutton and A. G. Barto, *Reinforcement Learning: An Introduction*, 2nd ed. MIT Press, 2018.
- [2] J. Moody and M. Saffell, “Learning to trade via direct reinforcement,” *IEEE Transactions on Neural Networks*, vol. 12, no. 4, pp. 875–889, 2001.
- [3] Y. Deng, F. Bao, Y. Kong, Z. Ren, and Q. Dai, “Deep direct reinforcement learning for financial signal representation and trading,” *IEEE Transactions on Neural Networks and Learning Systems*, vol. 28, no. 3, pp. 653–664, 2017.
- [4] T. G. Fischer, “Reinforcement learning in financial markets — a survey,” FAU Discussion Papers in Economics, Friedrich-Alexander University Erlangen-Nürnberg, Tech. Rep. 12/2018, 2018.
- [5] R. J. Williams, “Simple statistical gradient-following algorithms for connectionist reinforcement learning,” *Machine Learning*, vol. 8, no. 3–4, pp. 229–256, 1992.
- [6] V. Mnih, K. Kavukcuoglu, D. Silver, A. A. Rusu, J. Veness, M. G. Bellemare, A. Graves, M. Riedmiller, A. K. Fidjeland, G. Ostrovski, S. Petersen, C. Beattie, A. Sadik, I. Antonoglou, H. King, D. Kumaran, D. Wierstra, S. Legg, and D. Hassabis, “Human-level control through deep reinforcement learning,” *Nature*, vol. 518, no. 7540, pp. 529–533, 2015.
- [7] J. Schulman, F. Wolski, P. Dhariwal, A. Radford, and O. Klimov, “Proximal policy optimization algorithms,” *arXiv preprint arXiv:1707.06347*, 2017.
- [8] G. Tesauro, “Temporal difference learning and TD-Gammon,” *Communications of the ACM*, vol. 38, no. 3, pp. 58–68, 1995.
- [9] R. Bellman, *Dynamic Programming*. Princeton University Press, 1957.
- [10] R. S. Sutton, “Learning to predict by the methods of temporal differences,” *Machine Learning*, vol. 3, no. 1, pp. 9–44, 1988.
- [11] C. J. C. H. Watkins and P. Dayan, “Q-learning,” *Machine Learning*, vol. 8, no. 3–4, pp. 279–292, 1992.
- [12] J. N. Tsitsiklis and B. Van Roy, “An analysis of temporal-difference learning with function approximation,” *IEEE Transactions on Automatic Control*, vol. 42, no. 5, pp. 674–690, 1997.
- [13] L.-J. Lin, “Self-improving reactive agents based on reinforcement learning, planning and teaching,” *Machine Learning*, vol. 8, no. 3–4, pp. 293–321, 1992.
- [14] H. van Hasselt, A. Guez, and D. Silver, “Deep reinforcement learning with double Q-learning,” in *Proceedings of the 30th AAAI Conference on Artificial Intelligence*, 2016, pp. 2094–2100.
- [15] V. R. Konda and J. N. Tsitsiklis, “Actor-critic algorithms,” in *Advances in Neural Information Processing Systems 12*, 2000, pp. 1008–1014.

- [16] V. Mnih, A. P. Badia, M. Mirza, A. Graves, T. Lillicrap, T. Harley, D. Silver, and K. Kavukcuoglu, “Asynchronous methods for deep reinforcement learning,” in *Proceedings of the 33rd International Conference on Machine Learning (ICML)*, 2016, pp. 1928–1937.
- [17] J. Schulman, P. Moritz, S. Levine, M. Jordan, and P. Abbeel, “High-dimensional continuous control using generalized advantage estimation,” in *Proceedings of the 4th International Conference on Learning Representations (ICLR)*, 2016.
- [18] J. Schulman, S. Levine, P. Abbeel, M. Jordan, and P. Moritz, “Trust region policy optimization,” in *Proceedings of the 32nd International Conference on Machine Learning (ICML)*, 2015, pp. 1889–1897.
- [19] X.-Y. Liu, H. Yang, Q. Chen, R. Zhang, L. Yang, B. Xiao, and C. D. Wang, “FinRL: A deep reinforcement learning library for automated stock trading in quantitative finance,” *arXiv preprint arXiv:2011.09607*, 2020.
- [20] S. Huang and S. Ontañón, “A closer look at invalid action masking in policy gradient algorithms,” in *Proceedings of the 35th International FLAIRS Conference*, 2022.
- [21] Z. Jiang, D. Xu, and J. Liang, “A deep reinforcement learning framework for the financial portfolio management problem,” *arXiv preprint arXiv:1706.10059*, 2017.
- [22] T. Théate and D. Ernst, “An application of deep reinforcement learning to algorithmic trading,” *Expert Systems with Applications*, vol. 173, p. 114632, 2021.
- [23] M. Guan and X.-Y. Liu, “Explainable deep reinforcement learning for portfolio management: An empirical approach,” in *Proceedings of the 2nd ACM International Conference on AI in Finance (ICAIF)*, 2021.
- [24] D. Silver, A. Huang, C. J. Maddison, A. Guez, L. Sifre, G. van den Driessche, J. Schrittwieser, I. Antonoglou, V. Panneershelvam, M. Lanctot, S. Dieleman, D. Grewe, J. Nham, N. Kalchbrenner, I. Sutskever, T. Lillicrap, M. Leach, K. Kavukcuoglu, T. Graepel, and D. Hassabis, “Mastering the game of Go with deep neural networks and tree search,” *Nature*, vol. 529, no. 7587, pp. 484–489, 2016.
- [25] G. Brockman, V. Cheung, L. Pettersson, J. Schneider, J. Schulman, J. Tang, and W. Zaremba, “OpenAI Gym,” *arXiv preprint arXiv:1606.01540*, 2016.
- [26] L. R. Rabiner, “A tutorial on hidden Markov models and selected applications in speech recognition,” *Proceedings of the IEEE*, vol. 77, no. 2, pp. 257–286, 1989.
- [27] J. Yoon, D. Jarrett, and M. van der Schaar, “Time-series generative adversarial networks,” in *Advances in Neural Information Processing Systems*, vol. 32, 2019.
- [28] D. H. Bailey, J. M. Borwein, M. L. de Prado, and Q. J. Zhu, “Pseudo-mathematics and financial charlatanism: The effects of backtest overfitting on out-of-sample performance,” *Notices of the American Mathematical Society*, vol. 61, no. 5, pp. 458–471, 2014.

- [29] M. L. de Prado, *Advances in Financial Machine Learning*. Wiley, 2018.
- [30] P. Henderson, R. Islam, P. Bachman, J. Pineau, D. Precup, and D. Meger, “Deep reinforcement learning that matters,” in *Proceedings of the 32nd AAAI Conference on Artificial Intelligence*, 2018, pp. 3207–3214.
- [31] E. F. Fama, “Efficient capital markets: A review of theory and empirical work,” *The Journal of Finance*, vol. 25, no. 2, pp. 383–417, 1970.
- [32] H. Markowitz, “Portfolio selection,” *The Journal of Finance*, vol. 7, no. 1, pp. 77–91, 1952.
- [33] W. F. Sharpe, “The Sharpe ratio,” *The Journal of Portfolio Management*, vol. 21, no. 1, pp. 49–58, 1994.
- [34] D. P. Kingma and J. Ba, “Adam: A method for stochastic optimization,” in *Proceedings of the 3rd International Conference on Learning Representations (ICLR)*, 2015.
- [35] A. Raffin, A. Hill, A. Gleave, A. Kanervisto, M. Ernestus, and N. Dormann, “Stable-baselines3: Reliable reinforcement learning implementations,” *Journal of Machine Learning Research*, vol. 22, no. 268, pp. 1–8, 2021.
- [36] A. Paszke, S. Gross, F. Massa, A. Lerer, J. Bradbury, G. Chanan, T. Killeen, Z. Lin, N. Gimeshein, L. Antiga, A. Desmaison, A. Köpf, E. Yang, Z. DeVito, M. Raison, A. Tejani, S. Chilamkurthy, B. Steiner, L. Fang, J. Bai, and S. Chintala, “PyTorch: An imperative style, high-performance deep learning library,” in *Advances in Neural Information Processing Systems 32 (NeurIPS)*, 2019.
- [37] M. Towers, A. Kwiatkowski, J. Terry, J. U. Balis, G. De Cola, T. Deleu, M. Goulão, A. Kallinteris, M. Krimmel, A. KG *et al.*, “Gymnasium: A standard interface for reinforcement learning environments,” *arXiv preprint arXiv:2407.17032*, 2024.

Appendix A

Verification Tests

This appendix lists the 21 automated verification tests referenced in Section 4.5. Every test was run and passed before the final experiments in Chapter 5 were accepted. The tests are grouped into the same four groups as Section 4.5. Each entry gives the test and the property it checks. The test names are shown in the dissertation’s current terminology; a few of the underlying functions in the code keep older working names from development, such as “ladder” for the staged state configurations and “slice” for the fixed trade amount.

Environment and portfolio tests

Test	Property checked
episode always ends flat	Any remaining position is sold on the final trading day, leaving no units held and no exposure.
wealth can be calculated from the current state	The portfolio wealth can be calculated from the cash and exposure features at every step. This is why wealth is not included in the state as a separate feature.
drawdown cannot be calculated from the current state alone	Two price paths can arrive at the same cash, exposure, and wealth but have different drawdown values, because drawdown depends on the highest wealth reached earlier. This is why drawdown must be added to the state as a separate feature.
share action model buys one unit	The implementation also supports an alternative action model in which a buy purchases exactly one whole unit of the asset. This test confirms that model behaves as defined; the main experiments use the fixed-amount model below.
trade action model spends a fixed sum	A buy spends exactly the configured fraction of the initial capital, excluding the fee.
true buy-and-hold beats gradual buying when prices rise	The corrected buy-and-hold benchmark invests once at the start of the month and earns more than gradual buying in a rising market. This is the error that led to the benchmark correction described in Chapter 4.

Test	Property checked
a month with no trades leaves cash unchanged	If no trades are made during a month, the wealth is unchanged and no fees are charged. This checks the environment's accounting when the agent stays out of the market.

Feature and data tests

Test	Property checked
features use no future data	Removing the later prices from the series does not change any earlier feature value. This confirms that no feature is calculated using future information.
rolling features survive the episode boundary	Momentum on the first day of a month is calculated from the actual previous trading days, reaching back into the previous month, rather than restarting from zero.
no NaNs or infinities in any episode	Every feature value in every episode is a finite number, and every price is positive.
canonical order is independent of input order	Requesting the same features in a different order still produces the same state layout.
state dimensions are as documented	Each state configuration contains exactly the number of features reported in Chapter 3: nine for the main state and ten when drawdown is added.
every state configuration is a subset of the next	Each larger state configuration contains all the features of the smaller ones. This means a comparison between two states isolates the effect of the added features.
split is chronological and disjoint	The training, validation, and test periods are ordered in time and do not share any months.

Reward and risk tests

Test	Property checked
accounting identity holds for every policy	The sum of all daily rewards in a month equals the month's true change in wealth. This holds for every policy and state configuration tested.

Test	Property checked
reward equals wealth change when $\lambda = 0$	When the risk penalty is switched off, the risk-aware reward gives exactly the same values as the wealth-change reward.
risk penalty requires drawdown in the state	Configuring the risk-aware reward without the draw-down feature in the state raises an error. This enforces that the reward can be calculated from the information the agent observes.
risk penalty never increases the reward	For any sequence of actions, the penalised reward is never larger than the unpenalised reward.

Action-masking tests

Test	Property checked
mask blocks infeasible trades	Selling is not allowed when no position is held, and buying is not allowed when there is not enough cash. Selling becomes available while a position is held.
illegal action is forced to hold and flagged	If an invalid action is supplied from outside the environment, it is executed as Hold, flagged in the step information, and no trade is recorded.
mask is a deterministic function of observed features	The legal-action mask can be predicted exactly from the cash and exposure features. This means masking constrains the agent without giving it extra information.

Appendix B

Algorithm Listings

This appendix lists the three learning procedures used in this dissertation as they were implemented. REINFORCE and Deep Q-learning were built from scratch, while MaskablePPO is the library implementation used for comparison. All three apply the same action mask and operate in the same trading environment, as described in Chapter 3.

Algorithm 1: REINFORCE for the trading MDP (as implemented)

```
1: initialise policy network  $\pi\theta$ 
2: for iteration = 1, 2, ... until step budget exhausted do
3:   sample a training month; reset environment ( $C_0$  cash, no position)
4:   for  $t = 0, \dots, T - 1$  do
5:     compute legal-action mask; set illegal logits to  $-\infty$ 
6:     sample  $a_t \sim \pi\theta(\cdot | s_t)$ ; execute; store  $(\log \pi\theta(a_t | s_t), r_t)$ 
7:   end for
8:   compute rewards-to-go  $G_t$ ; standardise within the episode
9:    $L(\theta) \leftarrow -\sum_t \log \pi\theta(a_t | s_t) G_t$ 
10:  one Adam step on  $\nabla \theta L$ 
11: end for
```

Algorithm 2: Deep Q-learning for the trading MDP (as implemented)

```
1: initialise  $Q\phi$ ; target  $Q\phi^- \leftarrow Q\phi$ ; replay buffer  $D$ 
2: for environment step = 1, 2, ... until step budget exhausted do
3:   with prob.  $\epsilon$ : random legal action; else  $a_t = \arg \max_a \text{legal } Q\phi(s_t, a)$ 
4:   execute  $a_t$ ; store  $(s_t, a_t, r_t, s_{t+1}, \text{mask}_{t+1}, d_t)$  in  $D$ 
5:   sample minibatch; form masked temporal-difference targets  $y$  (Section 3.8.5); Adam step on  $L(\phi)$ 
6:   every  $K$  steps:  $\phi^- \leftarrow \phi$ 
7: end for
```

Algorithm 3: MaskablePPO for the trading MDP (as implemented)

```

1: initialise actor  $\pi_\theta$  and critic  $V_\psi$  (MaskablePPO, MLP 128–128)
2: for iteration = 1, 2, ... until step budget exhausted do
3:   for step = 1, ...,  $N$  do
4:     if episode terminated: sample a training month; reset ( $C_0$  cash, no position)
5:     observe  $s_t$  and legal-action mask; sample  $a_t \sim \pi_\theta(\cdot | s_t)$  over legal actions
6:     execute; store  $(s_t, a_t, r_t, V_\psi(s_t), \log \pi_\theta(a_t | s_t))$ 
7:   end for
8:   estimate advantages  $\hat{A}_t$  by GAE; form returns
9:   for epoch = 1, ...,  $K$  do
10:    for each minibatch from the rollout do
11:       $\rho_t \leftarrow \pi_\theta(a_t | s_t) / \pi_{\theta\_old}(a_t | s_t)$ 
12:      Adam step on  $LCLIP(\theta) + \text{value loss} - \text{entropy bonus}$ 
13:    end for
14:  end for
15: end for

```

Notes. $N = 2048$ is the Stable-Baselines3 rollout length used here, and $K = 10$ is the number of training epochs over each rollout. Month sampling and action masking use the same environment interface as REINFORCE and DQN, so all three algorithms face the same action space at every step.

Appendix C

Per-Month Held-Out Test Results

This appendix reports the wealth change ΔW in dollars for each of the 24 held-out test months ($C_0 = \$10,000$, 0.25 C_0 trade size, 0.05% fee, wealth-change reward). The columns for the learned agents show the mean of the six training seeds, so they match the seed means reported in Chapter 5. The final row is the total over the 24 months.

Month	Buy-and-hold	Always-buy (0.25 C_0)	REINFORCE	DQN	PPO
2024-01	+206.23	+282.35	+271.95	+72.70	+196.83
2024-02	+375.56	+305.79	+290.50	+54.06	+219.51
2024-03	+220.74	+267.59	+249.89	+31.05	+187.59
2024-04	-396.08	-325.47	-311.74	-95.75	-233.48
2024-05	+529.49	+366.11	+350.33	+291.96	+266.14
2024-06	+334.03	+264.80	+243.42	+127.65	+190.45
2024-07	+90.21	+3.05	+9.27	+72.89	+5.79
2024-08	+370.28	+651.47	+626.02	+363.21	+449.74
2024-09	+414.15	+487.66	+464.38	+105.42	+342.37
2024-10	-9.64	-26.18	-54.57	+96.12	-17.86
2024-11	+541.25	+431.90	+414.48	+132.34	+310.49
2024-12	-267.80	-297.22	-297.55	-32.13	-209.30
2025-01	+283.57	+188.10	+203.74	+18.17	+137.22
2025-02	-69.99	-148.04	-126.82	+46.50	-101.61
2025-03	-398.41	-320.38	-302.86	-184.04	-230.19
2025-04	-124.50	+243.72	+297.84	+79.27	+157.30
2025-05	+543.11	+479.42	+414.44	+390.80	+342.25
2025-06	+444.52	+414.11	+408.93	+69.76	+294.59
2025-07	+223.40	+167.96	+150.83	+51.18	+121.28
2025-08	+364.88	+255.00	+215.99	+98.64	+185.21
2025-09	+423.09	+345.46	+328.23	+119.11	+247.93
2025-10	+193.41	+175.59	+189.11	-76.82	+125.12
2025-11	-9.26	+90.48	+73.42	-32.31	+59.94
2025-12	+43.83	+10.78	+7.10	-7.54	+9.01
Total	+4326.04	+4314.08	+4116.31	+1792.25	+3056.31

Appendix D

Glossary of Terms

This glossary explains the recurring technical terms of the dissertation in plain language. The definitions here are informal; the precise definitions are given in Chapter 3.

Agent The learning program that makes the trading decisions. In this dissertation, it is a small neural network that decides each day whether to hold, buy, or sell.

Environment The simulated world the agent interacts with. Here, it is a simulated trading month built from real historical prices. It executes the agent's trades, charges the transaction fees, and updates the portfolio.

State The information the agent observes before making a decision. Here, it is nine numbers covering recent market behaviour, the agent's own cash and holdings, and how far the month has progressed. A tenth number is added when the drawdown feature is used.

Action One of the choices available to the agent. There are exactly three: hold the current position, buy a fixed amount of the asset (a quarter of the initial capital, written $0.25 C_0$), or sell the same amount.

Reward The feedback the agent receives after each action. Here, it is the day's change in total wealth, with trading fees already deducted. The agent's objective is to make the month's total reward as large as possible.

Episode One complete round of the task, after which everything resets. Here, it is one calendar month of trading, ending with a forced sale of any remaining holdings.

Policy The agent's decision rule. Formally, it is a function from the state to action probabilities; informally, it describes what the agent would do in any situation.

Policy gradient / REINFORCE A family of methods that improve the policy directly. After each month, decisions that came before a good outcome are made more probable, and decisions that came before a poor outcome are made less probable. REINFORCE is the original and simplest member of this family.

Q-value / Deep Q-learning The value-based family of methods. The agent learns an estimate of how much future profit each action is worth in each situation, then selects the action with the highest estimate. "Deep" means the estimates come from a neural network rather than a lookup table.

PPO A more recent policy-gradient method that limits how much the policy can change in a single update. This makes training steadier, but the policy can still settle into the first behaviour that performs well.

Neural network A flexible mathematical function with many adjustable parameters. The parameters are tuned gradually during training so that the outputs improve against an objective.

Training / validation / test split The data is divided into three parts by date. The agent learns on the earliest part (2018–2022), every tuning decision is made using only the middle part (2023), and the final evaluation uses the latest part (2024–2025), which was not used for any design decision. This separation is what makes the test result an honest measure of performance on unseen data.

Benchmark A simple non-learning strategy the agents are compared against. The most important one here is *buy-and-hold*: invest everything at the start of the month and wait.

Buy-and-hold The standard passive strategy. Previous research shows it is genuinely difficult to beat using price history alone. This is why a result that beats it easily should be examined carefully rather than accepted immediately.

Always-buy ($0.25 C_0$) The strongest strategy available to a policy that ignores its observations: buy one fixed trade every day buying is allowed, until fully invested. It is used as the reference point for learned policies that reduce to unconditional buying.

Transaction fee / basis points (bps) The cost of trading, charged as a fraction of the traded amount. 5 bps = 0.05%, which is 50 cents per \$1,000 traded.

Sharpe ratio The return earned per unit of variation in returns. Higher is better. However, a strategy that barely invests can report a high Sharpe ratio simply because its wealth barely moves.

Drawdown The fall from the highest wealth reached so far in the episode, as a fraction of that peak. It measures the largest loss experienced during the episode, not the average one.

Exposure The proportion of the capital invested in the asset at a given moment, from 0% (all cash) to 100% (fully invested).

Seed The starting value of the random number generator. Running the same experiment with several seeds (six in this study) checks that a result does not depend on one particular random initialisation.

Overfitting Learning patterns that are specific to the training data rather than patterns that hold in general, which leads to poor performance on new data. The chronological split and the held-out test months guard against this.

Action masking Removing impossible actions from the agent’s available choices, such as selling with no position or buying with no cash. The agent can neither attempt nor learn from invalid actions.

State-dependent policy A policy whose actions change with the observed state even when the same actions are legal. The opposite is a mask-determined policy, which follows a fixed pattern such as buying whenever buying is allowed. In a rising market, that fixed pattern earns a strong return, which is why profit alone cannot reveal it. The behavioural analysis in Chapter 5 was designed to detect this difference.

Accounting identity The check that the sum of all daily rewards in a month equals the month's true change in wealth. It is verified in every experiment. If it failed, the rewards would not reflect the actual portfolio.

Markov property The requirement that the current state contains everything from the past that is relevant to the next decision. This is the reason drawdown had to be added to the state before it could be used in the reward.