


Actor-Critic architectural flow (Iteration 1)

How the MDP, policy π_θ and value estimator V_ϕ interact during training



Left path = decide & act · Right path = score how good that was · Bottom = learn

Matches Algorithm 1 (Iteration 1): discrete actions, binary holding, one stock